

# **CIS 631**

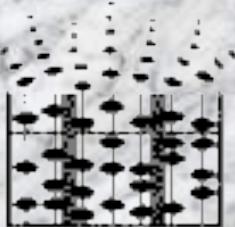
## *Parallel Processing*

### *Lecture 2: Parallel Architectures*

**Allen D. Malony**

[malony@cs.uoregon.edu](mailto:malony@cs.uoregon.edu)

Department of Computer and Information Science  
University of Oregon

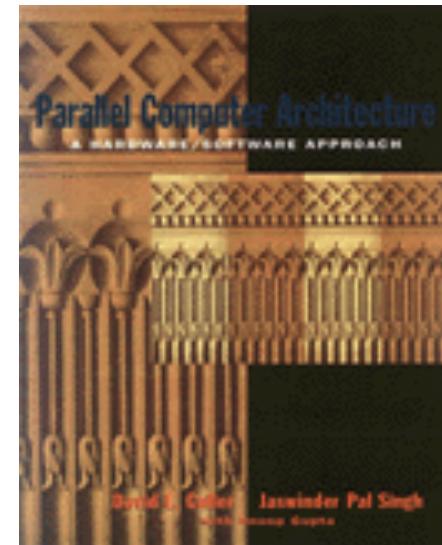


## *Acknowledgements*

- Portions of the lectures slides were adopted from:
  - David Patterson, CS252, University of California, Berkeley, 2001.
  - David E. Culler, CS 258, University of California, Berkeley, 1999.

<http://www.eecs.berkeley.edu/~culler/cs258-s99/>.

Parallel Computer Architecture:  
Hardware/Software Approach,  
D. Culler, J. Singh, A. Gupta,  
Morgan Kaufmann, 1998.



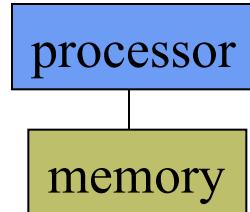
# *Outline*

- Parallel architecture types
- Instruction-level parallelism
- Vector processing
- SIMD
- Shared memory
  - Memory organization: UMA, NUMA
  - Coherency: CC-UMA, CC-NUMA
- Interconnection networks
- Distributed memory
- Clusters
- Clusters of SMPs
- Heterogeneous clusters of SMPs

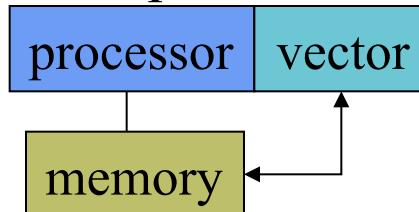
# Parallel Architecture Types

- Uniprocessor

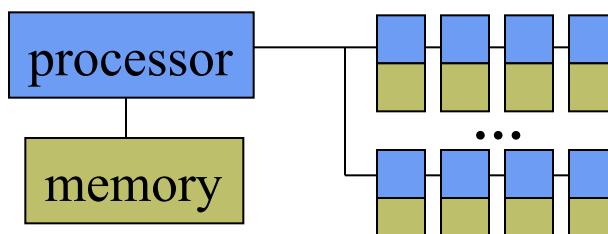
- Scalar processor



- Vector processor



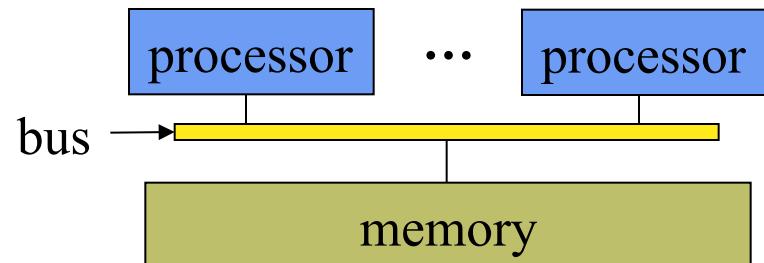
- Single Instruction Multiple Data (SIMD)



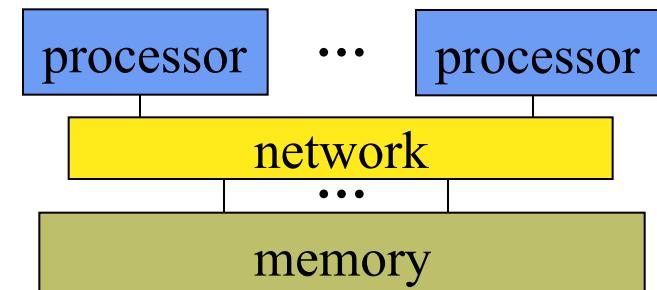
- Shared Memory Multiprocessor (SMP)

- Shared memory address space

- Bus-based memory system



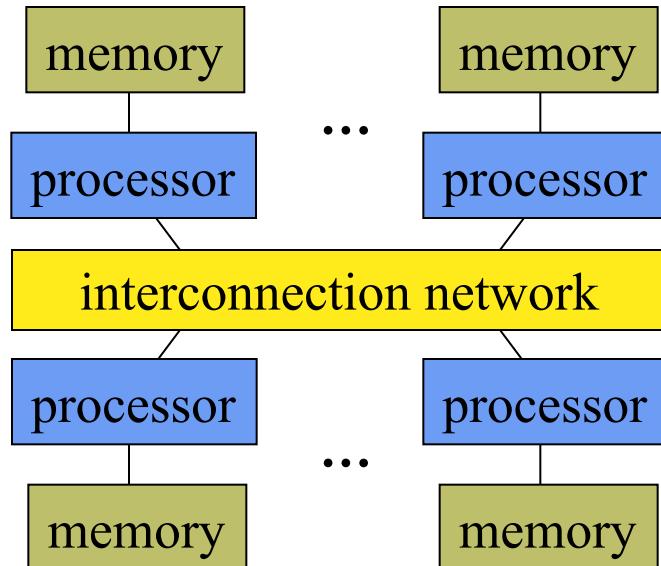
- Interconnection network



# *Introduction to Parallel Architectures*

- Distributed Memory Multiprocessor

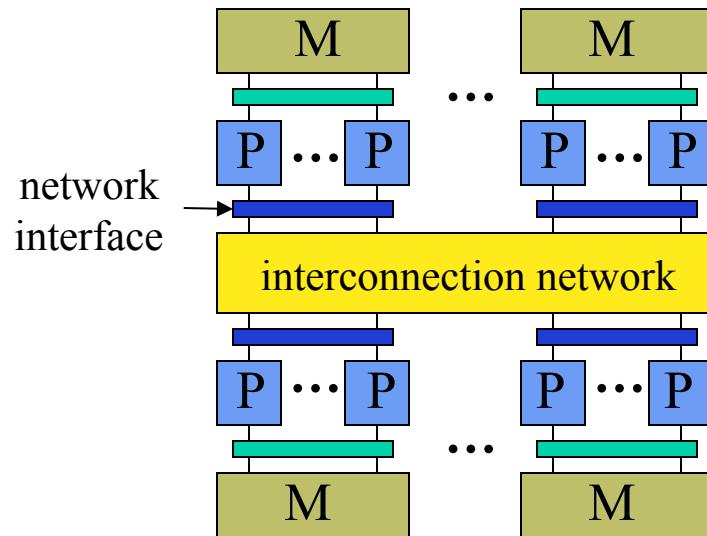
- Message passing between nodes



- Massively Parallel Processor (MPP)
    - Many, many processors

- Cluster of SMPs

- Shared memory addressing within SMP node
  - Message passing between SMP nodes

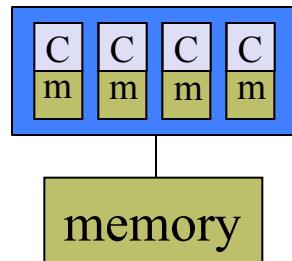


- Can also be regarded as MPP if processor number is large

# *Introduction to Parallel Architectures*

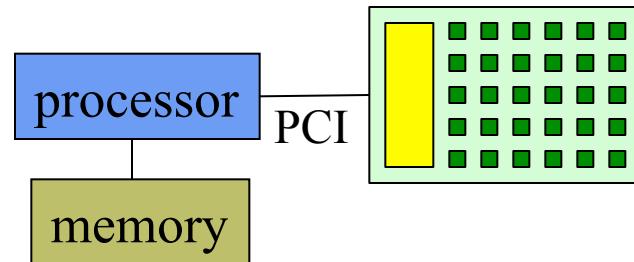
## Multicore

- Multicore processor

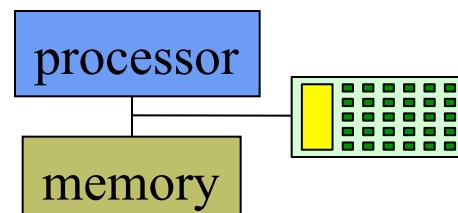


cores can be hardware multithreaded (hyperthread)

- GPU accelerator

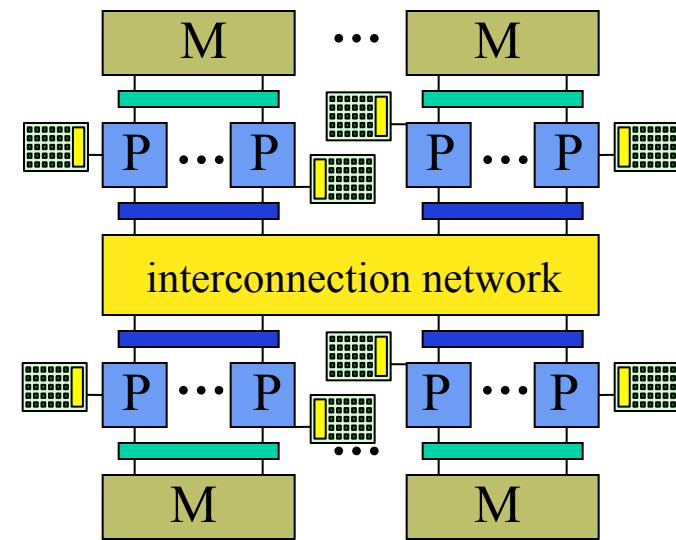


- “Fused” processor accelerator



## Multicore SMP+GPU Cluster

- Shared memory addressing within SMP node
- Message passing between SMP nodes
- GPU accelerators attached



# *How do you get parallelism in the hardware?*

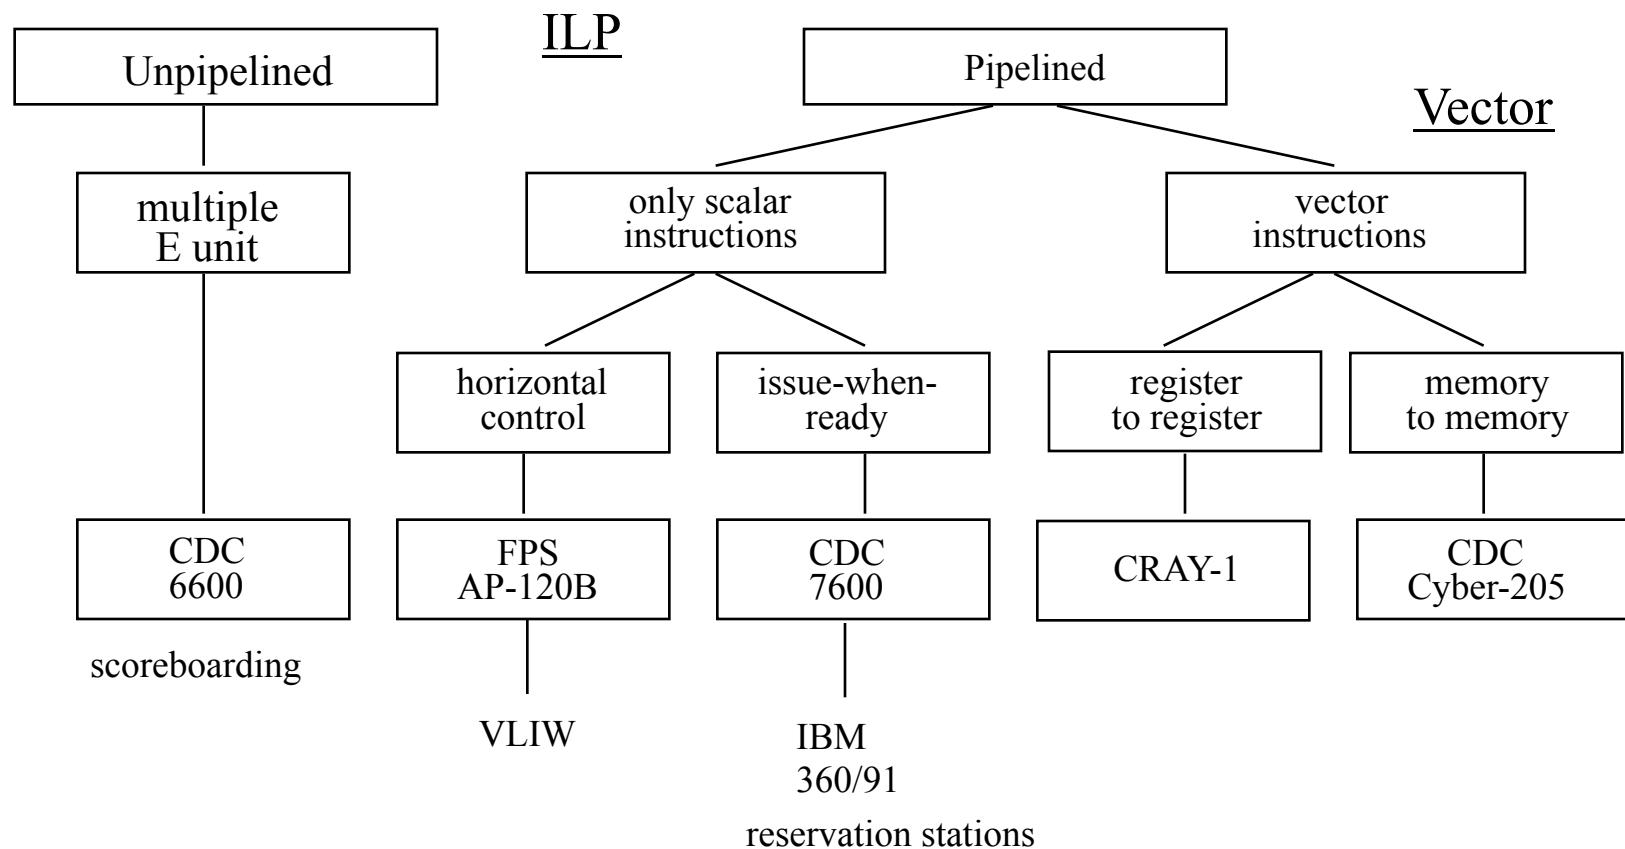
- Instruction-Level Parallelism (ILP)
- Data parallelism
  - Increase amount of data to be operated on at same time
- Processor parallelism
  - Increase number of processors
- Memory system parallelism
  - Increase number of memory units
  - Increase bandwidth to memory
- Communication parallelism
  - Increase amount of interconnection between elements
  - Increase communication bandwidth

# *Instruction-Level Parallelism*

- Opportunities for splitting up instruction processing
- Pipelining within instruction
- Pipelining between instructions
- Overlapped execution
- Multiple functional units
- Out of order execution
- Multi-issue execution
- Superscalar processing
- Superpipelining
- Very Long Instruction Word (VLIW)
- Hardware multithreading (hyperthreading)

# *Parallelism in Single Processor Computers*

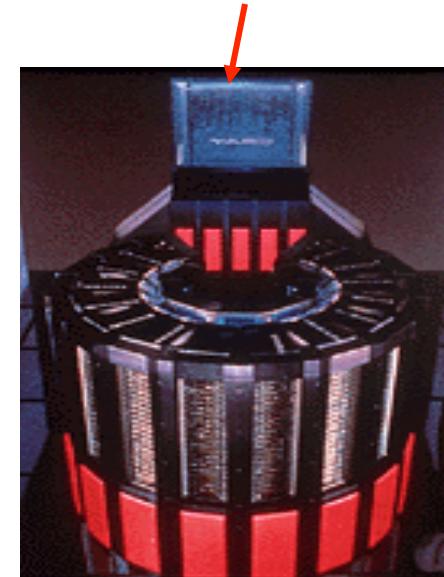
## □ History of processor architecture innovation



# *Vector Processing*

- Scalar processing
  - Processor instructions operate on scalar values
  - integer registers and floating point registers
- Vectors
  - Set of scalar data
  - Vector registers
    - integer, floating point (typically)
  - Vector instructions operate on vector registers (SIMD)
- Vector unit pipelining
- Multiple vector units
- Vector chaining

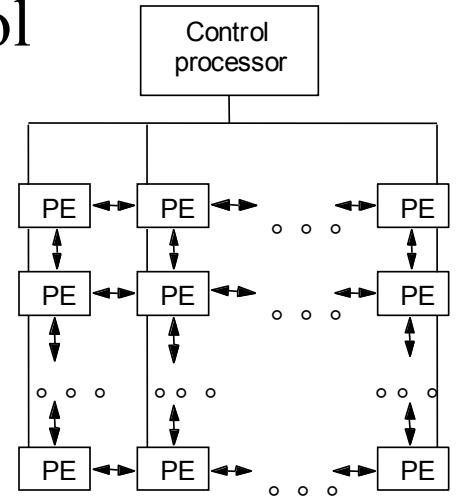
Liquid-cooled with inert fluorocarbon. (That's a waterfall fountain!!!)



Cray 2

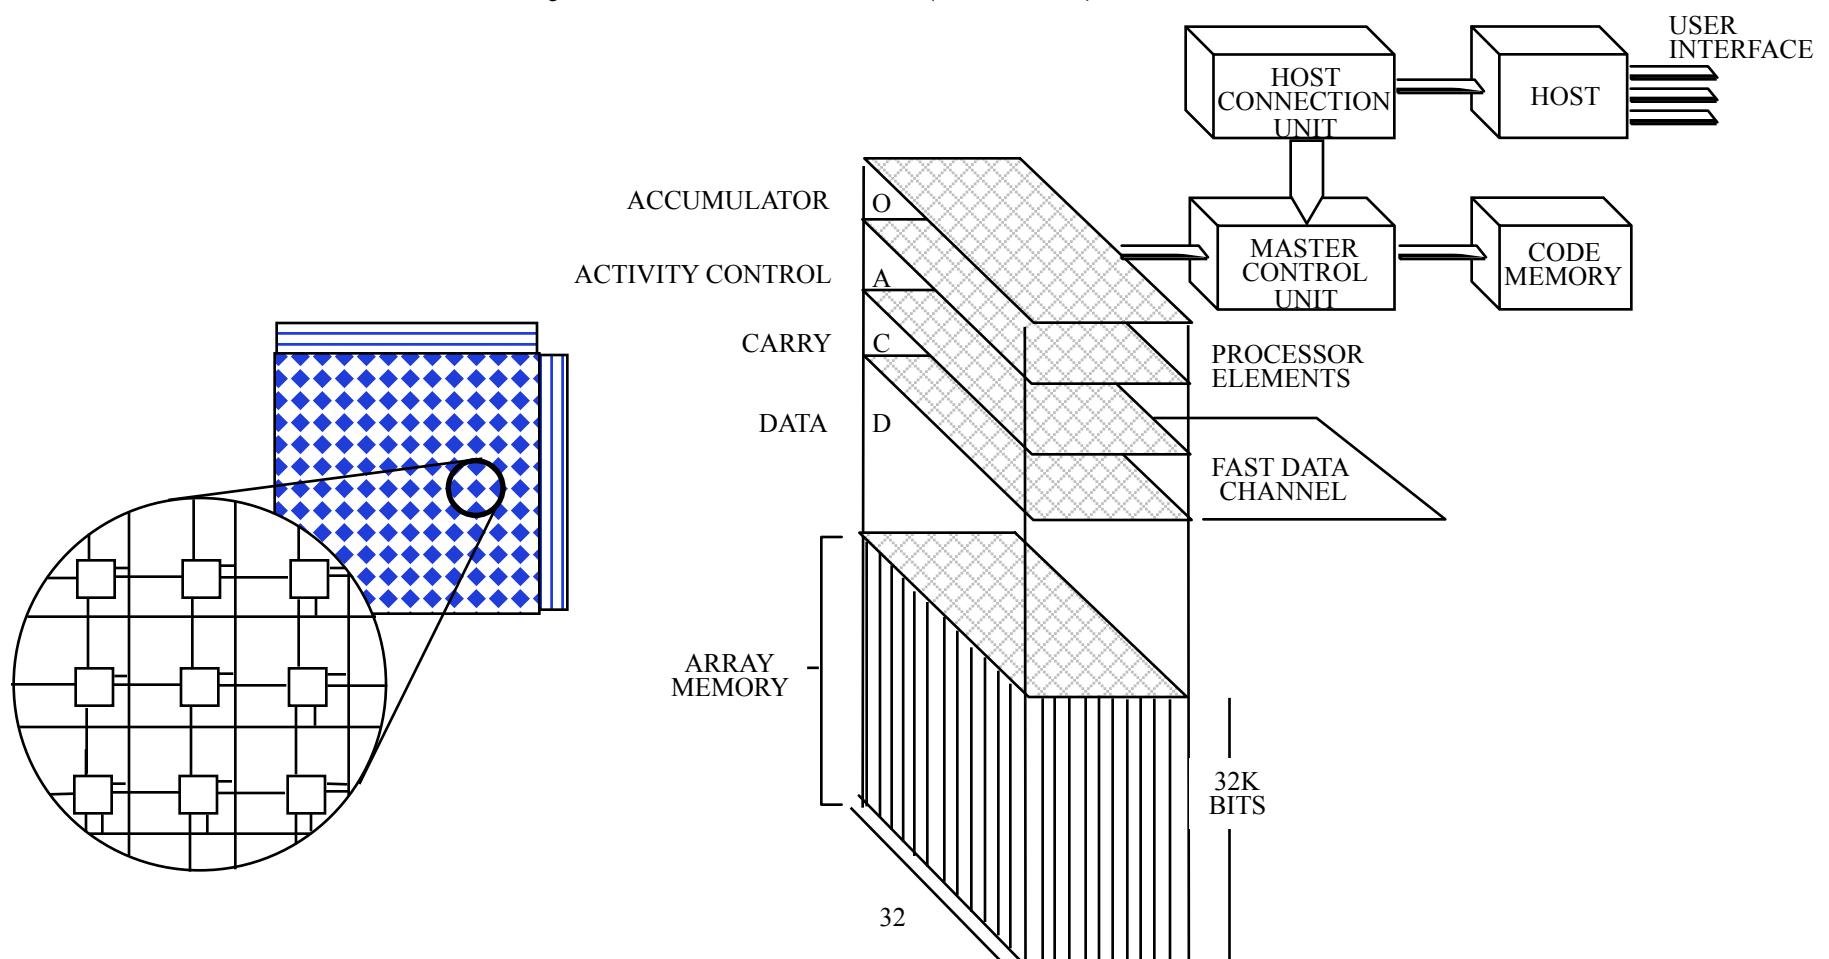
# *Data Parallel Architectures*

- SIMD (Single Instruction Multiple Data)
  - Logical single thread (instruction) of control
  - Processor associated with data elements
- Architecture
  - Array of simple processors with memory
  - Processors arranged in a regular topology
  - Control processor issues instructions
    - All processors execute same instruction (maybe disabled)
  - Specialized synchronization and communication
  - Specialized reduction operations
  - Array processing

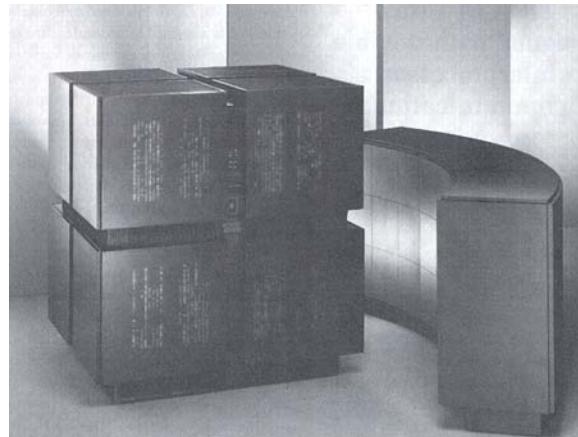


# *AMT DAP 500*

- Applied Memory Technology (AMT)
- Distributed Array Processor (DAP)

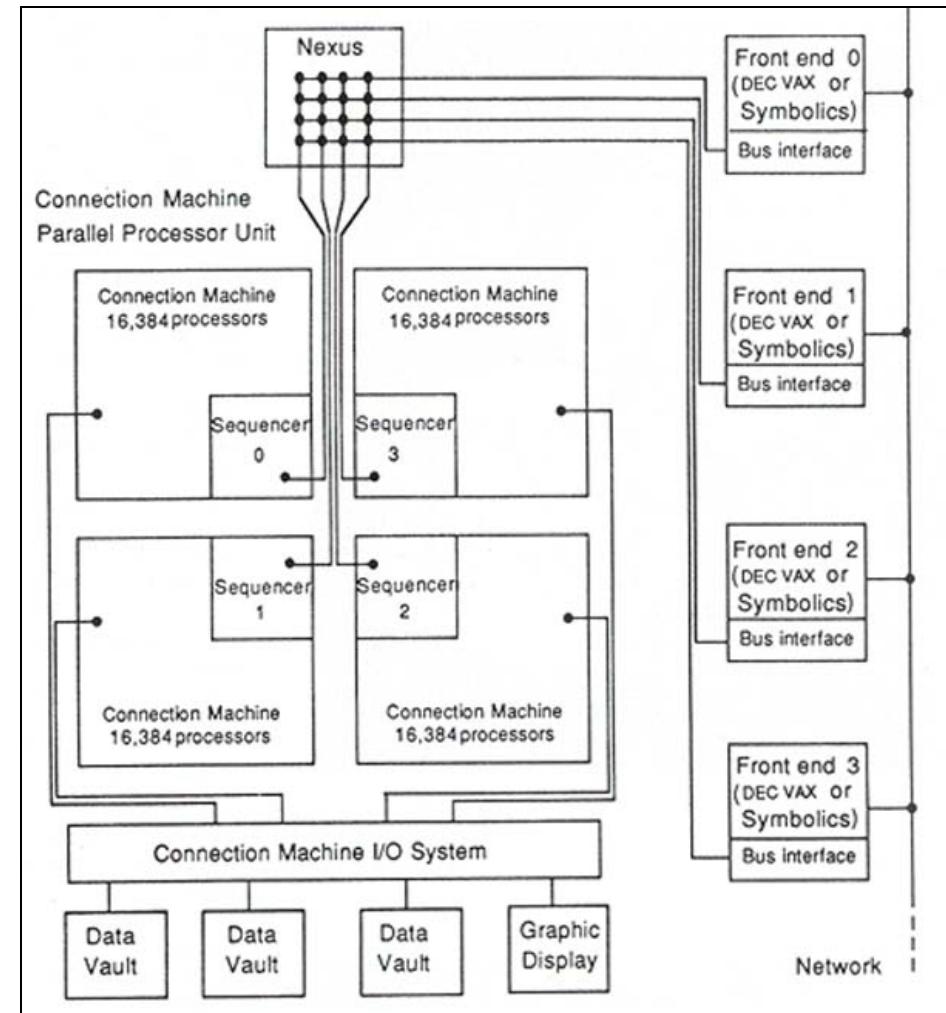
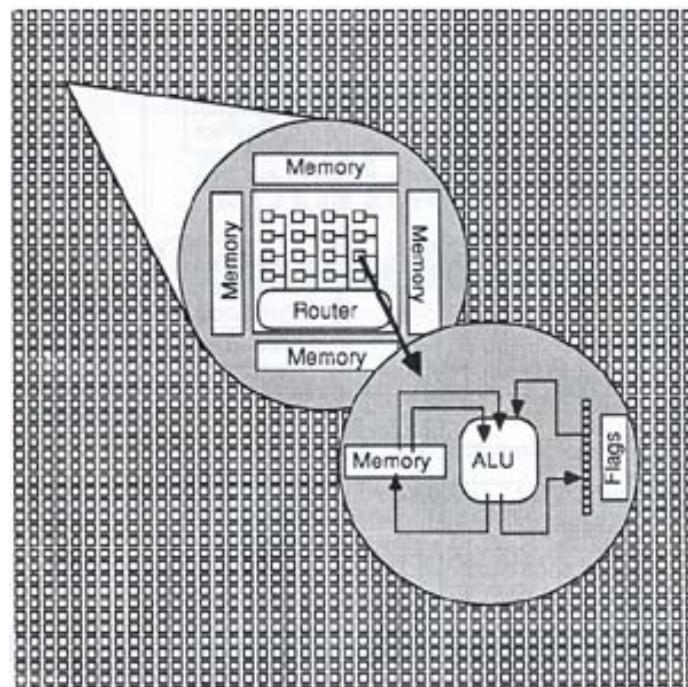


# *Thinking Machines Connection Machine*

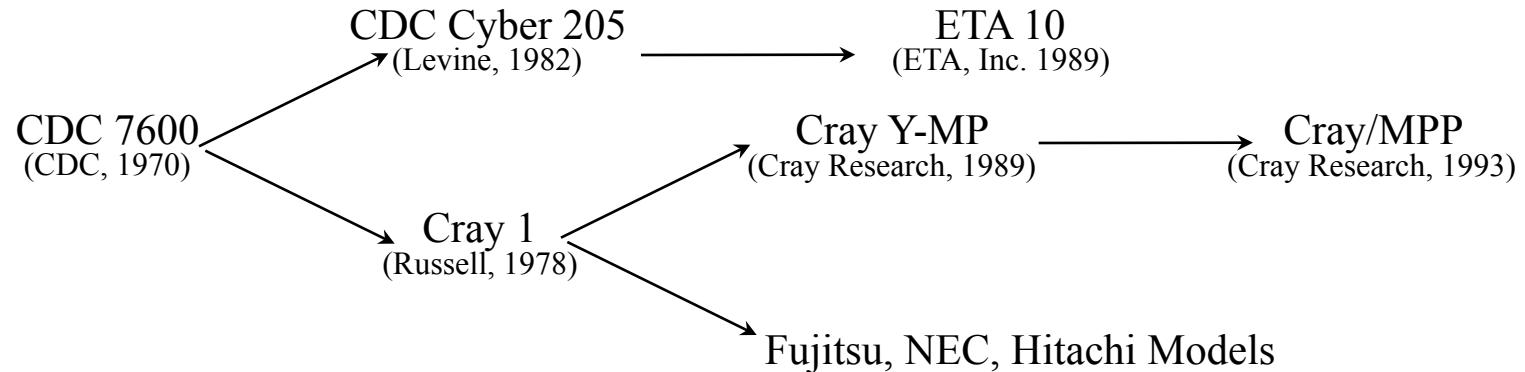


16,000 processors!!!

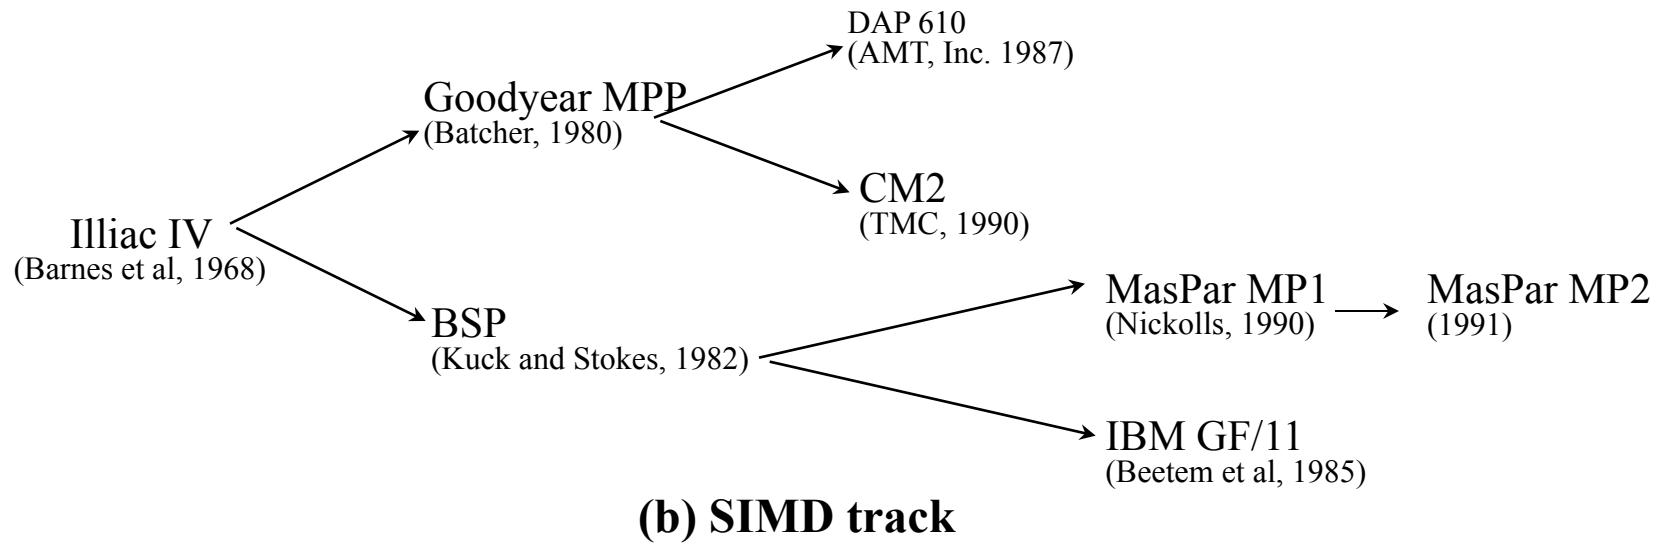
(Tucker, IEEE Computer, Aug. 1988)



# *Vector and SIMD Processing Timeline*



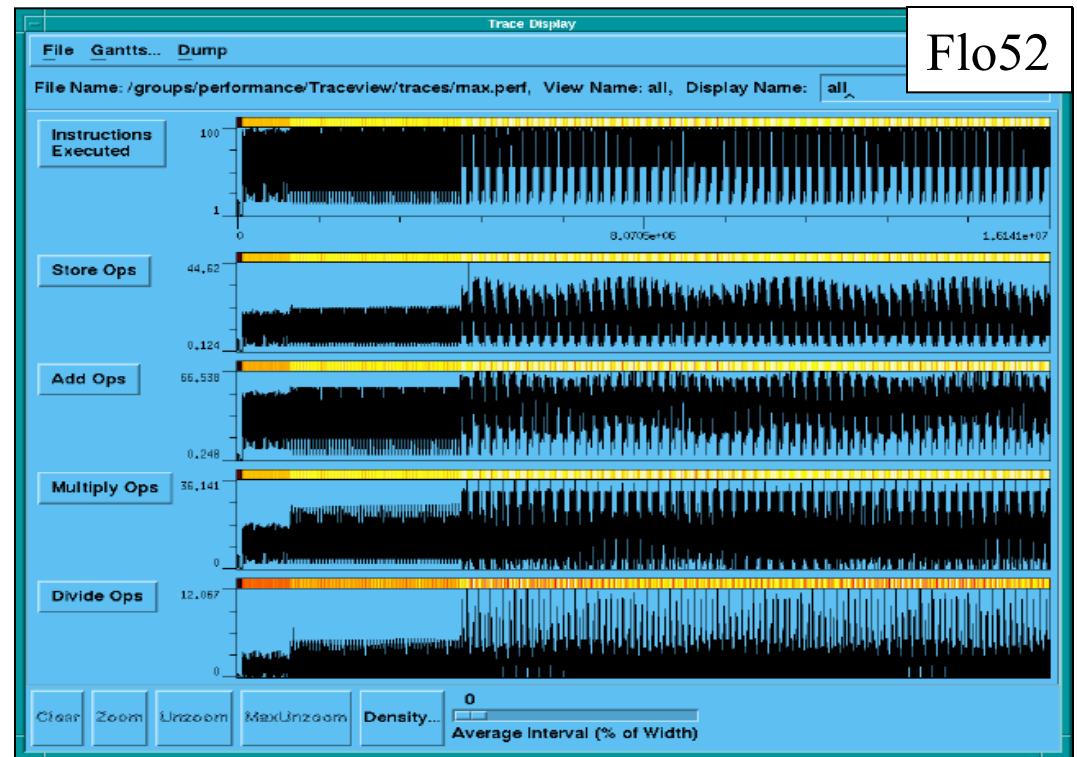
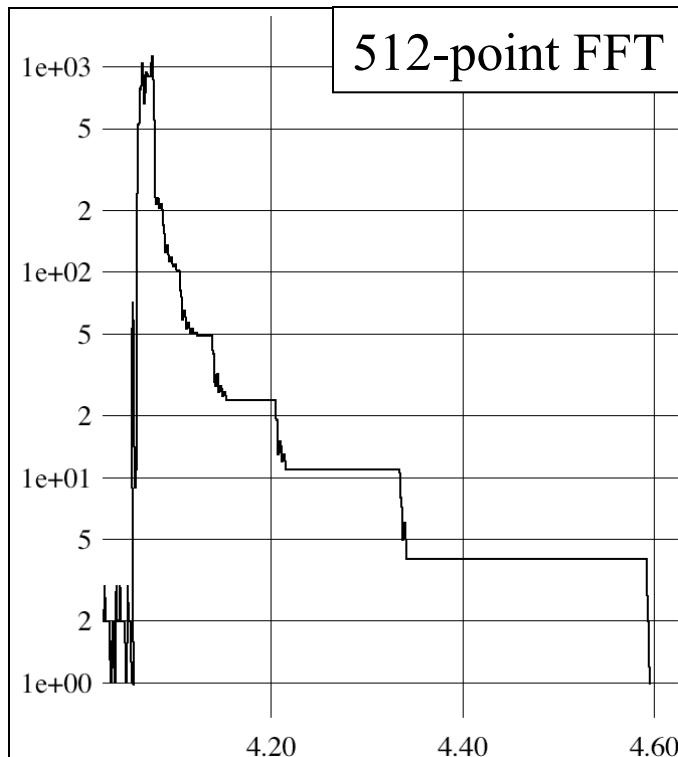
**(a) Multivector track**



**(b) SIMD track**

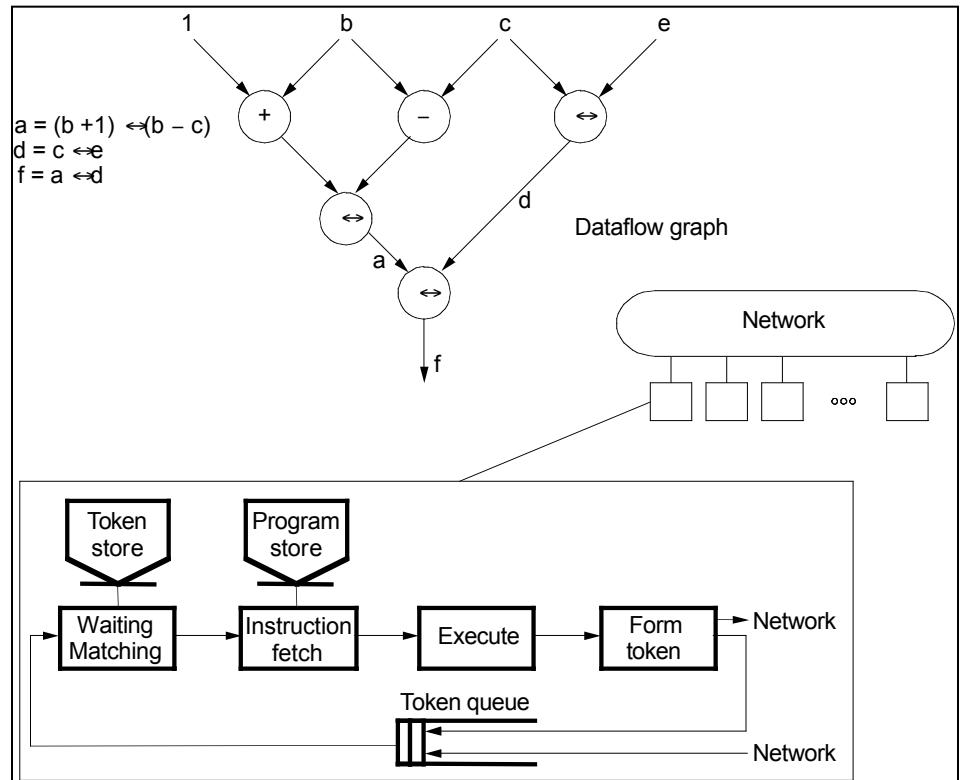
# *What's the maximum parallelism in a program?*

- “MaxPar: An Execution Driven Simulator for Studying Parallel Systems,” Ding-Kai Chen, M.S. Thesis, University of Illinois, Urbana-Champaign, 1989.
- Analyze the data dependencies in application execution



# Dataflow Architectures

- Represent computation as graph of dependencies
- Operations stored in memory until operands are ready
- Operations can be dispatched to processors
- Tokens carry tags of next instruction to processor
- Tag compared in matching store
- A match fires execution
- Machine does the hard parallelization work
- Hard to build right

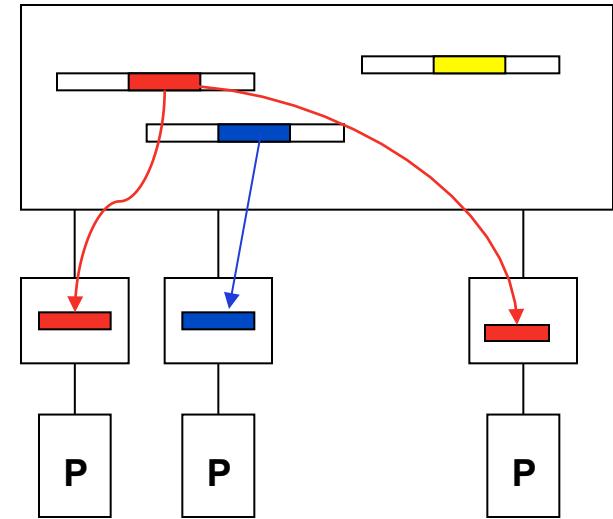


# *Shared Physical Memory*

- Add processors to single processor computer system
- Processors *share* computer system resources
  - Memory, storage, ...
- Sharing physical memory
  - Any processor can reference any memory location
  - Any I/O controller can reference any memory address
  - Single physical memory address space
- Operating system runs on any processor, or all
  - OS see single memory address space
  - Uses shared memory to coordinate
- Communication occurs as a result of loads and stores

# *Caching in Shared Memory Systems*

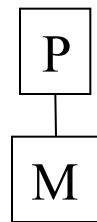
- Reduce average latency
  - automatic replication closer to processor
- Reduce average bandwidth
- Data is logically transferred from producer to consumer to memory
  - store reg → mem
  - load reg ← mem
- Processors can share data efficiently
- What happens when store and load are executed on different processors?
- Cache coherence problems



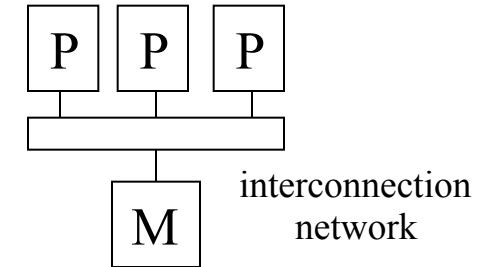
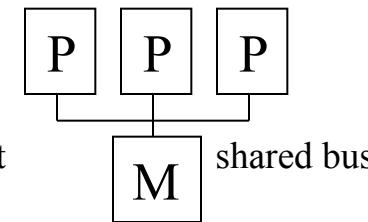
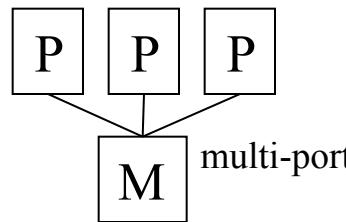
# *Shared Memory Multiprocessors (SMP)*

## □ Architecture types

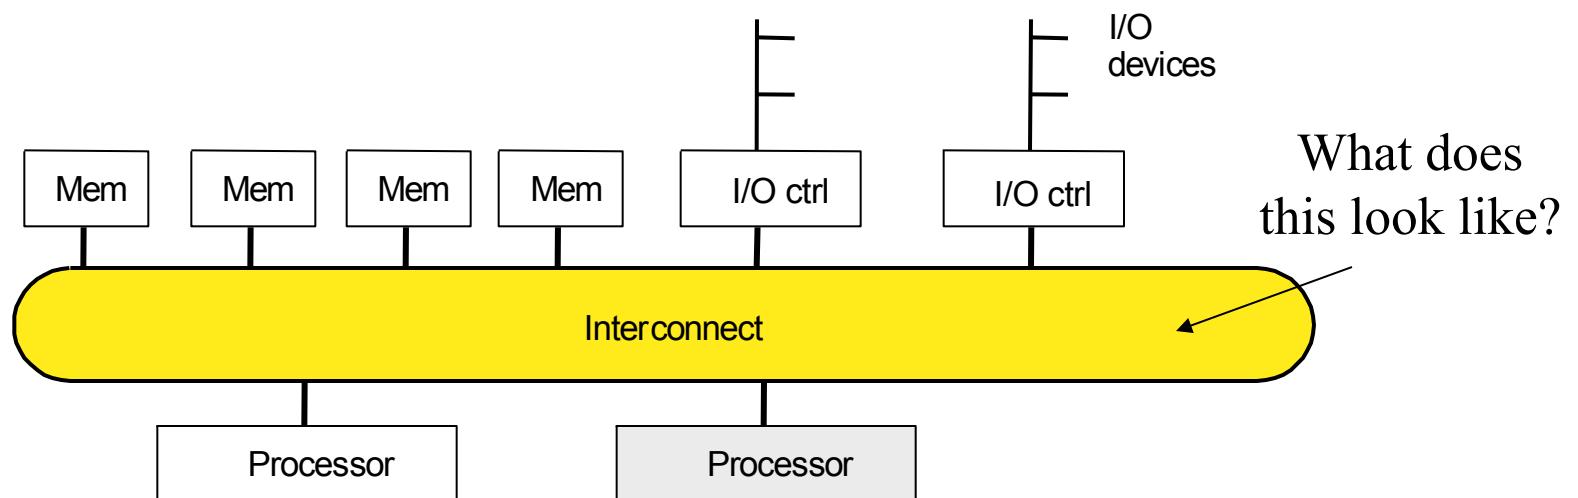
Single processor



Multiple processors

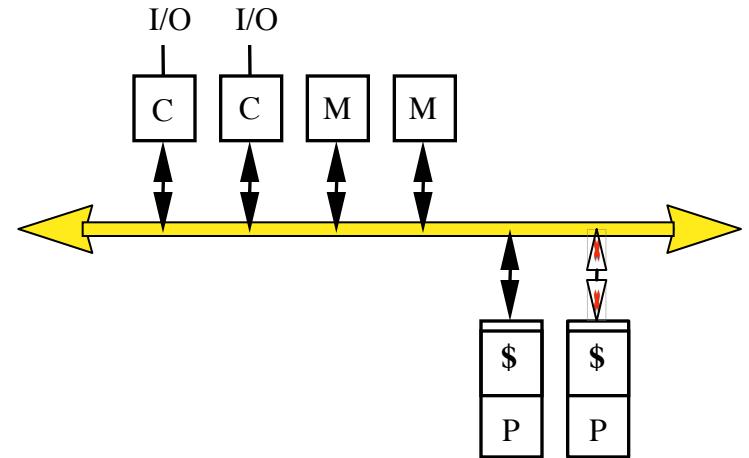


## □ Differences lie in memory system interconnection



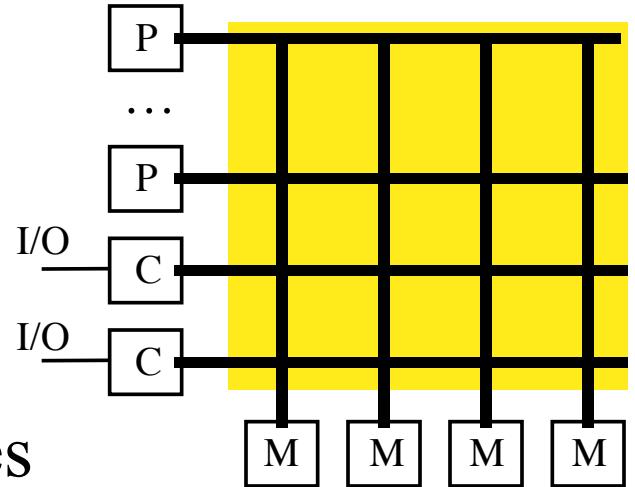
# *Bus-based SMP*

- Memory bus handles all memory read/write traffic
- Processors share bus
- *Uniform Memory Access* (UMA)
  - Memory (not cache) uniformly equidistant
  - Take same amount of time (generally) to complete
- May have multiple memory modules
  - Interleaving of physical address space
- Caches introduce memory hierarchy
  - Lead to data consistency problems
  - Cache coherency hardware necessary (CC-UMA)



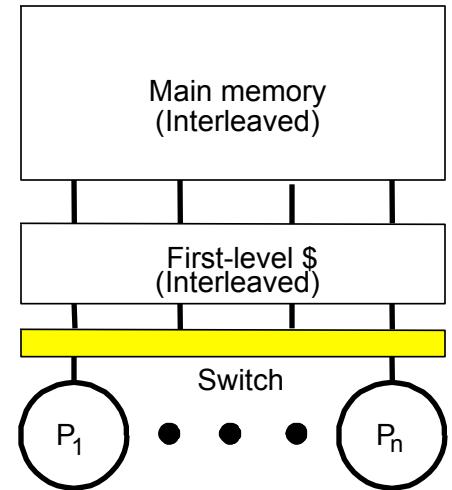
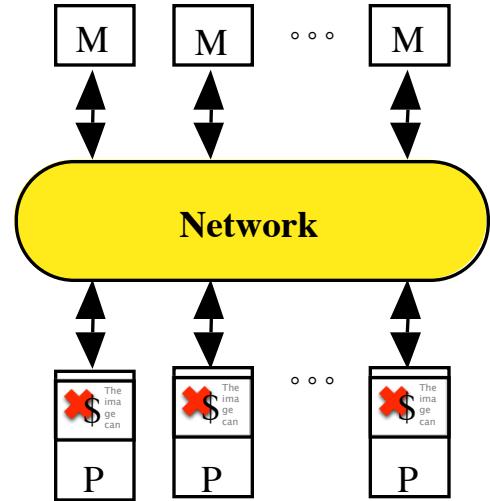
# *Crossbar SMP*

- Replicates memory bus for every processor and I/O controller
  - Every processor has direct path
- UMA SMP architecture
- Can still have cache coherency issues
- Multi-bank memory or interleaved memory
- Advantages
  - Bandwidth scales linearly (no shared links)
- Problems
  - High incremental cost (cannot afford for many processors)
  - Use switched multi-stage interconnection network



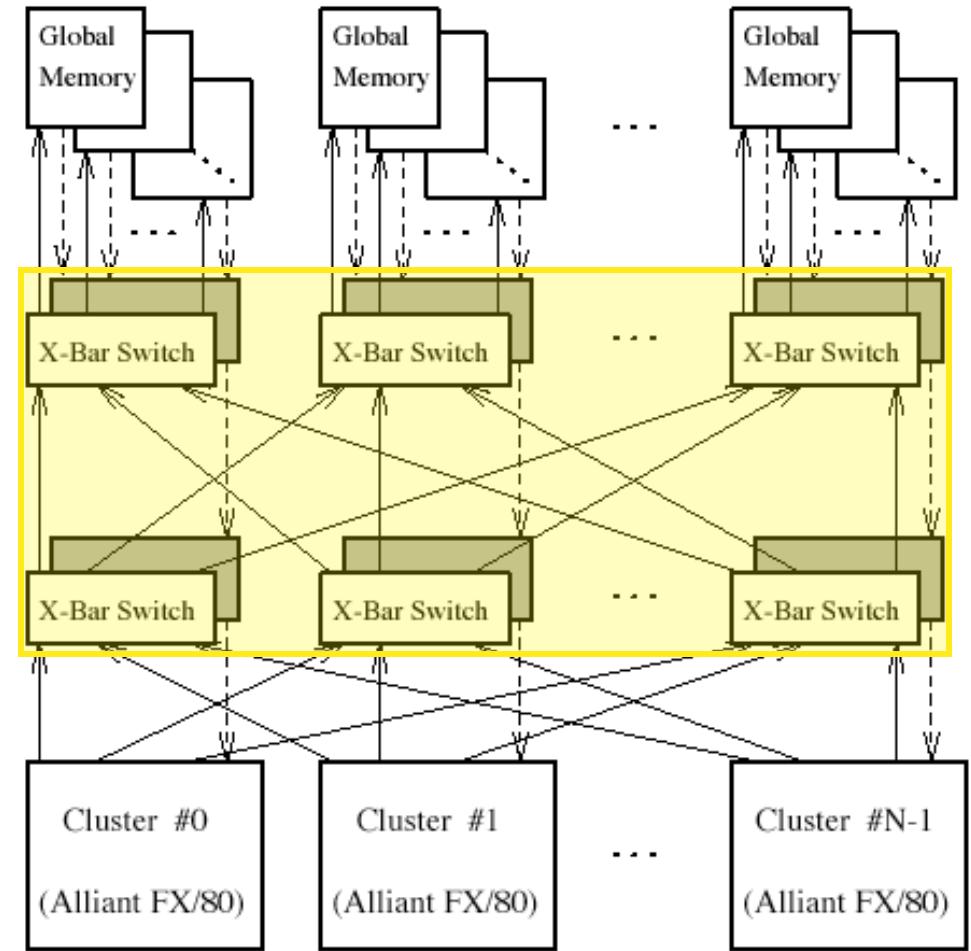
# *“Dance Hall” SMP and Shared Cache*

- Interconnection network connects processors to memory
- Centralized memory (UMA)
- Network determines performance
  - Continuum from bus to crossbar
  - Scalable memory bandwidth
- Memory is physically separated from processors
- Could have cache coherence problems
- Shared cache reduces coherence problem and provides fine grained data sharing

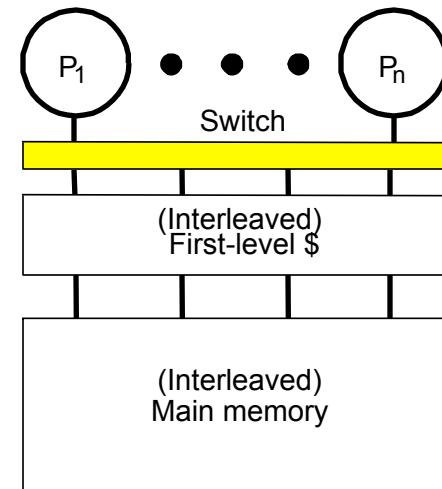


# *University of Illinois CSRD Cedar Machine*

- Center for Supercomputing Research and Development
- Multi-cluster scalable parallel computer
- Alliant FX/80
  - 8 processors w/ vectors
  - Shared cache
  - HW synchronization
- Omega switching network
- Shared global memory
- SW-based global memory coherency

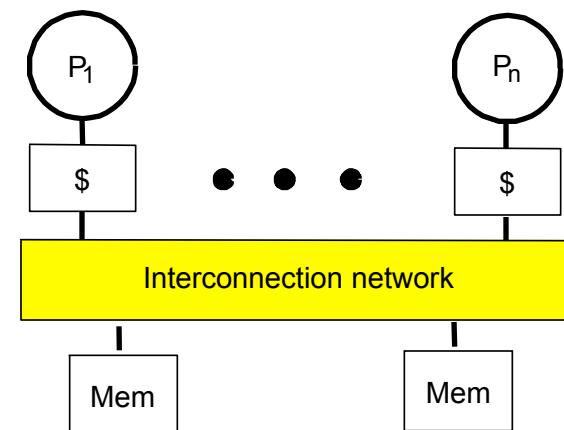


# *Natural Extensions of the Memory System*

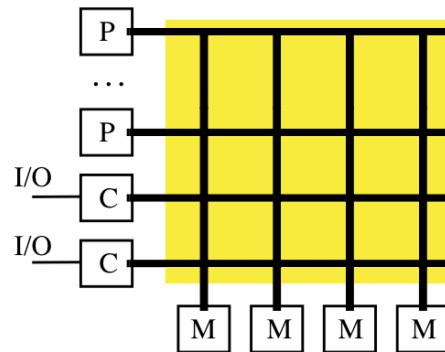


Shared Cache

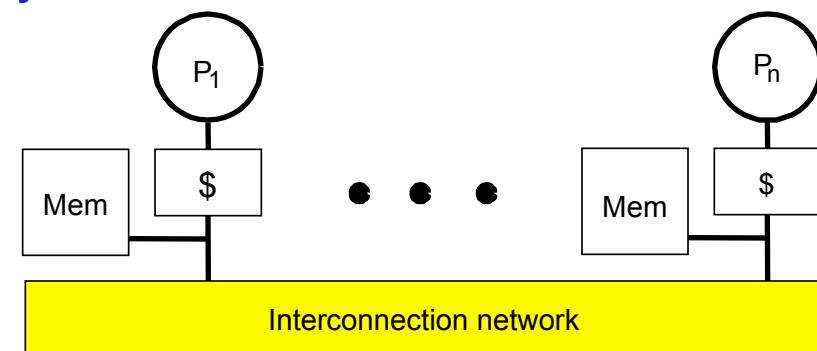
Scale →



Centralized Memory  
Dance Hall, UMA



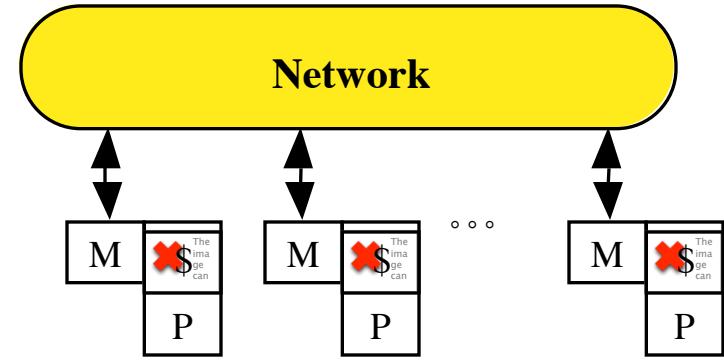
Crossbar, Interleaved



Distributed Memory (NUMA)

# *Non-Uniform Memory Access (NUMA) SMPs*

- Distributed memory
- Memory is physically resident close to each processor
- Memory is still shared
- *Non-Uniform Memory Access (NUMA)*
  - Local memory and remote memory
  - Access to local memory is faster, remote memory slower
  - Access is non-uniform
  - Performance will depend on data locality
- Cache coherency is still an issue (more serious)
- Interconnection network architecture is more scalable



# *Cache Coherency and SMPs*

- Caches play key role in SMP performance
  - Reduce average data access time
  - Reduce bandwidth demands placed on shared interconnect
- Private processor caches create a problem
  - Copies of a variable can be present in multiple caches
  - A write by one processor may not become visible to others
    - They'll keep accessing stale value in their caches

⇒ *Cache coherence* problem
- What do we do about it?
  - Organize the memory hierarchy to make it go away
  - Detect and take actions to eliminate the problem

# ***Definitions***

- Memory operation (load, store, read-modify-write, ...)
- Memory issue is operation presented to memory system
- Processor perspective
  - Write: subsequent reads return the value
  - Read: subsequent writes cannot affect the value
- *Coherent memory system*
  - There exists a serial order of memory operations on each location such that
    - operations issued by a process appear in order issued
    - value returned by each read is that written by previous write
  - ⇒ write propagation + write serialization

# *Motivation for Memory Consistency*

- Coherence implies that writes to a location become visible to all processors in the same order
- But when does a write become visible?
- How do we establish orders between a write and a read by different processors?
  - Use event synchronization
- Implement hardware protocol for cache coherency
- Protocol will be based on model of memory consistency

$P_1$	$P_2$
/* Assume initial value of A and flag is 0 */	
A = 1; flag = 1;	while (flag == 0); /* spin idly */ print A;

# *Memory Consistency*

- Specifies constraints on the order in which memory operations (from any process) can appear to execute with respect to each other
  - What orders are preserved?
  - Given a load, constrains the possible values returned by it
- Implications for both programmer and system designer
  - Programmer uses to reason about correctness
  - System designer can use to constrain how much accesses can be reordered by compiler or hardware
- Contract between programmer and system

# *Sequential Consistency*

- Total order achieved by interleaving accesses from different processes
  - Maintains *program order*
  - Memory operations (from all processes) appear to issue, execute, and complete atomically with respect to others
  - As if there was a single memory (no cache)

*“A multiprocessor is sequentially consistent if the result of any execution is the same as if the operations of all the processors were executed in some sequential order, and the operations of each individual processor appear in this sequence in the order specified by its program.” [Lamport, 1979]*

## ***Sequential Consistency (Sufficient Conditions)***

- There exist a total order consistent with the memory operations becoming visible in program order
- Sufficient Conditions
  - every process issues memory operations in program order
  - after write operation is issued, the issuing process waits for write to complete before issuing next memory operation (atomic writes)
  - after a read is issued, the issuing process waits for the read to complete and for the write whose value is being returned to complete (globally) before issuing its next memory operation
- Cache-coherent architectures implement consistency

# *Bus-based Cache-Coherent (CC) Architecture*

## □ Bus Transactions

- Single set of wires connect several devices
- Bus protocol: arbitration, command/addr, data
- Every device observes every transaction

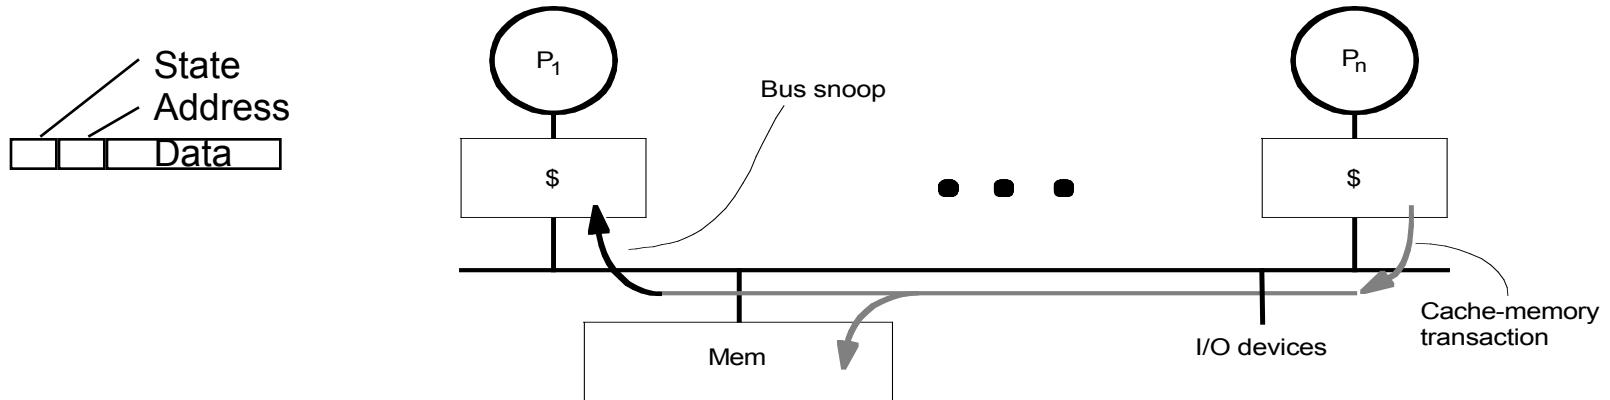
## □ Cache block state transition diagram

- FSM specifying how disposition of block changes
  - invalid, valid, dirty
- *Snoopy protocol*

## □ Basic Choices

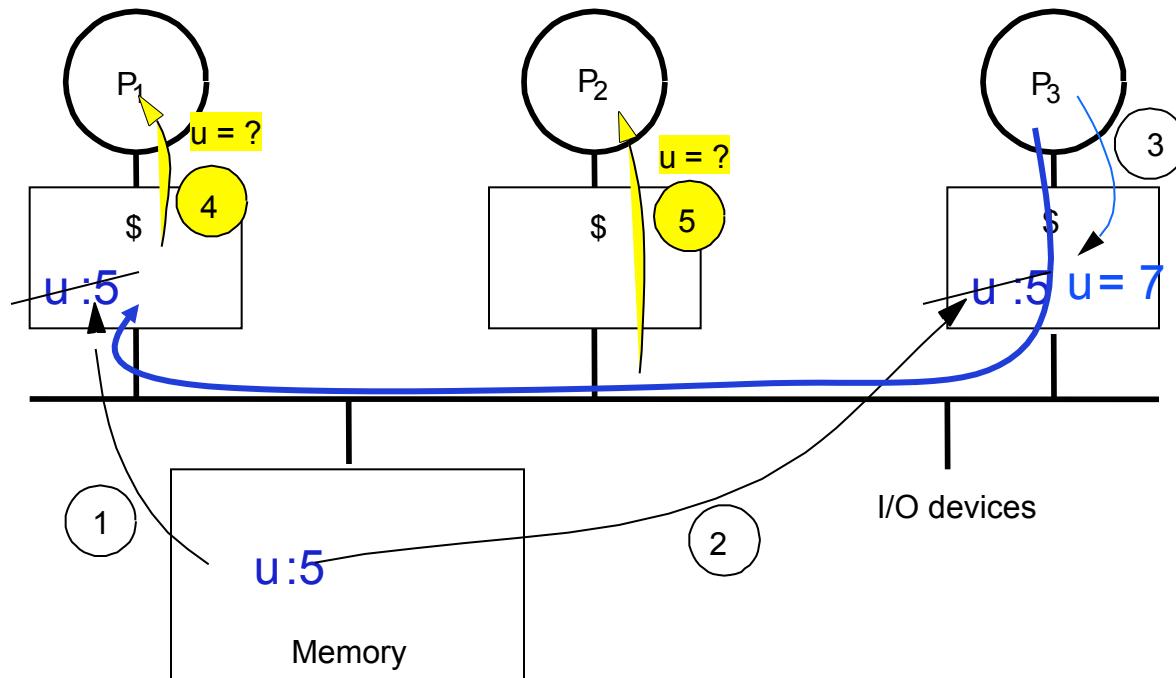
- Write-through vs Write-back
- Invalidate vs. Update

# *Snoopy Cache-Coherency Protocols*

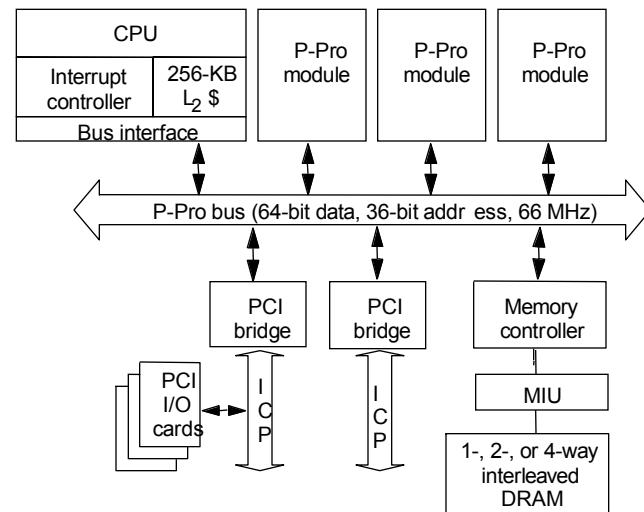
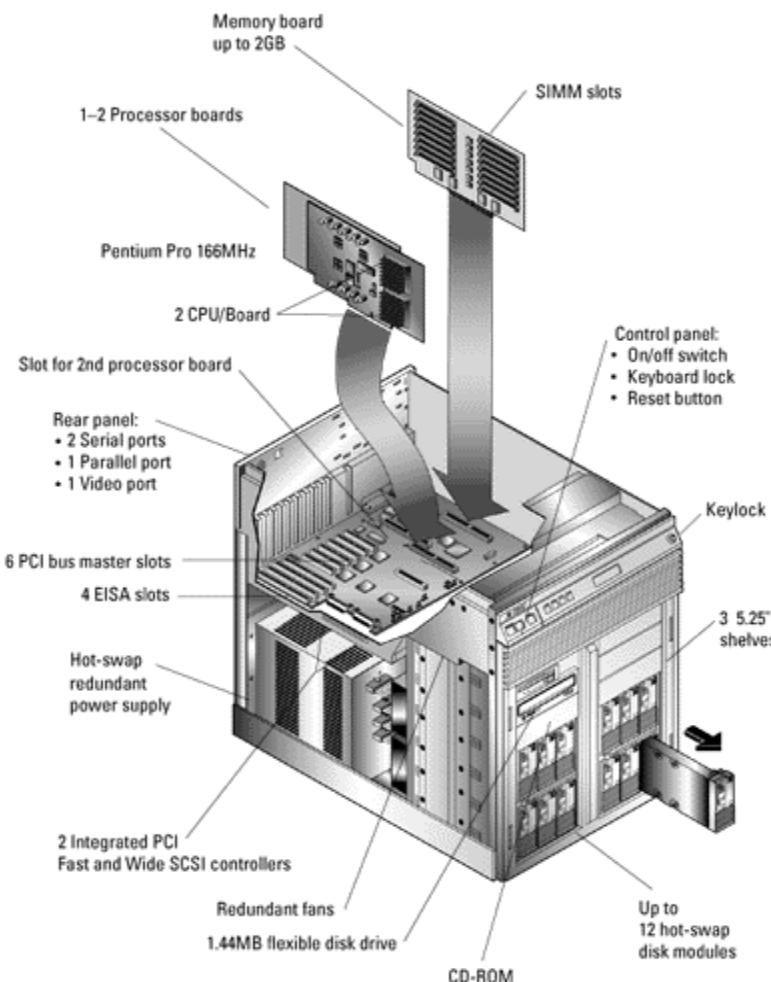


- Bus is a broadcast medium
- Caches know what they have
- Cache controller “snoops” all transactions on shared bus
  - relevant transaction if for a block its cache contains
  - take action to ensure coherence
    - invalidate, update, or supply value
  - depends on state of the block and the protocol

## *Example: Write-back Invalidate*

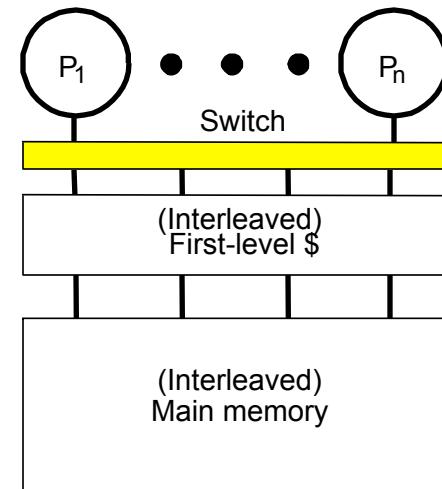


# Intel Pentium Pro Quad Processor



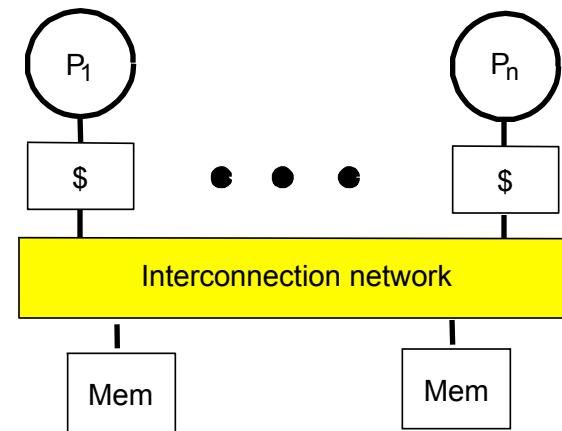
- All coherence and multiprocessing glue in processor module
- Highly integrated, targeted at high volume
- Low latency and bandwidth

# *Natural Extensions of the Memory System*

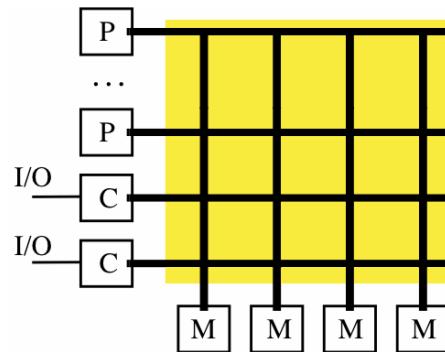


Shared Cache

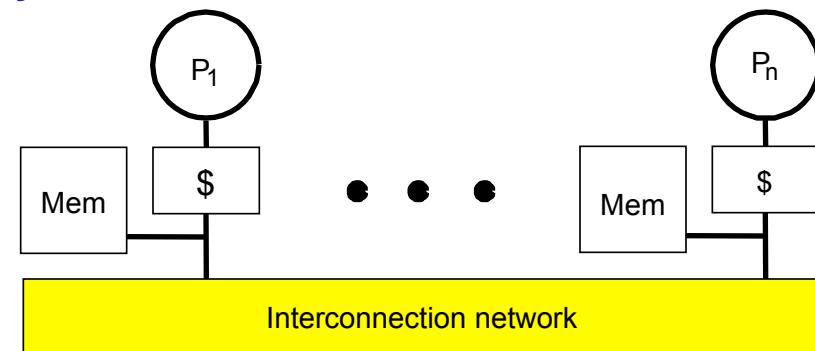
Scale →



Centralized Memory  
Dance Hall, UMA



Crossbar, Interleaved



Distributed Shared Memory (NUMA)

# *Memory Consistency*

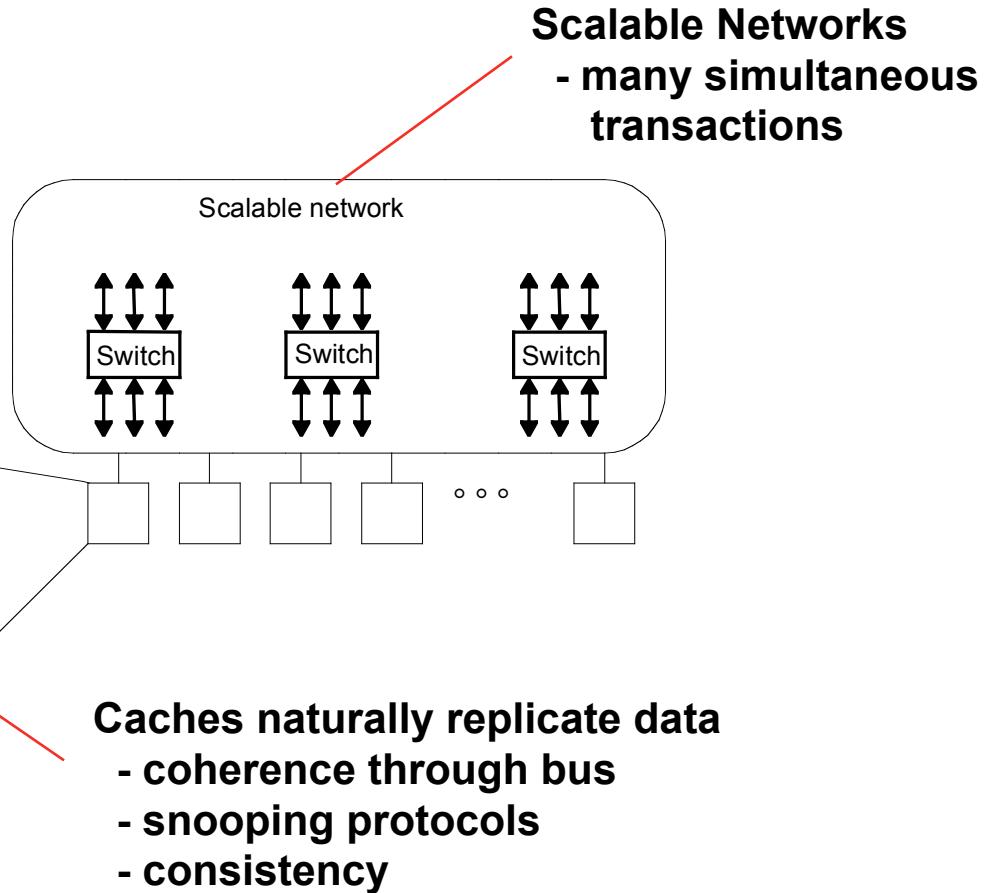
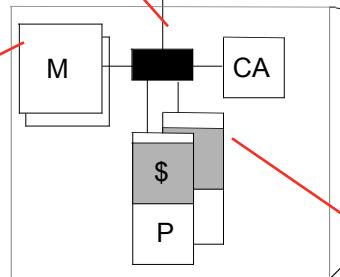
- Specifies constraints on the order in which memory operations (from any process) can appear to execute with respect to each other
  - What orders are preserved?
  - Given a load, constrains the possible values returned by it
- Implications for both programmer and system designer
  - Programmer uses to reason about correctness
  - System designer can use to constrain how much accesses can be reordered by compiler or hardware
- Contract between programmer and system
- Need coherency systems to enforce memory consistency

# *Context for Scalable Cache Coherence*

Realizing programming models through net transaction protocols

- efficient node-to-net interface
- interprets transactions

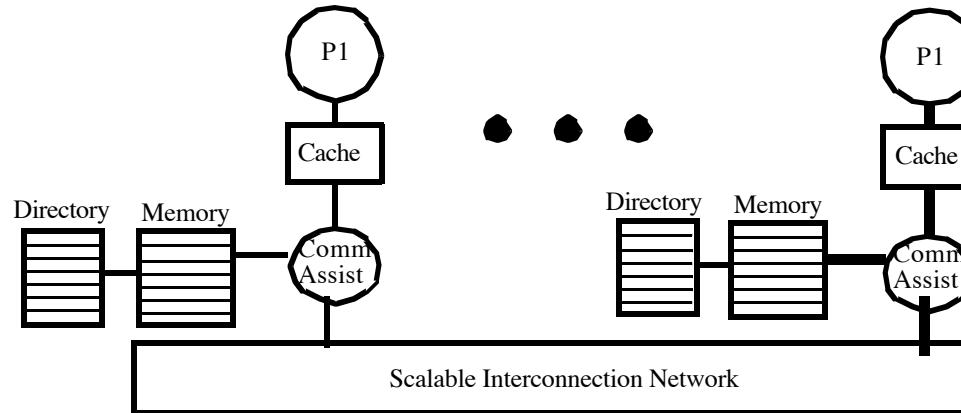
Scalable distributed memory



**Need cache coherence protocols that scale!**

- no broadcast or single point of order

# *Generic Solution: Directories*



- Maintain state vector explicitly
  - associate with memory block
  - records state of block in each cache
- On miss, communicate with directory
  - determine location of cached copies
  - determine action to take
  - conduct protocol to maintain coherence

# *Requirements of a Cache Coherent System*

- Provide set of states, state transition diagram, and actions
- Manage coherence protocol
  - (0) Determine when to invoke coherence protocol
  - (a) Find info about state of block in other caches to determine action
    - whether need to communicate with other cached copies
  - (b) Locate the other copies
  - (c) Communicate with those copies (inval/update)
- (0) is done the same way on all systems
  - state of the line is maintained in the cache
  - protocol is invoked if an “access fault” occurs on the line
- Different approaches distinguished by (a) to (c)

# *Bus-base Cache Coherence*

- All of (a), (b), (c) done through broadcast on bus
  - faulting processor sends out a “search”
  - others respond to the search probe and take necessary action
- Could do it in scalable network too
- Conceptually simple, but broadcast doesn’t scale
  - on bus, bus bandwidth doesn’t scale
  - on scalable network, every fault leads to at least  $p$  network transactions
- Scalable coherence:
  - can have same cache states and state transition diagram
  - different mechanisms to manage protocol

# *Basic Snoop Protocols*

- Write Invalidate :
  - Multiple readers, single writer
  - Write to shared data: an invalidate is sent to all caches
  - Read Miss:
    - Write-through: memory is always up-to-date
    - Write-back: snoop in caches to find most recent copy
- Write Broadcast (typically write through):
  - Write to shared data: broadcast on bus, snoop and update
- Write serialization: bus serializes requests!
- Write Invalidate versus Broadcast

# *Snooping Cache Variations*

<b>Basic Protocol</b>	<b>Berkeley Protocol</b>	<b>Illinois Protocol</b>	<b>MESI Protocol</b>
Exclusive	Owned Exclusive	Private Dirty	Modified (private,!=Memory)
Shared	Owned Shared	Private Clean	Exclusive (private,=Memory)
Invalid	Shared	Shared	Shared (shared,=Memory)
	Invalid	Invalid	Invalid

Owner can update via bus invalidate operation

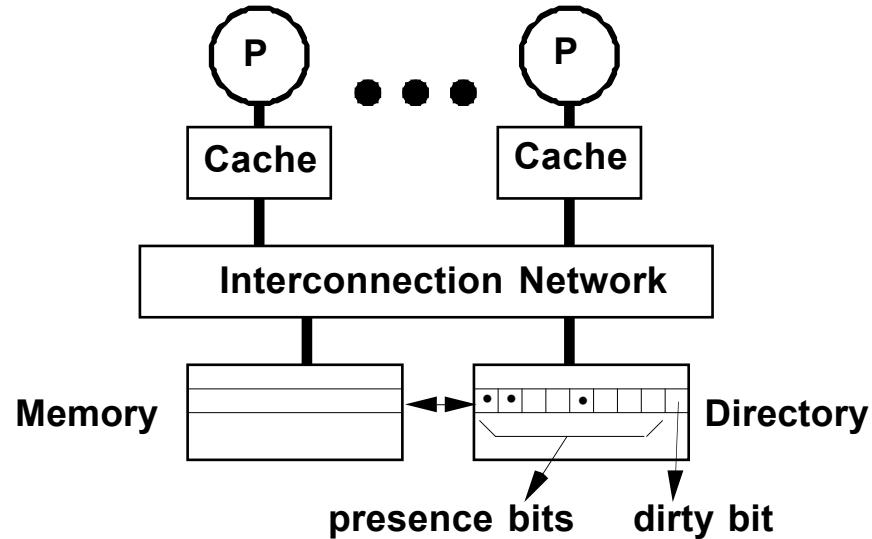
Owner must write back when replaced in cache

If read sourced from memory, then Private Clean  
if read sourced from other cache, then Shared  
Can write in cache if held private clean or dirty

## *Scalable Approach: Directories*

- Every memory block has associated directory information
  - keeps track of copies of cached blocks and their states
  - on a miss, find directory entry, look it up, and communicate only with the nodes that have copies if necessary
  - in scalable networks, communication with directory and copies is through network transactions
- Many alternatives for organizing directory information

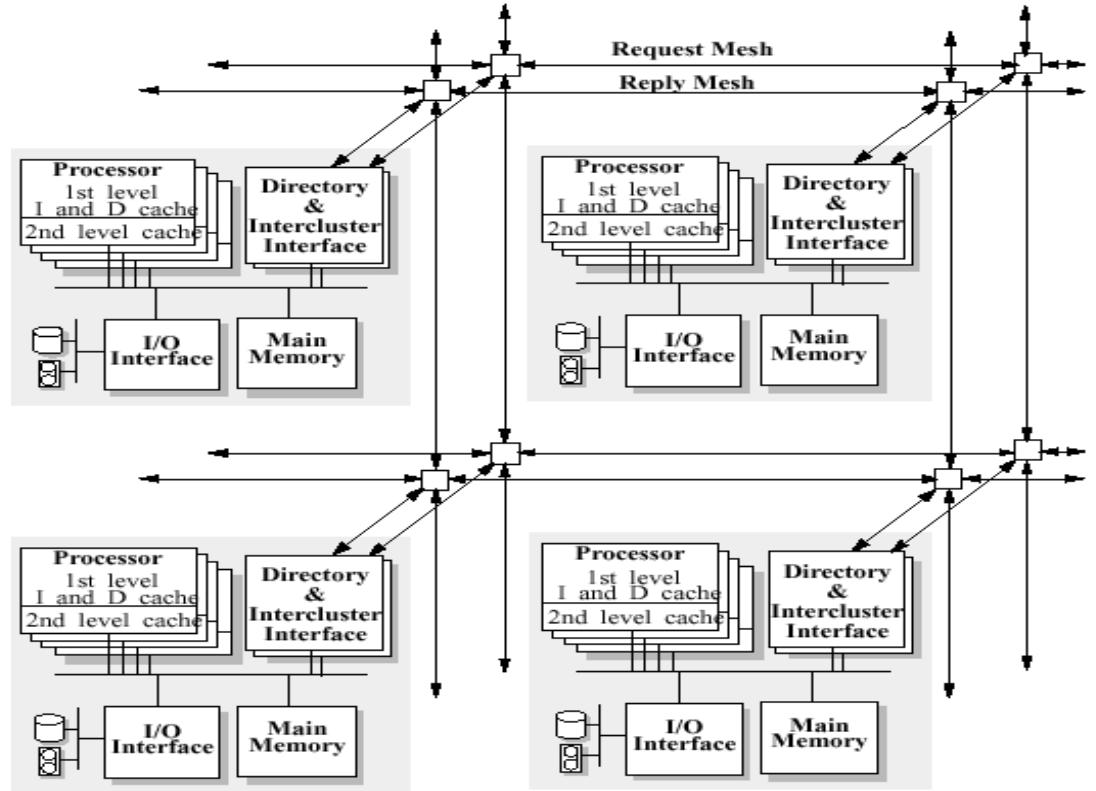
# *Basic Operation of Directory*



- k processors
- Each cache-block in memory
  - k presence bits and 1 dirty bit
- Each cache-block in cache
  - 1 valid bit and 1 dirty (owner) bit
- Read from memory
  - Dirty bit OFF
  - Dirty bit ON
- Write to memory
  - Dirty bit OFF

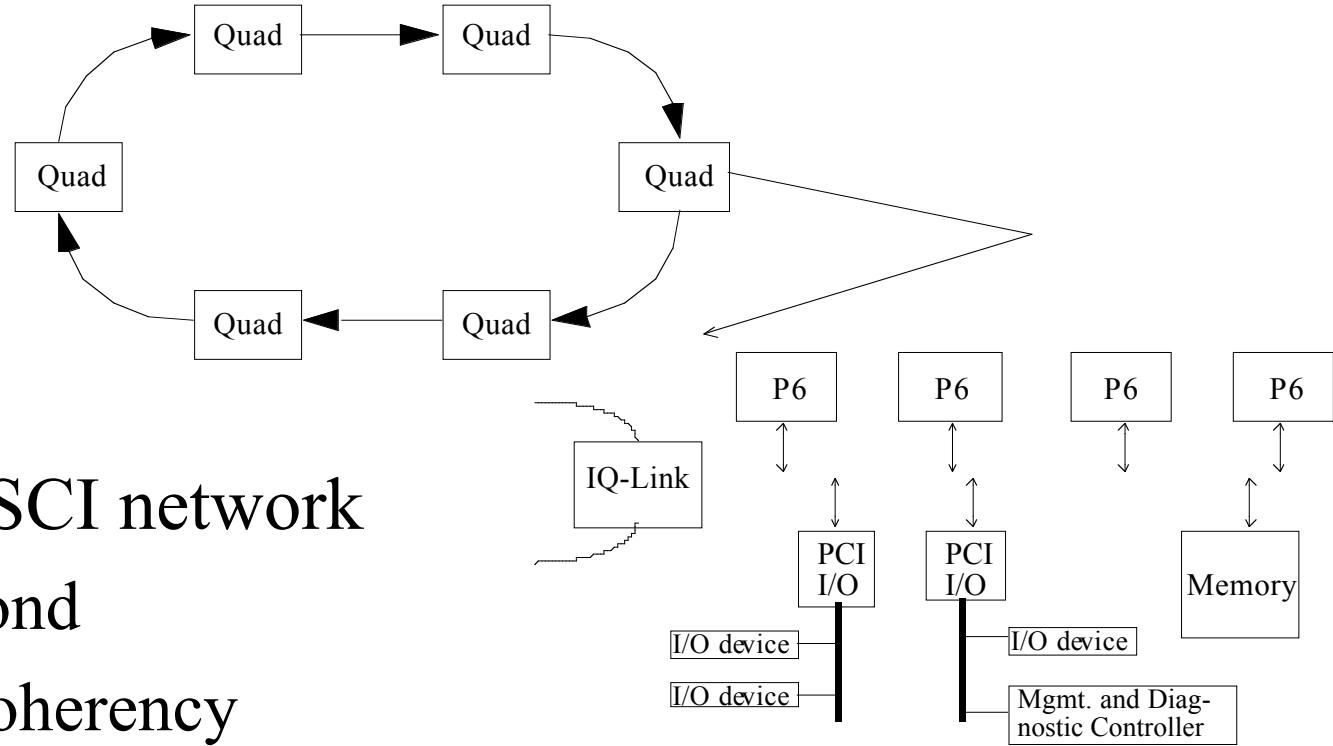
# *DASH Cache-Coherent SMP*

- Directory Architecture for Shared Memory
- Stanford research project (early 1990s) for studying how to build cache-coherent shared memory architectures
- Directory-based cache coherency
- IEEE Computer, Volume 25 Issue 3, March 1992



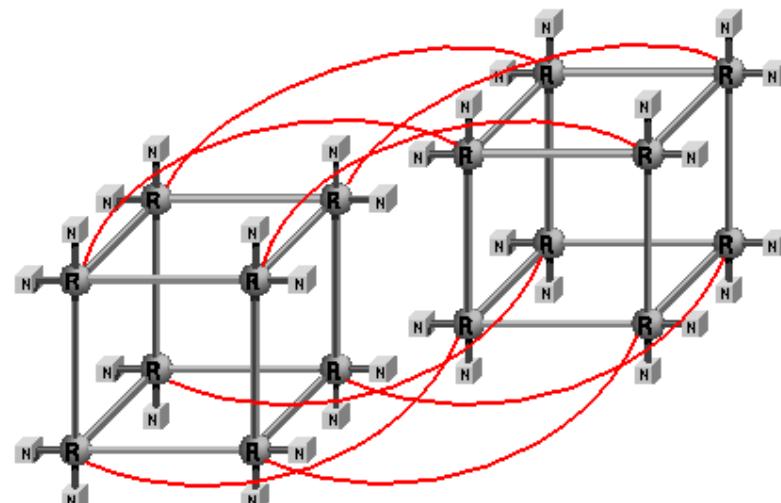
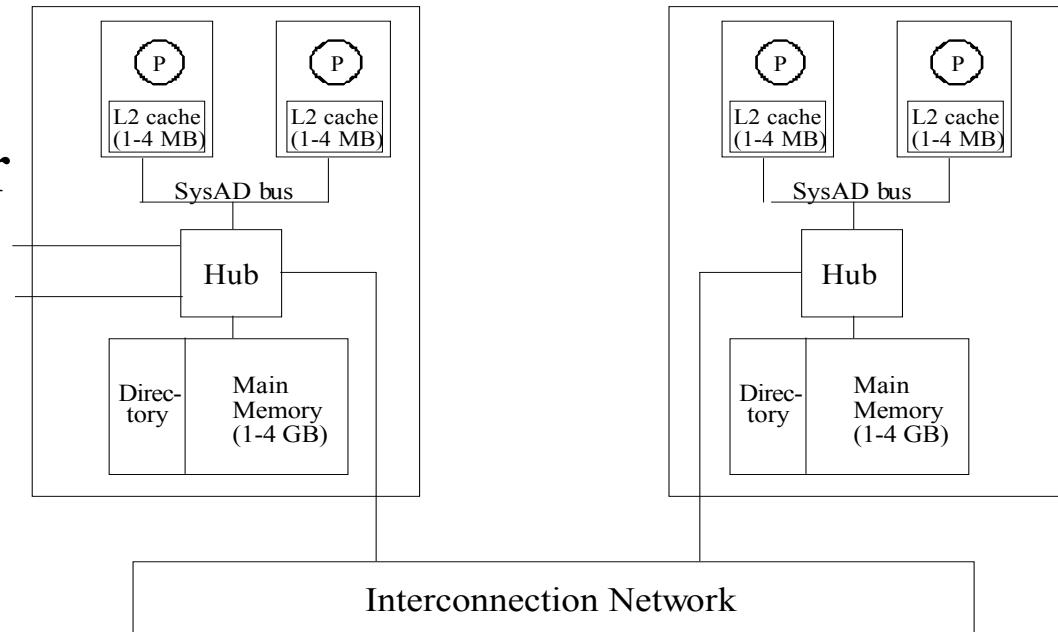
# *Sequent NUMA-Q*

- Ring-based SCI network
  - 1 GB/second
  - Built-in coherency
- Commodity SMPs as building blocks
  - Extend coherency mechanism
- Split transaction bus



# *SGI Origin 2000*

- Scalable shared memory multiprocessor
- MIPS R10000 CPU
- NUMAlink router
- Directory-based cache coherency (MESI)
- ASCI Blue Mountain

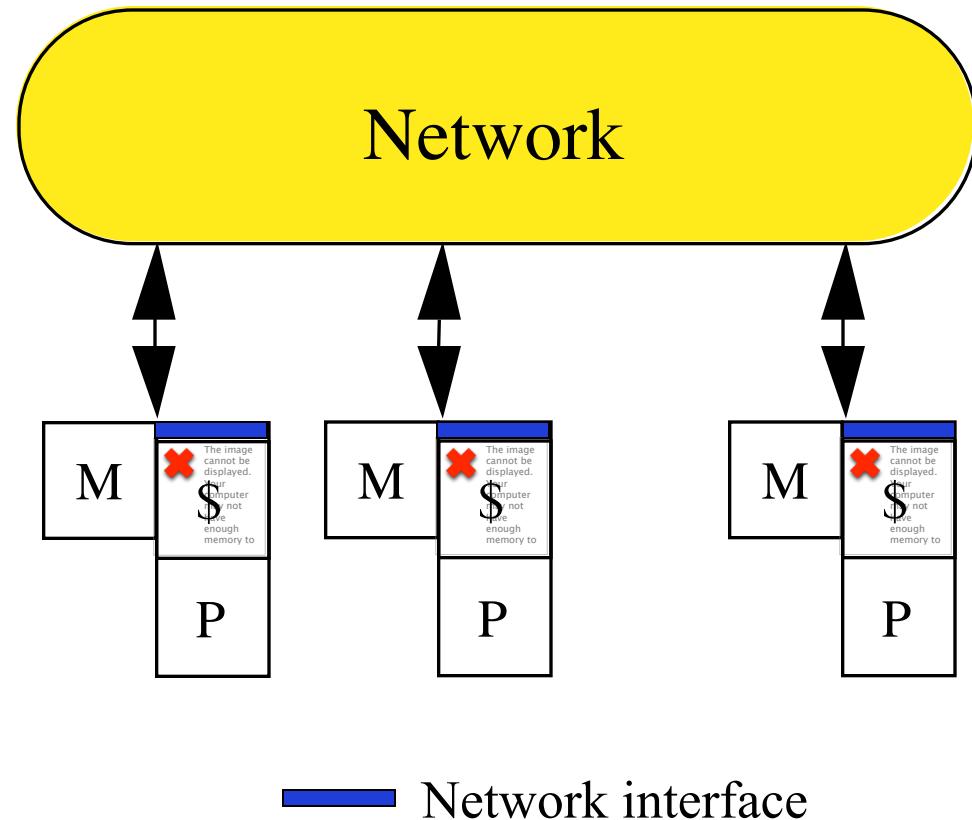


# *Distributed Memory Multiprocessors*

- Each processor has a local memory
  - Physically separated memory address space
- Processors must communicate to access non-local data
  - Message communication (message passing)
    - *Message passing architecture*
  - Processor interconnection network
- Parallel applications must be partitioned across
  - Processors: execution units
  - Memory: data partitioning
- Scalable architecture
  - Small incremental cost to add hardware (cost of node)

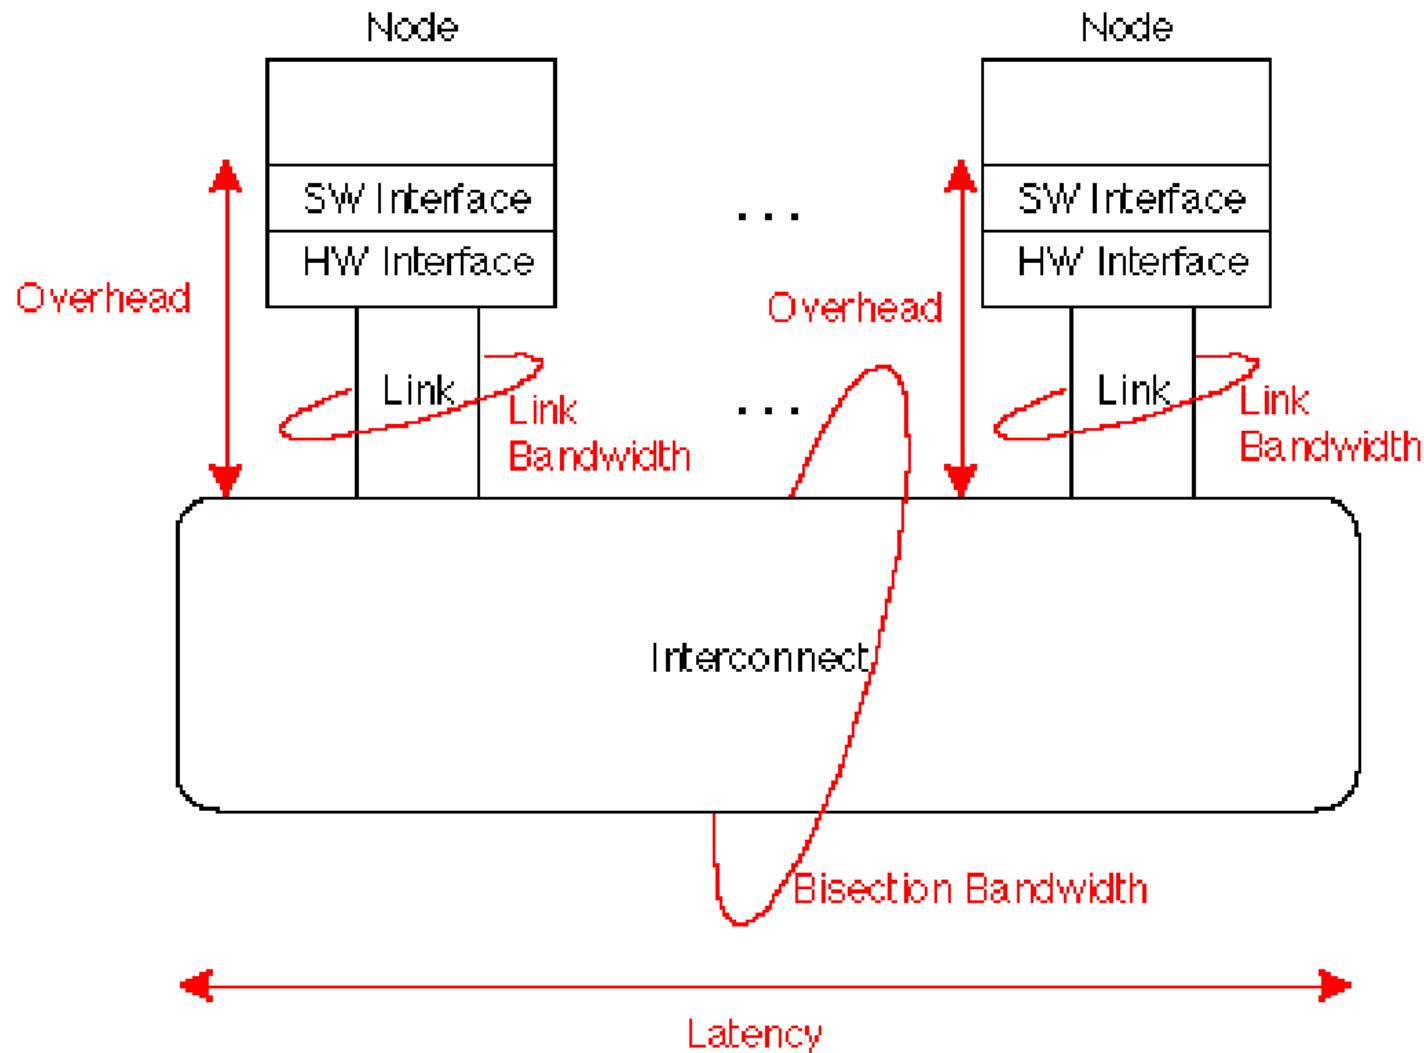
# *Distributed Memory (MP) Architecture*

- Nodes are complete computer systems
  - Including I/O
- Nodes communicate via interconnection network
  - Standard networks
  - Specialized networks
- Network interfaces
  - Communication integration
- Easier to build



— Network interface

# *Network Performance Measures*



**Overhead:** latency of interface vs. **Latency:** network

# *Performance Metrics: Latency and Bandwidth*

- Bandwidth
  - Need high bandwidth in communication
  - Match limits in network, memory, and processor
  - Network interface speed vs. network bisection bandwidth
- Latency
  - Performance affected since processor may have to wait
  - Harder to overlap communication and computation
  - Overhead to communicate is a problem in many machines
- Latency hiding
  - Increases programming system burden
  - Examples: communication/computation overlaps, prefetch

# *Scalable, High-Performance Interconnect*

- Interconnection network is core of parallel architecture
- Requirements and tradeoffs at many levels
  - Elegant mathematical structure
  - Deep relationship to algorithm structure
  - Hardware design sophistication
- Little consensus
  - Performance metrics?
  - Cost metrics?
  - Workload?
  - ...

# *What Characterizes an Interconnection Network?*

- Topology (what)
  - Interconnection structure of the network graph
- Routing Algorithm (which)
  - Restricts the set of paths that messages may follow
  - Many algorithms with different properties
- Switching Strategy (how)
  - How data in a message traverses a route
  - *circuit switching* vs. *packet switching*
- Flow Control Mechanism (when)
  - When a message or portions of it traverse a route
  - What happens when traffic is encountered?

# *Topological Properties*

- Routing distance
  - Number of links on route from source to destination
- Diameter
  - Maximum routing distance
- Average distance
- Partitioned network
  - Removal of links resulting in disconnected graph
  - Minimal cut
- Scaling increment
  - What is needed to grow the network to next valid degree

# *Interconnection Network Types*

<b>Topology</b>	<b>Degree</b>	<b>Diameter</b>	<b>Ave Dist</b>	<b>Bisection</b>	<b><math>D(D \text{ ave}) @ P=1024</math></b>
<b>1D Array</b>	<b>2</b>	<b><math>N-1</math></b>	<b><math>N / 3</math></b>	<b>1</b>	<b>huge</b>
<b>1D Ring</b>	<b>2</b>	<b><math>N/2</math></b>	<b><math>N/4</math></b>	<b>2</b>	
<b>2D Mesh</b>	<b>4</b>	<b><math>2 (N^{1/2} - 1)</math></b>	<b><math>2/3 N^{1/2}</math></b>	<b><math>N^{1/2}</math></b>	<b>63 (21)</b>
<b>2D Torus</b>	<b>4</b>	<b><math>N^{1/2}</math></b>	<b><math>1/2 N^{1/2}</math></b>	<b><math>2N^{1/2}</math></b>	<b>32 (16)</b>
<b>k-ary n-cube</b>	<b><math>2n</math></b>	<b><math>nk/2</math></b>	<b><math>nk/4</math></b>	<b><math>nk/4</math></b>	<b>15 (7.5) @n=3</b>
<b>Hypercube</b>	<b><math>n = \log N</math></b>	<b><math>n</math></b>	<b><math>n/2</math></b>	<b><math>N/2</math></b>	<b>10 (5)</b>

$N = \# \text{ nodes}$

# *Communication Performance*

- **Time(n)<sub>s-d</sub> = overhead + routing delay + channel occupancy + contention delay**
- **occupancy = (n + n<sub>h</sub>) / b**

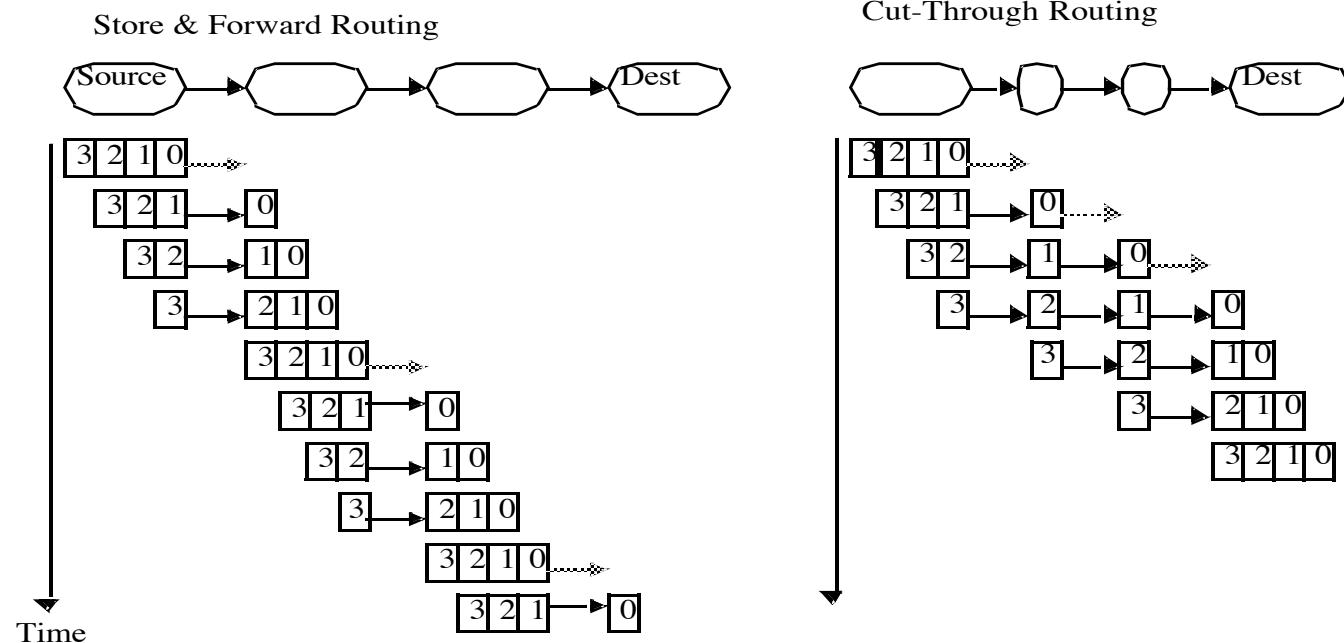
**n = message #bytes**

**n<sub>h</sub> = header #bytes**

**b = bitrate of communication link**

- What is the routing delay?
- Does contention occur and what is the cost?

# *Store-and-Forward vs. Cut-Through Routing*



- $h(n/b + \Delta)$   $n/b + h \Delta$
- What if message is fragmented?
- *Wormhole* vs. *Virtual cut-through*

# *Networks of Real Machines (circa 2000)*

Machine	Topology	Speed	Width	Delay	Flit
nCUBE/2	hypercube	25 ns	1	40 cycles	32
CM-5	fat-tree	25 ns	4	10 cycles	4
SP-2	banyan	25 ns	8	5 cycles	16
Paragon	2D mesh	11.5 ns	16	2 cycles	16
T3D	3D torus	6.67 ns	16	2 cycles	16
DASH	torus	30 ns	16	2 cycles	16
Origin	hypercube	2.5 ns	20	16 cycles	160
Myrinet	arbitrary	6.25 ns	16	50 cycles	16

A message is broken up into *flits* for transfer.

# *Message Passing Model*

- Hardware maintains send and receive message buffers
- Send message (synchronous)
  - Build message in local message send buffer
  - Specify receive location (processor id)
  - Initiate send and wait for receive acknowledge
- Receive message (synchronous)
  - Allocate local message receive buffer
  - Receive message byte stream into buffer
  - Verify message (e.g., checksum) and send acknowledge
- Memory to memory copy with acknowledgement and pairwise synchronization

## *Advantages of Shared Memory Architectures*

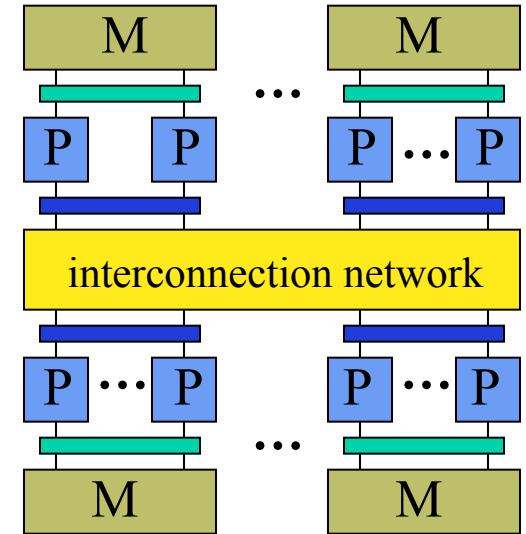
- Compatibility with SMP hardware
- Ease of programming when communication patterns are complex or vary dynamically during execution
- Ability to develop applications using familiar SMP model, attention only on performance critical accesses
- Lower communication overhead, better use of BW for small items, due to implicit communication and memory mapping to implement protection in hardware, rather than through I/O system
- HW-controlled caching to reduce remote communication by caching of all data, both shared and private

## *Advantages of Distributed Memory Architectures*

- The hardware can be simpler (especially versus NUMA) and is more scalable
- Communication is explicit and simpler to understand
- Explicit communication focuses attention on costly aspect of parallel computation
- Synchronization is naturally associated with sending messages, reducing the possibility for errors introduced by incorrect synchronization
- Easier to use sender-initiated communication, which may have some advantages in performance

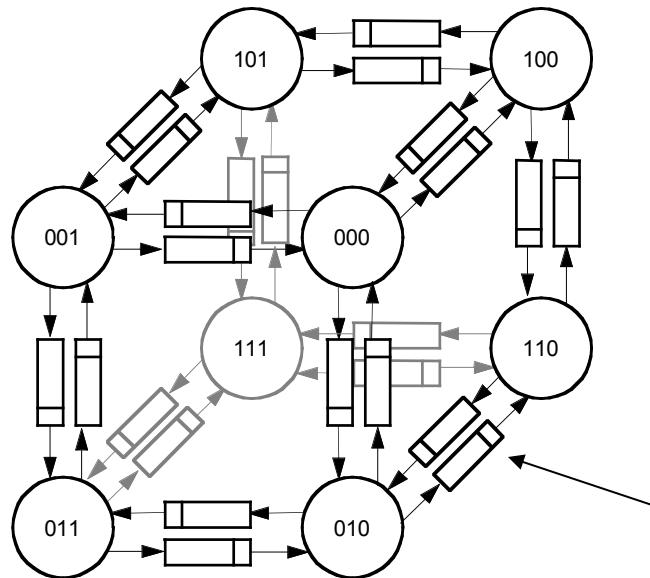
# *Clusters of SMPs*

- Clustering
  - Integrated packaging of nodes
- Motivation
  - Ammortize node costs by sharing packaging and resources
  - Reduce network costs
  - Reduce communications bandwidth requirements
  - Reduce overall latency
  - More parallelism in a smaller space
  - Increase node performance
- Scalable parallel systems today are built as SMP clusters

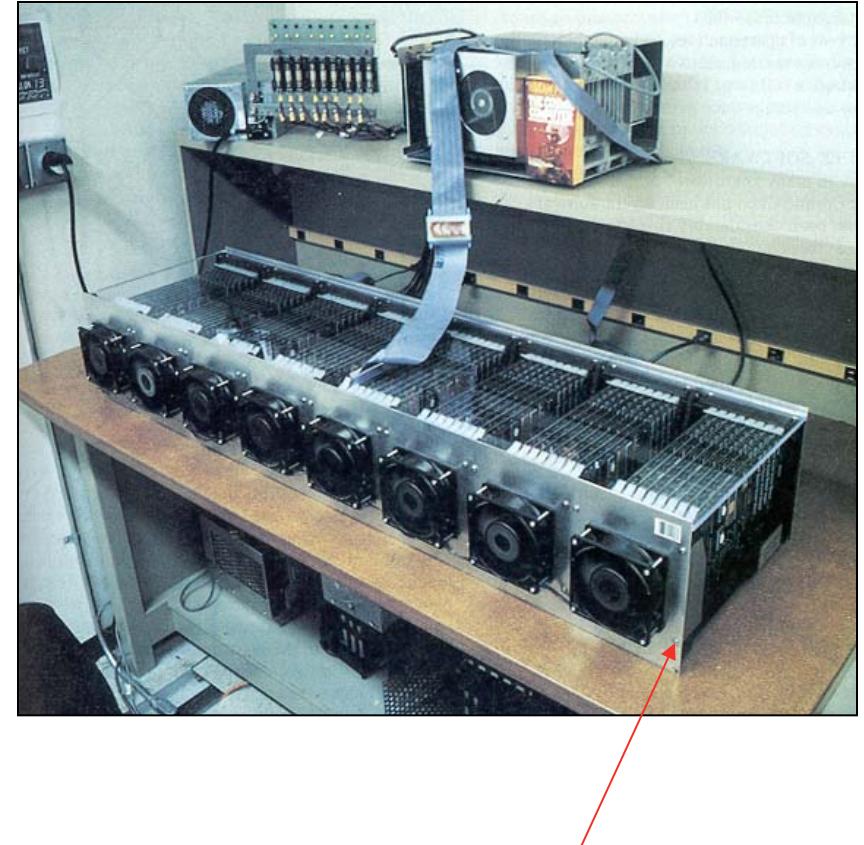


# *CalTech Cosmic Cube*

- First distributed memory message passing system
- Hypercube-based communications network



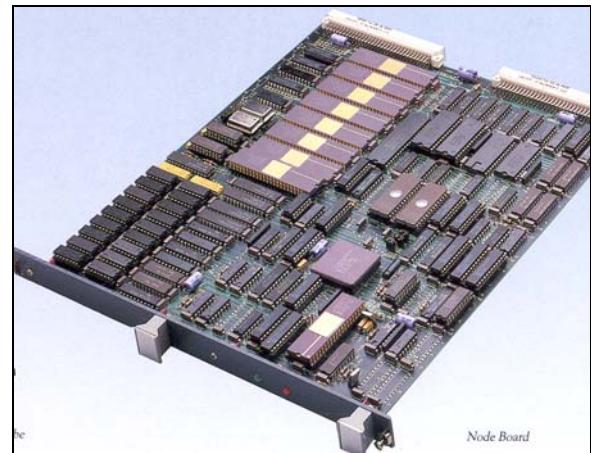
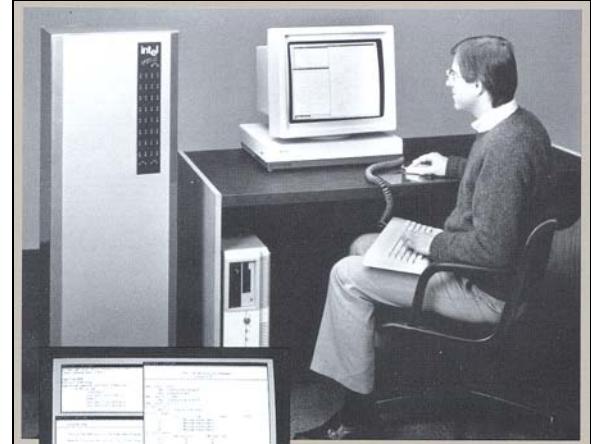
FIFO on each link  
- store and forward



- Chuck Seitz, Geoffrey Fox

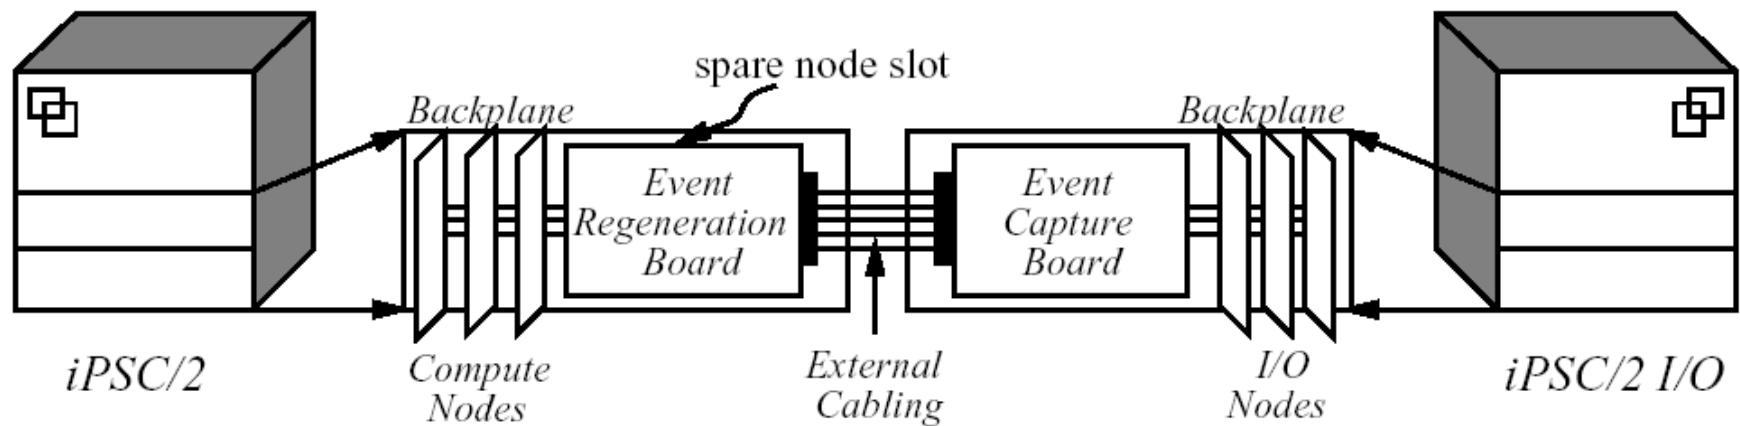
# *Intel iPSC/1, iPSC/2, iPSC/860*

- Shift to general links
  - DMA, enabling non-blocking ops
    - Buffered by system at destination until recv
  - Store&forward routing
- Diminishing role of topology
  - Any-to-any pipelined routing
  - node-network interface dominates communication time
  - Simplifies programming
  - Allows richer design space



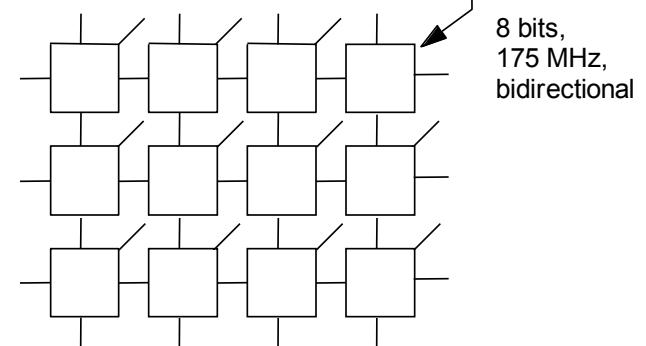
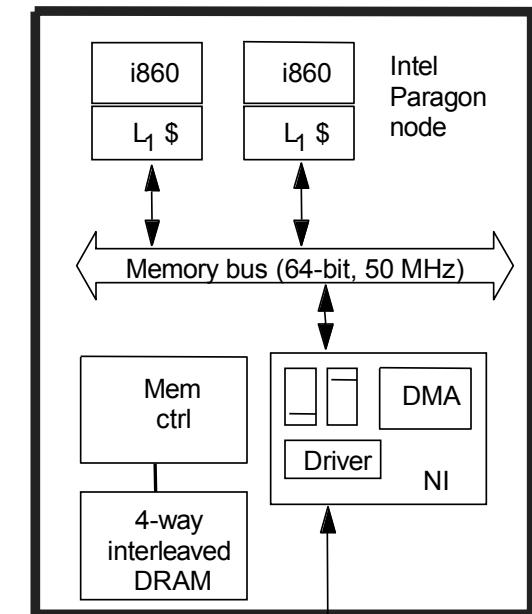
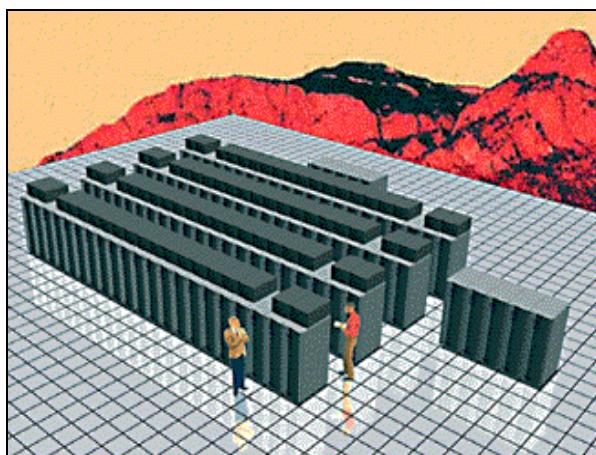
# *HyperMon Architecture (Who built this?)*

- Develop hardware support for tracing
  - Reduces intrusion trace buffering and I/O
- Hardware design
  - Memory-mapped interface
  - Synchronized timers and automatic timestamping
  - Support for event bursts and off-line streaming



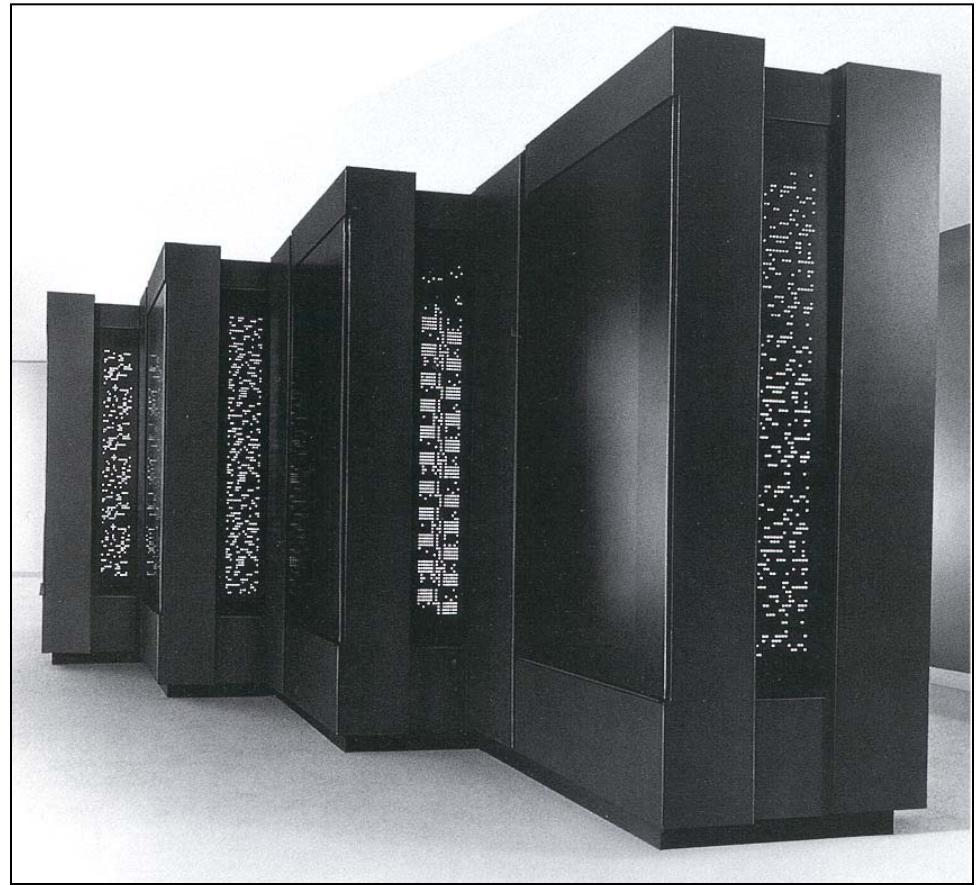
# *Intel Paragon and ASCI Red*

- DARPA project machine
  - Intel i860 processor
  - 2D grid network with processor node attached to every switch
  - 8bit, 175 MHz bidirectional links
- Forerunner design for ASCI Red
  - First Teraflop computer



# *Thinking Machine CM-5*

- Repackaged SparcStation
  - 4 per board
- Fat-Tree network
- Control network for global synchronization
- Suffered from hardware design and installation problems



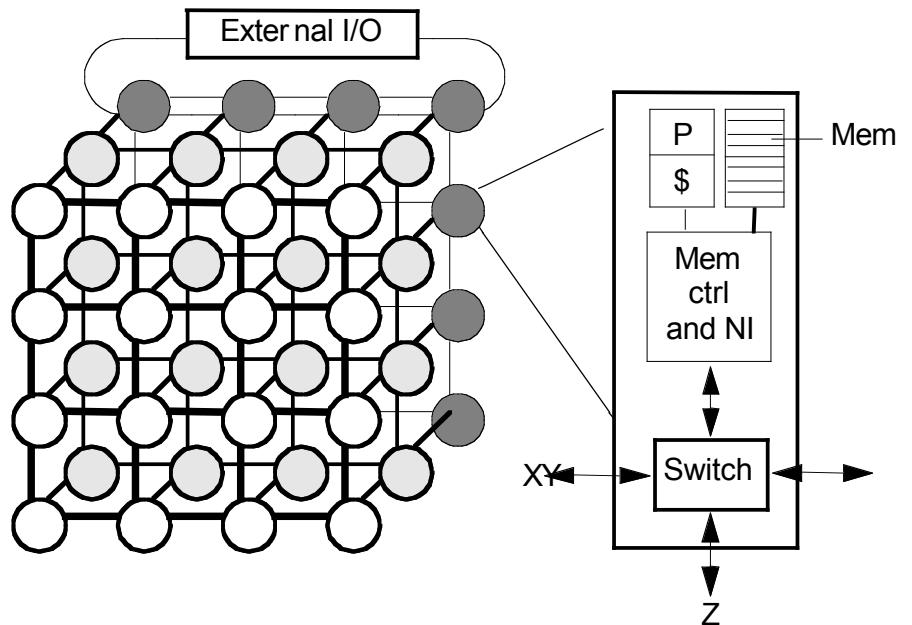
# *Berkeley Network Of Workstations (NOW)*

- 100 Sun Ultra2 workstations
- Intelligent network interface
  - proc + mem
- Myrinet network
  - 160 MB/s per link
  - 300 ns per hop



# Cray T3E

- Up to 1024 nodes
- 3D torus network
  - 480 MB/s links
- No memory coherence
- Access remote memory
  - Converted to messages
  - SHared MEMory communication
    - *put / get* operations
- Very successful machine

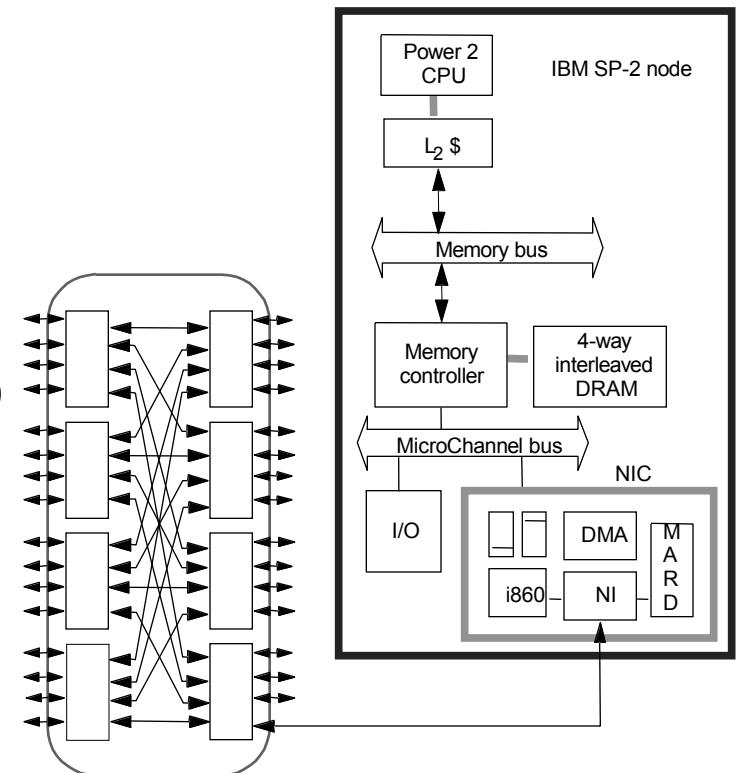


# IBM SP-2

- Made out of essentially complete RS6000 workstations
- Network interface integrated in I/O bus
- SP network very advanced
  - Formed from 8-port switches
- Predecessor design to
  - ASCI Blue Pacific (5856 CPUs)
  - ASCI White (8192 CPUs)



○ ASCI White (8192 CPUs)



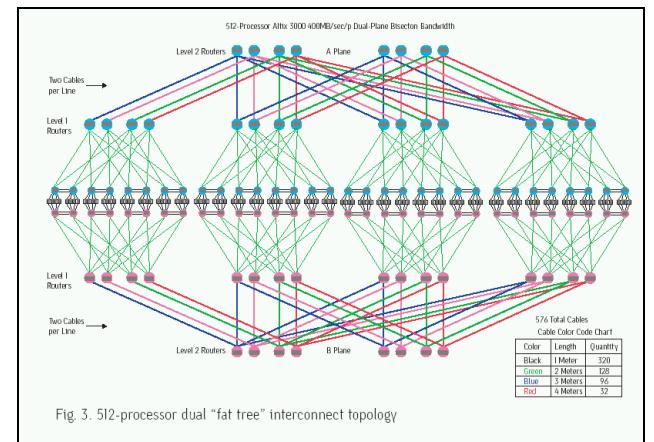
# NASA Columbia

## □ System hardware

- 20 SGI Altix 3700 superclusters
  - 512 Itanium2 processors (1.5 GHz)
  - 1 TB memory
- 10,240 processors (now 13,312)
- NUMAflex architecture
- NUMAlink “fat tree” network
- Fully shared memory!!!

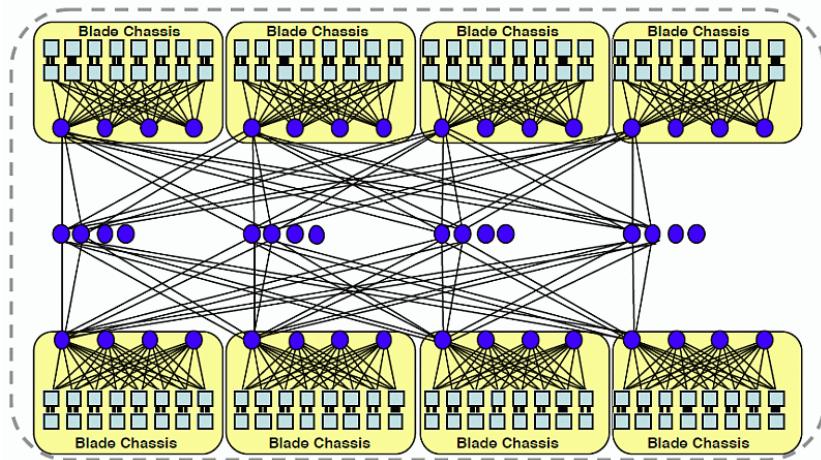
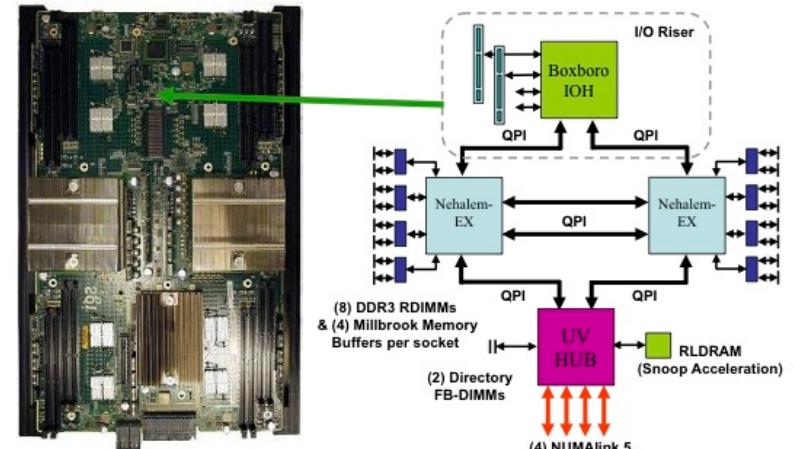
## □ Software

- Linux with PBS Pro job scheduling
- Intel Fortran/C/C++ compilers



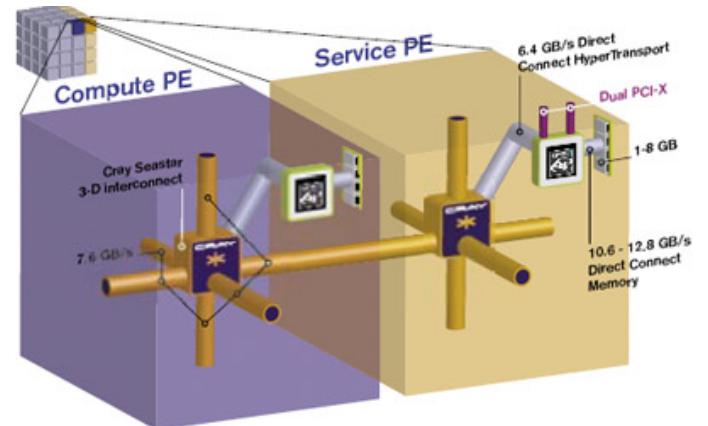
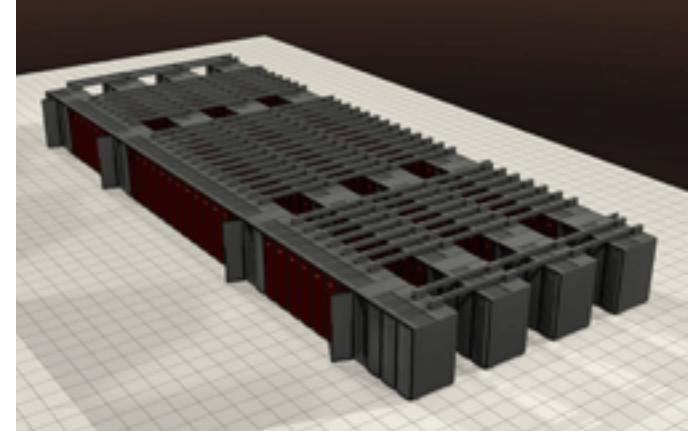
# *SGI Altix UV*

- Latest generation scalable shared memory architecture
- Scaling from 32 to 2,048 cores
  - Intel Nehalem EX
- Architectural provisioning for up to 262,144 cores
- Up to 16 terabytes of global shared memory in a single system image (SSI)
- High-speed 15GB per second interconnect NUMAlink 5



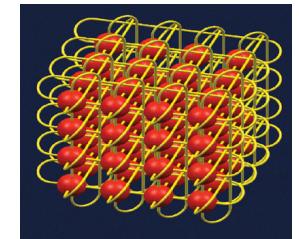
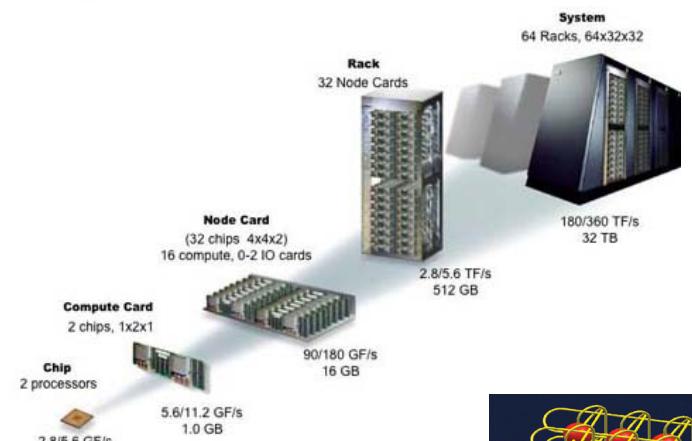
# *Sandia Red Storm*

- System hardware
  - Cray XT3
  - 135 compute node cabinets
  - 12,960 processors
    - AMD Opteron dual-core
  - 320 / 320 service / I/O node processors
  - 40 TB memory
  - 340 TB disk
  - 3D mesh interconnect
- Software
  - Catamount compute node kernel
  - Linux I/O node kernel
  - MPI



# LLNL BG/L

- System hardware
  - IBM BG/L (BlueGene)
  - 65,536 dual-processor compute nodes
    - PowerPC processors
    - “double hummer” floating point
  - I/O node per 32 compute nodes
  - 32x32x64 3D torus network
  - Global reduction tree
  - Global barrier and interrupt networks
  - Scalable tree network for I/O
- Software
  - Compute node kernel (CNK)
  - Linux I/O node kernel (ION)
  - MPI
  - Different operating modes



# *Tokyo Institute of Technology TSUBAME*

- System hardware
  - 655 Sun Fire X4600 servers
  - 11,088 processors
    - AMD Opteron dual-core
  - ClearSpeed accelerator
  - InfiniBand network
  - 21 TB memory
  - 42 Sun Fire X4500 servers
  - 1 PB of storage space

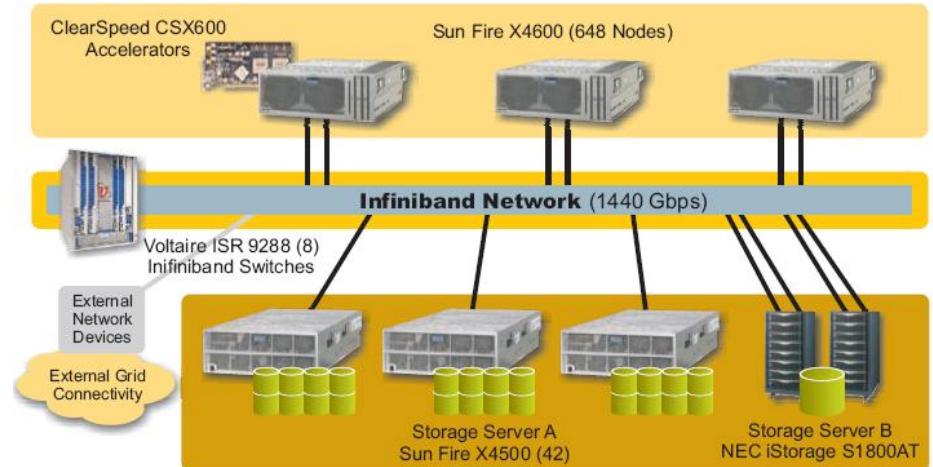
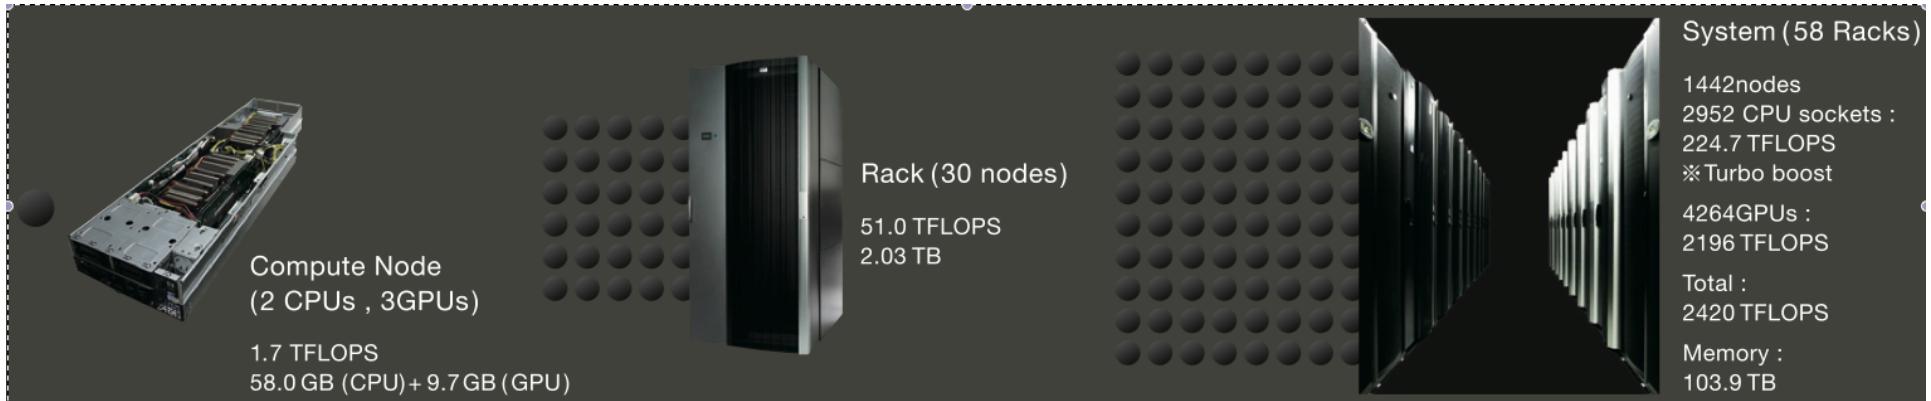


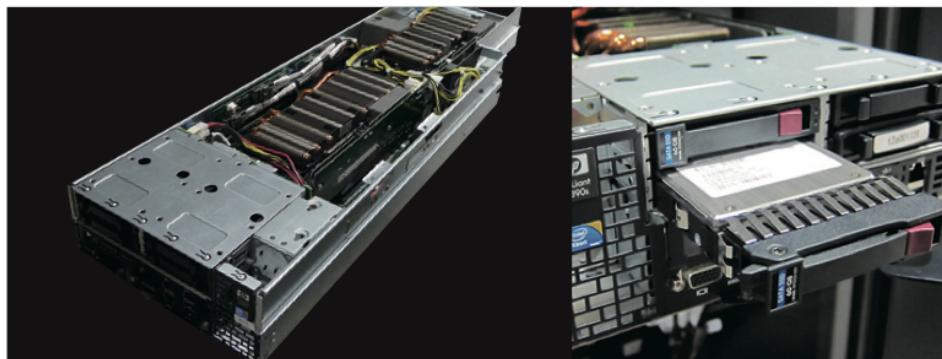
Figure 1. The TSUBAME grid system architecture

- Software
  - SuSE Linux Enterprise Server 9 SP3
  - Sun N1 Grid Engine 6.0
  - Lustre Client Software

# *Tokyo Institute of Technology TSUBAME2*



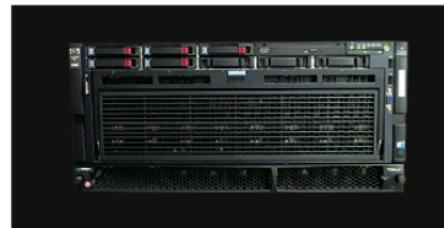
Thin Node 1408 nodes



HP ProLiant SL390s

GPU : NVIDIA Tesla M2050 (Fermi Core) x3 515GFLOPS VRAM 3GB/GPU  
CPU : Intel Xeon X5670 2.93GHz x2  
6 core/socket 76.7 GFLOPS (12cores/node) ※Turbo boost: 3.196GHz  
Memory : 58GB DDR3 1333MHz (partly 103GB)  
SSD : 60GB x2 (120GB/node) (partly 120GB x2 (240GB/node))

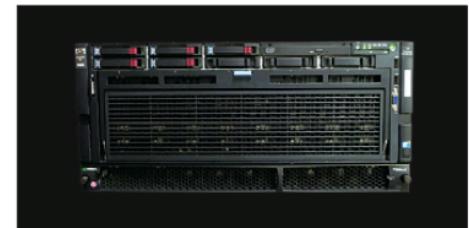
Medium Node 24 nodes



HP ProLiant DL580 G7

CPU: Intel Xeon X7550 (Nehalem-EX)  
2.0 GHz x4 sockets (32cores/node)  
GPU: NVIDIA Tesla S1070  
Memory: 137 GB (DDR3 1066MHz)  
SSD: 120GB x4 (480GB/node)  
Infiniband: QDR

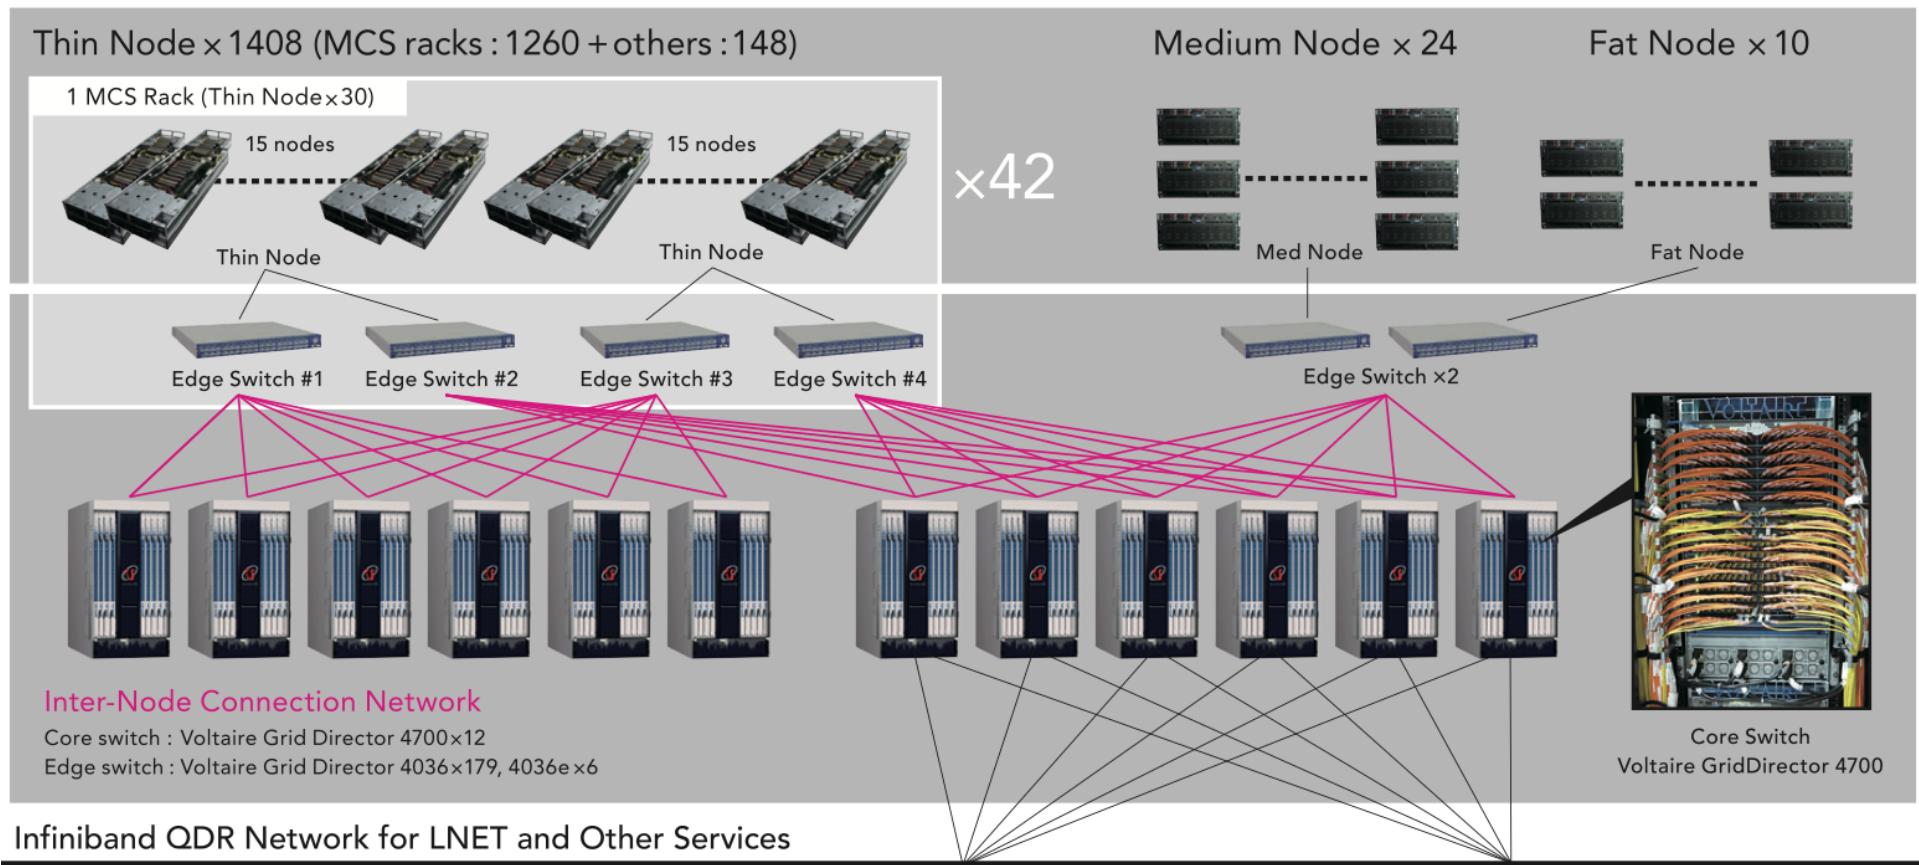
Fat Node 10 nodes



HP ProLiant DL580 G7

CPU: Intel Xeon X7550 (Nehalem-EX)  
2.0 GHz x4 sockets (32cores/node)  
GPU: NVIDIA Tesla S1070  
Memory: 274 GB (8 nodes),  
548 GB (2 nodes)  
DDR3 1066MHz  
SSD: 120GB x5 (600GB/node)  
Infiniband: QDR

# *TSUBAME2 – Interconnect*

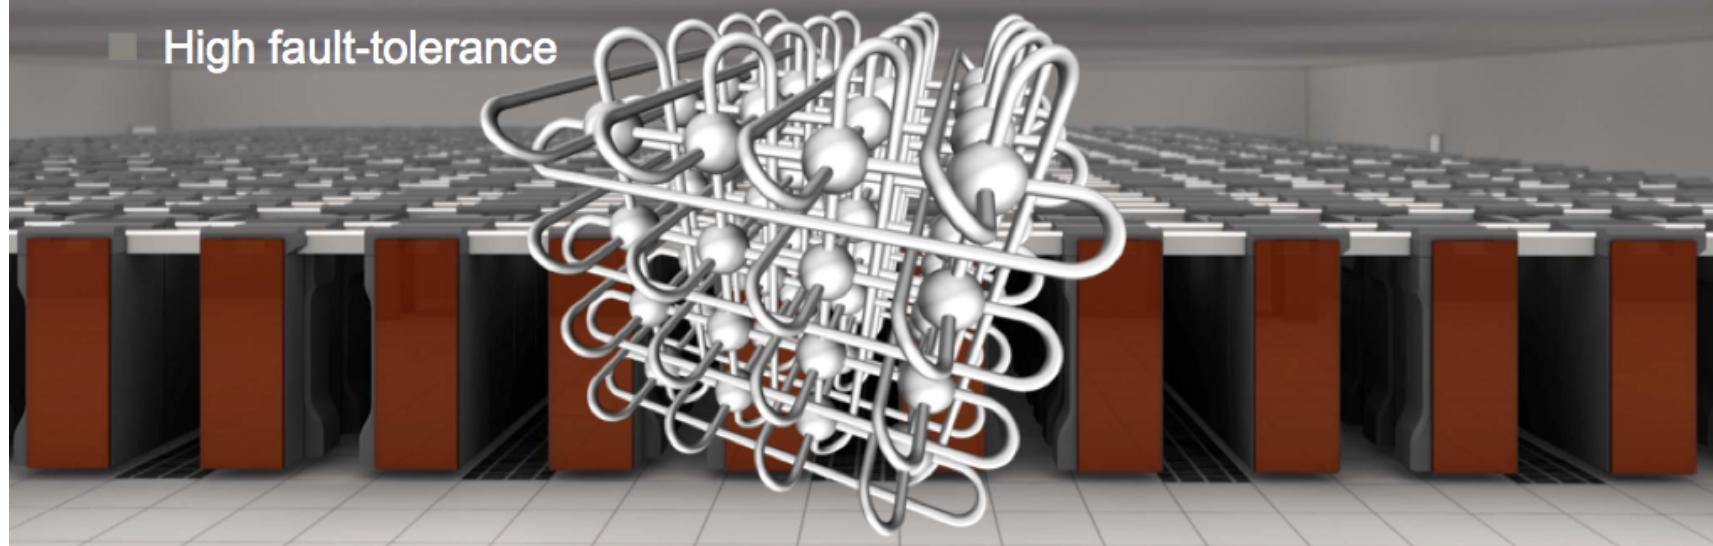


# *Japanese K Computer – Interconnect*

- 80,000 CPUs (SPARC64 VIIIfx), 640,000 cores
- 800 racks
- 8.6 Petaflops (Linpack)

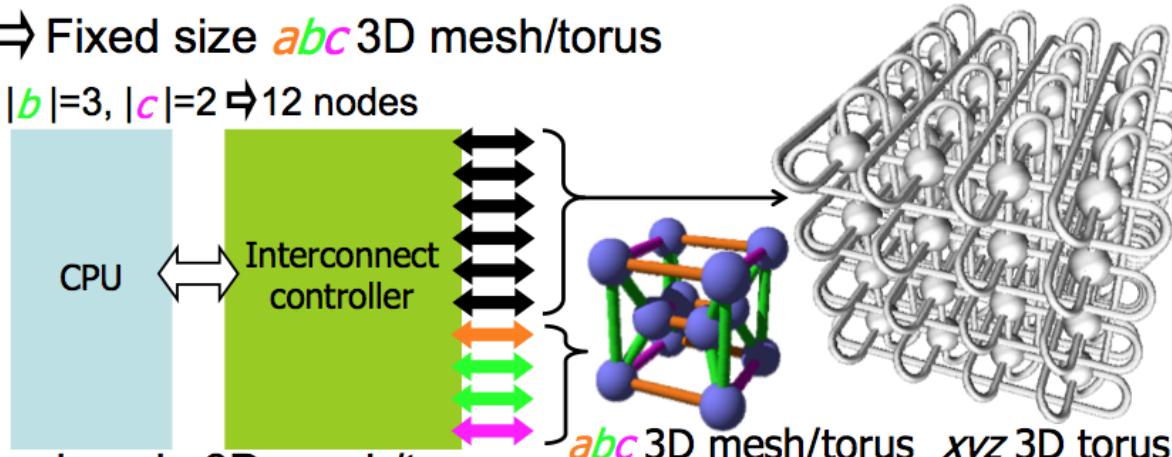
## ■ Tofu: Fujitsu's original 6D mesh/torus interconnect

- High communication performance
- High system scalability
- High fault-tolerance



# Japanese K Computer – Interconnect

- 6 links  $\Rightarrow$  Scalable  $xyz$  3D torus
- 4 links  $\Rightarrow$  Fixed size  $abc$  3D mesh/torus
  - $|a|=2, |b|=3, |c|=2 \Rightarrow 12$  nodes

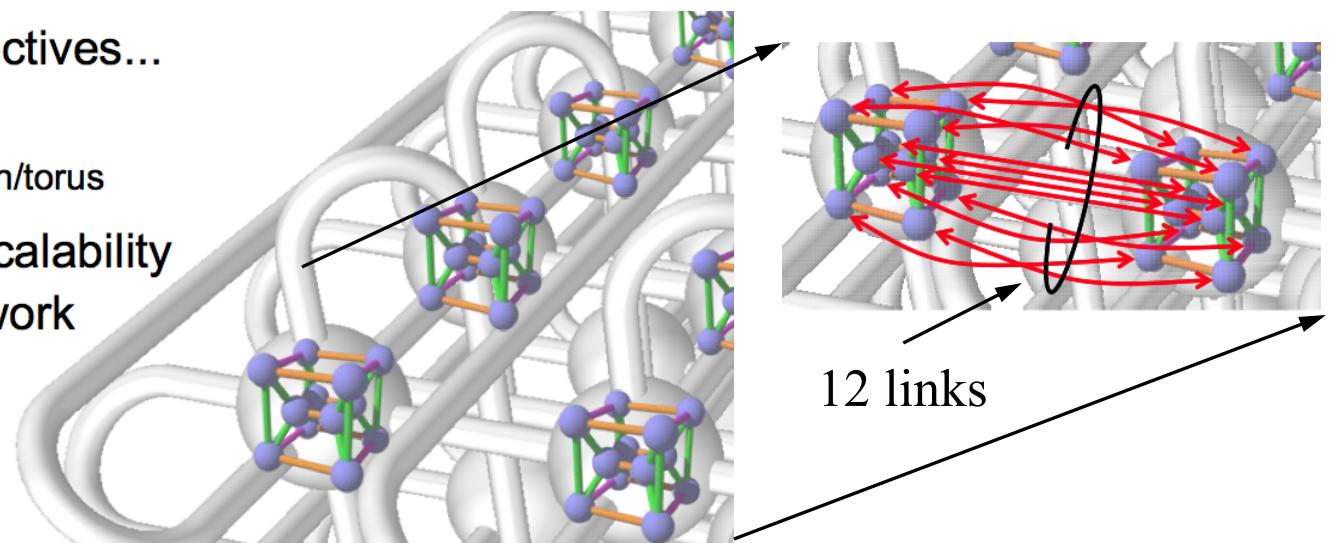


- Total topology is 6D mesh/torus
  - Cartesian product of  $xyz$  and  $abc$  mesh/torus

- From the other perspectives...

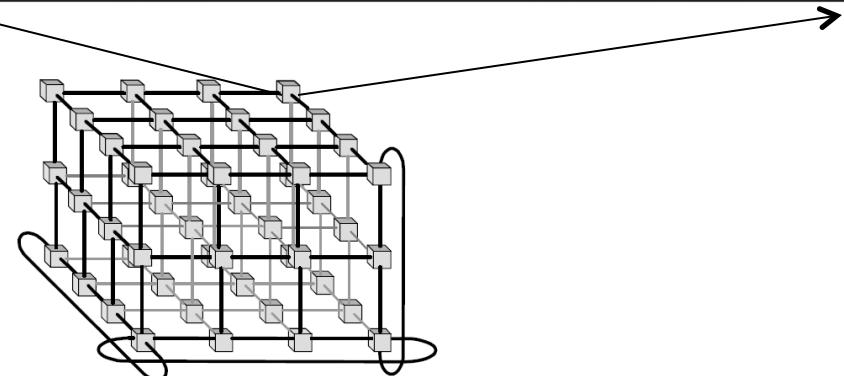
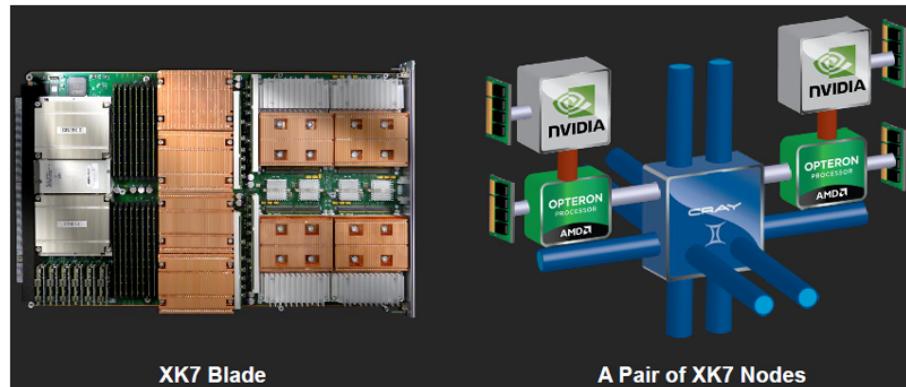
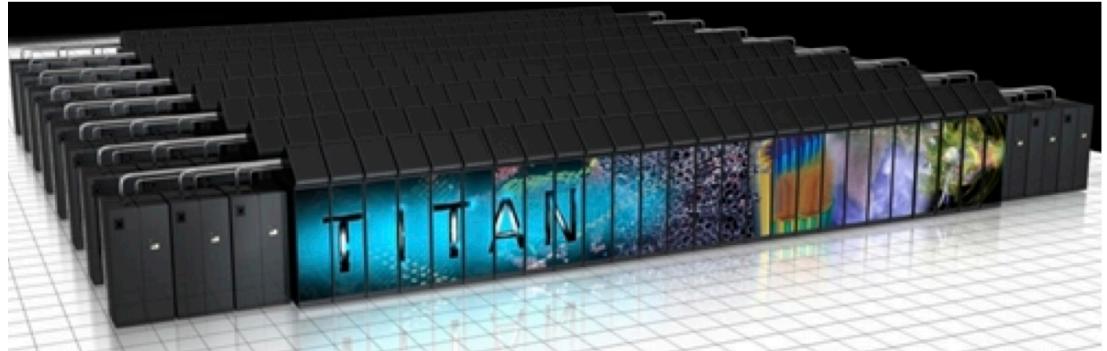
- Overlaid twelve  $xyz$  torus
  - X x Y x Z array of  $abc$  mesh/torus

- Twelve times higher scalability than the 3D torus network



# *ORNL Titan (<http://www.olcf.ornl.gov/titan>)*

- Cray XK7
  - 18,688 nodes
  - AMD Opteron
    - 16-core Interlagos
    - 299,008 Opteron cores
  - NVIDIA K20x
    - 18,688 GPUs
    - 50,233,344 GPU cores
- Gemini interconnect
  - 3D torus
- 20+ petaflops



## *Next Class*

- Parallel algorithm design
- Parallel programming
- or
- Parallel performance models