

# ***CIS 631***

## ***Parallel Processing***

### ***Lecture 5: Parallel Programming***

**Allen D. Malony**  
malony@cs.uoregon.edu

Department of Computer and Information Science  
University of Oregon



## *Acknowledgements*

- Portions of the lectures slides were adopted from:
  - I. Foster, “Designing and Building Parallel Programs,” 1995.
  - John Mellor-Crummey, COMP 422, “Parallel Computing,” Rice University, Spring 2010.

## *Term Project*

- ❑ Work on a code for parallel particle advection
  - Contribute by Prof. Hank Childs
  - Start by reading “GPU Acceleration of Particle Advection Workloads in a Parallel, Distributed Memory Setting”
  - Skype with Prof. Childs on Wednesday during class meeting to discuss
  - Verbal agreement needed to keep confidential the paper, the code, the ideas contained in each
- ❑ Work with David Ozog’s Framework for Parallel Task Operation (FRAMPTON)
  - Extend how FRAMPTON can process task DAGs
  - Develop new tasks adaptors

# *Outline*

- ❑ Quick look at parallel models
- ❑ Parallelism
  - Where can you find parallelism in a computation?
  - Dependencies
- ❑ Different types of parallelism
  - data parallelism
  - task parallelism
- ❑ Parallel programming
  - Creating parallel programs
- ❑ Standard models of parallelism and parallel programs

# Parallel Models 101

## ❑ Sequential models


- von Neumann (RAM) model

## ❑ Parallel model

- A parallel computer is simply a collection of *processors interconnected* in some manner to *coordinate* activities and *exchange data*
- Models that can be used as general frameworks for describing and analyzing parallel algorithms
  - *Simplicity*: description, analysis, architecture independence
  - *Implementability*: able to be realized, reflect performance

## ❑ Three common parallel models

- Directed acyclic graphs, shared-memory, network



# *Directed Acyclic Graphs (DAG)*

- ❑ Captures data flow parallelism
- ❑ Nodes represent operations to be performed
  - Inputs are nodes with no incoming arcs
  - Output are nodes with no outgoing arcs
  - Think of nodes as tasks
- ❑ Arcs are paths for flow of data results
- ❑ DAG represents the operations of the algorithm and implies precedent constraints on their order

for (i=1; i<100; i++)


$a[i] = a[i-1] + 100;$



# *Shared Memory Model*

## ❑ Parallel extension of RAM model (PRAM)

- Memory size is infinite
- Number of processors is unbounded
- Processors communicate via the memory
- Every processor accesses any memory location in 1 cycle
- Synchronous
  - All processors execute same algorithm synchronously
    - READ phase
    - COMPUTE phase
    - WRITE phase
  - Some subset of the processors can stay idle
- Asynchronous




## *Memory Access in PRAM*

- ❑ Exclusive Read (ER):  $p$  processors can simultaneously read the content of  $p$  distinct memory locations
- ❑ Concurrent Read (CR):  $p$  processors can simultaneously read the content of  $p'$  memory locations, where  $p' < p$
- ❑ Exclusive Write (EW):  $p$  processors can simultaneously write the content of  $p$  distinct memory locations
- ❑ Concurrent Write (CW):  $p$  processors can simultaneously write the content of  $p'$  memory locations, where  $p' < p$
- ❑ EREW and ERCW (weird)
- ❑ CREW and CRCW



# Network Model

- ❑  $G = (N, E)$ 
  - $N$  are processing nodes
  - $E$  are bidirectional communication links
- ❑ Each processor has its own memory
- ❑ No shared memory is available
- ❑ Network operation may be synchronous or asynchronous
- ❑ Requires communication primitives
  - Send ( $X, i$ )
  - Receive ( $Y, j$ )
- ❑ Captures message passing model for algorithm design



# *Parallelism*

- ❑ Ability to execute different parts of a computation concurrently on different machines
- ❑ Why do you want parallelism?
  - Shorter running time or handling more work
- ❑ What is being parallelized?
  - Task: instruction, statement, procedure, ...
  - Data: data flow, size, replication
  - Parallelism granularity
    - Coarse-grain versus fine-grained
- ❑ Thinking about parallelism
- ❑ Evaluation

## *Why is parallel programming important today?*

- ❑ Parallel programming has matured
  - Standard programming models
  - Common machine architectures
  - Programmer can focus on computation and use suitable programming model for implementation
- ❑ Increasing portability between models and architectures
- ❑ Reasonable hope of portability across platforms
- ❑ Problem
  - Performance optimization is still platform-dependent
  - Performance portability is a problem
  - Parallel programming methods are still evolving

# *Parallel Algorithm*

- ❑ Recipe to solve a problem “in parallel” on multiple processing elements
- ❑ Standard steps for constructing a parallel algorithm
  - Identify work that can be performed concurrently
  - Partition the concurrent work on separate processors
  - Properly manage input, output, and intermediate data
  - Coordinate data accesses and work to satisfy dependencies
- ❑ Which are hard to do?

# *Parallelism Views*

- ❑ Where can we find parallelism?
- ❑ Program (task) view
  - Statement level
    - Between program statements
    - Which statements can be executed at the same time?
  - Block level / Loop level / Routine level / Process level
    - Larger-grained program statements
- ❑ Data view
  - How is data operated on?
  - Where does data reside?
- ❑ Resource view

## *Parallelism, Correctness, and Dependence*

- ❑ Parallel execution, from any point of view, will be constrained by the sequence of operations needed to be performed for a correct result
- ❑ Parallel execution must address control, data, and system dependences
- ❑ A *dependency* arises when one operation depends on an earlier operation to complete and produce a result before this later operation can be performed
- ❑ We extend this notion of dependency to resources since some operations may depend on certain resources
  - For example, due to where data is located

## *Executing Two Statements in Parallel*

- ❑ Want to execute two statements in parallel
- ❑ On one processor:
  - Statement 1;
  - Statement 2;
- ❑ On two processors:

Processor 1:	Processor 2:
Statement 1;	Statement 2;
- ❑ Fundamental (*concurrent*) execution assumption
  - Processors execute independent of each other
  - No assumptions made about speed of processor execution

# *Sequential Consistency in Parallel Execution*

## ❑ Case 1:

Processor 1:	Processor 2:	time
statement 1;		↓
	statement 2;	

## ❑ Case 2:

Processor 1:	Processor 2:	time
	statement 2;	↓
statement 1;		

## ❑ Sequential consistency

- Statements execution does not interfere with each other
- Computation results are the same (independent of order)



## *Independent versus Dependent*

- ❑ In other words the execution of  
statement1;  
statement2;  
must be equivalent to  
statement2;  
statement1;
- ❑ Their order of execution must not matter!
- ❑ If true, the statements are *independent* of each other
- ❑ Two statements are *dependent* when the order of their execution affects the computation outcome

## Examples

### ❑ Example 1

S1: a=1;

S2: b=1;

### ❑ Example 2

S1: a=1;

S2: b=a;

### ❑ Example 3

S1: a=f(x);

S2: a=b;

### ❑ Example 4

S1: a=b;

S2: b=1;

### ❑ Statements are independent

### ❑ Dependent (*true (flow) dependence*)

○ Second is dependent on first

○ Can you remove dependency?

### ❑ Dependent (*output dependence*)

○ Second is dependent on first

○ Can you remove dependency? How?

### ❑ Dependent (*anti-dependence*)

○ First is dependent on second

○ Can you remove dependency? How?

# True Dependence and Anti-Dependence

- Given statements S1 and S2,  
     S1;  
     S2;
- S2 has a *true (flow) dependence* on S1  
     if and only if  
     S2 reads a value written by S1
- S2 has a *anti-dependence* on S1  
     if and only if  
     S2 writes a value read by S1


$$\begin{array}{c} X = \\ \vdots \\ \quad \swarrow \delta \\ = X \end{array}$$

$$\begin{array}{c} = X \\ \vdots \\ \quad \swarrow \delta^{-1} \\ X = \end{array}$$

# *Output Dependence*

- Given statements S1 and S2,  
    S1;  
    S2;
- S2 has an *output dependence* on S1  
    if and only if  
    S2 writes a variable written by S1

X =  
  :  
X =



The diagram shows a vertical sequence of three lines: 'X =', ':', and 'X ='. To the right of the colon is a curved arrow pointing from the 'X =' line above to the 'X =' line below, with the label  $\delta^0$  next to it.

- Anti- and output dependences are “name” dependencies
  - Are they “true” dependences?
- How can you get rid of output dependences?
  - Are there cases where you can not?

# Statement Dependency Graphs

□ Can use graphs to show dependence relationships


□ Example

S1: a=1;

S2: b=a;

S3: a=b+1;

S4: c=a;



□  $S_2 \delta S_3$  :  $S_3$  is flow-dependent on  $S_2$

□  $S_1 \delta^0 S_3$  :  $S_3$  is output-dependent on  $S_1$

□  $S_2 \delta^{-1} S_3$  :  $S_3$  is anti-dependent on  $S_2$

## *When can two statements execute in parallel?*

- ❑ Statements S1 and S2 can execute in parallel if and only if there are *no dependences* between S1 and S2
  - True dependences
  - Anti-dependences
  - Output dependences
- ❑ Some dependences can be removed by modifying the program
  - Rearranging statements
  - Eliminating statements

## *How do you compute dependence?*

- ❑ Data dependence relations can be found by comparing the IN and OUT sets of each node
- ❑ The IN and OUT sets of a statement **S** are defined as:
  - **IN(S)** : set of memory locations (variables) that may be used in **S**
  - **OUT(S)** : set of memory locations (variables) that may be modified by **S**
- ❑ Note that these sets include all memory locations that may be fetched or modified
- ❑ As such, the sets can be conservatively large

## *IN and OUT Sets and Computing Dependence*

- Assuming that there is a path from **S1** to **S2** , the following shows how to intersect the IN and OUT sets to test for data dependence

$out(S_1) \cap in(S_2) \neq \emptyset$        $S_1 \delta S_2$       flow dependence

$in(S_1) \cap out(S_2) \neq \emptyset$        $S_1 \delta^{-1} S_2$       anti - dependence

$out(S_1) \cap out(S_2) \neq \emptyset$        $S_1 \delta^0 S_2$       output dependence



## *Loop-Level Parallelism*

- ❑ Significant parallelism can be identified within loops

```
for (i=0; i<100; i++)  
    S1: a[i] = i;
```

```
for (i=0; i<100; i++) {  
    S1: a[i] = i;  
    S2: b[i] = 2*i;  
}
```

- ❑ Dependencies? What about  $i$ , the loop index?
- ❑ *DOALL* loop
  - All iterations are independent of each other
  - All statements be executed in parallel at the same time
    - Is this really true?


## *Iteration Space*

- ❑ Unroll loop into separate statements / iterations
- ❑ Show dependences between iterations

```
for (i=0; i<100; i++)  
    S1: a[i] = i;
```



```
for (i=0; i<100; i++) {  
    S1: a[i] = i;  
    S2: b[i] = 2*i;  
}
```



## ***Multi-Loop Parallelism***

- ❑ Significant parallelism can be identified between loops

for (i=0; i<100; i++) a[i] = i;

for (i=0; i<100; i++) b[i] = i;



- ❑ Dependencies?
- ❑ How much parallelism is available?
- ❑ Given 4 processors, how much parallelism is possible?
- ❑ What parallelism is achievable with 50 processors?

# Loops with Dependencies

Case 1:

```
for (i=1; i<100; i++)  
    a[i] = a[i-1] + 100;
```



❑ Dependencies?


○ What type?

❑ Is the Case 1 loop parallelizable?

❑ Is the Case 2 loop parallelizable?

Case 2:

```
for (i=5; i<100; i++)  
    a[i-5] = a[i] + 100;
```



## *Another Loop Example*

```
for (i=1; i<100; i++)  
    a[i] = f(a[i-1]);
```

- ❑ Dependencies?
  - What type?
- ❑ Loop iterations are not parallelizable
  - Why not?

# Loop Dependencies

- ❑ A *loop-carried* dependence is a dependence that is present only if the statements are part of the execution of a loop (i.e., between two statements instances in two different iterations of a loop)
- ❑ Otherwise, it is *loop-independent*, including between two statements instances in the same loop iteration
- ❑ Loop-carried dependences can prevent loop iteration parallelization
- ❑ The dependence is *lexically forward* if the source comes before the target or *lexically backward* otherwise
  - Unroll the loop to see


## *Loop Dependence Example*

```
for (i=0; i<100; i++)  
    a[i+10] = f(a[i]);
```

- ❑ Dependencies?
  - Between a[10], a[20], ...
  - Between a[11], a[21], ...
- ❑ Some parallel execution is possible
  - How much?

# Iteration Dependence and Pipelining

```
for (i=1; i<100; i++) {  
    S1: a[i] = ...;  
    S2: ... = a[i-1];  
}
```



## ❑ Dependencies?

- Between  $a[i]$  and  $a[i-1]$

## ❑ Is parallelism possible?

- Statements can be executed in pipelined parallel



## *Another Loop Dependence Example*

```
for (i=0; i<100; i++)  
    for (j=1; j<100; j++)  
        a[i][j] = f(a[i][j-1]);
```

- ❑ Dependencies?
  - Loop-independent dependence on i
  - Loop-carried dependence on j
- ❑ Which loop can be parallelized?
  - Outer loop parallelizable
  - Inner loop cannot be parallelized

## *Still Another Loop Dependence Example*

```
for (j=1; j<100; j++)  
    for (i=0; i<100; i++)  
        a[i][j] = f(a[i][j-1]);
```

- ❑ Dependencies?
  - Loop-independent dependence on i
  - Loop-carried dependence on j
- ❑ Which loop can be parallelized?
  - Inner loop parallelizable
  - Outer loop cannot be parallelized
  - Less desirable (why?)

## *Indirect Indexing and Dependences*

```
for (i=0; i<100; i++)  
    a[i] = f(a[index[i]]);
```

- ❑ Dependencies?
  - Cannot tell for sure
- ❑ Parallelization depends on knowledge of index values
  - User may know
  - Compiler does not know
  - User could inform the compiler

## *Hidden Dependencies – Printing*

```
printf("a");  
printf("b");
```

- ❑ Statements have a hidden output dependence
  - Due to the serial output stream

## *Hidden Dependences – Functions*

$a = f(x);$

$b = g(x);$

- Statements could have hidden dependence if  $f()$  and  $g()$  update the same variable through side effects

## *Parallelizing Compilers*

- ❑ Parallelizing compilers analyze program dependences to decide parallelization
- ❑ In parallelization by hand, user does the same analysis.
- ❑ Compiler more convenient and more correct
- ❑ User more knowledgeable
  - Can analyze more patterns

## *Key Ideas for Dependency Analysis*

- ❑ To execute in parallel:
  - Statement order must not matter
  - Statements must not have dependences
- ❑ Some dependences can be removed
- ❑ Some dependences may not be obvious

# *Dependencies and Synchronization*

- ❑ How is parallelism achieved when have dependencies?
  - Think about concurrency
  - Some parts of the execution are independent
  - Some parts of the execution are dependent
- ❑ Must control ordering of events on different processors
  - Dependencies pose constraints on parallel event ordering
  - Partial ordering of execution action
- ❑ Use synchronization mechanisms
  - Need for concurrent execution too
  - Maintains partial order




# Synchronization Primitives

- ❑ Suppose we had a set of primitives, **signal**(x) and **wait**(x)
- ❑ **wait**(x) blocks unless a **signal**(x) has occurred.
- ❑ **signal**(x) does not block, but causes a **wait**(x) to unblock, or causes a future **wait**(x) not to block



```
f() {  
    a=1; b=2; c=3;  
}  
g() {  
    d=4; e=5; a=6;  
}  
main() { f(); g(); }
```

```
f() {  
    a=1; signal(e_a); b=2; c=3;  
}  
g() {  
    d=4; e=5; wait(e_a); a=6;  
}  
main() { f(); g(); }
```



# Synchronization in Loops

```
for (i=0; i<100; i++) {  
    a[i] = ...;  
    ...;  
    ... = a[i-1];  
}
```




```
for (i=0; i<100; i++) {  
    a[i] = ...;  
    signal(e_a[i]);  
    ...;  
    wait(e_a[i-1]);  
    ... = a[i-1];  
}
```

- ❑ Loop cannot be parallelized unless have synchronization!
- ❑ Does it matters which processors get which iterations?
- ❑ This is called a *DOACROSS* loop
- ❑ How could you parallelize this without synchronization?

# *Fork-Join Parallelism*

```
x = g(a);  
for( i=0; i<100; i++ ) a[i] = f(i);  
y = h(a);  
for( i=0; i<100; i++ ) b[i] = x + h( a[i]);
```



- ❑ First loop is a DOALL loop
- ❑ Middle statement is sequential
- ❑ Second loop is a DOALL loop
- ❑ Execution moves between sequential and parallel phases
- ❑ Call this *fork-join* parallelism
- ❑ Fork-join, loop-level parallelism is basis for OpenMP

## *Fork-Join and Barrier Synchronization*

- ❑ **fork()** causes a number of processes to be created and to be run in parallel
- ❑ **join()** causes all these processes to wait until all of them have executed a **join()** (*barrier* synchronization)

**fork()**;

for( i=0; i<100; i++ ) a[i] = f(i);

**join()**;

y = h(a);

**fork()**;

for( i=0; i<100; i++ ) b[i] = x + h(a[i]);

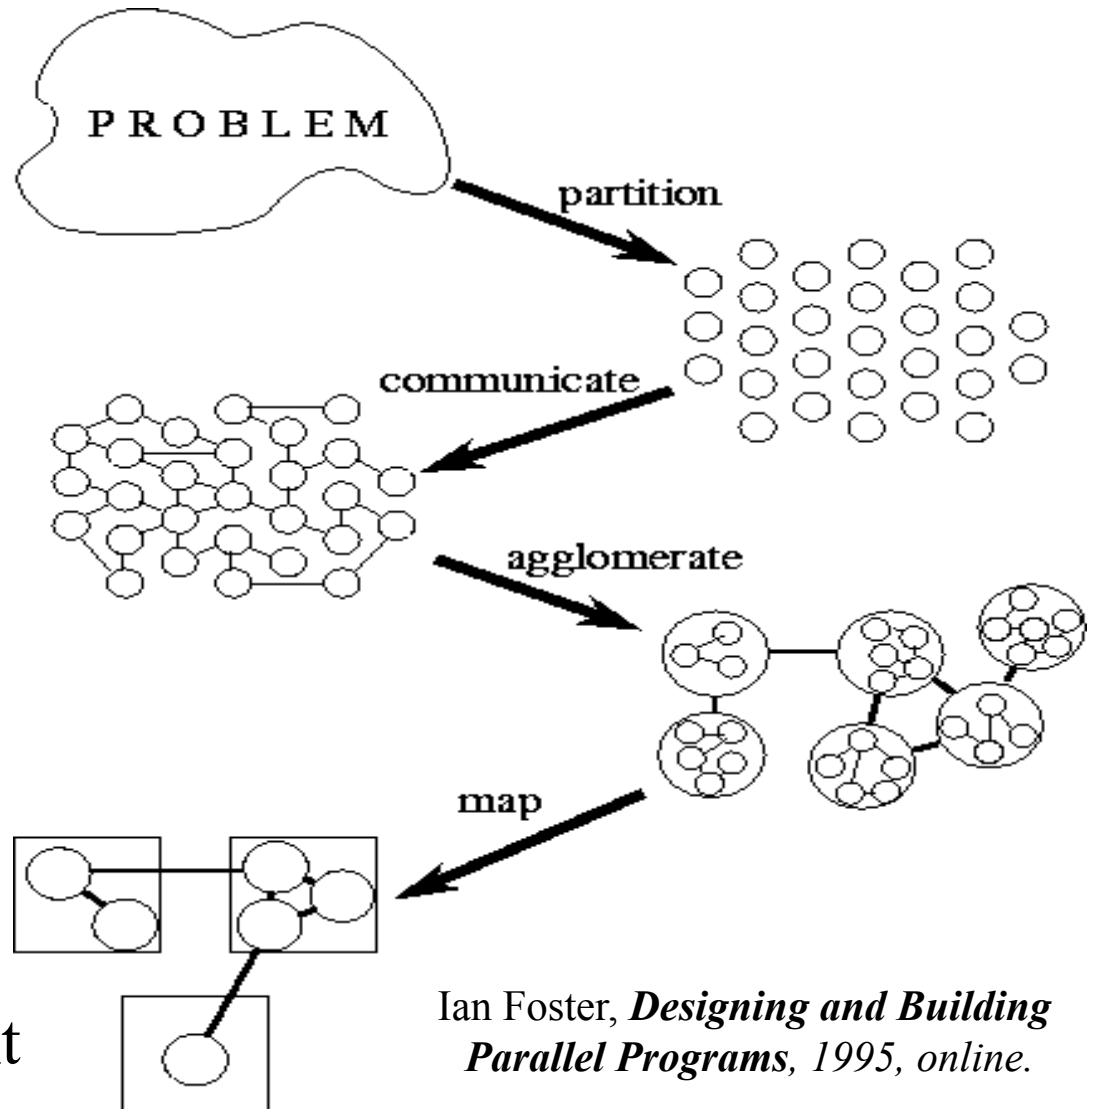
**join()**;

# *Synchronization Issues*

- ❑ Synchronization is necessary to make some programs execute correctly in parallel
- ❑ Dependences have to be “covered” by appropriate synchronization operations
- ❑ Different synchronization constructs exist in different parallel programming models
- ❑ However, synchronization is expensive
- ❑ To reduce synchronization
  - May need to limit parallelization
  - Look for opportunities to increase parallelism granularity

# Methodological Design

- ❑ Partition:
  - Task/data decomposition
- ❑ Communication
  - Task execution coordination
- ❑ Agglomeration
  - Evaluation of the structure
- ❑ Mapping
  - Resource assignment



Ian Foster, *Designing and Building Parallel Programs*, 1995, online.

## *Next Class*

- ❑ Parallel programming models
- ❑ Introduction to HPC Linux and LiveDVD