

CIS 631

Parallel Processing

Lecture 6: Parallel Programming

Allen D. Malony
malony@cs.uoregon.edu

Department of Computer and Information Science
University of Oregon



Acknowledgements


- Portions of the lectures slides were adopted from:
 - I. Foster, “Designing and Building Parallel Programs,” 1995.
 - Vijay Pai, COMP 422, “Parallel Programming,” Rice University, 2002.
 - A. Grama, A. Gupta, G. Karypis, and V. Kumar, “Introduction to Parallel Computing,” 2003.

Outline

- ❑ Dependency and Synchronization
- ❑ Methodological design of parallel programs
- ❑ Types of parallel programs
 - Data parallel vs. task parallel
 - Pipelining
 - Task graphs
 - Master-slave
 - Producer-consumer
 - Divide-and-conquer
 - SPMD
 - Loop scheduling

Fork-Join Parallelism

```
x = g(a);  
for( i=0; i<100; i++ ) a[i] = f(i);  
y = h(a);  
for( i=0; i<100; i++ ) b[i] = x + h( a[i]);
```



- ❑ First loop is a DOALL loop
- ❑ Middle statement is sequential
- ❑ Second loop is a DOALL loop
- ❑ Execution moves between sequential and parallel phases
- ❑ Call this *fork-join* parallelism

Fork-Join and Barrier Synchronization

- ❑ **fork()** causes a number of processes to be created and to be run in parallel
- ❑ **join()** causes all these processes to wait until all of them have executed a **join()** (*barrier* synchronization)

fork();

for(i=0; i<100; i++) a[i] = f(i);

join();

y = h(a);

fork();

for(i=0; i<100; i++) b[i] = x + h(a[i]);


join();

Synchronization Issues

- ❑ Synchronization is necessary to make some programs execute correctly in parallel
- ❑ Dependences have to be “covered” by appropriate synchronization operations
- ❑ Different synchronization constructs exist in different parallel programming models
- ❑ However, synchronization is expensive
- ❑ To reduce synchronization
 - May need to limit parallelization
 - Look for opportunities to increase parallelism granularity


Methodological Design

- ❑ Partition:
 - Task/data decomposition
- ❑ Communication
 - Task execution coordination
- ❑ Agglomeration
 - Evaluation of the structure
- ❑ Mapping
 - Resource assignment




Partitioning

- ❑ Partitioning stage is intended to expose opportunities for parallel execution
- ❑ Focus on defining large number of small task to yield a fine-grained decomposition of the problem
- ❑ A good partition divides into small pieces both the *computation* associated with a problem and the *data* on which this computation operates
- ❑ *Domain decomposition* focuses on computation data
- ❑ *Functional decomposition* focuses on computation tasks
- ❑ Mixing domain/functional decomposition is possible





Domain and Functional Decomposition

□ Domain decomposition of two / three-dimensional grid



□ Functional decomposition of a climate model




Partitioning Checklist

- ❑ Does your partition define at least an order of magnitude more tasks than there are processors in your target computer? If not, may lose design flexibility.
- ❑ Does your partition avoid redundant computation and storage requirements? If not, may not be scalable.
- ❑ Are tasks of comparable size? If not, it may be hard to allocate each processor equal amounts of work.
- ❑ Does the number of tasks scale with problem size? If not may not be able to solve larger problems with more processors
- ❑ Have you identified several alternative partitions?


Communication

- ❑ Tasks generated by a partition must interact to allow the computation to proceed
 - Information flow: data and control
- ❑ Types of communication
 - *Local* vs. *Global*: locality of communication
 - *Structured* vs. *Unstructured*: communication patterns
 - *Static* vs. *Dynamic*: determined by runtime conditions
 - *Synchronous* vs. *Asynchronous*: coordination degree
- ❑ Granularity and frequency of communication
 - Size of data exchange
- ❑ Communication as control



Types of Communication

- ❑ Point-to-point
- ❑ Group-based
- ❑ Hierarchical
- ❑ Collective




Communication Design Checklist

- ❑ Is the distribution of communications equal?
 - Unbalanced communication may limit scalability
- ❑ What is the communication locality?
 - Wider communication locales are more expensive
- ❑ What is the degree of communication concurrency?
 - Communication operations may be parallelized
- ❑ Is computation associated with different tasks able to proceed concurrently? Can communication be overlapped with computation?
 - Try to reorder computation and communication to expose opportunities for parallelism

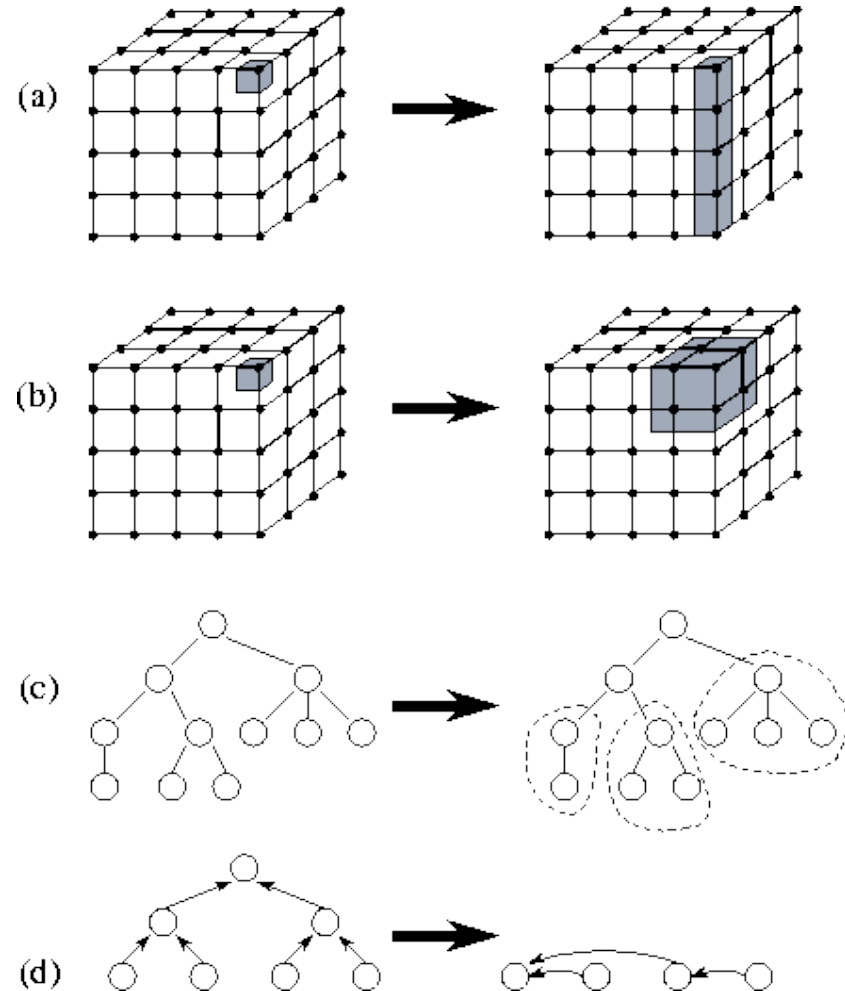
Agglomeration

- ❑ Move from parallel abstractions to real implementation
- ❑ Revisit partitioning and communication
 - View to efficient algorithm execution
- ❑ Is it useful to *agglomerate* (combine) tasks?
- ❑ Is it useful to *replicate* data and/or computation?
- ❑ Changes important algorithm and performance ratios
 - *Surface-to-volume*: reduction in communication at the expense of decreasing parallelism
 - *Communication/computation*: which cost dominates
- ❑ Replication may allow reduction in communication
- ❑ Maintain flexibility to allow overlap



Types of Agglomeration

- ❑ Element to column
- ❑ Element to block
 - Better surface to volume
- ❑ Task merging
- ❑ Task reduction
 - Reduces communication




Agglomeration Design Checklist

- ☐ Has increased locality reduced communication costs?
- ☐ Is replicated computation worth it?
- ☐ Does data replication compromise scalability?
- ☐ Is the computation still balanced?
- ☐ Is scalability in problem size still possible?
- ☐ Is there still sufficient concurrency?
- ☐ Is there room for more agglomeration?
- ☐ Fine-grained vs. coarse-grained?

Mapping

- ❑ Specify where each task is to execute
 - Less concern on shared-memory computers
- ❑ Attempt to minimize execution time
 - Place concurrent tasks on different processors to enhance physical concurrency
 - Place communicating tasks on same processor, or on processors close to each other, to increase locality
 - Strategies can conflict!
- ❑ Mapping problem is *NP-complete*
 - Use problem classifications and heuristics
- ❑ Static and dynamic load balancing



Mapping Algorithms

- ❑ Load balancing (partitioning) algorithms
- ❑ Data-based algorithms
 - Think of computational load with respect to amount of data being operated on
 - Assign data (i.e., work) in some known manner to balance
 - Take into account data interactions
- ❑ Task-based (task scheduling) algorithms
 - Used when functional decomposition yields many tasks with weak locality requirements
 - Use task assignment to keep processors busy computing
 - Consider centralized and decentralize schemes

Mapping Design Checklist

- ❑ Is static mapping too restrictive and non-responsive?
- ❑ Is dynamic mapping too costly in overhead?
- ❑ Does centralized scheduling lead to bottlenecks?
- ❑ Do dynamic load-balancing schemes require too much coordination to re-balance the load?
- ❑ What is the tradeoff of dynamic scheduling complexity versus performance improvement?
- ❑ Are there enough tasks to achieve high levels of concurrency? If not, processors may idle.

Types of Parallel Programs

- ❑ Flavors of parallelism
 - Data parallelism
 - All processors do same thing on different data
 - Task parallelism
 - Processors are assigned tasks that do different things
- ❑ Parallel execution models
 - Data parallel
 - Pipelining (Producer-Consumer)
 - Task graph
 - Work pool
 - Master-Worker

Data Parallel


- ❑ Data is decomposed (mapped) onto processors
- ❑ Processors performance similar (identical) tasks on data
- ❑ Tasks are applied concurrently
- ❑ Load balance is obtained through data partitioning
 - Equal amounts of work assigned
- ❑ Certainly may have interactions between processors
- ❑ Data parallelism scalability
 - Degree of parallelism tends to increase with problem size
 - Makes data parallel algorithms more efficient
- ❑ *Single Program Multiple Data (SPMD)*
 - Convenient way to implement data parallel computation

Matrix - Vector Multiplication

- ❑ $A \times b = y$
- ❑ Allocate tasks to rows of A


$$y[i] = \sum_j A[i,j] * b[j]$$

- ❑ Dependencies?
- ❑ Speedup?
- ❑ Computing each element of y can be done independently



Matrix-Vector Multiplication with Limited Tasks


- ❑ Suppose we only have 4 tasks
- ❑ Dependencies?
- ❑ Speedup?



Matrix Multiplication


□ $A \times B = C$

□ $A[i,:] \cdot B[:,j] = C[i,j]$



□ Row partitioning


○ N tasks



(a)

□ Block partitioning

○ $N*N/B$ tasks



(b)


□ Shading shows data sharing in B matrix

Granularity of Task and Data Decompositions



- ❑ Granularity can be with respect to tasks and data
- ❑ Task granularity
 - Equivalent to choosing the number of tasks
 - Fine-grained decomposition results in large # tasks
 - Large-grained decomposition has smaller # tasks
 - Translates to data granularity after # tasks chosen
 - consider matrix multiplication
- ❑ Data granularity
 - Think of in terms of amount of data needed in operation
 - Relative to data as a whole
 - Decomposition decisions based on input, output, input-output, or intermediate data

Mesh Allocation to Processors

- ❑ Mesh model of Lake Superior
- ❑ How to assign mesh elements to processors
- ❑ Distribute onto 8 processors
randomly



graph partitioning for
minimum edge cut



Pipeline Model


- ❑ Stream of data operated on by succession of tasks

 Task 1
  Task 2
  Task 3
  Task 4


- Tasks are assigned to processors

- ❑ Consider N data units


- ❑ Sequential



- ❑ Parallel (each task assigned to a processor)




4-way parallel




Pipeline Performance

- ❑ N data and T tasks
- ❑ Each task takes unit time t
- ❑ Sequential time = $N * T * t$
- ❑ Parallel pipeline time = $start + finish + (N - 2T) / T * t$
 $= O(N/T) \quad (\text{for } N \gg T)$
- ❑ Try to find a lot of data to pipeline
- ❑ Try to divide computation in a lot of pipeline tasks
 - More tasks to do (longer pipelines)
 - Shorter tasks to do
- ❑ Pipeline computation special form of *producer-consumer*
 - Producer tasks output data input by consumer tasks




Tasks Graphs

- ❑ Computations in any parallel algorithms can be viewed as a task dependency graph
- ❑ Task dependency graphs may be simple or non-trivial
 - Pipeline Task 1 → Task 2 → Task 3 → Task 4
 - Arbitrary (represents the algorithm dependencies)



Numbers are
time taken to
perform task


(a)



(b)


Task Graph Performance

- ❑ Determined by the *critical path*
 - Sequence of dependent tasks that takes the longest time



(a)

Min time = 27




(b)

Min time = 34

- *Critical path length* bounds parallel execution time


Task Assignment (Mapping) to Processors

- ❑ Given a set of tasks and number of processors
- ❑ How to assign tasks to processors?
- ❑ Should take dependencies into account
- ❑ Task mapping will determine execution time



(a)

Total time = ?





(b)

Total time = ?

Task Graphs in Action

- ❑ Uintah task graph scheduler
 - C-SAFE: Center for Simulation of Accidental Fires and Explosions, University of Utah
 - Large granularity tasks
- ❑ PLASMA
 - DAG-based parallel linear algebra
 - DAGuE: A generic distributed DAG engine for HPC

Task graph for PDE solver




DAG of QR for a 4×4 tiles matrix on a 2×2 grid of processors.

Bag o' Tasks Model and Worker Pool

- ❑ Set of tasks to be performed
- ❑ How do we schedule them?
 - Find independent tasks
 - Assign tasks to available processors

- ❑ Bag o' Tasks approach

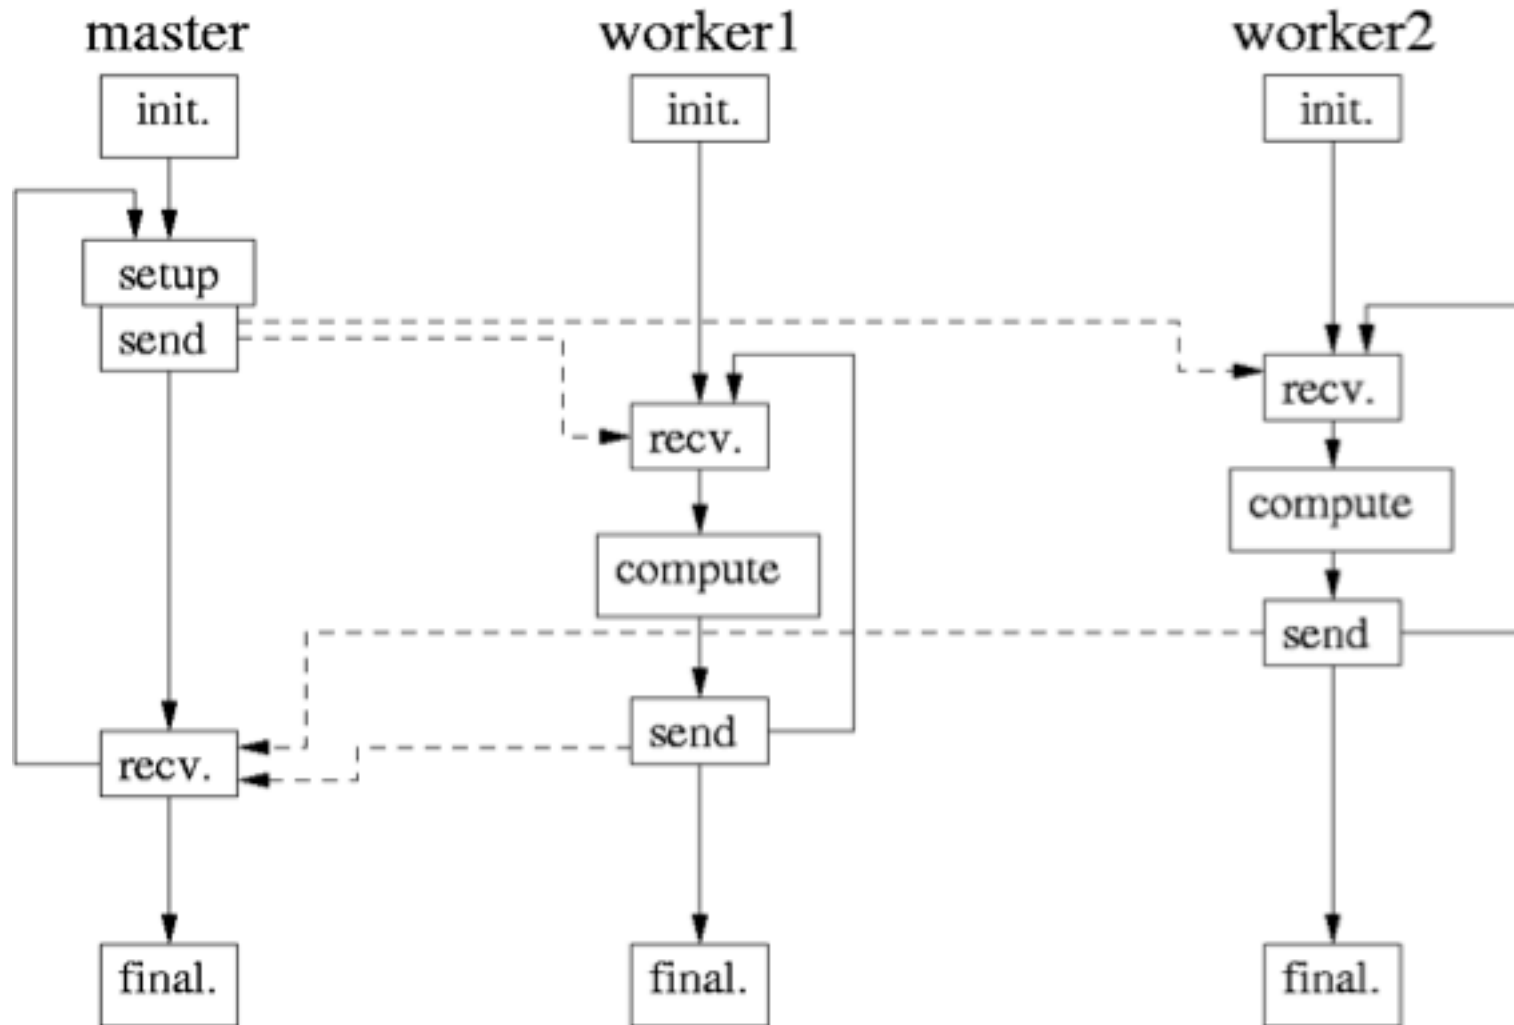
- Tasks are stored in a bag waiting to run
- If all dependencies are satisfied, it is moved to a ready to run queue
- Scheduler assigns a task to a free processor selected from a pool of (worker) processors



Master-Worker Parallelism


- ❑ One or more master processes generate work
- ❑ Masters allocate work to worker processes
- ❑ Workers idle if have nothing to do
- ❑ Workers are mostly stupid and must be told what to do
 - Execute independently
 - May need to synchronize, but must be told to do so
- ❑ Master may become the bottleneck if not careful
- ❑ What are the performance factors and expected performance behavior
 - Consider task granularity and asynchrony
 - How do they interact?

Master-Worker Execution Model (Li Li)



Li Li, "Model-based Automatics Performance Diagnosis of Parallel Computations," Ph.D. thesis, 2007.

M-W Execution Trace (Li Li)



Search-Based (Exploratory) Decomposition

- ❑ 15-puzzle problem
- ❑ 15 tiles numbered 1 through 15 placed in 4x4 grid
 - Blank tile located somewhere in grid
 - Initial configuration is out of order
 - Find shortest sequence of moves to put in order

1	2	3	4
5	6	7	8
9	10	7	11
13	14	15	12

(a)

1	2	3	4
5	6	7	8
9	10	11	11
13	14	15	12

(b)

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	12

(c)

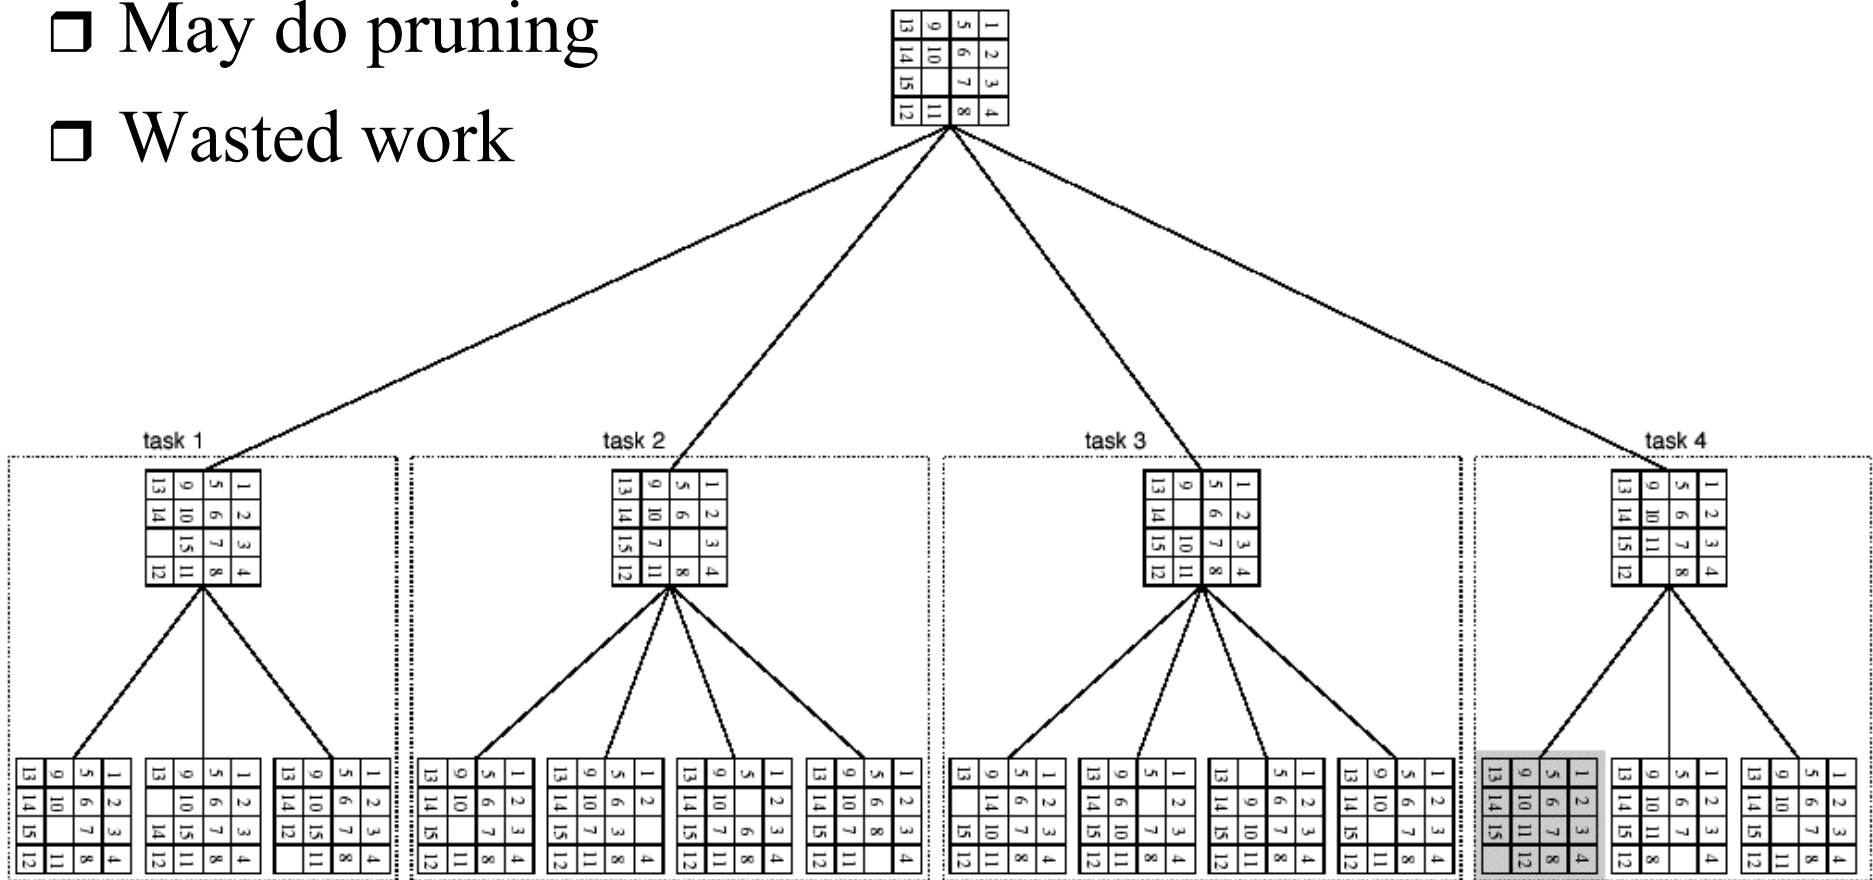
1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	

(d)

- ❑ Sequential search across space of solutions
 - May involve some heuristics

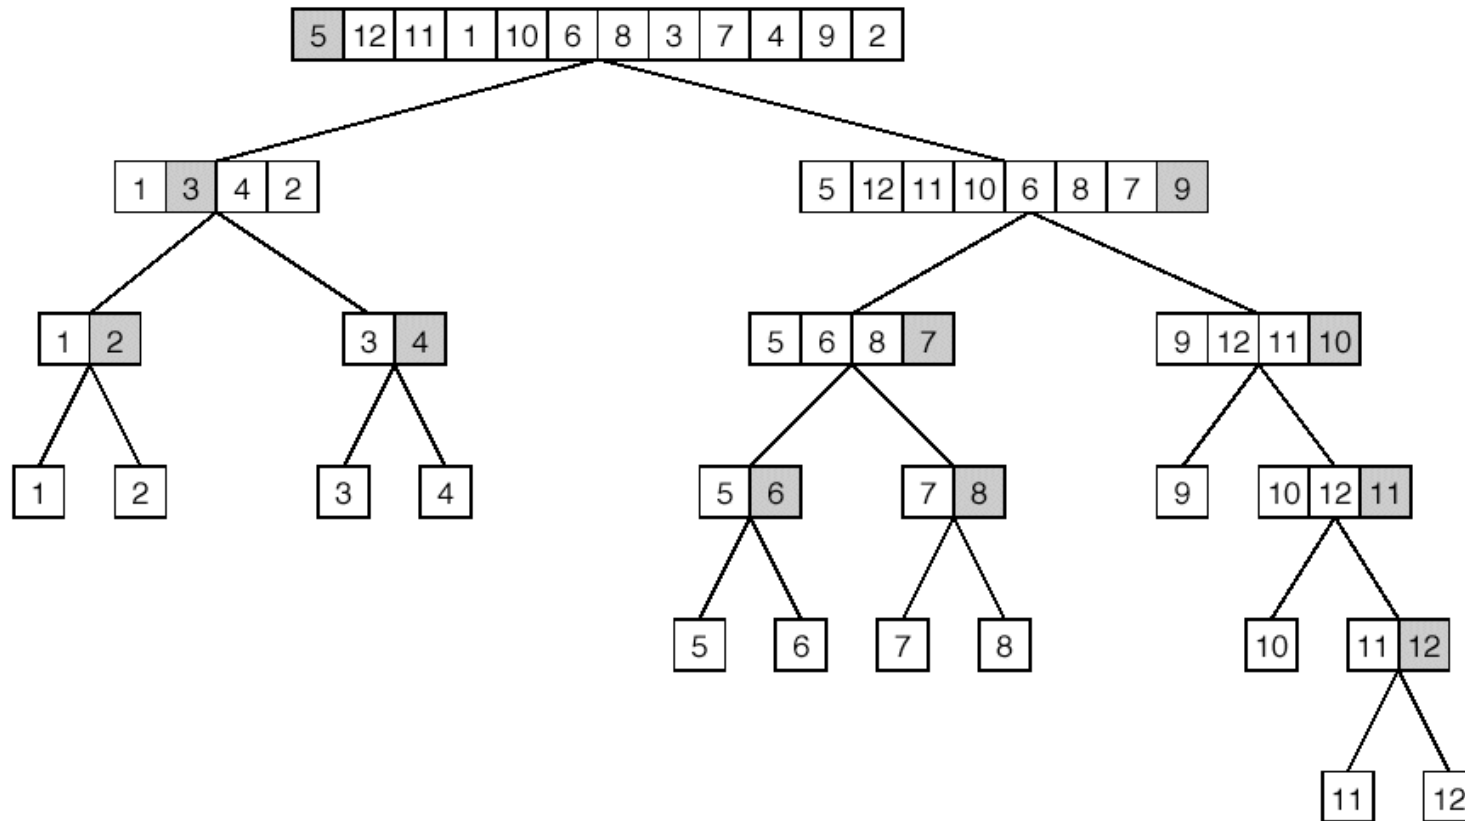
Parallelizing the 15-Puzzle Problem

- ❑ Enumerate move choices at each stage
- ❑ Assign to processors
- ❑ May do pruning
- ❑ Wasted work



Divide-and-Conquer Parallelism

- ❑ Break problem up in orderly manner into smaller, more manageable chunks and solve
- ❑ Quicksort example



Next Class

- ❑ Programming models
- ❑ Standard parallel programming techniques
 - shared memory (Pthreads)
 - message passing (MPI)
 - data parallelism (Fortran 90, CUDA)
 - shared memory + data parallelism (OpenMP)
 - object-oriented parallelism (?)