

CIS 631

Parallel Processing

Lecture 9: Shared Memory Parallel Programming

Allen D. Malony
malony@cs.uoregon.edu

Department of Computer and Information Science
University of Oregon



Acknowledgements

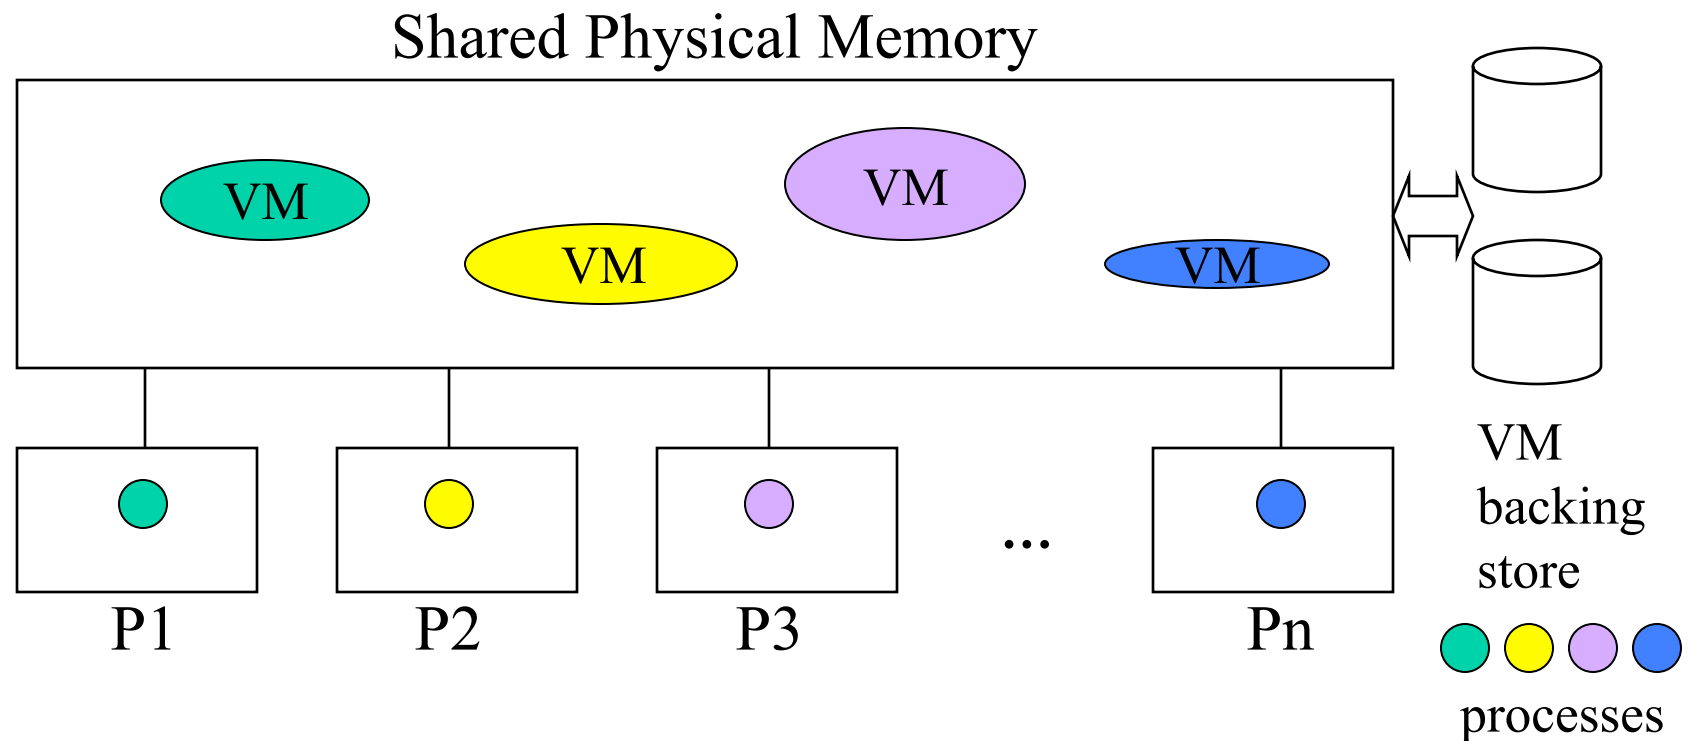
- ❑ Portions of the lectures slides were adopted from:
 - I. Foster, “Designing and Building Parallel Programs,” 1995
 - Vijay Pai, COMP 422, “Parallel Programming,” Rice University, 2002
 - John Mellor-Crummey, COMP 422, “Parallel Programming,” Rice University, 2010
 - A. Grama, A. Gupta, G. Karypis, and V. Kumar, “Introduction to Parallel Computing,” 2003

Outline

- ❑ Shared memory parallelism and programming
- ❑ Process-based vs. thread-based programming
- ❑ Pthreads programming

Shared Memory Multiprocessors

- ❑ All processes share the same **physical memory** space
- ❑ Processes generally do not share **virtual memory** space



- ❑ How to program for parallel execution?

Shared Address Space Programming Taxonomy


- ❑ Process model
 - Each process's data is private, but can create shared space
 - Example: Linux shared memory segments
- ❑ Lightweight process (LWP) / thread model
 - Virtual memory is global, shared between LWP / threads
 - Example: Pthreads, Cilk (lazy, lightweight threads)
- ❑ Language-based model
 - Threads have shared /private data, built on runtime system
 - Example: OpenMP, Java
- ❑ Global address space (for distributed memory)
 - Language or library supported

Shared Memory Programming – IPC (Really!?)

- ❑ Multiple processes can run in parallel concurrently
- ❑ Multiple processes can be used in a parallel execution
- ❑ How?
 - Inter-process communication (IPC)
 - Sockets
 - Message passing (MPI)
 - IPC
 - Memory mapping implementation
 - Through file system (yucky, but possible)
- ❑ Effectively programming as a multi-computer system
- ❑ BUT, get advantage of shared file system, devices, ... !!!

Shared Memory Programming – Process

- ❑ Shared memory parallel programming possible only if processes can access same virtual memory space (some)
- ❑ One way is to use OS VM address sharing mechanisms
 - Linux (*shm*) memory mapping and synchronization
 - Incur overheads of using OS support
 - Less than going through IPC and network interfaces though!

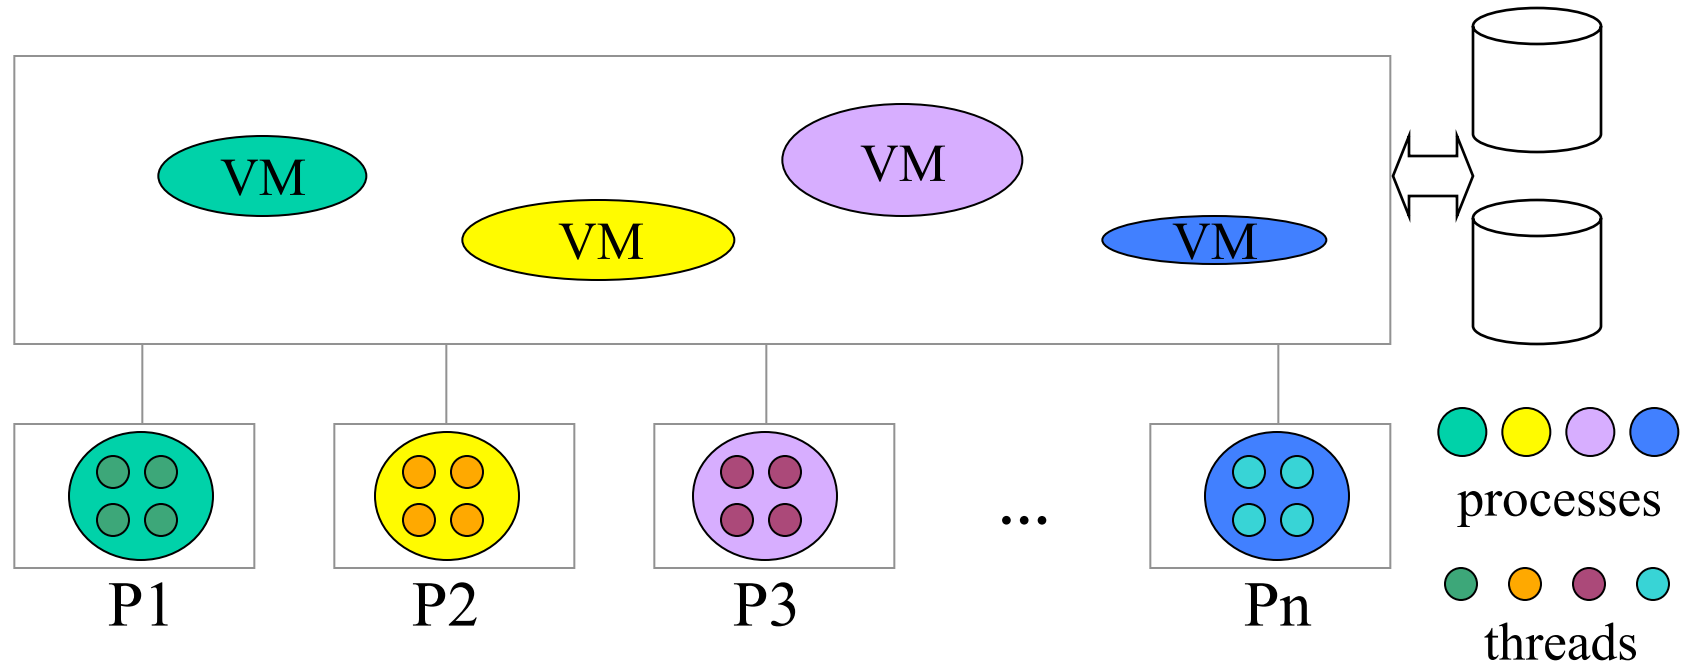


Shared Memory Programming – Multithreading

- ❑ Process-level parallel programming via virtual memory space sharing suffers from overheads
 - Processes are heavy-weight and cost more to switch
 - Interaction with OS mechanisms not very clean
- ❑ Desire way to:
 - Create lighter-weight units of computation
 - Program light-weight computation interaction
- ❑ Must be able to address shared data directly
 - Object x is the same object x no matter which computation unit references it
 - The virtual memory space (or some portion) is shared

Multi-threading

- ❑ Creation of “threads of execution” that share process VM



- ❑ Threads inherit process' s VM address space
- ❑ All global data AND instructions are shared
- ❑ Thread-private data possible
 - Through memory allocation and local variables

Multi-threading: Advantages and Disadvantages

❑ Advantages

- Light-weight computation
- Fast thread switching
- Shared data through memory addressing
- Synchronization support through shared memory
- Low overhead
- User has control over threads of execution

❑ Disadvantages

- User has control over threads of execution and memory
- Prone to concurrent programming (synchronization) errors
- Performance of memory system may be harder to optimize

Shared Memory Parallel Programming

- ❑ What does the user have to do?
 - Decide how to decompose computation into parallel parts
 - Create (and destroy) threads to support decomposition
 - Add synchronization to satisfy dependences
- ❑ In some sense, similar to message passing programming
 - Think distributed memory, then program share memory
- ❑ Execution on a shared memory multiprocessor
 - OS can run threads on available processors
 - Threads run concurrently and can run in parallel
 - Cache coherency kicks in to support consistency model
- ❑ Synchronization programming is the hard part!

What's a thread?

- ❑ Control perspective
 - A thread is a single stream of control in a program
 - A thread can execute an instruction stream
- ❑ State perspective
 - A thread embodies a state of execution
 - It contains an instruction pointer, a stack, registers, ...
- ❑ “Thread of execution”
 - Control plus state
- ❑ Threads inherit from parent process
 - Virtual address space, file descriptors, ...
 - Threads are peers, only have parent relationship to process

General Thread Structure

- ❑ Typically, a thread is the execution of a piece of code
 - Represents a portion of the program (light-weight)
 - Given a well-defined entry point (e.g., routine)
 - Inherits process' s symbol table (e.g., to get to libraries)
- ❑ Task-based parallelism
 - Parallel parts form separate procedures or functions
 - Kick off thread with specific routine
 - Kick off thread with a common routine and then specialize
- ❑ Data-based parallelism
 - Invoke threads to work on different data
 - Automatic with shared memory support

Why threads?

- ❑ Portable, widely-available programming model
 - Requires OS support, but practically all OSes do
- ❑ Easier to program (some say)
- ❑ Efficiencies in scheduling
 - Versus processes (Why?)
- ❑ Efficiencies in latency hiding
 - I/O, communication (How?)
- ❑ More dynamic concurrency
- ❑ Requires code to be *thread-safe*

POSIX Pthreads

- ❑ POSIX standard multi-threading interface
 - For general multi-threaded concurrent programming
 - Largely independent across implementations
 - Broadly supported on different platforms
 - Common target for library and language implementation
- ❑ Provides primitives for
 - Thread creation and management
 - Synchronization
- ❑ We will only look at a subset of Pthreads
- ❑ See webpage for links to Pthreads programming

Thread Creation


```
#include <pthread.h>
int pthread_create(
    pthread_t *thread_id,
    const pthread_attr_t *attribute,
    void *(*thread_function) (void *),
    void *arg);
```

- ❑ thread_id
 - thread's unique identifier
- ❑ attribute
 - contain details on scheduling policy, priority, stack, ...
- ❑ thread_function
 - function to be run in parallel (entry point)
- ❑ arg
 - arguments for function func

Example of Thread Creation

```
void *func(void *arg) {  
    int *I=arg;  
    ...  
}
```

```
void main()  
{  
    int X;  
    pthread_t id;  
    ...  
    pthread_create(&id, NULL, func, &X);  
    ...  
}
```



Pthread Termination

```
void pthread_exit(void *status)
```

- ❑ Terminates the currently running thread
- ❑ Implicitly called when function called in `pthread_create` returns

Thread Joining

```
int pthread_join(  
    pthread_t thread_id,  
    void **status);
```

- ❑ Waits for thread `thread_id` to terminate
 - Either by returning
 - Or by calling `pthread_exit()`
- ❑ Status receives the return value or the value given as argument to `pthread_exit()`

Thread Joining Example

```
void *func(void *) {  
    ...  
}
```

```
pthread_t    id;  
int X;
```

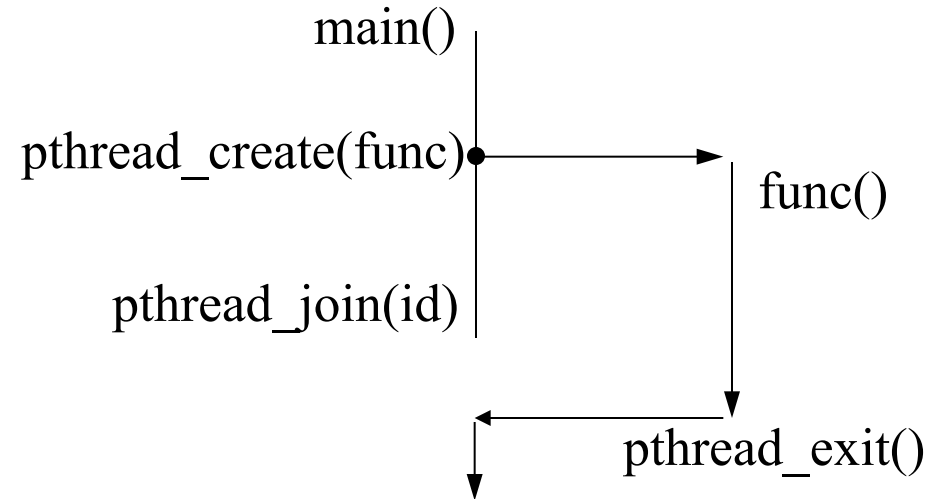
...

```
pthread_create(&id, NULL, func, &X);
```

...

```
pthread_join(id, NULL);
```

...




General Program Structure

- ❑ Encapsulate parallel parts in functions
- ❑ Use function arguments to parameterize thread behavior
- ❑ Call `pthread_create()` with the function
- ❑ Call `pthread_join()` for each thread created
- ❑ Need to take care to make program “thread safe”

Matrix Multiply

$$\square A \times B = C$$


$$\square A[i,:] \cdot B[:,j] = C[i,j]$$



```
for( i=0; i<n; i++ )  
    for( j=0; j<n; j++ ) {  
        c[i][j] = 0.0;  
        for( k=0; k<n; k++ )  
            c[i][j] += a[i][k]*b[k][j];  
    }
```

Parallel Matrix Multiply

- ❑ All i- or j-iterations can be run in parallel
- ❑ If we have p processors, n/p rows to each processor
- ❑ Corresponds to partitioning i-loop



Matrix Multiply: Parallel Part

```
void mmult(void* s)
{
    int slice = (int) s;
    int from = (slice*n)/p;
    int to = ((slice+1)*n)/p;
    for(i=from; i<to; i++)
        for(j=0; j<n; j++) {
            c[i][j] = 0.0;
            for(k=0; k<n; k++)
                c[i][j] += a[i][k]*b[k][j];
        }
}
```


Matrix Multiply: Main

```
int main()
{
    pthread_t thrd[p];

    for( i=0; i<p; i++ )
        pthread_create(&thrd[i], NULL,
                        mmult, (void*) i);

    for( i=0; i<p; i++ )
        pthread_join(thrd[i], NULL);
}
```

Pthread Process Management

- ❑ `pthread_create()`
 - Creates a parallel thread executing a given function
 - Passes function arguments
 - Returns thread identifier
- ❑ `pthread_exit()`
 - terminates thread.
- ❑ `pthread_join()`
 - waits for particular thread to terminate

Pthreads Synchronization

- ❑ Create/exit/join
 - Provide some coarse form of synchronization
 - “Fork-join” parallelism
 - Requires thread creation/destruction
- ❑ Need for finer-grain synchronization
 - Mutex locks
 - Condition variables

Pthreads “Hello World”

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>

#define NUM_THREADS      5

void *TaskCode(void *argument)
{
    int tid;

    tid = *((int *) argument);
    printf("Hello World! It's me, thread %d!\n", tid);

    /* optionally: insert more useful stuff here */

    return NULL;
}

int main(void)
{
    pthread_t threads[NUM_THREADS];
    int thread_args[NUM_THREADS];
    int rc, i;

    /* create all threads */
    for (i=0; i<NUM_THREADS; ++i) {
        thread_args[i] = i;
        printf("In main: creating thread %d\n", i);
        rc = pthread_create(&threads[i], NULL, TaskCode, (void *) &thread_args[i]);
        assert(0 == rc);
    }

    /* wait for all threads to complete */
    for (i=0; i<NUM_THREADS; ++i) {
        rc = pthread_join(threads[i], NULL);
        assert(0 == rc);
    }

    exit(EXIT_SUCCESS);
}
```

Mutex Locks – Create, Destroy

- ❑ Creates a new mutex lock `mutex`

```
pthread_mutex_init(  
    pthread_mutex_t * mutex,  
    const pthread_mutex_attr *attr);
```

- ❑ Destroys the mutex specified by `mutex`.

```
pthread_mutex_destroy(  
    pthread_mutex_t *mutex);
```

Mutex Locks – Lock

- ❑ Tries to acquire the lock specified by mutex

```
pthread_mutex_lock(  
    pthread_mutex_t *mutex)
```

- ❑ If mutex is already locked
 - Calling thread blocks until mutex is unlocked
- ❑ If mutex is not locked
 - Mutex is locked and calling thread returned
- ❑ Mutually exclusive

Mutex Locks – Unlock

- ❑ Unlock mutex lock

```
pthread_mutex_unlock(  
    pthread_mutex_t *mutex) ;
```

- ❑ If calling thread has mutex currently locked
 - Mutex will be unlocked
 - If other threads are blocked waiting on this mutex
 - one will unblock and acquire mutex
 - which one is determined by the scheduler

Use of Mutex Locks

- ❑ Pthreads provides only exclusive locks
- ❑ Other systems allow other types of locks
 - Shared-read, exclusive-write locks
- ❑ Critical sections
 - Code sections only to be executed by one thread
- ❑ Applications
 - Update of shared variables
 - Queues
 - Stacks
- ❑ Problem is that mutex locks can be inefficient
 - Excessive polling

Condition Variables

- ❑ Condition variables are objects for thread synchronization
- ❑ Allows a thread to block itself until specified data reaches a predefined state
- ❑ A condition variable is associated with a predicate
 - When predicate become true, the condition variable is used to signal thread(s) waiting on the condition
- ❑ A condition variable always has an associated mutex
- ❑ A thread locks the mutex and tests the predicate
 - If not true, threads waits on condition variable
- ❑ A blocked thread is released on a signal
 - It then acquires the mutex before resuming

Condition variables – Init, Destroy

- ❑ Creates a new condition variable `cond`

```
pthread_cond_init(  
    pthread_cond_t *cond,  
    pthread_cond_attr *attr)
```

- ❑ Destroys the condition variable `cond`

```
pthread_cond_destroy(  
    pthread_cond_t *cond)
```

Condition Variables – Wait

- ❑ Wait on cond

```
pthread_cond_wait(  
    pthread_cond_t *cond,  
    pthread_mutex_t *mutex)
```

- ❑ Blocks the calling thread
- ❑ Unlocks the mutex on success

Condition Variables – Signal

- ❑ Unblocks one thread waiting on cond

```
pthread_cond_signal(  
    pthread_cond_t *cond)
```

- ❑ Which one is determined by scheduler
- ❑ If no thread waiting, then signal is a no-op

Condition Variables – Broadcast

- Unblocks all threads waiting on cond

```
pthread_cond_broadcast(  
    pthread_cond_t *cond)
```

- If no thread waiting, then broadcast is a no-op

Use of Condition Variables

- ❑ To implement signal-wait synchronization
- ❑ Be careful
 - A signal is “forgotten” if there is no corresponding wait that has already happened.

PIPE Example

- ❑ Send picture images into a pipeline
- ❑ Transformations performed at each pipeline stage
- ❑ Separate flag for each picture image

```
P1: for( i=0; i<num_pics, read(in_pic); i++ ) {  
    int_pic_1[i] = trans1( in_pic );  
    signal( event_1_2[i] );  
}
```

```
P2: for( i=0; i<num_pics; i++ ) {  
    wait( event_1_2[i] );  
    int_pic_2[i] = trans2( int_pic_1[i] );  
    signal( event_2_3[i] );  
}
```

PIPE Using Pthreads

- ❑ Replace by Pthreads condition variable wait/signal
 - Will not work
 - Signals before a wait are forgotten
 - Need to remember a signal
- ❑ Create a signal semaphore

```
semaphore_signal(i) {  
    pthread_mutex_lock(&mutex_rem[i]);  
    arrived [i]= 1;  
    pthread_cond_signal(&cond[i]);  
    pthread_mutex_unlock(&mutex_rem[i]);  
}
```


PIPE Using Pthreads

```
semaphore_wait(i) {  
    pthreads_mutex_lock(&mutex_rem[i]);  
    if( arrived[i] == 0 ) {  
        pthreads_cond_wait(&cond[i],  
                           mutex_rem  
                           [i]);  
    }  
    arrived[i] = 0;  
    pthreads_mutex_unlock(&mutex_rem[i]);  
}
```

PIPE with Pthreads

P1:

```
for( i=0; i<num_pics, read(in_pic); i++ ) {  
    int_pic_1[i] = trans1( in_pic );  
    semaphore_signal( event_1_2[i] );  
}
```

P2:

```
for( i=0; i<num_pics; i++ ) {  
    semaphore_wait( event_1_2[i] );  
    int_pic_2[i] = trans2( int_pic_1[i] );  
    semaphore_signal( event_2_3[i] );  
}
```

Note on Semaphores

- ❑ Many shared memory programming systems (other than Pthreads) have semaphores as basic primitive
- ❑ If they do, you should use it, not construct it yourself
- ❑ Implementation may be more efficient than what you can do yourself

Reality Bites ...

- ❑ Thread create/exit/join is not so cheap
- ❑ More efficient if could have a parallel program where
 - Create/exit/join would happen rarely (once!)
 - Cheaper synchronization were used
- ❑ We need something that makes all threads wait
 - Barrier synchronization

Barrier Synchronization

- ❑ A wait at a barrier causes a thread to wait until all threads have performed a wait at the barrier
- ❑ At that point, they all proceed

Implementing Barriers in Pthreads

- ❑ Count the number of arrivals at the barrier
- ❑ Wait if this is not the last arrival
- ❑ Make everyone unblock if this is the last arrival
- ❑ Since the arrival count is a shared variable, enclose the whole operation in a mutex lock-unlock

Implementing Barriers in Pthreads

```
void barrier()
{
    pthread_mutex_lock(&mutex_arr);
    arrived++;
    if (arrived < N) {
        pthread_cond_wait(
            &cond, &mutex_arr);
    }
    else {
        pthread_cond_broadcast(&cond);
        arrived = 0; /* next barrier */
    }
    pthread_mutex_unlock(&mutex_arr);
}
```

Note on Barriers

- ❑ Many shared memory programming systems (other than Pthreads) have barriers as basic primitive
- ❑ If they do, you should use it, not construct it yourself
- ❑ Implementation may be more efficient than what you can do yourself

Busy Waiting

- ❑ Not an explicit part of the API
- ❑ Available in shared memory programming environment

```
initially:  flag = 0;
```

```
P1:    produce data;  
       flag = 1;
```

```
P2:    while( !flag ) ;  
       consume data;
```

Use of Busy Waiting

- ❑ On the surface, simple and efficient
- ❑ In general, not a recommended practice
- ❑ Often leads to messy and unreadable code
 - Blurs data/synchronization distinction
- ❑ On some architectures, may be inefficient
- ❑ May not even work as intended
 - Depending on consistency model

Private Data in Pthreads

- ❑ To make a variable private in Pthreads, you need to make an array out of it
- ❑ Index the array by thread identifier, which you can get by the `pthread_self()` call
- ❑ Not very elegant or efficient

Other Primitives in Pthreads

- ❑ Set the attributes of a thread
- ❑ Set the attributes of a mutex lock
- ❑ Set scheduling parameters

Next Class

- ❑ Shared memory parallel programming
- ❑ OpenMP