

# **CIS 631**

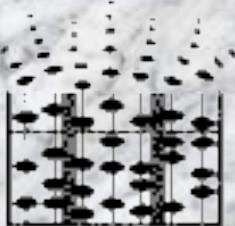
## *Parallel Processing*

### *Lecture 10: Shared Memory Parallel Programming and OpenMP*

**Allen D. Malony**

[malony@cs.uoregon.edu](mailto:malony@cs.uoregon.edu)

Department of Computer and Information Science  
University of Oregon



## *Acknowledgements*

- Portions of the lectures slides were adopted from:
  - John Mellor-Crummey, COMP 422, “Parallel Programming,” Rice University, 2002
  - Vijay Pai, COMP 422, “Parallel Programming,” Rice University, 2002
  - A. Grama, A. Gupta, G. Karypis, and V. Kumar, “Introduction to Parallel Computing,” 2003
  - L. Kale, CS 320, “Parallel Computing,” University of Illinois, Urbana-Champaign, 2003

# *Outline*

- OpenMP

# *References*

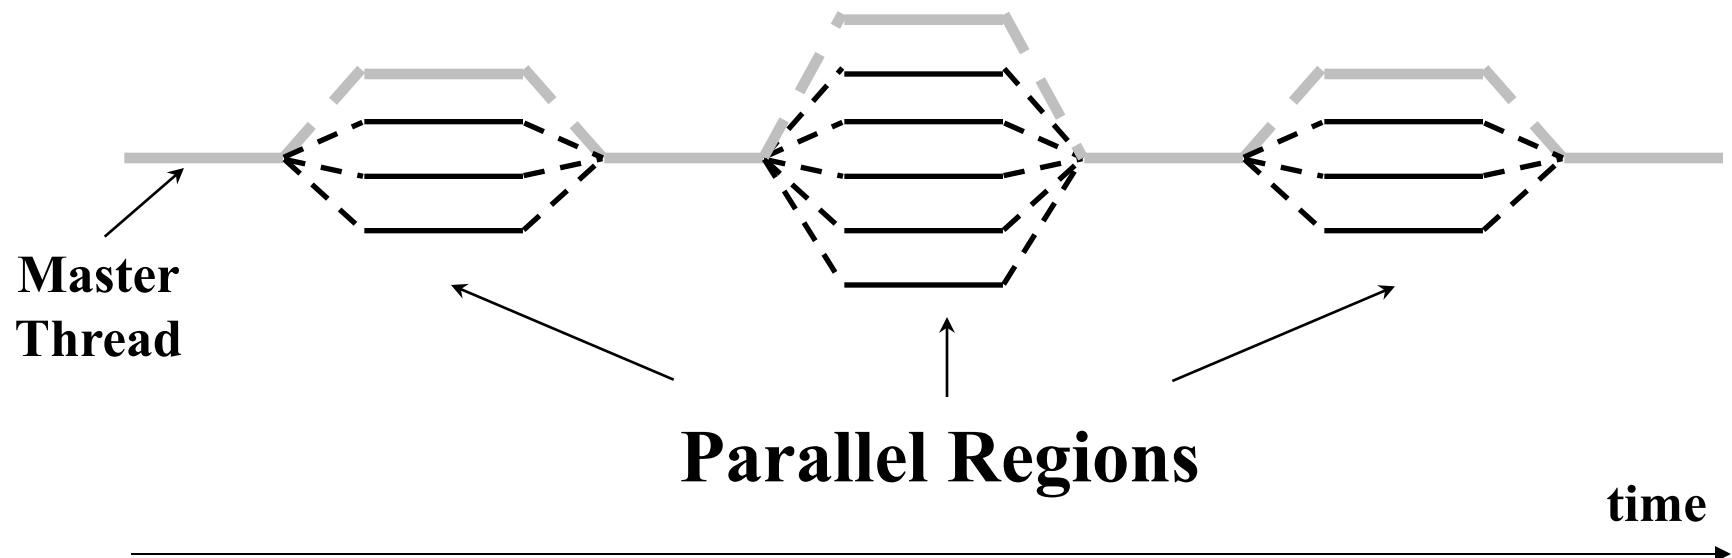
- <http://www.openmp.org>
  - History of OpenMP
  - Current status and specifications
  - Tutorials
  - Everything you need to know!!!

# *OpenMP – What Is It?*

- An API for writing multithreaded applications
- A set of compiler directives and library routines for parallel application programmers
- Makes it easy to create shared memory parallel programs
  - Using Fortran, C, and C++
- Standardizes last 20+ years of SMP practice
- Supported by many hardware and software vendors
  - Intel, Microsoft, Cray, PGI, ...
- Specs (<http://openmp.org/wp/openmp-specifications/>)
  - OpenMP 1.0 (1998) – OpenMP 3.1 (2011, current)
  - OpenMP 4.0 public review release candidate (Nov. 2012)

# *OpenMP – Programming Model*

- Fork-Join Parallelism
  - Master thread spawns slave threads as needed
  - Parallelism is added incrementally
    - sequential program evolves into a parallel program



# *OpenMP – Thread Interaction*

- This is shared memory parallel programming
  - Threads communicate by sharing variables
  - Unintended sharing of data can lead to *race conditions*
- Race condition
  - Program's outcome depends on thread ordering
  - Typically not desired
  - *Non-deterministic* execution
- To control race conditions
  - Use synchronization to protect data conflicts
  - Synchronization is expensive
  - Change data storage to minimize need for synchronization

# *OpenMP – General Rules*

- Most OpenMP constructs are *compiler directives*
- Directives inform the compiler
  - Provide compiler with knowledge
  - Usage assumptions
- Directives are ignored by non-OpenMP compilers!
  - Essentially act as comment for backward compatibility
- For C and C++, the syntax is:  
`#pragma omp construct [clause [clause] ...]`
- For Fortran, the syntax is:  
`C$OMP construct [clause [clause] ...]`  
`! $OMP construct [clause [clause] ...]`  
`*$OMP construct [clause [clause] ...]`

# *Structured Blocks*

- Most OpenMP constructs apply to structured blocks
- Structured block
  - A block of code with one point of entry at the top and one point of exit at the bottom
  - Only other branches allowed out of the block are STOP statements in Fortran and exit() in C/C++
- Loops are a common example of structured blocks
  - Excellent source of parallelism

## ***PARALLEL Directive***

- Specifies the following should be executed in parallel:
  - A program section (structured block)
  - If applied to a loop, what happens is:
    - iterations are executed in parallel
    - do loop (Fortran)
    - for loop (C/C++)
- PARALLEL DO is called a “worksharing” directive
  - Causes work to be shared across threads
  - More on this later

# *PARALLEL DO: Syntax*

- Fortran

```
!$omp parallel do [clause [,] [clause ...]]  
  do index = first, last [, stride]  
    body of the loop  
  enddo  
[ !$omp end parallel do]
```

- C/C++

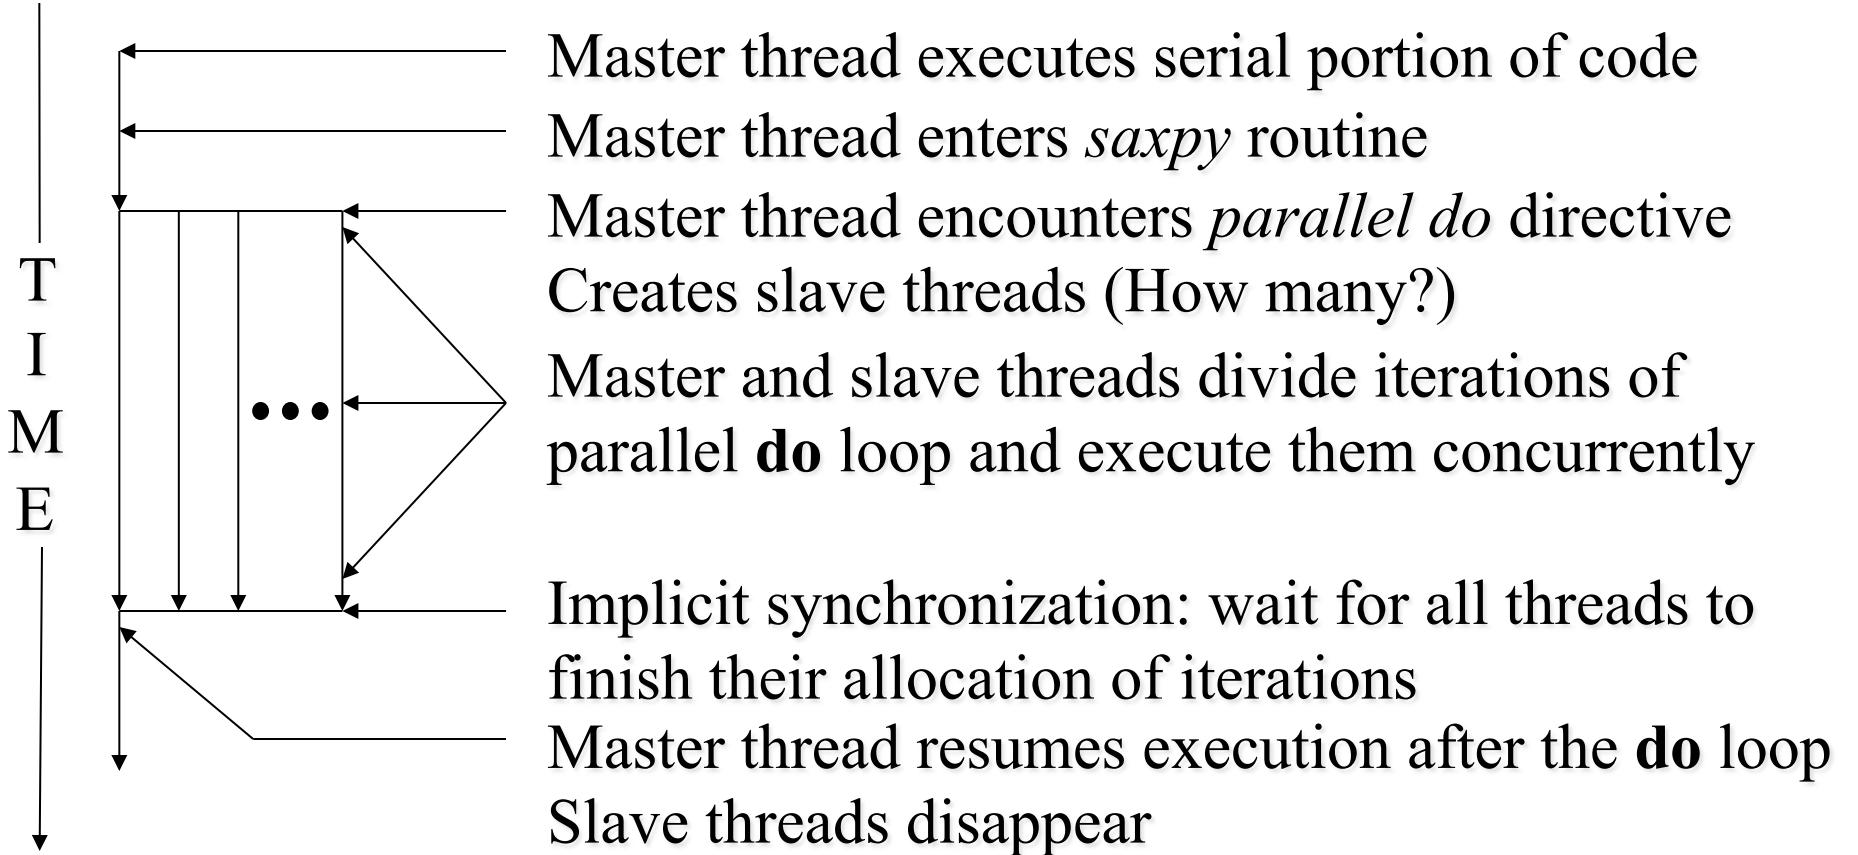
```
#pragma omp parallel for [clause [clause ...]]  
  for (index = first; text_expr;  
       increment_expr) {  
    body of the loop  
  }
```

## *Example: PARALLEL DO*

- Single precision  $a^*x + y$  (*saxpy*)

```
subroutine saxpy (z, a, x, y, n)
integer i, n
real z(n), a, x(n), y(n)
!$omp parallel do
do i = 1, n
    z(i) = a * x(i) + y(i)
enddo
return
end
```

# *Execution Model of PARALLEL DO*

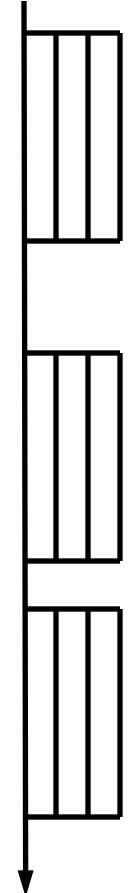


- Abstract execution model

# *Loop-level Parallelization Paradigm*

- Execute each loop in parallel
  - Where possible
- Easy to parallelize code
- Similar to automatic parallelization
- Incremental parallelization
  - One loop at a time
  - Does not break code (Really?)
- Fine-grain overhead
  - Frequent synchronization
- Performance determined by sequential part (Why?)

```
C$OMP PARALLEL DO  
do i=1,n  
.....  
enddo  
alpha = xnorm/sum  
C$OMP PARALLEL DO  
do i=1,n  
.....  
enddo  
C$OMP PARALLEL DO  
do i=1,n  
.....  
enddo
```



## *Example: PARALLEL DO – Bad saxpy*

- Single precision  $a^*x + y$  (*saxpy*)

```
subroutine saxpy (z, a, x, y, n)
integer i, n
real z(n), a, x(n), y(n)
!$omp parallel do
do i = 1, n
    y(i) = a * x(i+1) + y(i+1)
enddo
return
end
```

What happens here?

# *How Many Threads?*

- Use environment variable
  - `setenv OMP_NUM_THREADS 8` (Unix machines)
- Use `omp_set_num_threads()` function

```
subroutine saxpy (z, a, x, y, n)
integer i, n
real z(n), a, x(n), y(n)
XS call omp_set_num_threads(4)
!$omp parallel do
    do i = 1, n
        z(i) = a * x(i) + y(i)
    enddo
    return
end
```

Not a directive,  
but a call to the  
OpenMP library

## *Assigning Iterations to Threads*

- A parallel loop in OpenMP is a worksharing directive
- The manner in which iterations of a parallel loop are assigned to threads is called the loop's *schedule*
- Default schedule assigns iterations to threads as evenly as possible (good enough for `saxpy`)
- Alternative user-specified schedules possible
- More on scheduling later

# *Communication Between Threads*

- Six clause types allow the programmer to specify how data (i.e., each variable) is shared between threads executing a parallel do:
  - `private(list of variable/array names)`
  - `shared(list)`
  - `default(private | shared | none)`
  - `reduction(intrinsic operator : list)`
  - `firstprivate(list)`
  - `lastprivate(list)`
- Data scope clauses

# *Default Data Sharing Rules*

- Most variables are shared by default
  - This means there is only one copy shared by all threads
  - Global variables are *shared* among threads
    - Fortran: COMMON blocks, SAVE variables, MODULE variables
    - C/C++: File scope variables, static
- Stack variables in sub-programs called from parallel regions are *private*
  - Threads have their own private stack
- Automatic variables within a statement block are *private*

## *Private Clause*

- Each thread has copy of all variables declared private
- Private variables (implicit or explicit) are uninitialized when a thread starts
  - This is the thread's responsibility
- The value of a private variable is unavailable to the master thread after a parallel loop terminates
- Is there a way to pass private data back?

# *Firstprivate and Lastprivate*

- `firstprivate (list)`
  - Initializes each thread's copy of a private variable to the value of the master thread's copy, for all variables in list
- `lastprivate (list)`
  - Writes back to the master's copy the value contained in the private copy belonging to the thread that executed the sequentially *last* iteration of the loop, for all variables in list

# *Communication Between Threads*

- Unless one of the data scope clauses is present, most data/variables are shared by default

```
subroutine saxpy (z, a, x, y, n)
integer i, n
real z(n), a, x(n), y(n)
!$omp parallel do
do i = 1, n
    z(i) = a * x(i) + y(i)
enddo
return
end
```

- Is there a problem here?
- OpenMP will implicitly take care of the index variable

## *Example: Data Scope Clauses*

```
double x, y;
int i, j, m, n, maxiter;
int depth[300][200];
extern int mandel_val();
n = 300;
m = 200;
maxiter = 200;
#pragma omp parallel for private(j, x, y)
for (i = 1; i <= m; i++)
    for (j = 1; j <= n; j++) {
        x = i/ (double) m;
        y = j/ (double) n;
        depth[j][i] =
            mandel_val(x, y, maxiter);
    }
```

This is the parallel loop

## *Example: Private Clause*

- What is wrong with this example?

```
program wrong
IS = 0
C$OMP PARALLEL DO PRIVATE (IS)
DO J=1,100
...
... = IS
...
100 CONTINUE
print *, IS
```

## *Corrected Example*

```
program wrong right
IS = 0
C$OMP PARALLEL DO FIRSTPRIVATE(IS)
C$OMP+ LASTPRIVATE(IS)
DO J=1,100
...
... = IS
...
100 CONTINUE
print *, IS
```

## *Example: firstprivate, lastprivate*

```
common /mycom/ x, c, y, z
real x(n, n), c(n, n), y(n), z(n)
x(1, 1) = ...
x(2, 1) = ...
!$omp parallel do firstprivate(x) lastprivate(i, x)
do i = 1, n-1
    x(1, 2) = c(i, 1) * x(1, 1)
    x(2, 2) = c(i, 2) * x(2, 1) ** 2
    y(i) = x(2, 2) + x(1, 2)
    z(i) = x(2, 2) - x(1, 2)
enddo
y(i+1) = x(1, 2) + x(2, 2)
```

- What if you did not have lastprivate?

## *Example: reduction*

```
subroutine sum (values, n, s)
integer n, s
real values(n)
s = 0
!$omp parallel do reduction(+ : s)
    do i = 1, n
        s = s + values(i)
    enddo
    return
end
```

- If *s* were shared, you'd need to protect it with locks
- If it were private, how do you get the global sum?
- Reduction clause simplifies your code

## *More on Reductions*

- Operator must be commutative and associative
- Caution with floating point numbers and roundoff errors
  - $A + (B+C)$  may not be the same as  $(A+B) +C$

## *Default clause*

- Default storage attribute is DEFAULT (SHARED)
  - No need to specify
- DEFAULT (PRIVATE)
  - Each variable in static extent of the parallel region is made private as if specified in a private clause
  - Mostly saves typing
- DEFAULT (NONE)
  - No default for variables in static extent
  - Must list storage attribute for each variable in static extent
- C/C++ do not support DEFAULT (PRIVATE)

## *Example: DEFAULT Clause*

```
    itotal = 1000
C$OMP PARALLEL DO PRIVATE(np, each)
do i = 1, 100
    np = omp_get_num_threads()
    each = itotal/np
    ...
enddo
```

---

```
    itotal = 1000
C$OMP PARALLEL DO DEFAULT(PRIVATE) SHARED(itotal)
do i = 1, 100
    np = omp_get_num_threads()
    each = itotal/np
    ...
enddo
```

# *Assigning Iterations to Threads*

- Motivation
  - Balance the work per thread
  - When the work per iteration is inherently unbalanced

```
!$omp parallel do private(xkind)
    do i = 1, n
        xkind = f(i)
        if (xkind .lt. 10 then)
            call smallwork(x[i])
        else
            call bigwork(x[i])
        endif
    enddo
```

What happens if iterations are statically assigned to threads?

# ***Schedule Clause***

- `schedule (type[, chunk])`
- `type` = static, dynamic, guided, runtime
- `chunk` = scalar integer value
- static
  - Iterations are divided as evenly as possible among all threads
  - *Simple static*
- static, `chunk`
  - Iterations are divided into chunks of size `chunk`
  - Chunks are then assigned in round robin fashion to threads
  - *Interleaved*

# *Schedule Clause*

- `dynamic, chunk`
  - Iterations are divided into chunks of size `chunk` (1 if unspecified) and are assigned to threads dynamically after an initial round robin assignment
  - *Simple dynamic* (chunk size of 1)
- `guided, chunk`
  - Chunk size decreases exponentially from an implementation dependent value (usually N/P) to `chunk` (1 if unspecified)
  - Chunks are assigned dynamically
  - *Guided self scheduling*

# *Schedule Clause*

- *runtime*
  - The schedule type is chosen at runtime based on the environmental variable OMP\_SCHEDULE

```
setenv OMP_SCHEDULE "dynamic, 3"
```
- The choice of schedule is a tradeoff between load balancing and schedule overhead
  - Simple static has the minimum overhead, but has poor load balancing capability
  - Guided is expensive, but balances load better

## *Schedule Clause*

- The best schedule varies based on the structure of each loop and can also vary based on the input data set
- There was a lot of research work on loop scheduling in parallelizing compilers that motivated this
  - C. Polychronopolous, “Guided Self-Scheduling: A Practical Scheduling Scheme for Parallel Supercomputers,” IEEE Transactions on Computers, December 1987.
- *Caution*
  - The schedule clause is a tool for performance enhancement,  
*not for ensuring program correctness*

## *PARALLEL DO: The Small Print*

- The programmer has to make sure that the iterations can in fact be executed in parallel
  - No automatic verification by the compiler

```
subroutine noparallel (z, a, x, y, n)
integer i, n
real z(n), a, x(n), y(n)
!$omp parallel do
do i = 2, n
    z(i) = a * x(i) + y(i) + z(i-1)
enddo
return
end
```

## *PARALLEL DO: Restrictions*

- Number of times that the loop body is executed (*trip-count*) must be available at runtime before the loop is executed
  - Fortran:

```
Do index = lowerbound, upperbound [, stride]
```

➤ trip count computable from bounds and stride
  - C/C++:

```
for (index=start; index op end; incr_expr)
```

➤ index must be an integer variable  
➤ op must be <, <=, > or >=  
➤ expressions start and end must not change during execution

## ***PARALLEL DO: Restrictions (continued)***

- `increment_expr` must change the value of `index` by the same amount after each iteration.
- `increment_expr` must be of the form `index++`, `++index`, `index--`, `--index`, `index += incr`, `index -= incr`, `index = index + incr`, `index = incr +index`, or `index = index - incr`, where `incr` is an expression that does not change during the loop
- Loop body must be able to complete all iterations:
  - Fortran: no `exit` or `goto` that branches outside the loop
  - C/C++: no `break` or `goto` that branches outside the loop
  - C++: no exception caught by a `try` block outside the loop

## *PARALLEL Directive*

- Fortran

```
!$omp parallel [clause [,] [clause ...]]  
    structured block  
 !$omp end parallel
```

- C/C++

```
#pragma omp parallel [clause [clause ...]]  
    structured block
```

## *Parallel Directive: Details*

- When a parallel directive is encountered, threads are spawned which execute the code of the enclosed structured block (i.e., the *parallel region*)
- The number of threads can be specified just like for the PARALLEL DO directive
- The parallel region is replicated and each thread executes a copy of the replicated region

## *Example: Parallel Region*

```
double A[1000];
omp_set_num_threads(4);
#pragma omp parallel
{
    int ID = omp_thread_num();
    pooh(ID, A);
}
printf("all done\n");
```

The diagram illustrates the execution flow of the provided C code. It begins with variable declarations and an `omp_set_num_threads` call. A vertical line leads to the start of a parallel block, indicated by the `#pragma omp parallel` directive. Inside this block, the thread ID is assigned using `ID = omp_thread_num()`. A horizontal line with vertical tick marks indicates the parallel region spans from the start of the block to the end of the block. Within this region, there is a loop with three iterations, each calling `pooh` with arguments `(0, A)`, `(1, A)`, `(2, A)`, and `(3, A)`. The loop is represented by three vertical tick marks on the horizontal line. Finally, the program prints "all done" using `printf("all done\n")`. A green star is placed next to the `printf` line.

# *Parallel versus Parallel Do*

- Arbitrary structured blocks versus loops
- Coarse grained versus fine grained
- Replication versus work division (work sharing)

```
!$omp parallel do
do I = 1,10
    print *, 'Hello world', I
enddo
```

PARALLEL DO is a  
work sharing directive

Output: 10 Hello world messages

```
!$omp parallel
do I = 1,10
    print *, 'Hello world', I
enddo
!$omp end parallel
```

Output:  $10*T$  Hello world messages  
where  $T$  = number of threads

## *Parallel Directive: Clauses*

- private (*list*)
- shared (*list*)
- default (private | shared | none)
- reduction(*intrinsic operator* : *list*)
- if(*logical\_expression*)
- copyin(*list*)

# Parallel: Back to Motivation

```
omp_set_num_threads(2);
#pragma omp parallel private(i, j, x, y, my_width,
    my_thread, i_start, i_end)
{
    my_width = m/2;
    my_thread = omp_get_thread_num();
    i_start = 1 + my_thread * my_width;
    i_end = i_start + my_width - 1;
    for (i = i_start; i <= i_end; i++)
        for (j = 1; j <= n; j++) {
            x = i/ (double) m;
            y = j/ (double) n;
            depth[j][i] = mandel_val(x, y, maxiter);
        }
    for (i = i_start; i <= i_end; i++)
        for (j = 1; j <= n; j++)
            dith_depth[j][i] = 0.5*depth[j][i]
                + 0.25*(depth[j-1][i] + depth[j+1][i])
}
```

What is going  
on here?

# *Work Sharing in Parallel Regions*

- Manual division of work (previous example)
- OMP Work sharing constructs
  - Simplify the programmers job in dividing work among the threads that execute a parallel region
    - **do** directive  
have different threads perform different iterations of a loop
    - **sections** directive  
identify sections of work to be assigned to different threads
    - **single** directive  
specify that a section of code is to be executed by one thread only (remember default is replicated)

# *DO Directive*

## □ Fortran

```
!$omp parallel [clause [,] [clause ...]]  
...  
!$omp do [clause [,] [clause ...]]  
    do loop  
!$omp enddo [nowait]  
...  
!$omp end parallel
```

## □ C/C++

```
#pragma omp parallel [clause [clause ...]]  
{  
...  
#pragma omp for [clause [clause] ... ]  
    for-loop  
}
```

## ***DO Directive: Details***

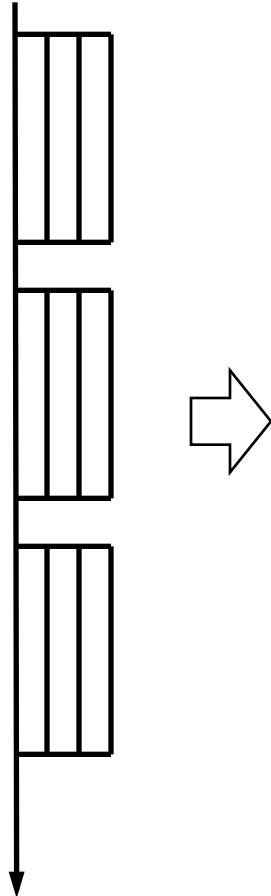
- The DO directive does not spawn new threads!
  - It just assigns work to the threads already spawned by the PARALLEL directive
- The work↔thread assignment is identical to that in the PARALLEL DO directive

```
!$omp parallel do
do I = 1,10
  print *, 'Hello world', I
enddo
```

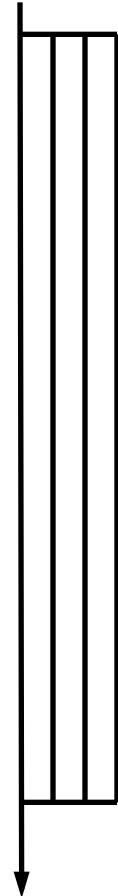
```
!$omp parallel
 !$omp do
do I = 1,10
  print *, 'Hello world', I
enddo
 !$omp enddo
 !$omp end parallel
```

# *Coarser-Grain Parallelism*

```
C$OMP PARALLEL DO  
  do i=1,n  
  .....  
  enddo  
C$OMP PARALLEL DO  
  do i=1,n  
  .....  
  enddo  
C$OMP PARALLEL DO  
  do i=1,n  
  .....  
  enddo
```



```
C$OMP PARALLEL  
C$OMP DO  
  do i=1,n  
  .....  
  enddo  
C$OMP DO  
  do i=1,n  
  .....  
  enddo  
C$OMP DO  
  do i=1,n  
  .....  
  enddo  
C$OMP PARALLEL
```



- What's going on here? Is this possible? When?
- Is this better? Why?

## ***DO/FOR Directive: Clauses***

- private (*list*)
- firstprivate(*list*)
- lastprivate(*list*)
- reduction(*intrinsic operator* : *list*)
- schedule(*type*[, *chunk*] )
- ordered
- nowait

# *Parallel + for: Back to Motivation*

```
omp_set_num_threads(2);
#pragma omp parallel private(i, j)
{
#pragma omp for private (x,y)
for (i = 1; i <= m; i++)
    for (j = 1; j <= n; j++) {
        x = i/ (double) m;
        y = j/ (double) n;
        depth[j][i] = mandel_val(x, y, maxiter);
    }
#pragma omp for
    for (i = 1; i <= m; i++)
        for (j = 1; j <= n; j++)
            dith_depth[j][i] = 0.5*depth[j][i]
                + 0.25*(depth[j-1][i] + depth[j+1][i])
}
```

# *SECTIONS Directive*

- Fortran

```
!$omp sections [clause [,] [clause ...]]  
[ !$omp section  
    code for section 1  
[ !$omp section  
    code for section 2]  
...  
!$omp end sections [nowait]
```

- C/C++

```
#pragma omp sections [clause [clause ...]]  
{  
    [#pragma omp section]  
        block  
    ...  
}
```

## ***SECTIONS Directive: Details***

- Sections are assigned to threads
  - Each section executes once
  - Each thread executes zero or more sections
- Sections are not guaranteed to execute in any order

```
#pragma omp parallel
#pragma omp sections
{
    x_calculation();
#pragma omp section
    y_calculation();
#pragma omp section
    z_calculation();
}
```

## ***SECTIONS Directive: Clauses***

- private (*list*)
- firstprivate(*list*)
- lastprivate(*list*)
- reduction(*intrinsic operator* : *list*)
- nowait

# *SINGLE Directive: Syntax*

## □ Fortran

```
!$omp single [clause [,] [clause ...]]  
    structured block  
 !$omp end single [nowait]
```

## □ C/C++

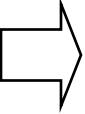
```
#pragma omp single [clause [clause ...]]  
    structured block
```

## □ Clauses:

- private (list)
- firstprivate(list)
- nowait

# *Statements Between Loops*

```
C$OMP PARALLEL DO
C$OMP& REDUCTION(+: sum)
    do i=1,n
        sum = sum + a[i]
    enddo
    alpha = sum/scale
C$OMP PARALLEL DO
    do i=1,n
        a[i] = alpha * a
    enddo
```



```
C$OMP PARALLEL
C$OMP DO REDUCTION(+: sum)
    do i=1,n
        sum = sum + a[i]
    enddo
    alpha = sum/scale
C$OMP SINGLE
    alpha = sum/scale
C$OMP END SINGLE
C$OMP DO
    do i=1,n
        a[i] = alpha * a[i]
    enddo
C$OMP END PARALLEL
```

# *Restrictions on Work Sharing*

- Structured blocks
  - Same nesting level
  - Restriction on branching as before
- If one thread reaches a work sharing construct, all threads reach the construct and in the same order
- Nesting of work sharing constructs is illegal

# *Parallel SECTIONS Directive*

## □ Fortran

```
!$omp parallel sections [clause [,] [clause ...]]  
[ !$omp section  
    code for section 1  
[ !$omp section  
    code for section 2]  
...  
!$omp end parallel sections [nowait]
```

## □ C/C++

```
#pragma omp parallel sections [clause [clause ...]]  
{  
    [#pragma omp section]  
        block  
    ...  
}
```

## *Parallel Sections : Clauses*

- private(*list*)
- shared(*list*)
- firstprivate(*list*)
- lastprivate(*list*)
- default (private | shared | none)
- reduction(*intrinsic operator* : *list*)
- if(*logical\_expression*)
- copyin(*list*)

# *Synchronization*

- Concurrent access to shared data may result in data inconsistency
- Mechanism required to maintain data consistency
  - Mutual exclusion
- Sometimes code sections executed by different threads need to be sequenced in some particular order
  - Event synchronization

# *Mutual Exclusion*

- Mechanisms for ensuring the consistency of data that is accessed concurrently by several threads
  - Critical Directive
  - Atomic Directive
  - Library Lock routines

## *Mutual Exclusion Features*

- Apply to critical, atomic as well as library routines
  - NO Fairness guarantee
  - Guarantee of Progress
  - Careful when nesting - lots of chances for deadlock

# *Event Synchronization*

- Mechanisms for controlling the relative order in which threads execute a section of code
  - Barriers
  - Ordered Sections
  - Master directive

# *Critical Section Directive*

## □ Fortran

```
!$omp critical [(name)]  
    structured block  
!$omp end critical [(name)]
```

## □ C/C++

```
#pragma omp critical [(name)]  
    structured block
```

## *Example: Critical Directive*

```
    cur_max = MINUS_INFINITY
 !$omp parallel do
 do i = 1, n
 ...
 !$omp critical
 if (a(i) .gt. cur_max) then
     cur_max = a(i)
 endif
 !$omp end critical
 ...
 enddo
```

# *OpenMP 3.0*

- Task parallelism is the big news!
  - Allows to parallelize irregular problems
    - unbounded loops
    - recursive algorithms
    - producer/consumer
- Loop parallelism improvements
  - STATIC schedule guarantees
  - Loop collapsing
  - New AUTO schedule
- OMP\_STACK\_SIZE environment variable

# *General Task Characteristics in OpenMP 3.0*

- A task has
  - Code to execute
  - A data environment (it owns its data)
  - An assigned thread that executes the code and uses the data
- Two activities: packaging and execution
  - Each encountering thread packages a new instance of a task (code and data)
  - Some thread in the team executes the task at some later time

# ***Definitions***

- Task construct
  - Task directive plus structured block
- Task
  - The package of code and instructions for allocating data created when a thread encounters a task construct
- Task region
  - The dynamic sequence of instructions produced by the execution of a task by a thread

# **Task Construct**

**#pragma omp task [clause[,]clause] ...]**

*structured-bloc*

*where clause can be one of:*

*if (expression)*

*untied*

*shared (list)*

*private (list)*

*firstprivate (list)*

*default( shared | none )*

# *Intel Thread Building Blocks (TBB)*

- Sutter: <http://www.go-parallel.com>, March 2008
  - Main site: <http://threadingbuildingblocks.org/>
- C++ Library
- Targets threading for performance (designed to parallelize computationally intensive work)
- Is compatible with other threading packages
- Emphasizes scalable data parallel programming
- Relies on generic programming
- Specifies templates and tasks instead of threads
  - library schedules tasks onto threads
  - manages load balancing

## Parallel\_for

- parallel\_for is a template function provided by library
- Example:
  - concurrently apply a function to each element in an array
- Serial version:

```
void SerialApplyFoo( float a[], size_t n ) {  
    for( size_t i=0; i<n; ++i )  
        Foo(a[i]);  
}
```
- Iteration space is 0...(n-1)

## *Parallel\_for (continued)*

- Parallel version requires two steps – Step 1:

```
#include "tbb / blocked_range.h"  
class ApplyFoo {  
    float *const my_a;  
public:  
    ApplyFoo( float a[] ) : my_a(a) {}  
    void operator()( const blocked_range<size_t>& r ) const {  
        float *a = my_a;  
        for( size_t i=r.begin(); i!=r.end(); ++i )  
            Foo(a[i]);  
    }  
};
```

## ***Parallel\_for (continued)***

- Parallel version requires two steps – Step 2:

```
#include "tbb/parallel_for.h"  
void ParallelApplyFoo( float a[], size_t n ) {  
    parallel_for(blocked_range<size_t>(0,n,IdealGrainSize),  
        ApplyFoo(a));  
}
```

- `parallel_for` breaks iteration space into chunks each of which are run on separate threads
- `blocked_range<T>(begin,end,grainsize)` - recursively divisible struct
- `blocked_range` is for 1D iteration spaces
  - TBB provides 2D and users can define their own
- *grainsize* specifies the number of iterations for a “reasonable size” chunk to deal out to a processor
  - If the iteration space has more than *grainsize* iterations, `parallel_for` splits it into separate subranges that are scheduled separately
- operator () processes a chunk

## ***TBB Library - Algorithms***

- parallel\_reduce
  - math ops with elements of an array in parallel
- parallel\_do
  - loops indeterminate length iteration spaces
- parallel\_\*
  - several others...

## ***TBB Library – Pipeline***

- TBB implements the pipeline pattern
- Data flows through a series of pipeline stages, and each stage processes the data in some way
- Given an incoming stream of data, some stages operate in parallel and others do not

## **TBB – Containers**

- concurrent\_hash\_map<Key,T,HashCompare>
- concurrent\_queue<T,Allocator>
- concurrent\_vector<T,Allocator>
- Supports concurrent access, concurrent operations and parallel iteration
- TBB library retains control over memory allocation

## ***TBB – Allocation***

- `tbb_allocator<T>`
  - Allocates and frees memory via the TBB malloc library if available, otherwise it reverts to using malloc and free
- `scalable_allocator<T>`
  - Allocates and frees memory in a way that scales with the number of processors
- others...

## ***TBB – Synchronization***

- Library provides several flavors of mutexes including:  
recursive, spin, queuing, reader/writer...
- atomic<T>
  - read, write
  - fetch-and-add, fetch-and-store, compare-and-swap

## *Next Class*

- Parallel algorithms