

CIS 631

Parallel Processing

Lecture 11: Parallel Algorithms

Allen D. Malony
malony@cs.uoregon.edu

Department of Computer and Information Science
University of Oregon



Acknowledgements

- Portions of the lectures slides were adopted from:
 - A. Grama, A. Gupta, G. Karypis, and V. Kumar, “Introduction to Parallel Computing,” 2003.
 - Chapters 8, 9, and 10

Outline

- ❑ Dense matrix algorithms
- ❑ Sorting algorithms
- ❑ Graph algorithms

Dense Matrix Algorithms

- ❑ Great deal of activity in algorithms and software for solving linear algebra problems
 - Solution of linear systems ($Ax = b$)
 - Least-squares solution of over- or under-determined systems ($\min ||Ax-b||$)
 - Computation of eigenvalues and eigenvectors ($Ax = \lambda x$)
 - Driven by numerical problem solving in scientific computation
- ❑ Solutions involves various forms of matrix computations
- ❑ Focus on high-performance matrix algorithms
 - Key insight is to maximize computation to communication

Solving a System of Linear Equations

□ $Ax=b$

$$a_{0,0}x_0 + a_{0,1}x_1 + \dots + a_{0,n-1}x_{n-1} = b_0$$

$$a_{1,0}x_0 + a_{1,1}x_1 + \dots + a_{1,n-1}x_{n-1} = b_1$$

...

$$a_{n-1,0}x_0 + a_{n-1,1}x_1 + \dots + a_{n-1,n-1}x_{n-1} = b_{n-1}$$

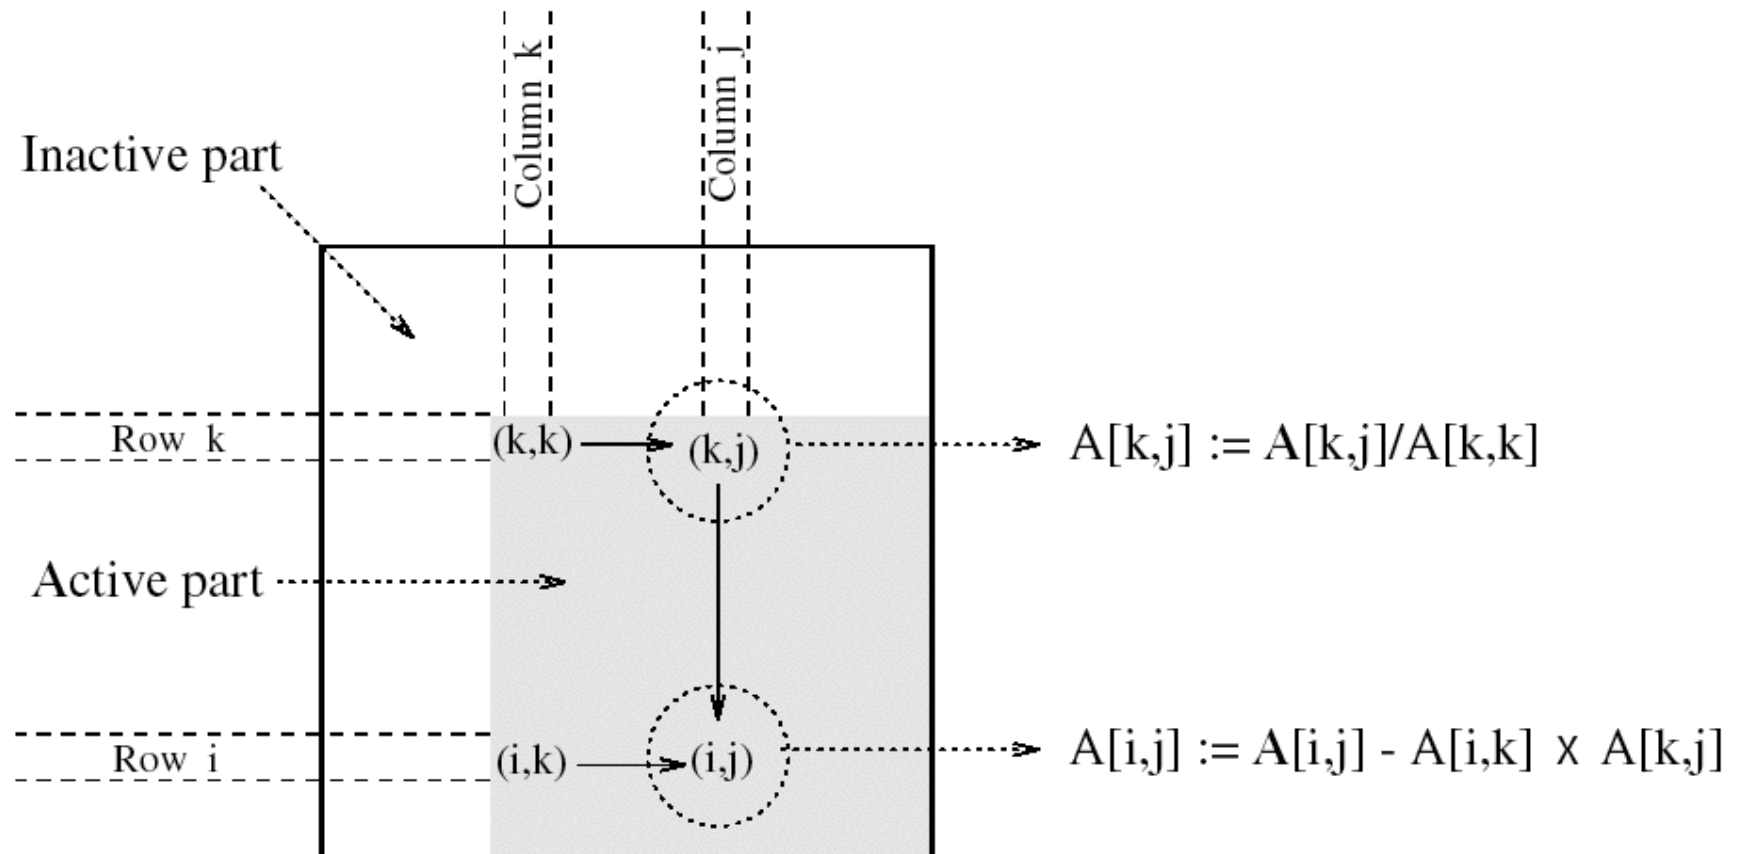
□ Gaussian elimination (classic algorithm)

- Forward elimination to $Ux=y$ (U is upper triangular)
 - Without or with partial pivoting
- Back substitution to solve for x
- Parallel algorithms based on partitioning of A

Sequential Gaussian Elimination

```
1.  procedure GAUSSIAN ELIMINATION (A, b, y)
2.  Begin
3.      for  $k := 0$  to  $n - 1$  do /* Outer loop */
4.      begin
5.          for  $j := k + 1$  to  $n - 1$  do
6.               $A[k, j] := A[k, j] / A[k, k]$ ; /* Division step */
7.               $y[k] := b[k] / A[k, k]$ ;
8.               $A[k, k] := 1$ ;
9.          for  $i := k + 1$  to  $n - 1$  do
10.         begin
11.             for  $j := k + 1$  to  $n - 1$  do
12.                  $A[i, j] := A[i, j] - A[i, k] \times A[k, j]$ ; /* Elimination step */
13.                  $b[i] := b[i] - A[i, k] \times y[k]$ ;
14.                  $A[i, k] := 0$ ;
15.             endfor; /*Line9*/
16.         endfor; /*Line3*/
17.     end GAUSSIAN ELIMINATION
```

Computation Step in Gaussian Elimination



$$5x + 3y = 22$$

$$8x + 2y = 13$$



$$x = (22 - 3y) / 5$$

$$8(22 - 3y)/5 + 2y = 13$$



$$x = (22 - 3y) / 5$$

$$y = (13 - 176/5) / (24/5 + 2)$$

Rowwise Partitioning on Eight Processes

P_0	1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
P_1	0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
P_2	0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
P_3	0	0	0	(3,3)	(3,4)	(3,5)	(3,6)	(3,7)
P_4	0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
P_5	0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
P_6	0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
P_7	0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

P_0	1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
P_1	0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
P_2	0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
P_3	0	0	0	1	(3,4)	(3,5)	(3,6)	(3,7)
P_4	0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
P_5	0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
P_6	0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
P_7	0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

(a) Computation:

(i) $A[k,j] := A[k,j]/A[k,k]$ for $k < j < n$

(ii) $A[k,k] := 1$

(b) Communication:

One-to-all broadcast of row $A[k,*]$

Rowwise Partitioning on Eight Processes

P_0	1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
P_1	0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
P_2	0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
P_3	0	0	0	1	(3,4)	(3,5)	(3,6)	(3,7)
P_4	0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
P_5	0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
P_6	0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
P_7	0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

(c) Computation:

(i) $A[i,j] := A[i,j] - A[i,k] \times A[k,j]$
for $k < i < n$ and $k < j < n$

(ii) $A[i,k] := 0$ for $k < i < n$

2D Mesh Partitioning on 64 Processes

1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
0	0	0	(3,3)	(3,4)	(3,5)	(3,6)	(3,7)
0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

(a) Rowwise broadcast of $A[i,k]$
for $(k-1) < i < n$

1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
0	0	0	(3,3)	(3,4)	(3,5)	(3,6)	(3,7)
0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

(b) $A[k,j] := A[k,j]/A[k,k]$
for $k < j < n$

1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
0	0	0	1	(3,4)	(3,5)	(3,6)	(3,7)
0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

(c) Columnwise broadcast of $A[k,j]$
for $k < j < n$

1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
0	0	0	1	(3,4)	(3,5)	(3,6)	(3,7)
0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

(d) $A[i,j] := A[i,j] - A[i,k] \times A[k,j]$
for $k < i < n$ and $k < j < n$

Back Substitution to Find Solution


```
1. procedure BACK SUBSTITUTION ( $U, x, y$ )
2. begin
3.   for  $k := n - 1$  downto 0 do /* Main loop */
4.   begin
5.      $x[k] := y[k];$ 
6.     for  $i := k - 1$  downto 0 do
7.        $y[i] := y[i] - x[k] \times U[i, k];$ 
8.   endfor;
9. end BACK SUBSTITUTION
```

Dense Linear Algebra (www.netlib.gov)


- ❑ Basic Linear Algebra Subroutines (BLAS)
 - Level 1 (*vector-vector*): vectorization
 - Level 2 (*matrix-vector*): vectorization, parallelization
 - Level 3 (*matrix-matrix*): parallelization
- ❑ LINPACK (Fortran)
 - Linear equations and linear least-squares
- ❑ EISPACK (Fortran)
 - Eigenvalues and eigenvectors for matrix classes
- ❑ LAPACK (Fortran, C) (LINPACK + EISPACK)
 - Use BLAS internally
- ❑ ScaLAPACK (Fortran, C, MPI) (scalable LAPACK)

Numerical Libraries


- ❑ PETSc (<http://www.mcs.anl.gov/petsc/petsc-as>)
 - data structures / routines for partial differential equations
 - MPI based



- ❑ SuperLU (<http://crd.lbl.gov/~xiaoye/SuperLU/>)
 - Large sparse nonsymmetric linear systems




- ❑ Hypre (<http://www.llnl.gov/CASC/hypre>)
 - Large sparse linear systems



- ❑ TAO (<http://www.mcs.anl.gov/research/projects/tao/>)
 - Toolkit for Advanced Optimization

- ❑ DOE ACTS (<http://acts.nerisc.gov/>)
 - Advanced CompuTational Software



Sorting Algorithms


- ❑ Task of arranging unordered collection into order
- ❑ Permutation of a sequence of elements
- ❑ Internal versus external sorting
 - External sorting uses auxiliary storage
- ❑ Comparison-based
 - Compare pairs of elements and exchange
 - $O(n \log n)$
- ❑ Noncomparison-based
 - Use known properties of elements
 - $O(n)$

Sorting on Parallel Computers

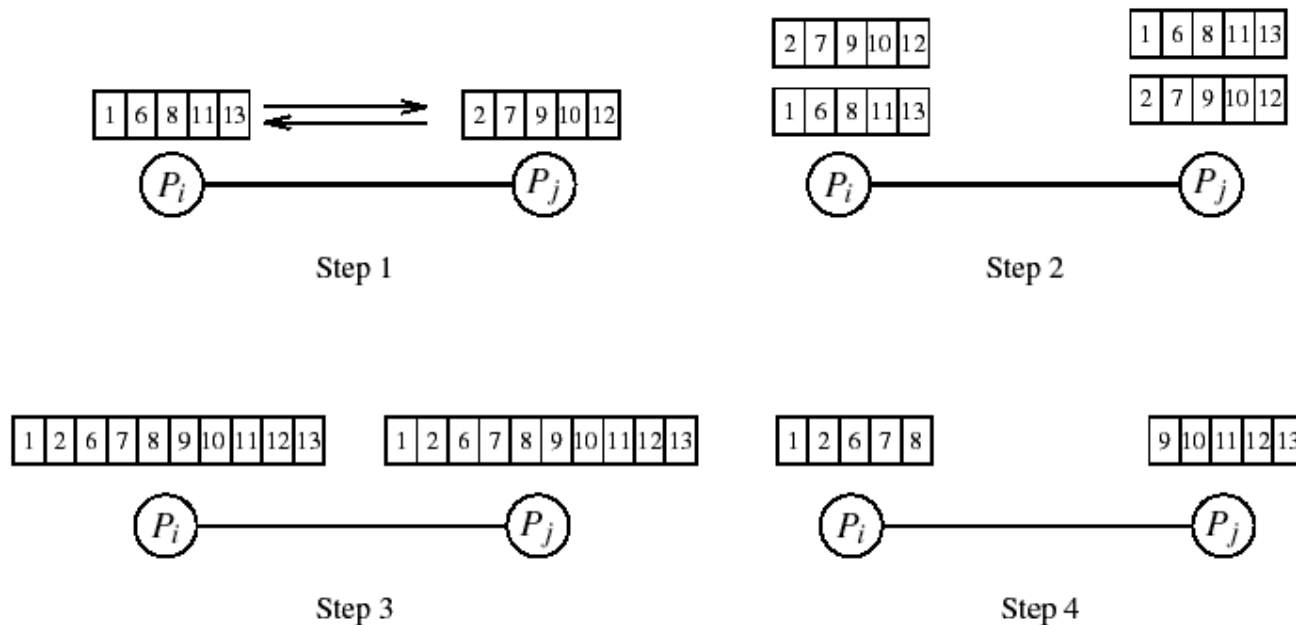
- ❑ Where are the elements stored?
 - Need to be distributed across processes
 - Sorted order will be with respect to process order
- ❑ How are comparisons performed?
 - One element per process
 - compare-exchange
 - interprocess communication will dominate execution time
 - More than one element per process
 - compare-split
- ❑ Sorting networks
 - Based on comparison network model
- ❑ Contrast with shared memory sorting algorithms

Single vs. Multi Element Comparision

❑ One element per processor

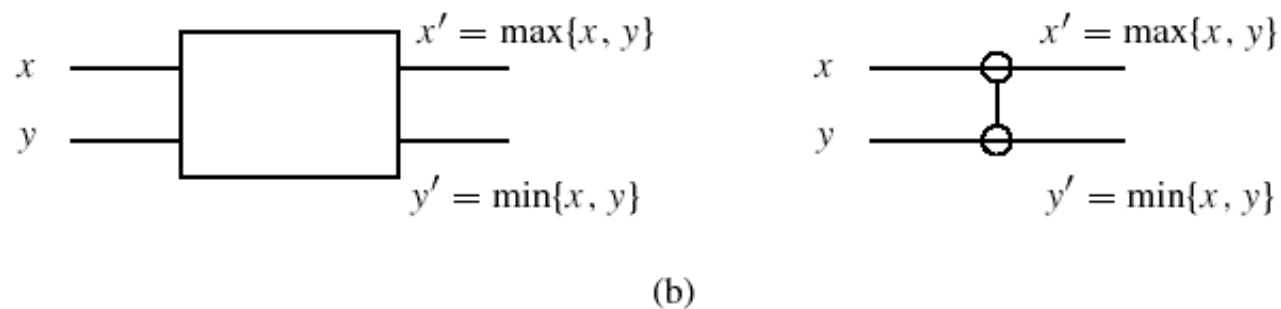
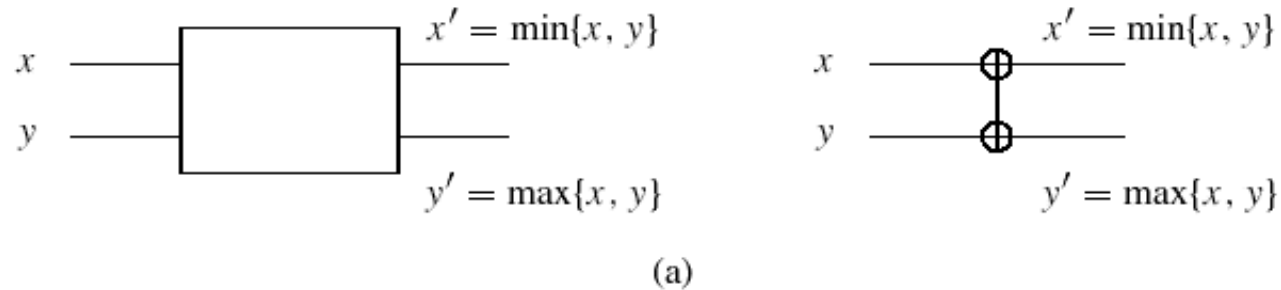


❑ Multiple elements per processor



Sorting Networks

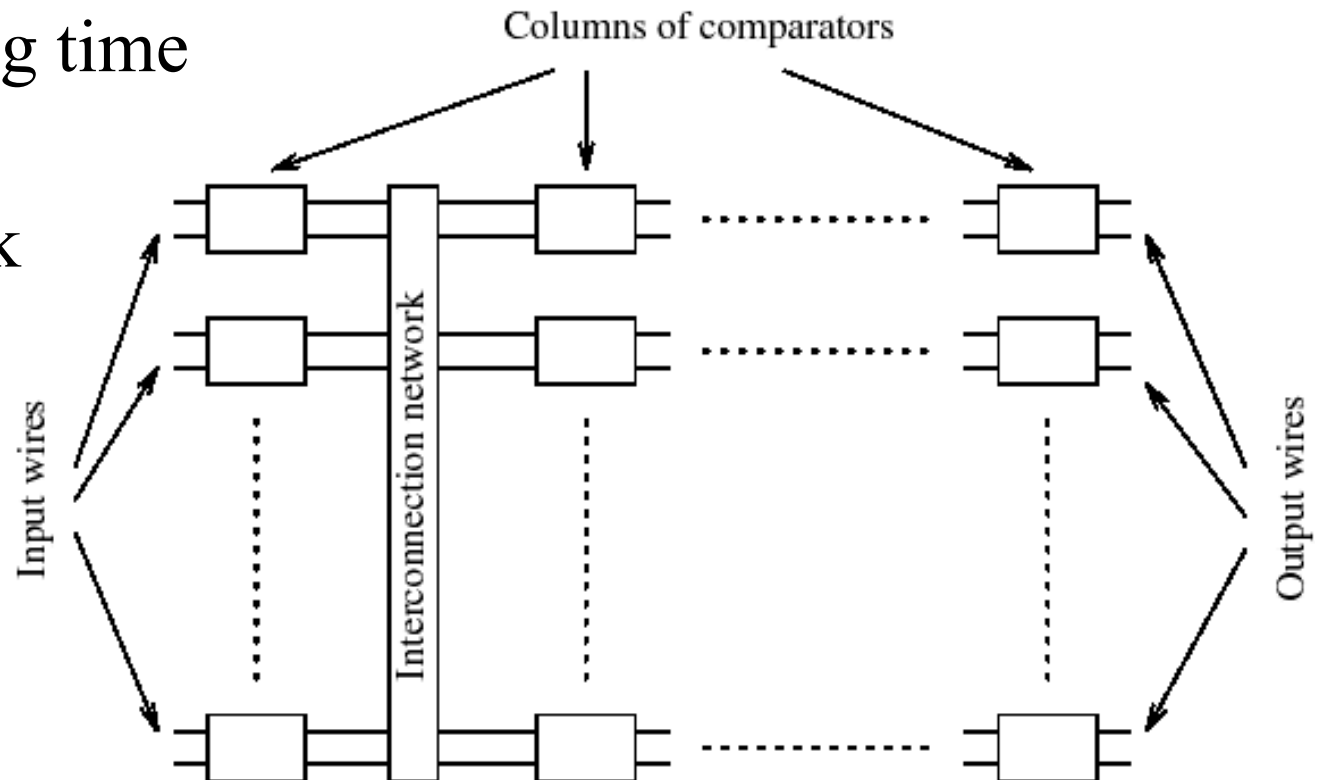
- ❑ Networks to sort n elements in less than $O(n \log n)$
- ❑ Key component in network is a comparator
 - Increasing or decreasing comparator



- ❑ Comparators connected in parallel and permute elements

Sorting Network Design


- ❑ Multiple comparator stages (# stages, # comparators)
- ❑ Connected together by interconnection network
- ❑ Output of last stage is the sorted list
- ❑ $O(\log^2 n)$ sorting time
- ❑ Convert any sorting network to sequential algorithm



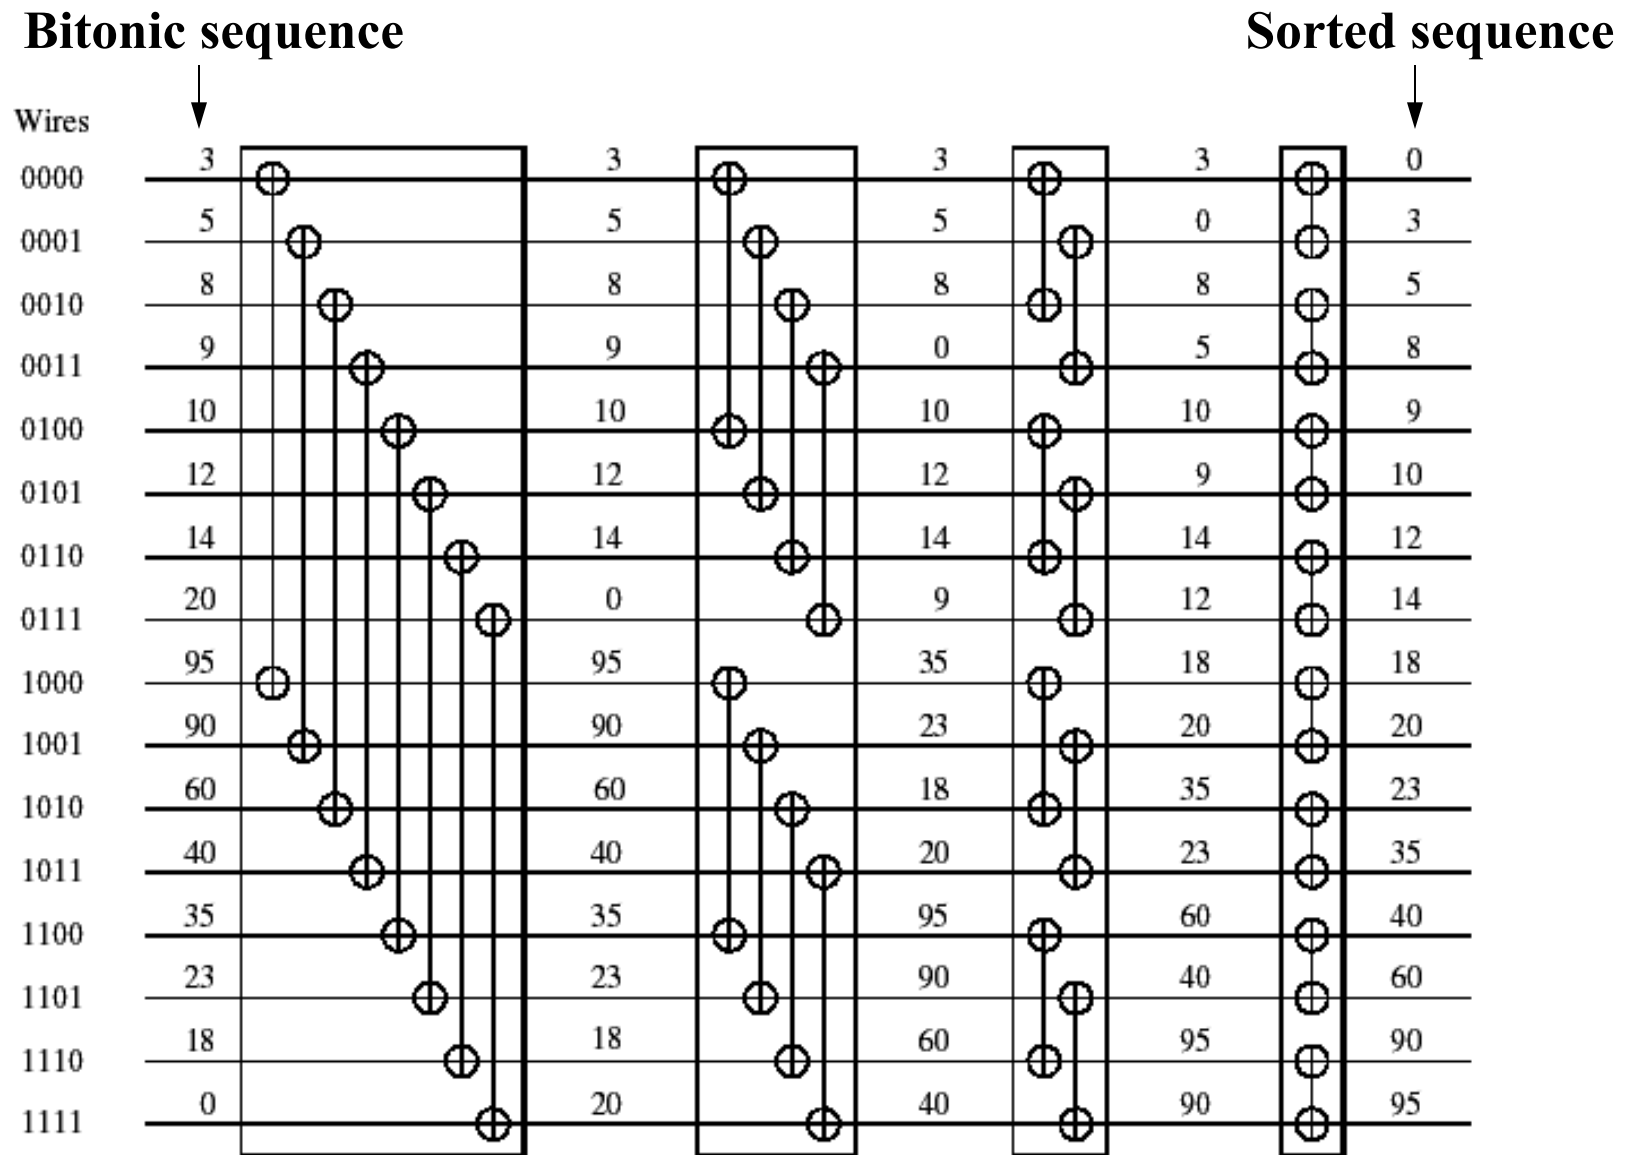
Bitonic Sort

- ❑ Create a *bitonic sequence* then sort the sequence
- ❑ Bitonic sequence
 - sequence of elements $\langle a_0, a_1, \dots, a_{n-1} \rangle$
 - $\langle a_0, a_1, \dots, a_i \rangle$ is monotonically increasing
 - $\langle a_i, a_{i+1}, \dots, a_{n-1} \rangle$ is monotonically decreasing
- ❑ Sorting using *bitonic splits* is called *bitonic merge*
- ❑ *Bitonic merge network* is a network of comparators
 - Implement bitonic merge
- ❑ Bitonic sequence is formed from unordered sequence
 - Bitonic sort creates a bitonic sequence
 - Start with sequence of size two (default bitonic)

Bitonic Sort Network




Bitonic Merge Network




Parallel Bitonic Sort on a Hypercube

```
1. procedure BITONIC SORT(label, d)
2. begin
3.   for i := 0 to d - 1 do
4.     for j := i downto 0 do
5.       if (i + 1)st bit of label = j th bit of label then
6.         comp exchange max(j);
7.       else
8.         comp exchange min(j);
9.   end BITONIC SORT
```


Parallel Bitonic Sort on a Hypercube (Last stage)




Step 1



Step 2



Step 3




Step 4

Bubble Sort and Variants

- ❑ Can easily parallelize sorting algorithms of $O(n^2)$
- ❑ *Bubble sort* compares and exchanges adjacent elements
 - $O(n)$ each pass
 - $O(n)$ passes
 - Available parallelism?
- ❑ *Odd-even transposition sort*
 - Compares and exchanges odd and even pairs
 - After n phases, elements are sorted
 - Available parallelism?

Odd-Even Transposition Sort



Parallel Odd-Even Transposition Sort on Ring

```
1.  procedure ODD-EVEN PAR( $n$ )
2.  begin
3.       $id :=$  process' s label
4.      for  $i := 1$  to  $n$  do
5.          begin
6.              if  $i$  is odd then
7.                  if  $id$  is odd then
8.                      compare-exchange  $\min(id + 1)$ ;
9.                  else
10.                     compare-exchange  $\max(id - 1)$ ;
11.             if  $i$  is even then
12.                 if  $id$  is even then
13.                     compare-exchange  $\min(id + 1)$ ;
14.                 else
15.                     compare-exchange  $\max(id - 1)$ ;
16.             end for
17. end ODD-EVEN PAR
```

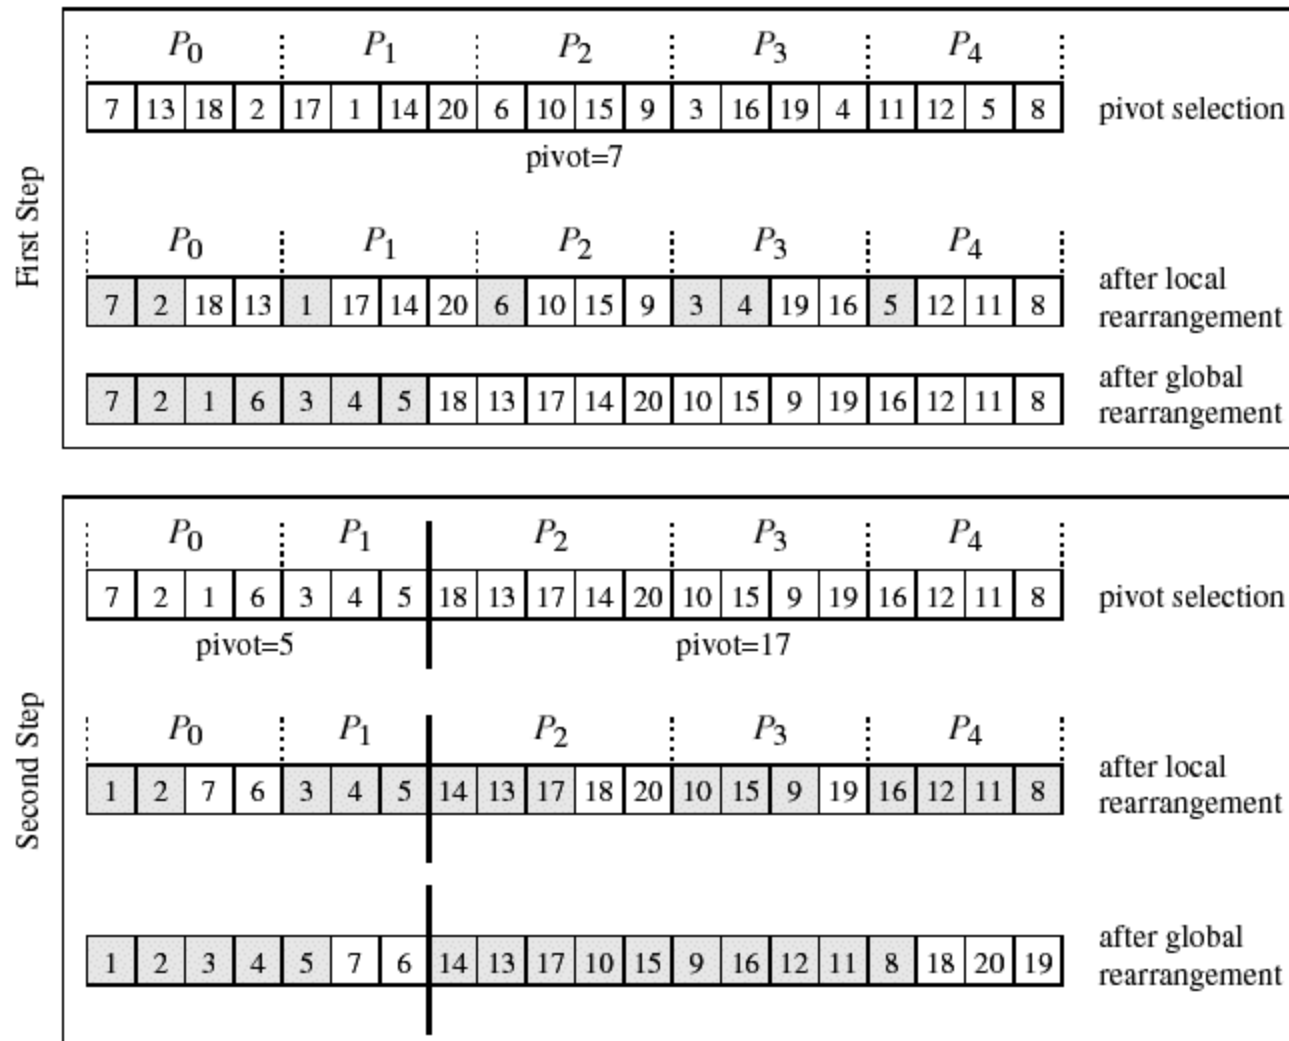
Quicksort

- ❑ *Quicksort* has average complexity of $O(n \log n)$
- ❑ Divide-and-conquer algorithm
 - Divide into subsequences where every element in first is less than or equal to every element in the second
 - Pivot is used to split the sequence
 - Conquer step recursively applies quicksort algorithm
- ❑ Available parallelism?

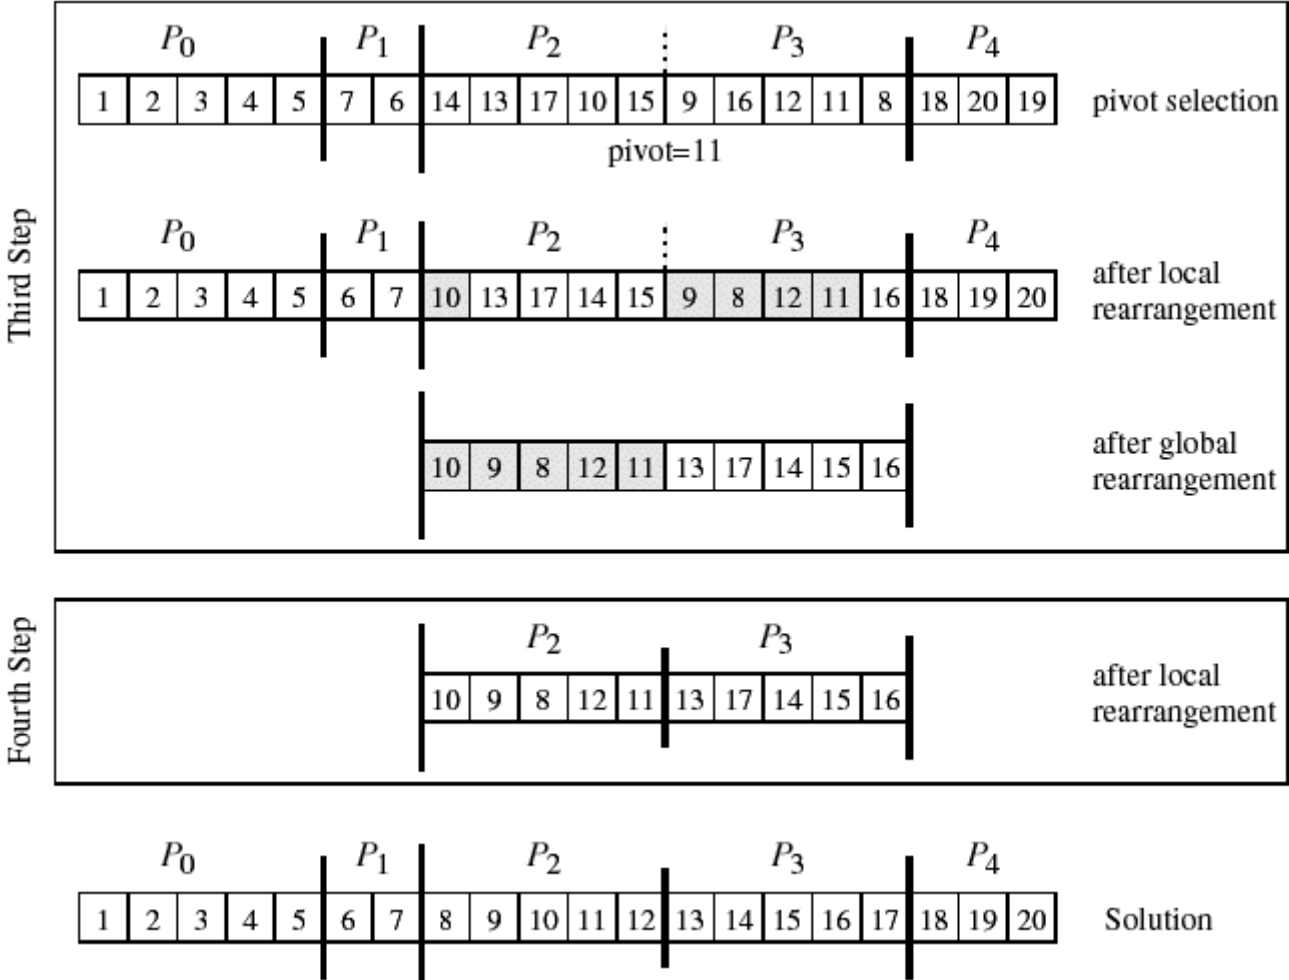
Sequential Quicksort

```
1.  procedure QUICKSORT (A, q, r )
2.  begin
3.      if  $q < r$  then
4.          begin
5.               $x := A[q]$ ;
6.               $s := q$ ;
7.              for  $i := q + 1$  to  $r$  do
8.                  if  $A[i] \leq x$  then
9.                      begin
10.                          $s := s + 1$ ;
11.                         swap( $A[s]$ ,  $A[i]$ );
12.                     end if
13.                 swap( $A[q]$ ,  $A[s]$ );
14.                 QUICKSORT (A, q, s);
15.                 QUICKSORT (A, s + 1, r );
16.             end if
17. end QUICKSORT
```

Parallel Shared Address Space Quicksort




Parallel Shared Address Space Quicksort




Bucket Sort and Sample Sort

- ❑ *Bucket sort* is popular when elements (values) are uniformly distributed over an interval
 - Create m buckets and place elements in appropriate bucket
 - $O(n \log(n/m))$
 - If $m=n$, can use value as index to achieve $O(n)$ time
- ❑ *Sample sort* is used when uniformly distributed assumption is not true
 - Distributed to m buckets and sort each with quicksort
 - Draw sample of size s
 - Sort samples and choose $m-1$ elements to be *splitters*
 - Split into m buckets and proceed with bucket sort


Parallel Sample Sort




Initial element distribution




Local sort & sample selection



Sample combining





Global splitter selection



Final element assignment

Graph Algorithms

- ❑ Graph theory important in computer science
- ❑ Many complex problems are graph problems
- ❑ $G = (V, E)$
 - V finite set of points called vertices
 - E finite set of edges
 - $e \in E$ is an pair (u, v) , where $u, v \in V$
 - Unordered and ordered graphs




Graph Terminology

- Vertex *adjacency* if (u,v) is an edge
- *Path* from u to v if there is an edge sequence starting at u and ending at v
- If there exists a path, v is *reachable* from u
- A graph is *connected* if all pairs of vertices are connected by a path
- A *weighted* graph associates weights with each edge
- *Adjacency matrix* is an $n \times n$ array A such that
 - $A_{i,j} = 1$ if $(v_i, v_j) \in E$; 0 otherwise
 - Can be modified for weighted graphs (∞ is no edge)
 - Can represent as *adjacency lists*



Graph Representations

□ Adjacency matrix



$$A = \begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 1 & 1 & 0 \end{bmatrix}$$

□ Adjacency list



Minimum Spanning Tree

- ❑ A *spanning tree* of an undirected graph G is a subgraph of G that is a tree containing all the vertices of G
- ❑ The *minimum spanning tree* (MST) for a weighted undirected graph is a spanning tree with minimum weight
- ❑ Prim's algorithm can be used
 - Greedy algorithm
 - Selects an arbitrary starting vertex
 - Chooses new vertex guaranteed to be in MST
 - $O(n^2)$
 - Prim's algorithm is iterative


Prim's Minimum Spanning Tree Algorithm

```
1. procedure PRIM MST( $V, E, w, r$  )
2. begin
3.    $VT := \{r\}$ ;
4.    $d[r] := 0$ ;
5.   for all  $v \in (V - VT)$  do
6.     if edge  $(r, v)$  exists set  $d[v] := w(r, v)$ ;
7.     else set  $d[v] := \infty$ ;
8.   while  $VT \neq V$  do
9.     begin
10.      find a vertex  $u$  such that  $d[u] := \min\{d[v] | v \in (V - VT)\}$ ;
11.       $VT := VT \cup \{u\}$ ;
12.      for all  $v \in (V - VT)$  do
13.         $d[v] := \min\{d[v], w(u, v)\}$ ;
14.      endwhile
15. end PRIM MST
```




Example: Prim's MST Algorithm

(a) Original graph



	a	b	c	d	e	f
d[]	1	0	5	1	∞	∞
a	0	1	3	∞	∞	3
b	1	0	5	1	∞	∞
c	3	5	0	2	1	∞
d	∞	1	2	0	4	∞
e	∞	∞	1	4	0	5
f	2	∞	∞	∞	5	0


(b) After the first edge has been selected



	a	b	c	d	e	f
d[]	1	0	2	1	4	∞
a	0	1	3	∞	∞	3
b	1	0	5	1	∞	∞
c	3	5	0	2	1	∞
d	∞	1	2	0	4	∞
e	∞	∞	1	4	0	5
f	2	∞	∞	∞	5	0


Example: Prim's MST Algorithm

(c) After the second edge has been selected



	a	b	c	d	e	f
d[]	1	0	2	1	4	3
a	0	1	3	∞	∞	3
b	1	0	5	1	∞	∞
c	3	5	0	2	1	∞
d	∞	1	2	0	4	∞
e	∞	∞	1	4	0	5
f	2	∞	∞	∞	5	0

(d) Final minimum spanning tree





	a	b	c	d	e	f
d[]	1	0	2	1	1	3
a	0	1	3	∞	∞	3
b	1	0	5	1	∞	∞
c	3	5	0	2	1	∞
d	∞	1	2	0	4	∞
e	∞	∞	1	4	0	5
f	2	∞	∞	∞	5	0

Parallel Formulation of Prim's Algorithm

- ❑ Difficult to perform different iterations of the **while** loop in parallel because $d[v]$ may change each time
- ❑ Can parallelize each iteration though
- ❑ Partition vertices into p subsets V_i , $i=0, \dots, p-1$
- ❑ Each process P_i computes
$$d_i[u] = \min\{d_i[v] \mid v \in (V - V_T) \cap V_i\}$$
- ❑ Global minimum is obtained using all-to-one reduction
- ❑ New vertex is added to V_T and broadcast to all processes
- ❑ New values of $d[v]$ are computed for local vertex
- ❑ $O(n^2/p) + O(n \log p)$ (computation + communication)

Partitioning in Prim's Algorithm



Single-Source Shortest Paths

- ❑ Find *shortest path* from a vertex v to all other vertices
- ❑ The shortest path in a weighted graph is the edge with the minimum weight
- ❑ Weights may represent time, cost, loss, or any other quantity that accumulates additively along a path
- ❑ Dijkstra's algorithm finds shortest paths from a vertex s
 - Similar to Prim's MST algorithm
 - MST with vertex v as starting vertex
 - Incrementally finds shortest paths in greedy manner
 - Keep track of minimum cost to reach a vertex from s
 - $O(n^2)$

Dijkstra's Single-Source Shortest Paths Algorithm

1. **procedure** DIJKSTRA SINGLE SOURCE SP(V, E, w, s)
2. **begin**
3. $V_T := \{s\};$
4. **for** all $v \in (V - V_T)$ **do**
5. **if** (s, v) exists set $l[v] := w(s, v);$
6. **else** set $l[v] := \infty;$
7. **while** $V_T \neq V$ **do**
8. **begin**
9. find a vertex u such that $l[u] := \min\{l[v] | v \in (V - V_T)\};$
10. $V_T := V_T \cup \{u\};$
11. **for** all $v \in (V - V_T)$ **do**
12. $l[v] := \min\{l[v], l[u] + w(u, v)\};$
13. **endwhile**
14. **end** DIJKSTRA SINGLE SOURCE SP

Parallel Formulation of Dijkstra's Algorithm

- ❑ Very similar to Prim's MST parallel formulation
- ❑ Use 1D block mapping as before
- ❑ All processes perform computation and communication similar to that performed in Prim's algorithm
- ❑ Parallel performance is the same
 - $O(n^2/p) + O(n \log p)$
 - Scalability
 - $O(n^2)$ is the sequential time
 - $O(n^2) / [O(n^2/p) + O(n \log p)]$

All Pairs Shortest Path

- ❑ Find the shortest path between all pairs of vertices
- ❑ Outcome is a $n \times n$ matrix $D = \{d_{i,j}\}$ such that $d_{i,j}$ is the cost of the shortest path from vertex v_i to vertex v_j
- ❑ Dijkstra's algorithm
 - Execute single-source algorithm on each process
 - $O(n^3)$
 - Source-partitioned formulation (use sequential algorithm)
 - Source-parallel formulation (use parallel algorithm)
- ❑ Floyd's algorithm
 - Builds up distance matrix from the bottom up

Floyd's All-Pairs Shortest Paths Algorithm


1. **procedure** FLOYD ALL PAIRS SP(A)
2. **begin**
3. $D^{(0)} = A;$
4. **for** $k := 1$ **to** n **do**
5. **for** $i := 1$ **to** n **do**
6. **for** $j := 1$ **to** n **do**
7. $d^{(k)}_{i,j} := \min d^{(k-1)}_{i,j}, d^{(k-1)}_{i,k} + d^{(k-1)}_{k,j};$
8. **end** FLOYD ALL PAIRS SP

Parallel Floyd's Algorithm

1. **procedure** FLOYD ALL PAIRS PARALLEL (A)
2. **begin**
3. $D^{(0)} = A;$
4. **for** $k := 1$ **to** n **do**
5. **forall** $P_{i,j}$, where $i, j \leq n$, **do in parallel**
6. $d^{(k)}_{i,j} := \min d^{(k-1)}_{i,j}, d^{(k-1)}_{i,k} + d^{(k-1)}_{k,j};$
7. **end** FLOYD ALL PAIRS PARALLEL

Parallel Graph Algorithm Library – Boost

- ❑ Parallel Boost Graph Library (Indiana University)
 - Generic C++ library for high-performance parallel and distributed graph computation
 - Builds on the Boost Graph Library (BGL)
 - offers similar data structures, algorithms, and syntax
 - Research platform for parallel graph algorithms
 - Provide solid implementations for solving large-scale graph problems
 - Boost Software License (BSD-like)



D. Gregor and A. Lumsdaine, “The Parallel BGL: A Generic Library for Distributed Graph Computations, Parallel Object-Oriented Scientific Computing (POOSC), July 2005.

Original BGL: Algorithms

- ❑ Searches (breadth-first, depth-first, A*)
- ❑ Single-source shortest paths (Dijkstra, Bellman-Ford, DAG)
- ❑ All-pairs shortest paths (Johnson, Floyd-Warshall)
- ❑ Minimum spanning tree (Kruskal, Prim)
- ❑ Components (connected, strongly connected, biconnected)
- ❑ Maximum cardinality matching
- ❑ Max-flow (Edmonds-Karp, push-relabel)
- ❑ Sparse matrix ordering (Cuthill-McKee, King, Sloan, minimum degree)
- ❑ Layout (Kamada-Kawai, Fruchterman-Reingold, Gursoy-Atun)
- ❑ Betweenness centrality
- ❑ PageRank
- ❑ Isomorphism
- ❑ Vertex coloring
- ❑ Transitive closure
- ❑ Dominator tree

Original BGL Summary

- ❑ The original BGL is large, stable, efficient
 - Lots of algorithms, graph types
 - Peer-reviewed code with many users, nightly regression testing, and so on
 - Performance comparable to FORTRAN.
- ❑ Who should use the BGL?
 - Programmers comfortable with C++
 - Users with graph sizes from tens of vertices to millions of vertices


BGL-Python

- ❑ Python is ideal for rapid prototyping:
 - It's a scripting language (no compiler)
 - Dynamically typed means less typing for you
 - Easy to use: you already know Python...
- ❑ BGL-Python provides access to the BGL from within Python
 - Similar interfaces to C++ BGL
 - Easier to learn than C++
 - Great for scripting, GUI applications
 - `help(bgl.dijkstra_shortest_paths)`


The Parallel BGL

- ❑ A version of the C++ BGL for computational clusters
 - Distributed memory for huge graphs
 - Parallel processing for improved performance
- ❑ An active research project
- ❑ Closely related to the original BGL
 - Parallelizing BGL programs should be “easy”

A simple, directed graph...



distributed across 3 processors



Parallel Graph Algorithms

- ❑ Breadth-first search
- ❑ Eager Dijkstra's single-source shortest paths
- ❑ Crauser et al. single-source shortest paths
- ❑ Depth-first search
- ❑ Minimum spanning tree (Boruvka, Dehne & Götz)
- ❑ Connected components
- ❑ Strongly connected components
- ❑ Biconnected components
- ❑ PageRank
- ❑ Graph coloring
- ❑ Fruchterman-Reingold layout
- ❑ Max-flow (Dinic's)

Parallel BGL in Python

- ❑ Preliminary support for the Parallel BGL in Python
 - Just `import boost.graph.distributed`
 - Similar interface to sequential BGL-Python
- ❑ Several options for usage with MPI:
 - Straight MPI: `mpirun -np 2 python script.py`
 - pyMPI: allows interactive use of the interpreter
- ❑ Initially used to prototype our distributed Fruchterman-Reingold implementation.

Parallel BGL Summary

- ❑ The Parallel BGL is built for huge graphs
 - Millions to hundreds of millions of nodes
 - Distributed-memory parallel processing on clusters
 - Future work will permit larger graphs...
- ❑ Parallel programming has a learning curve
 - Parallel graph algorithms much harder to write
 - Distributed graph manipulation can be tricky
- ❑ Parallel BGL is an active research library

<http://osl.iu.edu/research/pbgl/>

Next Class

- ❑ Algorithms for simulation
- ❑ Analytical modeling of parallel programs