

# Map Pattern

Parallel Computing

CIS 410/510

Department of Computer and Information Science



UNIVERSITY OF OREGON

# *Table of Contents*

- ❑ Map
- ❑ Optimizations
  - Sequences of Maps
  - Code Fusion
  - Cache Fusion
- ❑ Related Patterns
- ❑ Example Implementation: Scaled Vector Addition (SAXPY)
  - Problem Description
  - Various Implementations

# *Table of Contents*

## □ Map

## □ Optimizations

- Sequences of Maps
- Code Fusion
- Cache Fusion

## □ Related Patterns

## □ Example Implementation: Scaled Vector Addition (SAXPY)

- Problem Description
- Various Implementations

# Mapping

- “Do the same thing many times”

```
foreach i in foo:
```

```
    do something
```


- Well-known higher order function in languages like ML, Haskell, Scala

$$\text{map} : \forall ab.(a \rightarrow b) \text{List}\langle a \rangle \rightarrow \text{List}\langle b \rangle$$


applies a function each element in a list and returns a list of results

# Example Maps

Add 1 to every item in an array



Double every item in an array



**Key Point:** An operation is a map if it can be applied to each element without knowledge of neighbors.

# *Key Idea*

- Map is a “foreach loop”

# *Key Idea*

- Map is a “foreach loop” where each iteration is independent

# Key Idea

- Map is a “foreach loop” where each iteration is independent

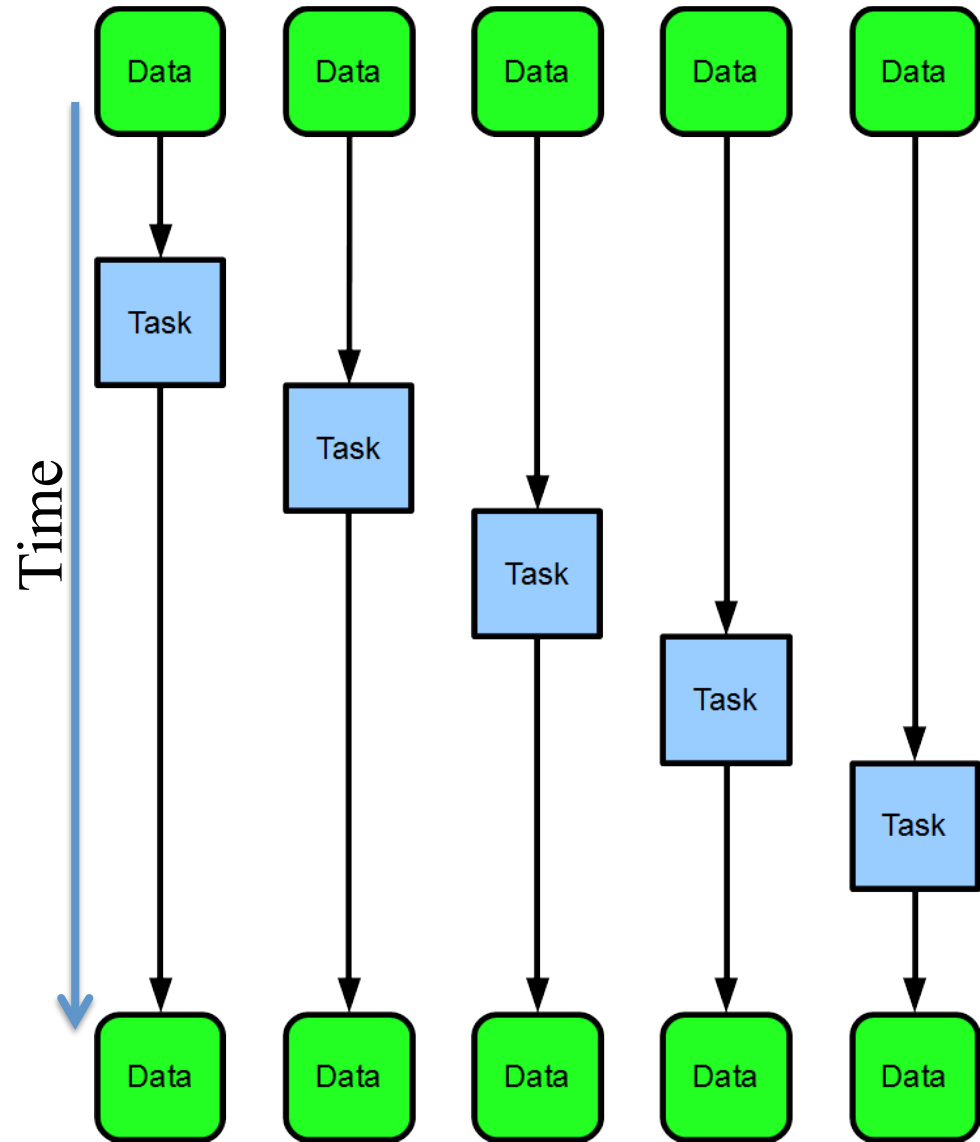
## Embarrassingly Parallel

Independence is a big win. We can run map completely in parallel.  
Significant speedups! More precisely:  $T(\infty)$  is  $O(1)$  plus implementation overhead that is  $O(\log n)$ ...so  $T(\infty) \in O(\log n)$ .




# *Sequential Map*

```
for(int n=0;  
    n< array.length;  
    ++n) {  
    process(array[n]);  
}
```



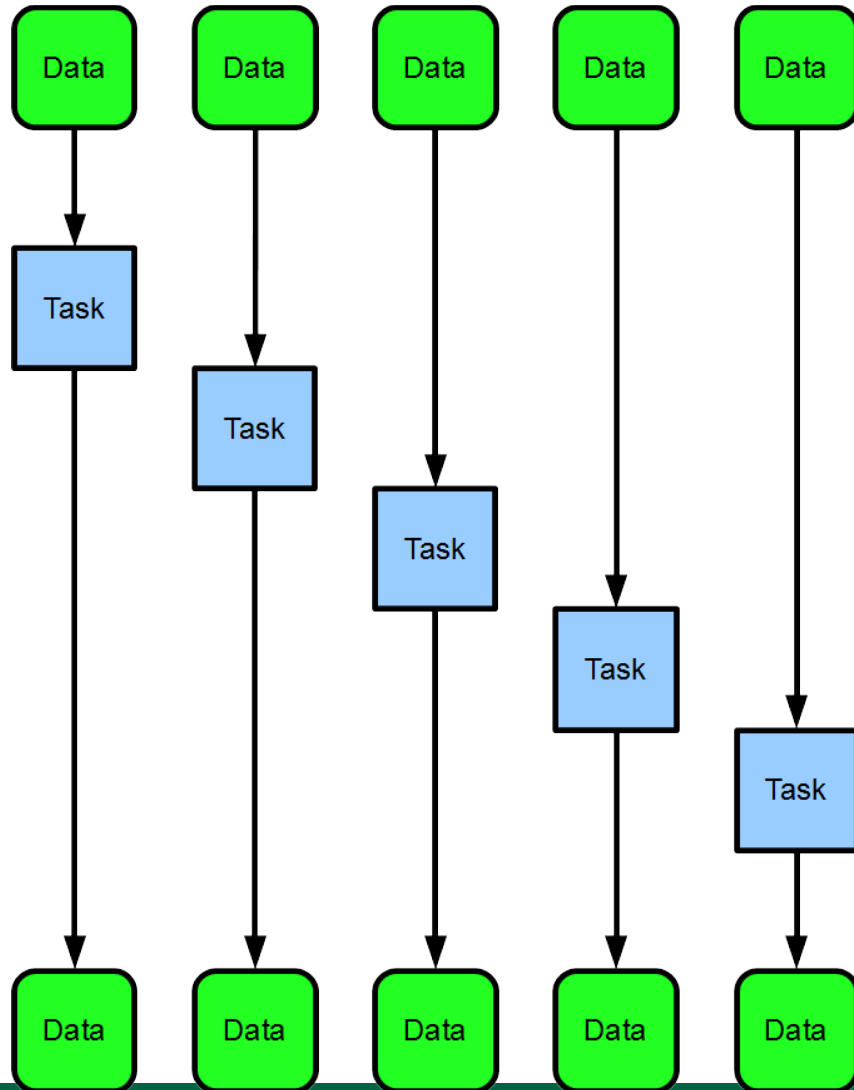
# *Parallel Map*

```
parallel_for_each(  
  x in array){  
    process(x);  
  }
```




# Comparing Maps

## Serial Map




## Parallel Map




# Comparing Maps

## Serial Map



## Parallel Map



## Speedup

The space here is speedup. With the parallel map, our program finished execution early, while the serial map is still running.

# *Independence*

- ❑ The key to (embarrassing) parallelism is independence

# *Independence*

- ❑ The key to (embarrassing) parallelism is independence

**Warning: No shared state!**

Map function should be “pure” (or “pure-ish”) and should not modify shared states

- ❑ Modifying shared state breaks perfect independence
- ❑ Results of accidentally violating independence:
  - non-determinism
  - data-races
  - undefined behavior
  - segfaults

# *Implementation and API*

- ❑ OpenMP and CilkPlus contain a parallel ***for*** language construct
- ❑ Map is a mode of use of parallel ***for***
- ❑ TBB uses **higher order functions** with lambda expressions/“functors”
- ❑ Some languages (CilkPlus, Matlab, Fortran) provide **array notation** which makes some maps more concise

## Array Notation

$A[:]$  =  $A[:] * 5$ ;

is CilkPlus array notation for “multiply every element in  $A$  by 5”

# *Unary Maps*

## Unary Maps

So far we have only dealt with mapping over a single collection...



# Map with 1 Input, 1 Output

	0	1	2	3	4	5	6	7	8	9	10	11
x	3	7	0	1	4	0	0	4	5	3	1	0
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
result	6	14	0	2	8	0	0	8	10	6	2	0

```
int oneToOne ( int x[11] ) {  
    return x*2;  
}
```

# *N-ary Maps*

## N-ary Maps

But, sometimes it makes sense to map over multiple collections at once...

# Map with 2 Inputs, 1 Output

	0	1	2	3	4	5	6	7	8	9	10	11
x	3	7	0	1	4	0	0	4	5	3	1	0
y	2	4	2	1	8	3	9	5	5	1	2	1
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
result	5	11	2	2	12	3	9	9	10	4	3	1

```
int twoToOne ( int x[11], int y[11] ) {  
    return x+y;  
}
```

# *Table of Contents*

## □ Map

## □ Optimizations

- Sequences of Maps
- Code Fusion
- Cache Fusion

## □ Related Patterns


## □ Example Implementation: Scaled Vector Addition (SAXPY)

- Problem Description
- Various Implementations

# *Table of Contents*

- ❑ Map
- ❑ Optimizations
  - Sequences of Maps
  - Code Fusion
  - Cache Fusion
- ❑ Related Patterns
- ❑ Example Implementation: Scaled Vector Addition (SAXPY)
  - Problem Description
  - Various Implementations

# Sequences of Maps




- ❑ Often several map operations occur in sequence
  - Vector math consists of many small operations such as additions and multiplications applied as maps
- ❑ A naïve implementation may write each intermediate result to memory, wasting memory BW and likely overwhelming the cache

# *Table of Contents*

- ❑ Map
- ❑ Optimizations
  - Sequences of Maps
  - Code Fusion
  - Cache Fusion
- ❑ Related Patterns
- ❑ Example Implementation: Scaled Vector Addition (SAXPY)
  - Problem Description
  - Various Implementations

# Code Fusion




- ❑ Can sometimes “fuse” together the operations to perform them at once
- ❑ Adds arithmetic intensity, reduces memory/cache usage
- ❑ Ideally, operations can be performed using registers alone



# *Table of Contents*

- ❑ Map
- ❑ Optimizations
  - Sequences of Maps
  - Code Fusion
  - Cache Fusion
- ❑ Related Patterns
- ❑ Example Implementation: Scaled Vector Addition (SAXPY)
  - Problem Description
  - Various Implementations

# Cache Fusion



- ❑ Sometimes impractical to fuse together the map operations
- ❑ Can instead break the work into blocks, giving each CPU one block at a time
- ❑ Hopefully, operations use cache alone

# *Table of Contents*

- ❑ Map
- ❑ Optimizations
  - Sequences of Maps
  - Code Fusion
  - Cache Fusion
- ❑ Related Patterns
- ❑ Example Implementation: Scaled Vector Addition (SAXPY)
  - Problem Description
  - Various Implementations

# *Overview*


Three patterns related to map are discussed here:

- Stencil
- Workpile
- Divide-and-Conquer

They will be discussed more in detail in a later lecture.

# Stencil


- ❑ Each instance of the map function accesses neighbors of its input, offset from its usual input
- ❑ Common in imaging and PDE solvers



# *Workpile*

- ❑ Work items can be added to the map while it is in progress, from inside map function instances
- ❑ Work grows and is consumed by the map
- ❑ Workpile pattern terminates when no more work is available

# *Divide-and-Conquer*



- Applies if a problem can be divided into smaller subproblems recursively until a base case is reached that can be solved serially

# *Table of Contents*

- ❑ Map
- ❑ Optimizations
  - Sequences of Maps
  - Code Fusion
  - Cache Fusion
- ❑ Related Patterns
- ❑ Example Implementation: Scaled Vector Addition (SAXPY)
  - Problem Description
  - Various Implementations



# *Table of Contents*

- ❑ Map
- ❑ Optimizations
  - Sequences of Maps
  - Code Fusion
  - Cache Fusion
- ❑ Related Patterns
- ❑ Example Implementation: Scaled Vector Addition (SAXPY)
  - Problem Description
  - Various Implementations

# *Problem Description*

- ❑  $y \leftarrow ax + y$ 
  - Scales vector  $x$  by  $a$  and adds it to vector  $y$
  - Result is stored in input vector  $y$
- ❑ Comes from the BLAS (Basic Linear Algebra Subprograms) library
- ❑ **Every element in vector  $x$  and vector  $y$  are independent**

*Visual:*  $y \leftarrow ax + y$

	0	1	2	3	4	5	6	7	8	9	10	11
a * x + y	4	4	4	4	4	4	4	4	4	4	4	4
	2	4	2	1	8	3	9	5	5	1	2	1
	3	7	0	1	4	0	0	4	5	3	1	0
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
y	11	23	8	5	36	12	36	49	50	7	9	4

*Visual:*  $y \leftarrow ax + y$

	0	1	2	3	4	5	6	7	8	9	10	11
a	4	4	4	4	4	4	4	4	4	4	4	4
*												
x	2	4	2	1	8	3	9	5	5	1	2	1
+												
y	3	7	0	1	4	0	0	4	5	3	1	0
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
y	11	23	8	5	36	12	36	49	50	7	9	4

Twelve processors used  $\rightarrow$  one for each element in the vector

*Visual:*  $y \leftarrow ax + y$

	0	1	2	3	4	5	6	7	8	9	10	11
a	4	4	4	4	4	4	4	4	4	4	4	4
*	2	4	2	1	8	3	9	5	5	1	2	1
+												
y	3	7	0	1	4	0	0	4	5	3	1	0
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
y	11	23	8	5	36	12	36	49	50	7	9	4

Six processors used  $\rightarrow$  one for every two elements in the vector

*Visual:*  $y \leftarrow ax + y$

	0	1	2	3	4	5	6	7	8	9	10	11
a	4	4	4	4	4	4	4	4	4	4	4	4
*	2	4	2	1	8	3	9	5	5	1	2	1
+												
y	3	7	0	1	4	0	0	4	5	3	1	0
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
y	11	23	8	5	36	12	36	49	50	7	9	4

Two processors used  $\rightarrow$  one for every six elements in the vector

# *Table of Contents*

- ❑ Map
- ❑ Optimizations
  - Sequences of Maps
  - Code Fusion
  - Cache Fusion
- ❑ Related Patterns
- ❑ Example Implementation: Scaled Vector Addition (SAXPY)
  - Problem Description
  - Various Implementations

# ***Serial SAXPY Implementation***

```
1 void saxpy_serial(  
2     size_t n,           // the number of elements in the vectors  
3     float a,            // scale factor  
4     const float x[],    // the first input vector  
5     float y[]           // the output vector and second input vector  
6 ) {  
7     for (size_t i = 0; i < n; ++i)  
8         y[i] = a * x[i] + y[i];  
9 }
```



# ***TBB SAXPY Implementation***

```
1 void saxpy_tbb(  
2     int n,          // the number of elements in the vectors  
3     float a,        // scale factor  
4     float x[],      // the first input vector  
5     float y[]       // the output vector and second input vector  
6 ) {  
7     tbb::parallel_for(  
8         tbb::blocked_range<int>(0, n),  
9         [&](tbb::blocked_range<int> r) {  
10         for (size_t i = r.begin(); i != r.end(); ++i)  
11             y[i] = a * x[i] + y[i];  
12         }  
13     );  
14 }
```

# *Cilk Plus SAXPY Implementation*

```
1 void saxpy_cilk(  
2     int n,          // the number of elements in the vectors  
3     float a,        // scale factor  
4     float x[],      // the first input vector  
5     float y[]       // the output vector and second input vector  
6 ) {  
7     cilk_for (int i = 0; i < n; ++i)  
8         y[i] = a * x[i] + y[i];  
9 }
```

# *OpenMP SAXPY Implentation*

```
1 void saxpy_openmp(  
2     int n,          // the number of elements in the vectors  
3     float a,        // scale factor  
4     float x[],      // the first input vector  
5     float y[]       // the output vector and second input vector  
6 ) {  
7     #pragma omp parallel for  
8     for (int i = 0; i < n; ++i)  
9         y[i] = a * x[i] + y[i];  
10 }
```