



# Data Reorganization Pattern

Parallel Computing

CIS 410/510

Department of Computer and Information Science



UNIVERSITY OF OREGON

# *Table of Contents*

- Gather Pattern
  - Shifts, Unzip, Zip
- Scatter Pattern
  - Collision Rules: atomic, permutation, merge, priority
- Pack Pattern
  - Split, Unsplits, Bin
  - Fusing Map and Pack
  - Expand
- Partitioning Data
- AoS vs. SoA

# Gather: Serial Implementation

```
1 template<typename Data, typename Idx>
2 void gather(
3     size_t n, //number of elements in data collection
4     size_t m, //number of elements in index collection
5     Data a[], //input data collection (n elements)
6     Data A[], //output data collection (m elements)
7     Idx idx[] //input index collection (m elements)
8 ) {
9     for (size_t i = 0; i < m; ++i) {
10         size_t j = idx[i]; //get ith index
11         assert(0 <= j && j < n); //check array bounds
12         A[i] = a[j]; //perform random read
13     }
14 }
```

## Serial implementation of gather in pseudocode

# Gather: Serial Implementation

```
1 template<typename Data, typename Idx>
2 void gather(
3     size_t n, //number of elements in data collection
4     size_t m, //number of elements in index collection
5     Data a[], //input data collection (n elements)
6     Data A[], //output data collection (m elements)
7     Idx idx[] //input index collection (m elements)
8 ) {
9     for (size_t i = 0; i < m; ++i) {
10         size_t j = idx[i]; //get ith index
11         assert(0 <= i && i < n); //check array bounds
12         A[i] = a[j]; //perform random read
13     }
14 }
```

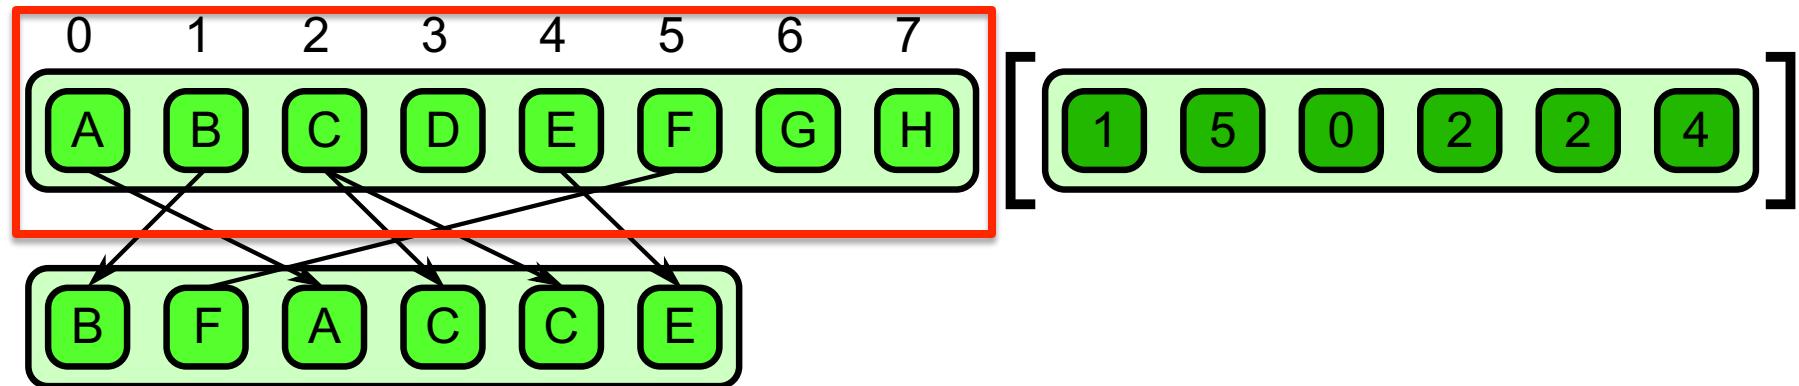
Parallelize over  
for loop to  
perform random  
read

## Serial implementation of gather in pseudocode

# *Gather: Defined*

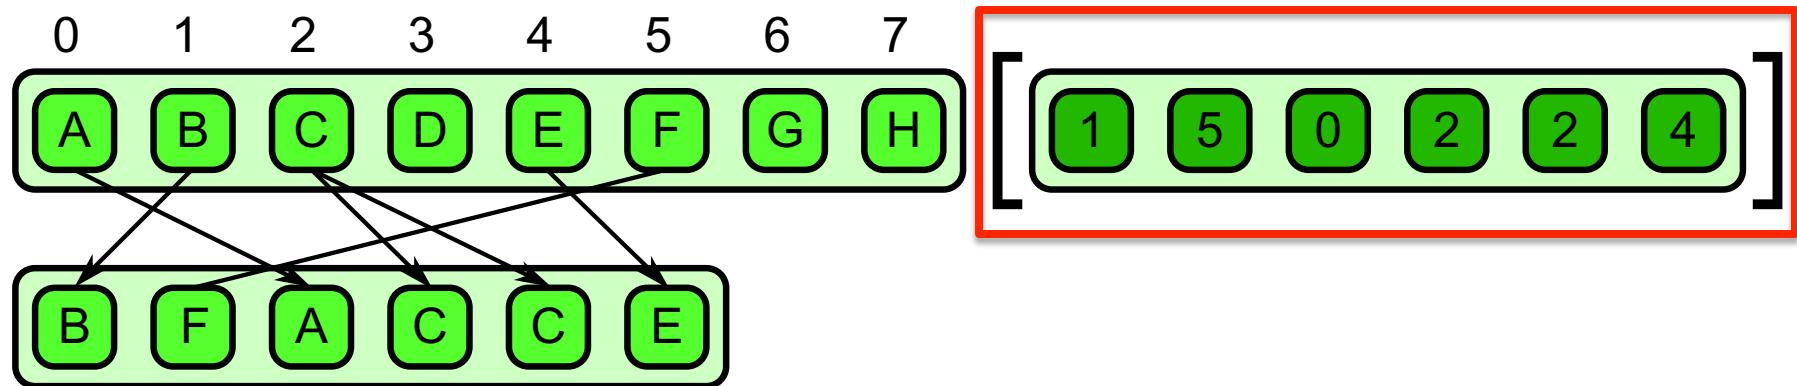
- Results from the combination of a map with a random read

# *Gather: Defined*



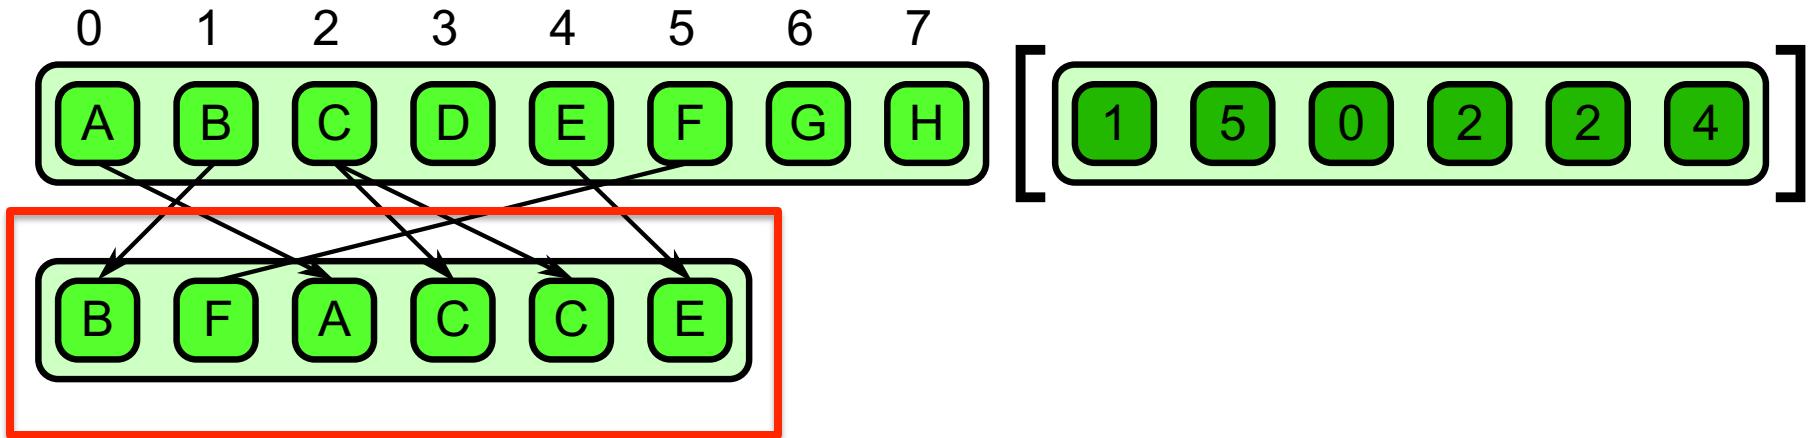
Given a **collection of locations** (address or array indices)

# Gather: Defined



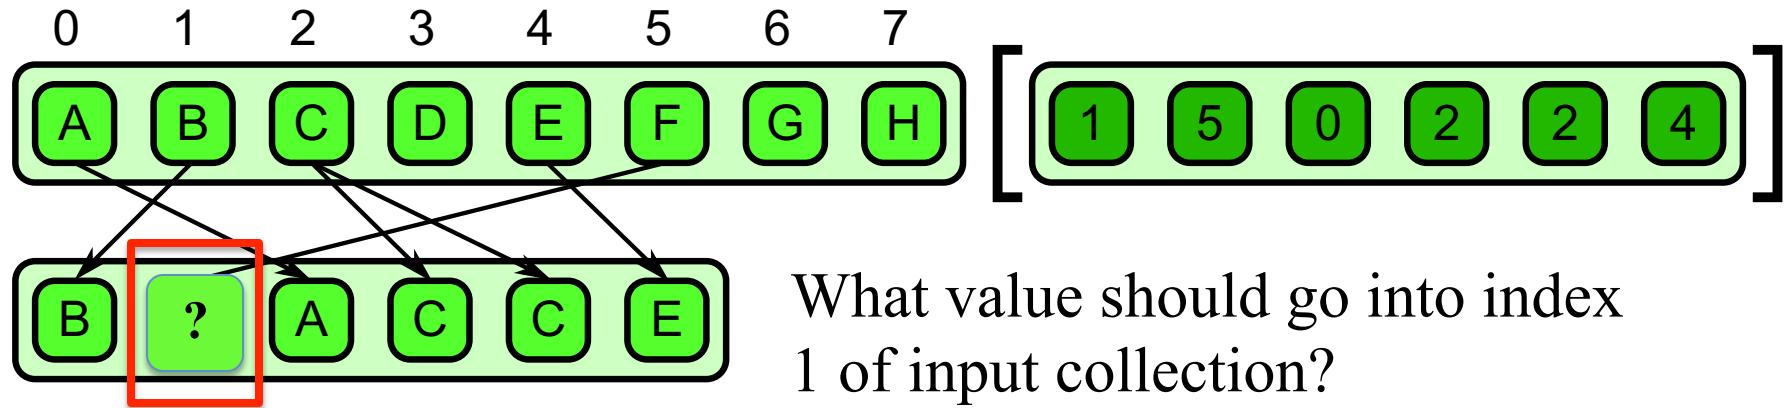
Given a collection of locations (address or array indices) and a **source array**

# *Gather: Defined*



Given a collection of locations (address or array indices) and a source array, gather collects all the data from the source array at the given locations and places them into an **output collection**

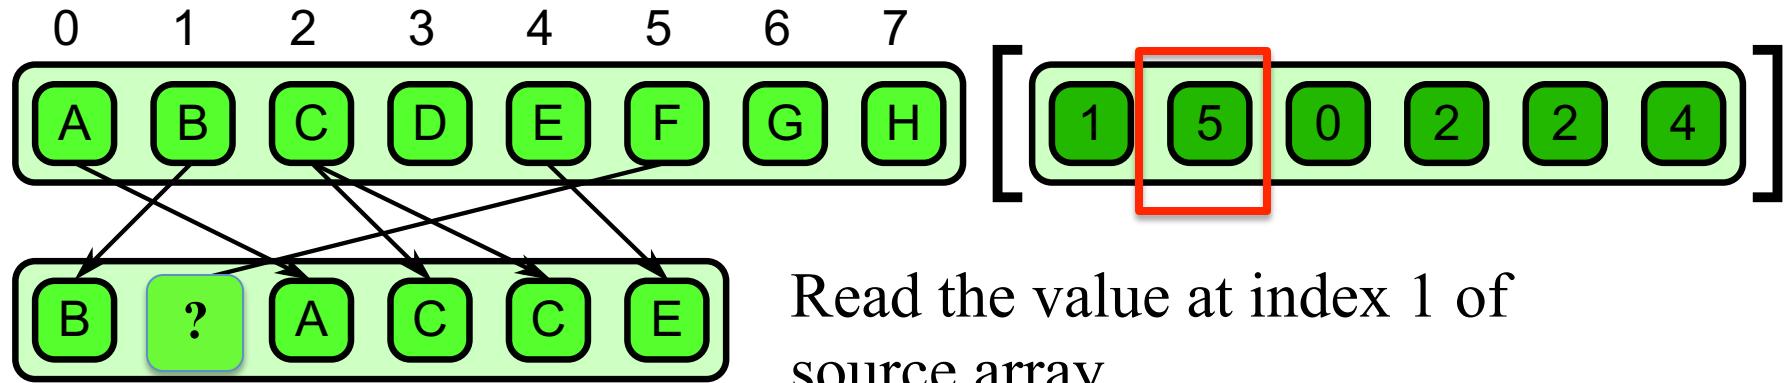
# Gather: Defined



What value should go into index 1 of input collection?

Given a collection of locations (address or array indices) and a source array, gather collects all the data from the source array at the given locations and places them into an **output collection**

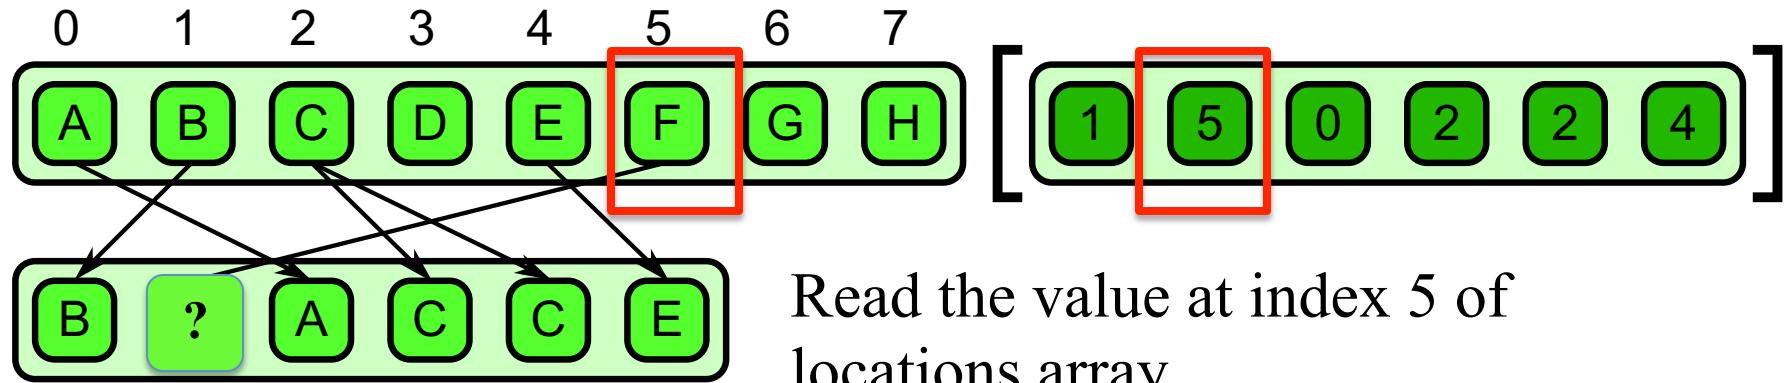
# Gather: Defined



Read the value at index 1 of source array.

Given a collection of locations (address or array indices) and a source array, gather collects all the data from the source array at the given locations and places them into an **output collection**

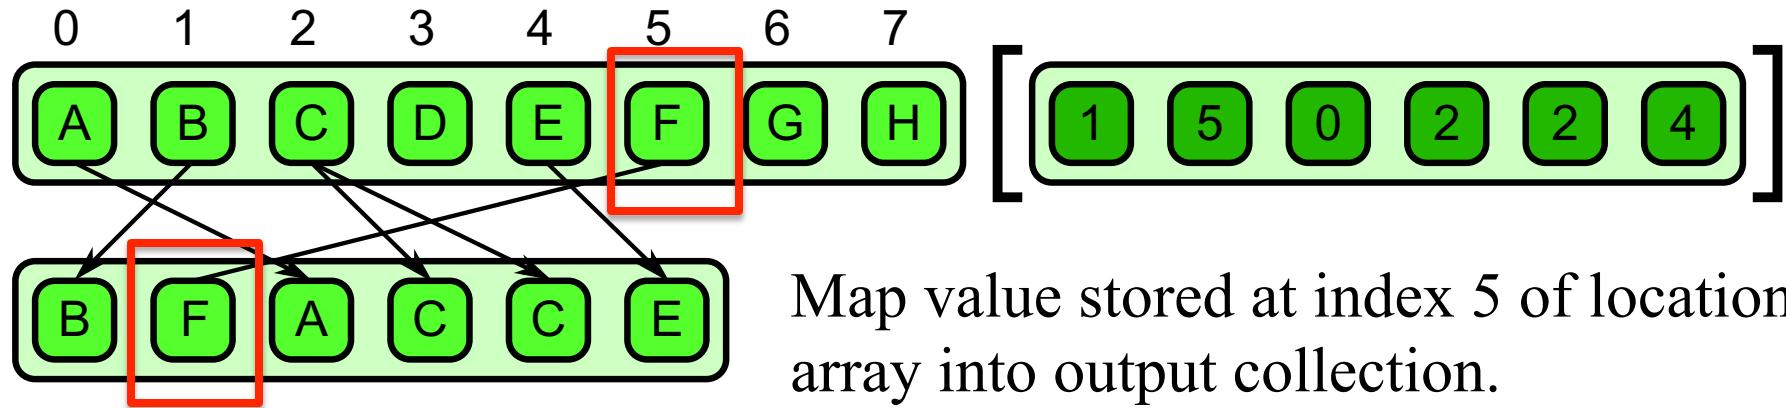
# Gather: Defined



Read the value at index 5 of locations array.

Given a collection of locations (address or array indices) and a source array, gather collects all the data from the source array at the given locations and places them into an **output collection**

# Gather: Defined



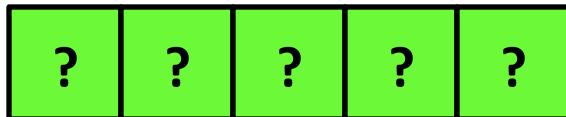
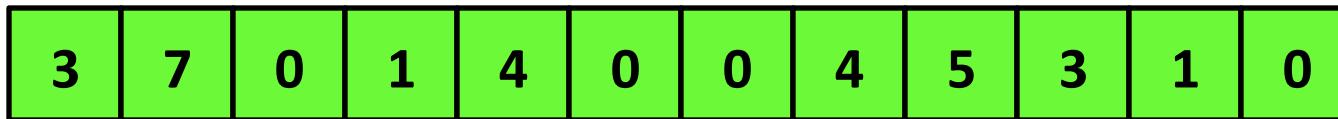
Map value stored at index 5 of locations array into output collection.

Given a collection of locations (address or array indices) and a source array, gather collects all the data from the source array at the given locations and places them into an **output collection**

# *Quiz 1*

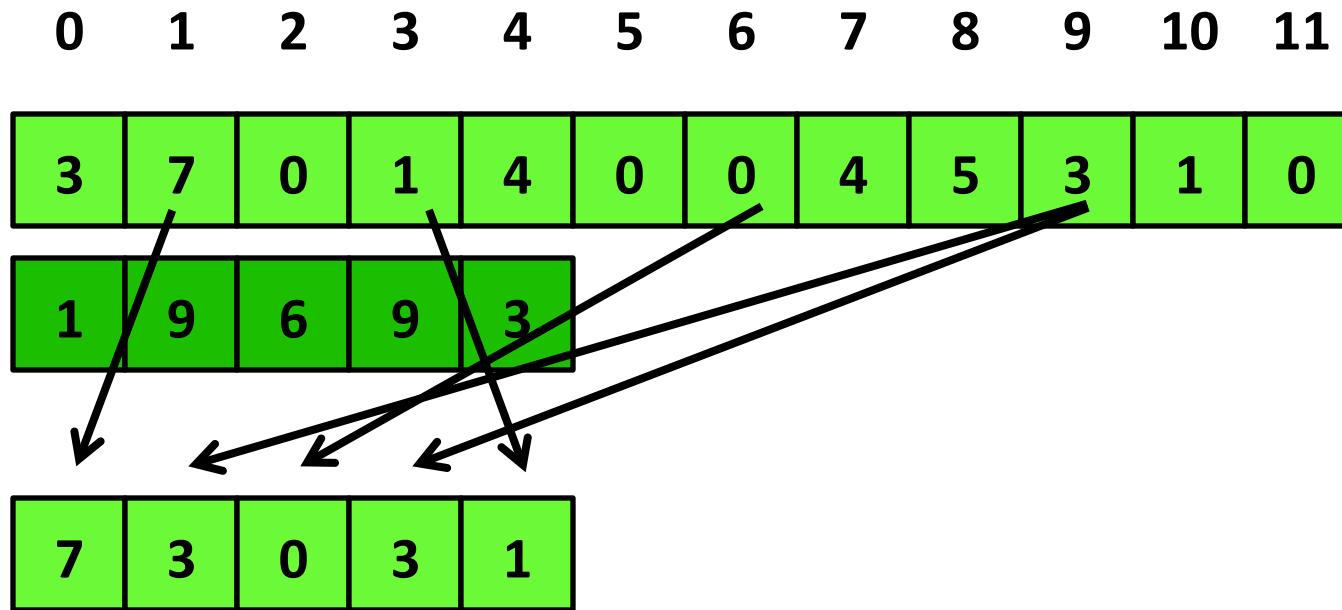
Given the following locations and source array, use a gather to determine what values should go into the output collection:

0    1    2    3    4    5    6    7    8    9    10    11

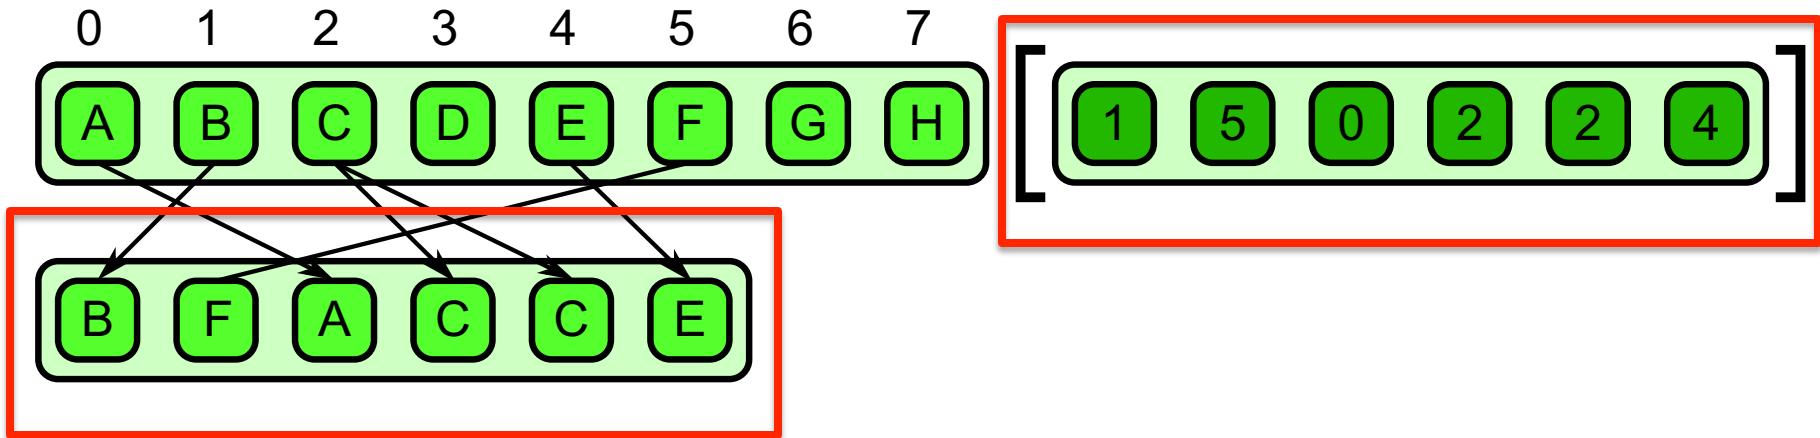


# Quiz 1

Given the following locations and source array, use a gather to determine what values should go into the output collection:

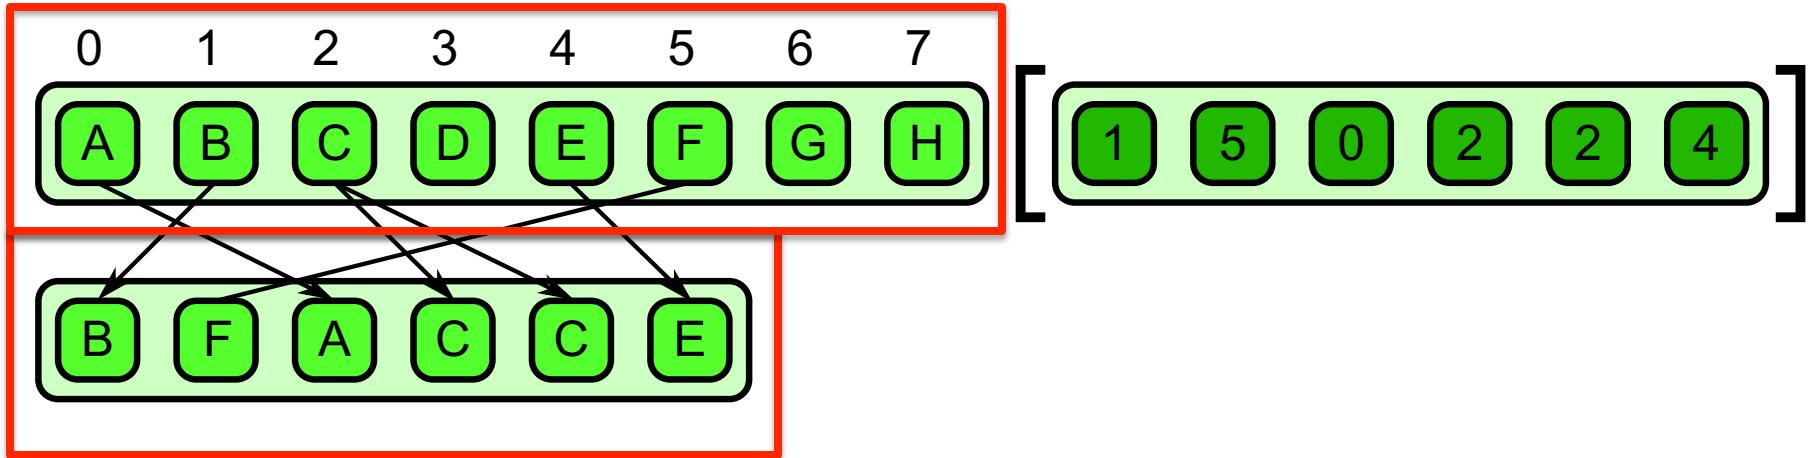


# Gather: Array Size



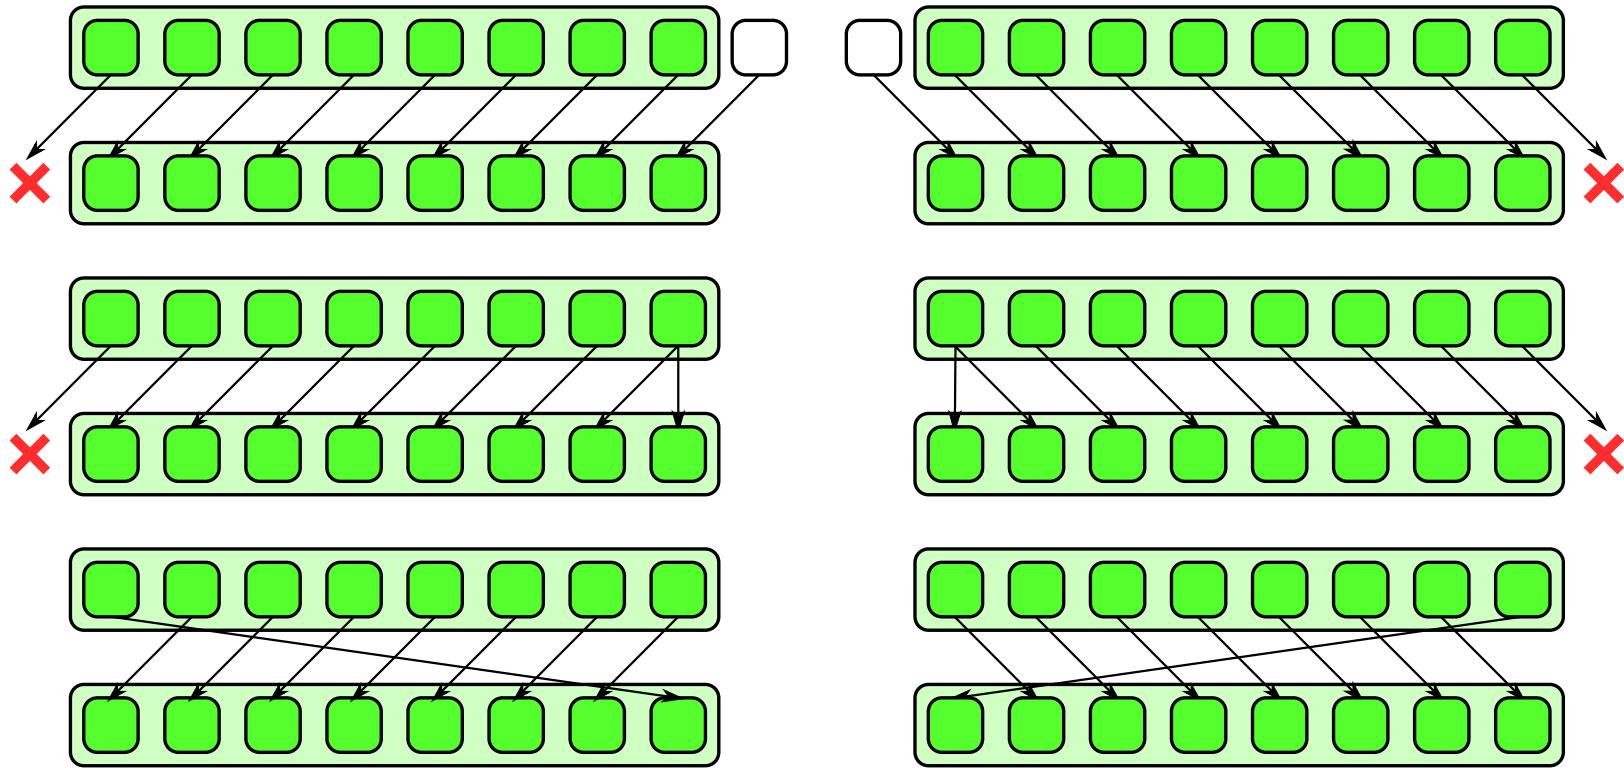
- Output data collection has the same number of elements as the number of indices in the index collection

# Gather: Array Size



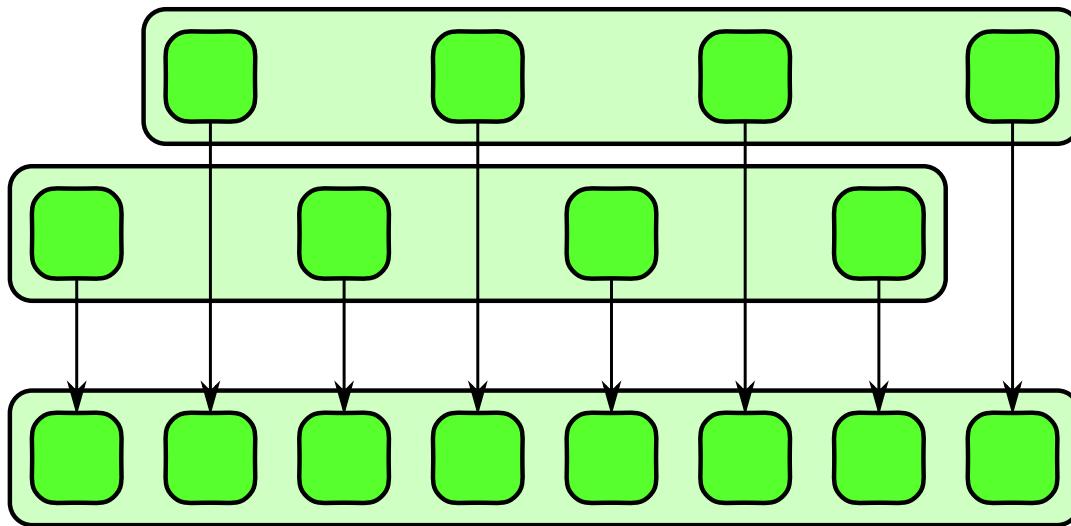
- Output data collection has the same number of elements as the number of indices in the index collection
- Elements of the output collection are the same type as the input data collection

# *Special Case of Gather: Shifts*



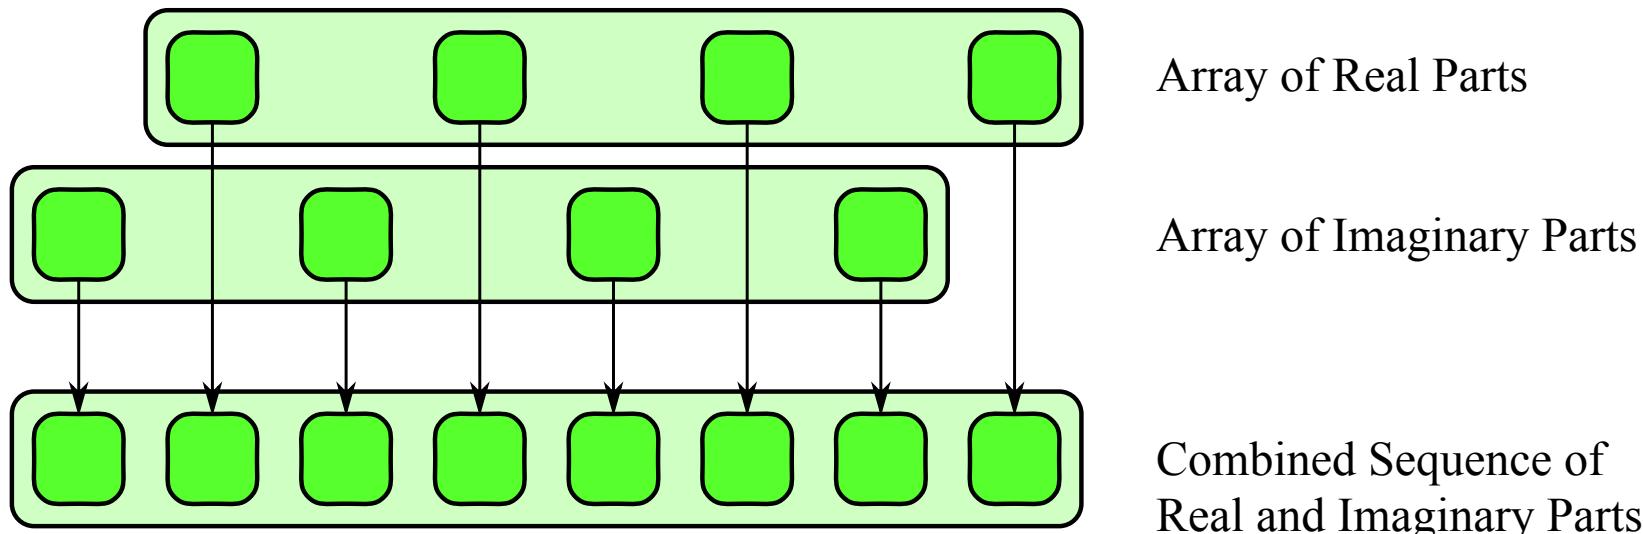
- Moves data to the left or right in memory
- Data accesses are offset by fixed distances

# *Special Case of Gather: Zip*



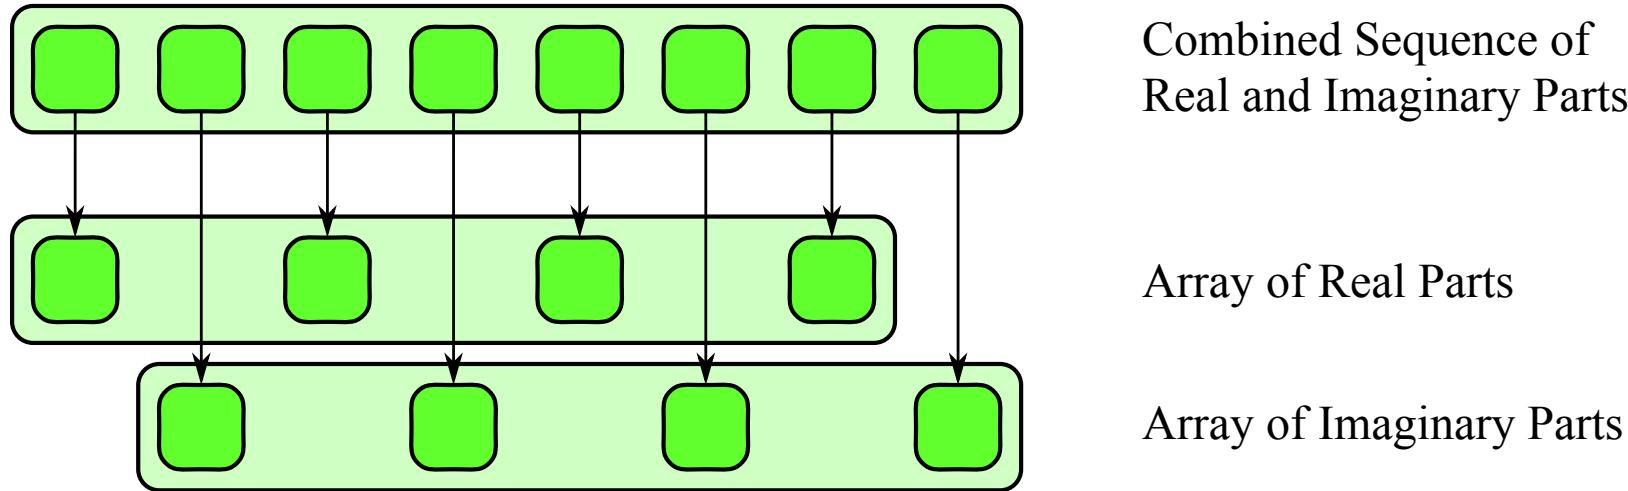
- Interleaves data

# *Special Case of Gather: Zip*



- Interleaves data
- Ex: Given two separate arrays of real parts and imaginary parts, use zip to combine them into a sequence of real and imaginary pairs.

# *Special Case of Gather: Unzip*



- Reverses a zip, extracting subarrays at certain offsets and strides from an input array
- Ex: Given a sequence of complex numbers organized as pairs, extract real and imaginary parts into separate arrays

# *Gather vs. Scatter*

## Gather

- Combination of map with random **reads**
- Read locations provided as input

## Scatter

- Combination of map with random **writes**
- Write locations provided as input
- Race conditions

# *Scatter: Serial Implementation*

```
1 template<typename Data, typename Idx>
2 void scatter(
3     size_t n, //number of elements in output data collection
4     size_t m, //number of elements in input data and index collection
5     Data a[], //input data collection (m elements)
6     Data A[], //output data collection (n elements)
7     Idx idx[] //input index collection (m elements)
8 ) {
9     for (size_t i = 0; i < m; ++i) {
10         size_t j = idx[i]; //get ith index
11         assert(0 <= j && j < n); //check output array bounds
12         A[j] = a[i]; //perform random write
13     }
14 }
```

## Serial implementation of scatter in pseudocode

# *Scatter: Serial Implementation*

```
1 template<typename Data, typename Idx>
2 void scatter(
3     size_t n, //number of elements in output data collection
4     size_t m, //number of elements in input data and index collection
5     Data a[], //input data collection (m elements)
6     Data A[], //output data collection (n elements)
7     Idx idx[] //input index collection (m elements)
8 ) {
9     for (size_t i = 0; i < m; ++i) {
10         size_t j = idx[i]; //get ith index
11         assert(0 <= j && j < n); //check output array bounds
12         A[j] = a[i]; //perform random write
13     }
14 }
```

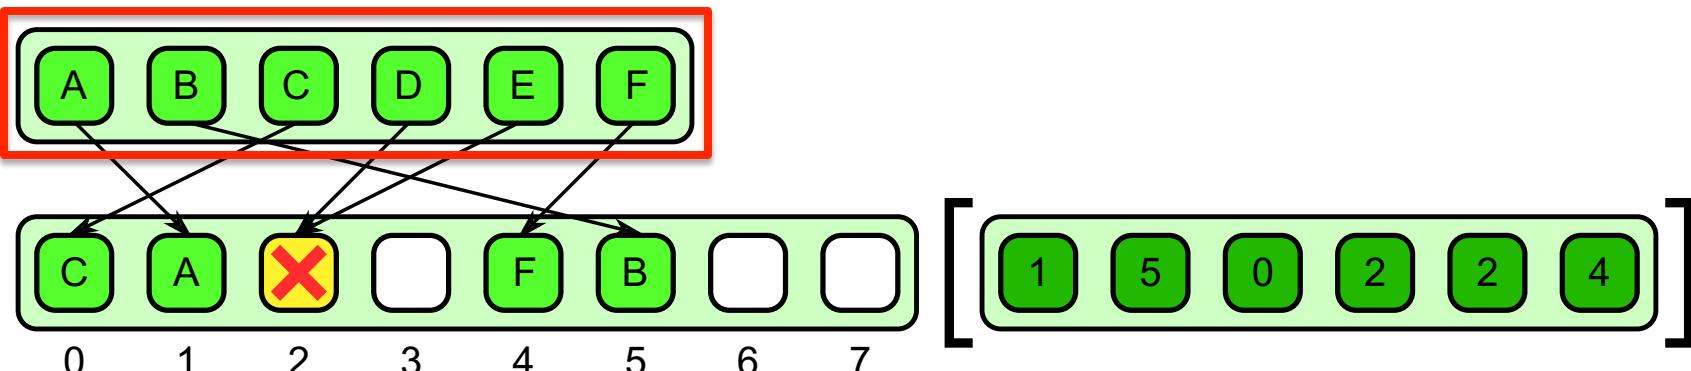
Parallelize over  
for loop to  
perform random  
write

## Serial implementation of scatter in pseudocode

# *Scatter: Defined*

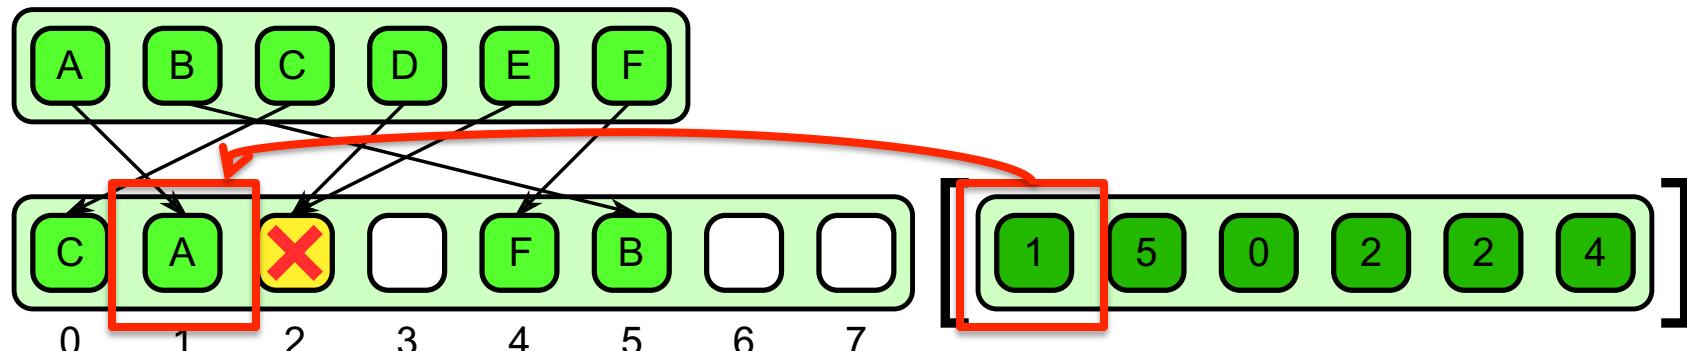
- Results from the combination of a map with a random write

# *Scatter: Defined*



Collection of **input data**

# *Scatter: Defined*



Collection of input data written in parallel to the **write locations** specified.

# *Quiz 2*

Given the following locations and source array, what values should go into the input collection:

0    1    2    3    4    5    6    7    8    9    10    11

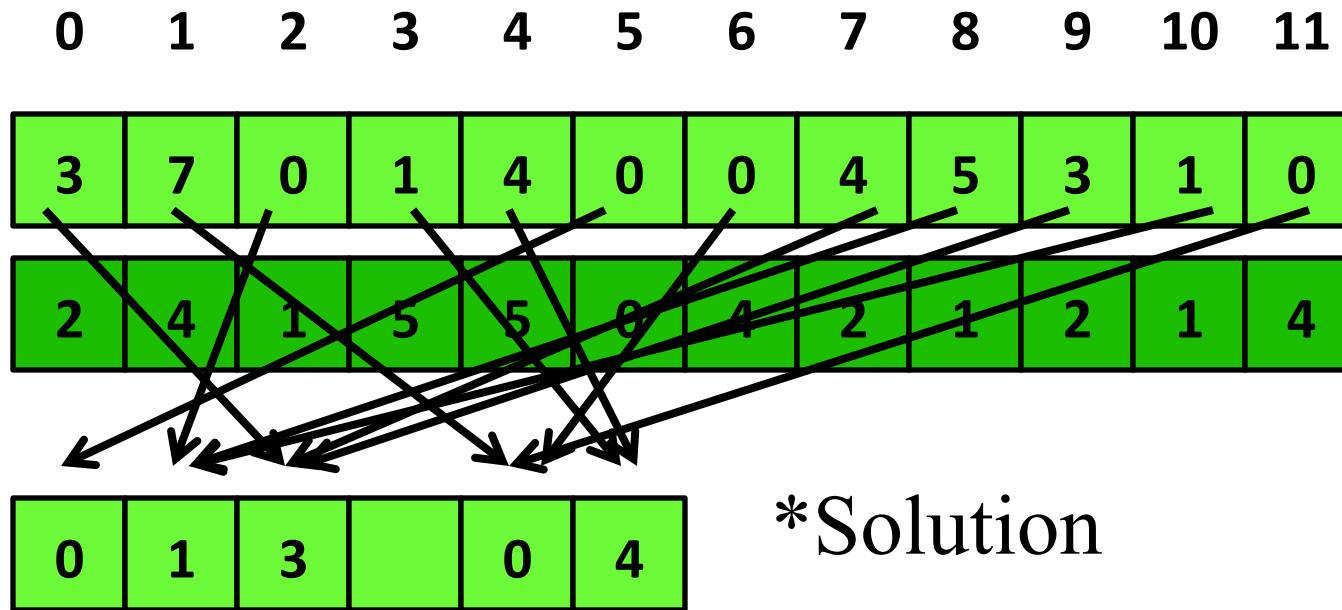
3	7	0	1	4	0	0	4	5	3	1	0
---	---	---	---	---	---	---	---	---	---	---	---

2	4	1	5	5	0	4	2	1	2	1	4
---	---	---	---	---	---	---	---	---	---	---	---

?	?	?		?	?
---	---	---	--	---	---

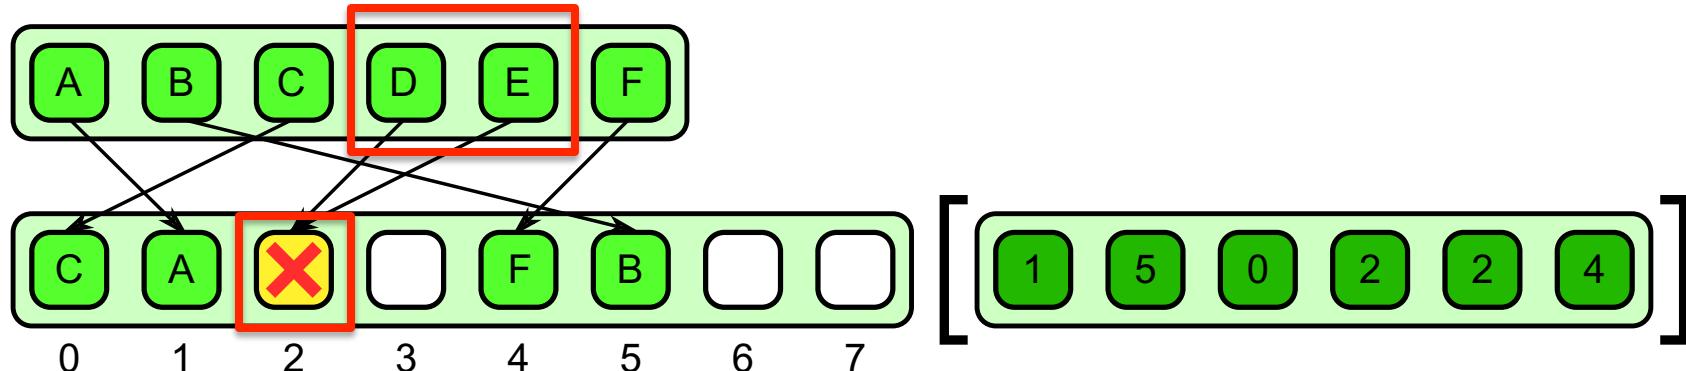
# Quiz 2

Given the following locations and source array, what values should go into the input collection:



\*Solution

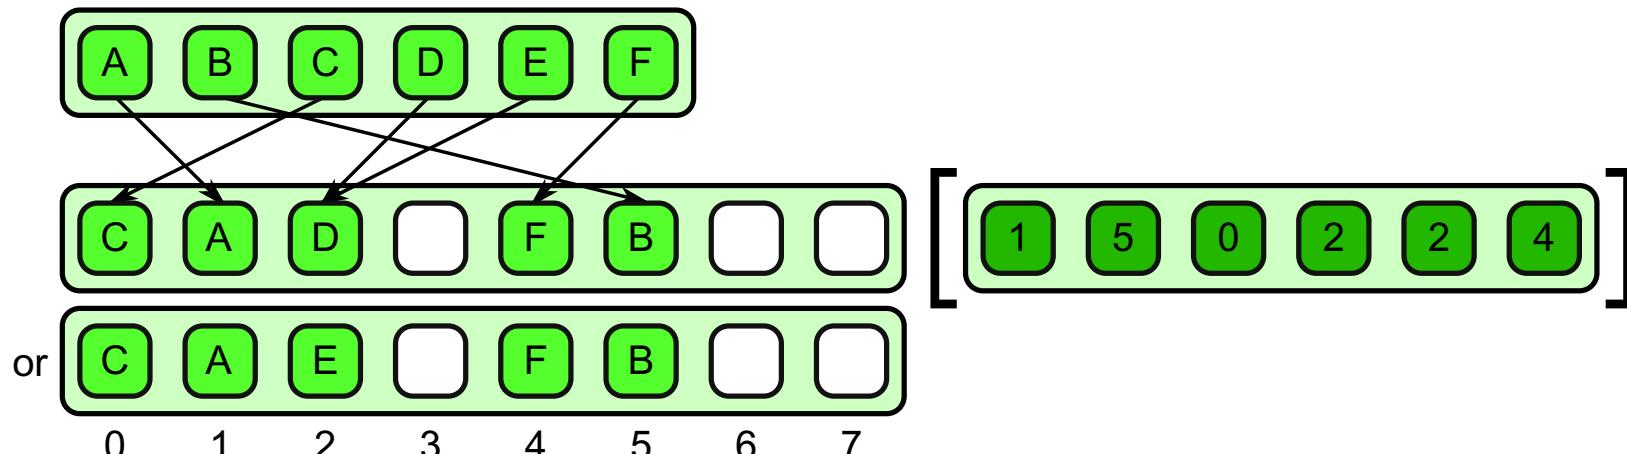
# *Scatter: Race Conditions*



Collection of input data written in parallel to the **write locations** specified.

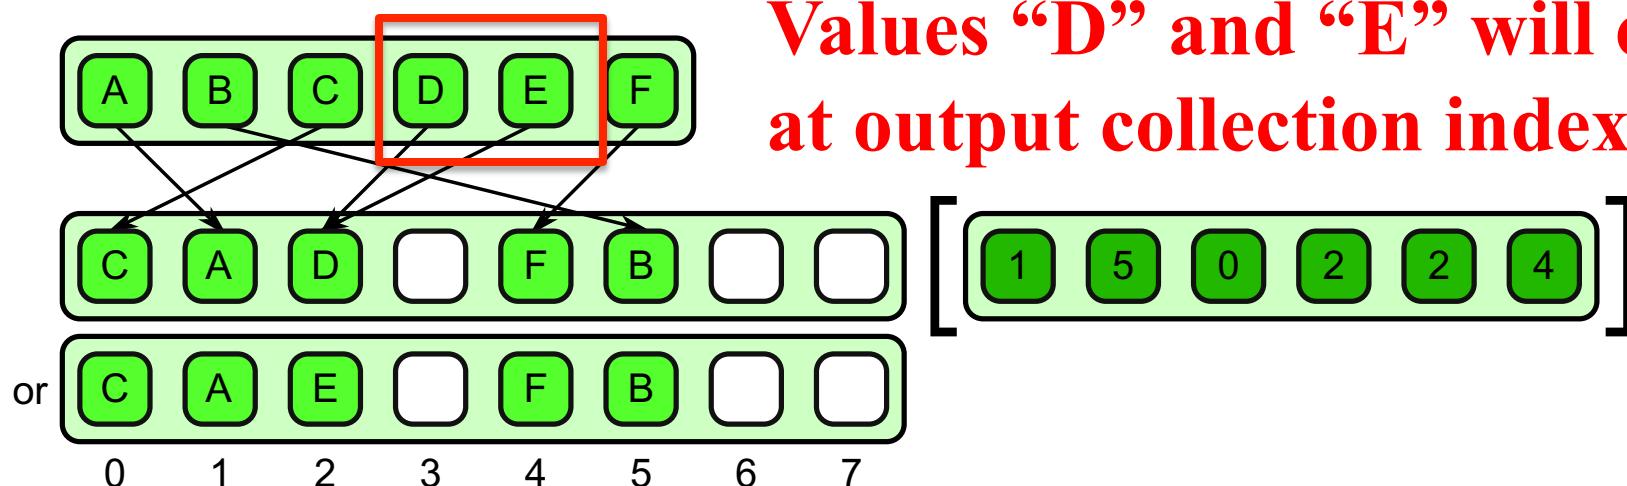
Race Condition: Two (or more) values being written to the same location in output collection.  
Unclear what the result should be.  
**Need rules to resolve collisions!**

# *Collision Resolution: Atomic Scatter*



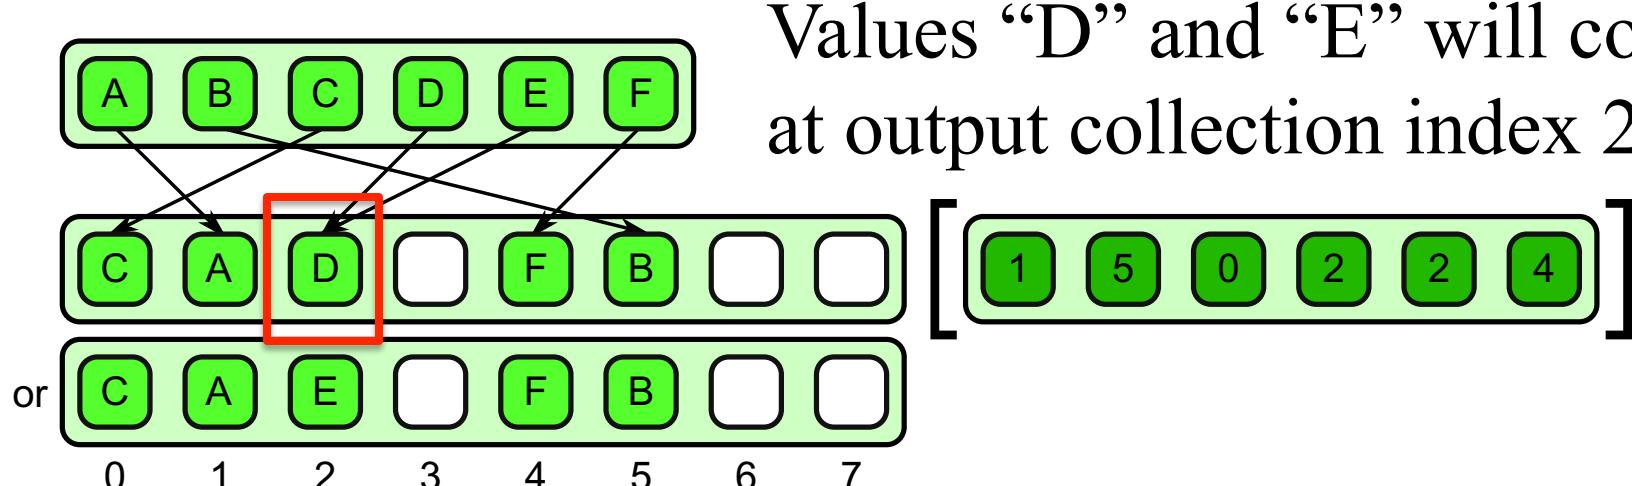
- Non-deterministic approach
- Upon collision, one and only one of the values written to a location will be written in its entirety

# *Collision Resolution: Atomic Scatter*



- Non-deterministic approach
- Upon collision, one and only one of the values written to a location will be written in its entirety

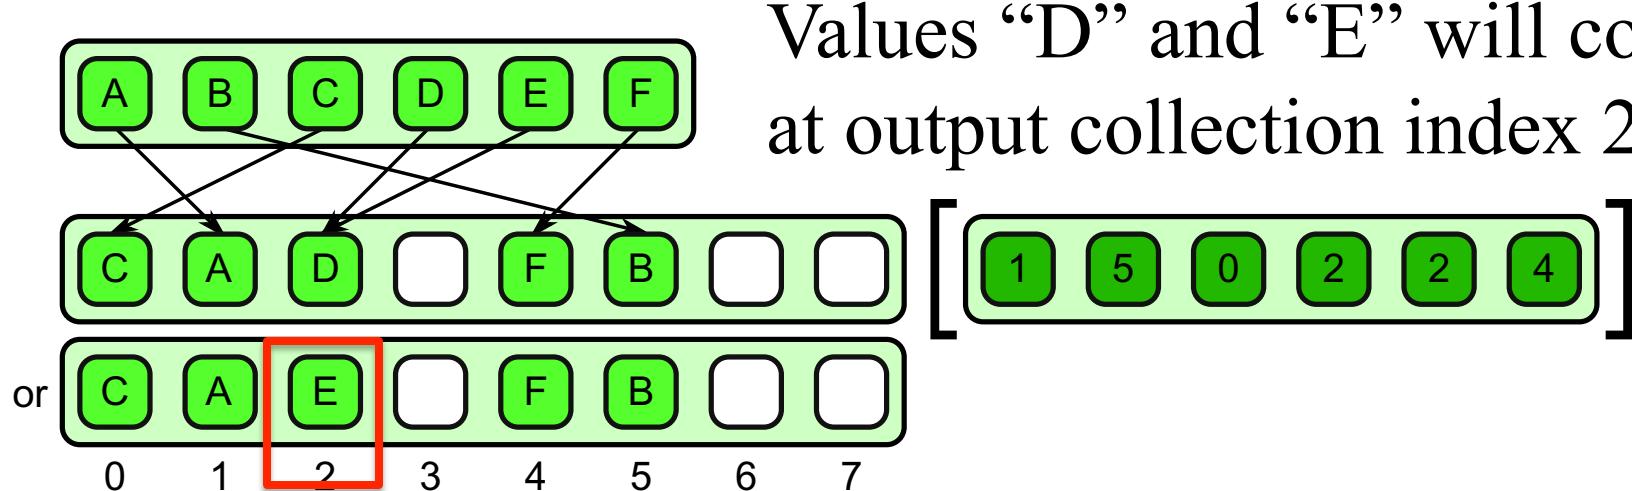
# *Collision Resolution: Atomic Scatter*



- Non-deterministic approach
- Upon collision, one and only one of the values written to a location will be written in its entirety
- No rule determines which of the input items will be retained

Either “D”...

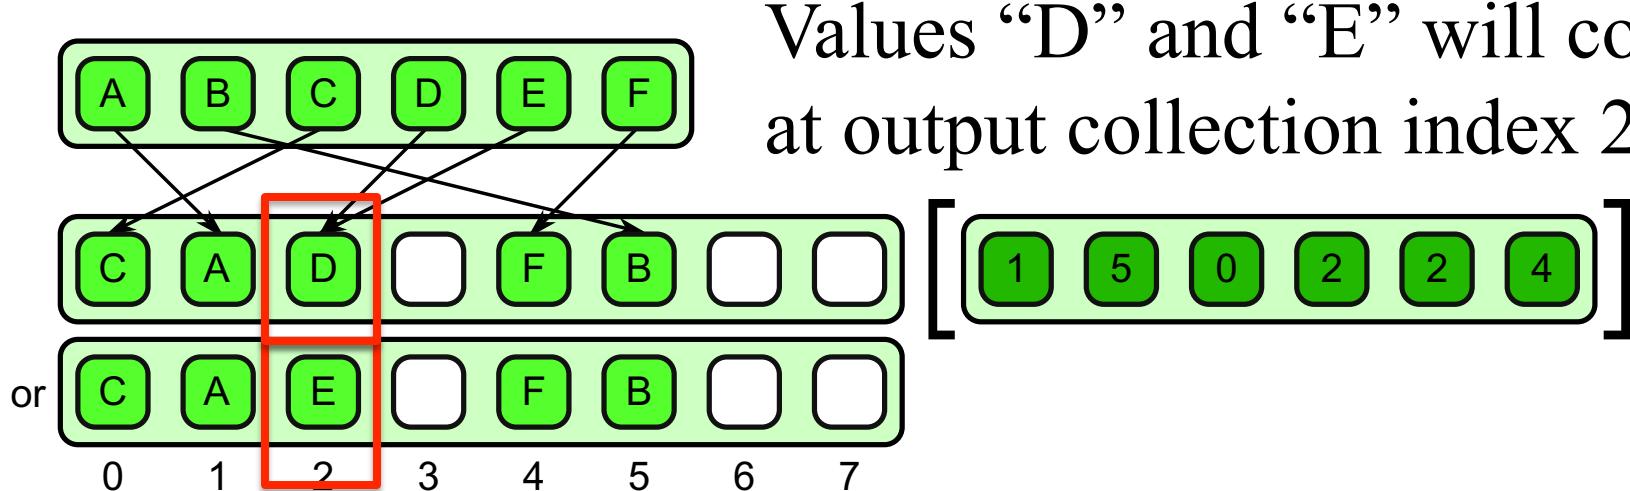
# *Collision Resolution: Atomic Scatter*



- Non-deterministic approach
- Upon collision, one and only one of the values written to a location will be written in its entirety
- No rule determines which of the input items will be retained

Either “D” or “E”...

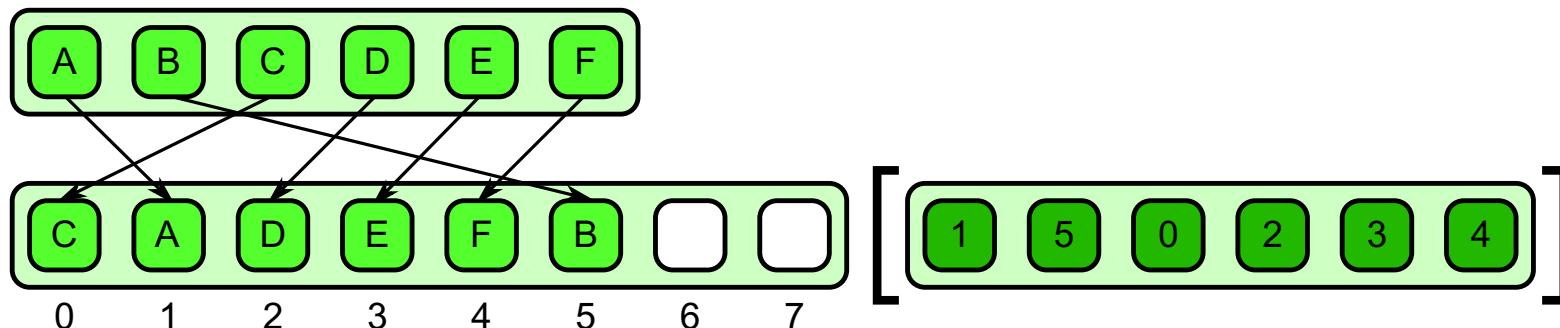
# *Collision Resolution: Atomic Scatter*



- Non-deterministic approach
- Upon collision, one and only one of the values written to a location will be written in its entirety
- No rule determines which of the input items will be retained

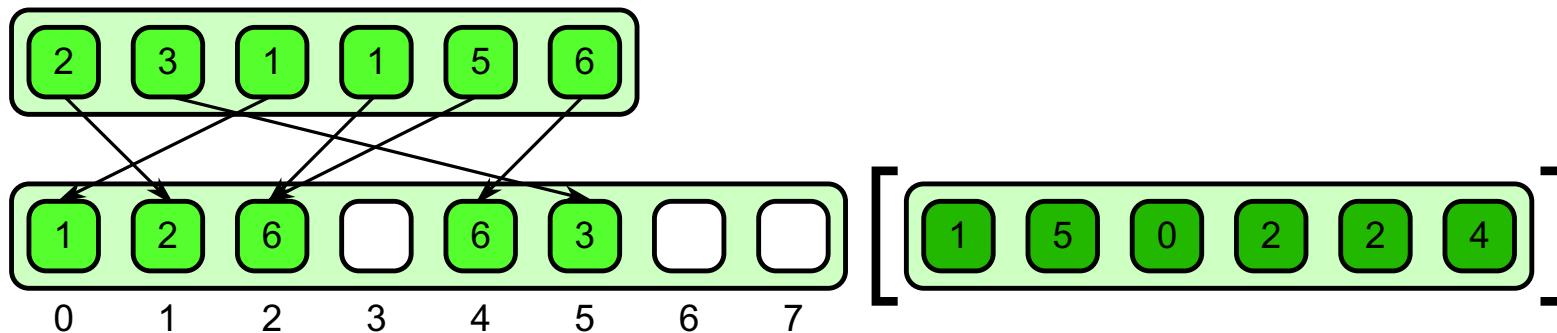
Either “D” or “E” will be written to index 2 of output collection

# *Collision Resolution: Permutation Scatter*



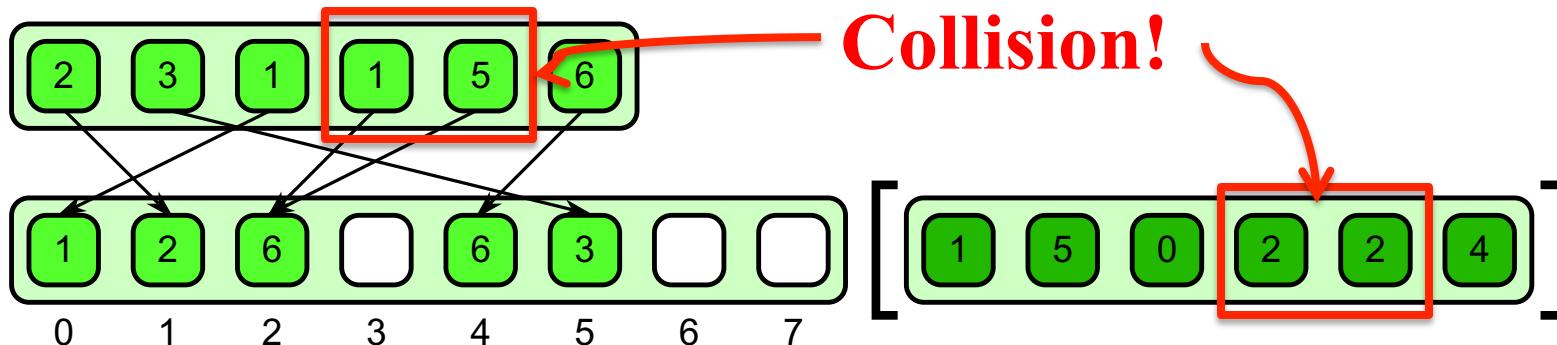
- Collisions are **illegal**
- Check for collisions in advance → turn scatter into gather
- Ex: FFT scrambling, matrix/image transpose, unpacking

# *Collision Resolution: Merge Scatter*



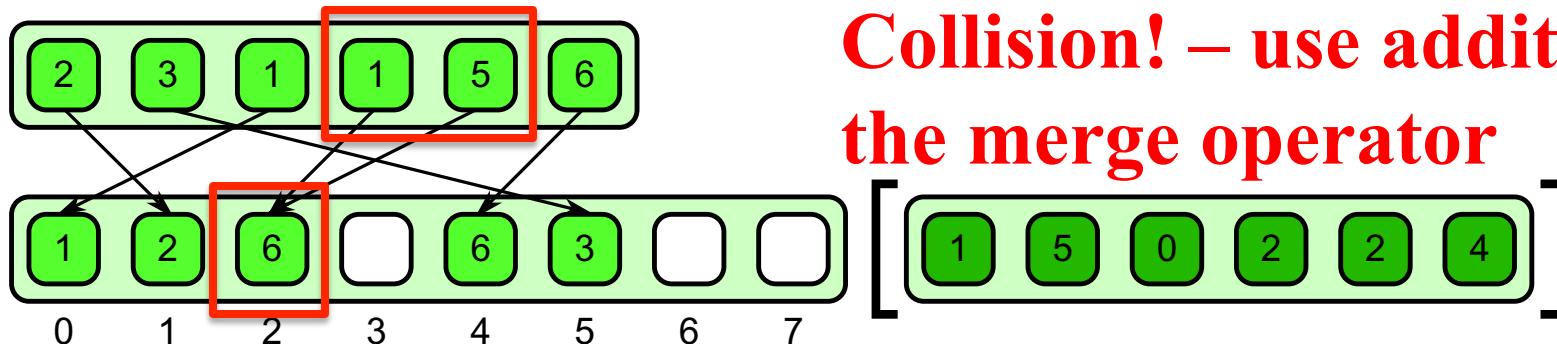
- ❑ Associative and commutative operators are provided to merge elements in case of a collision

# *Collision Resolution: Merge Scatter*



- ❑ Associative and commutative operators are provided to merge elements in case of a collision

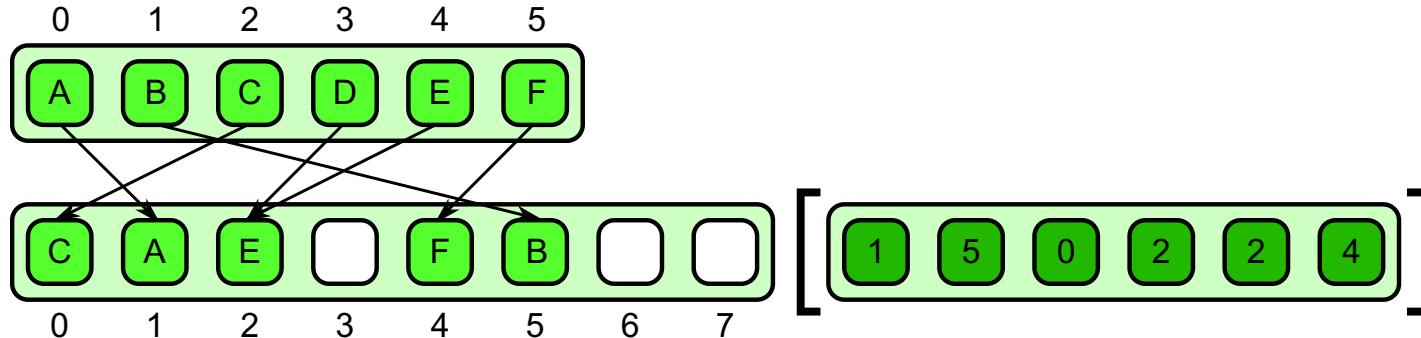
# *Collision Resolution: Merge Scatter*



**Collision! – use addition as the merge operator**

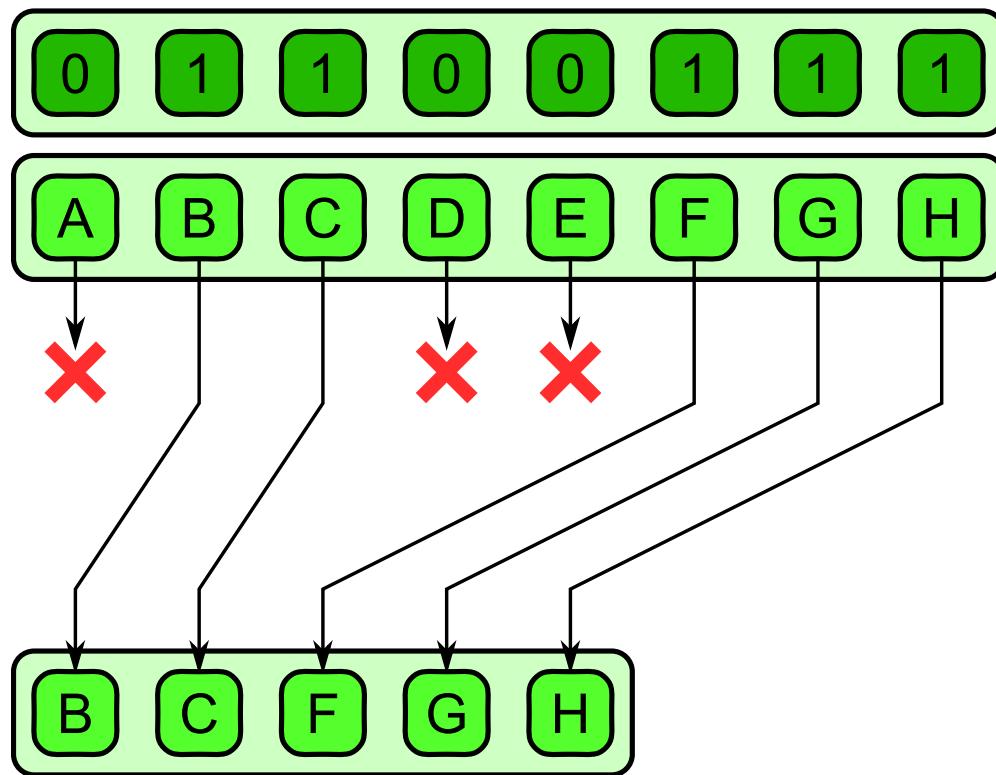
- ❑ Associative and commutative operators are provided to merge elements in case of a collision

# *Collision Resolution: Priority Scatter*



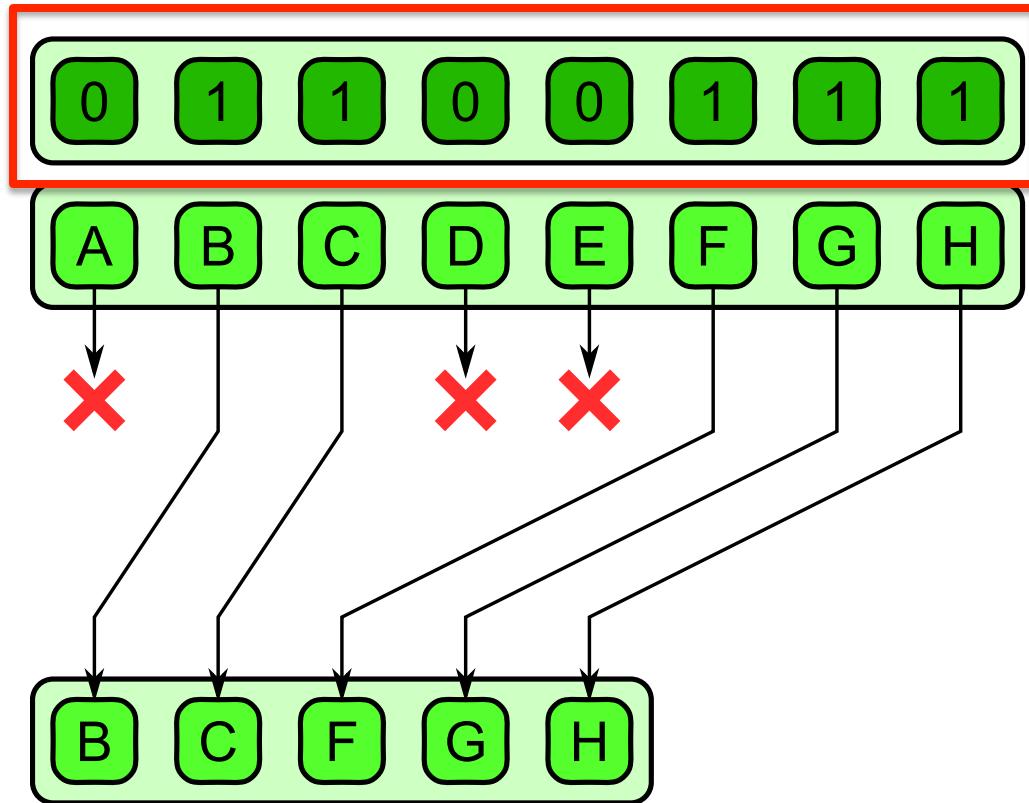
- Every element in the input array is assigned a priority based on its position
- Priority is used to decide which element is written in case of a collision
- Ex: 3D graphics rendering

# *Pack: Defined*



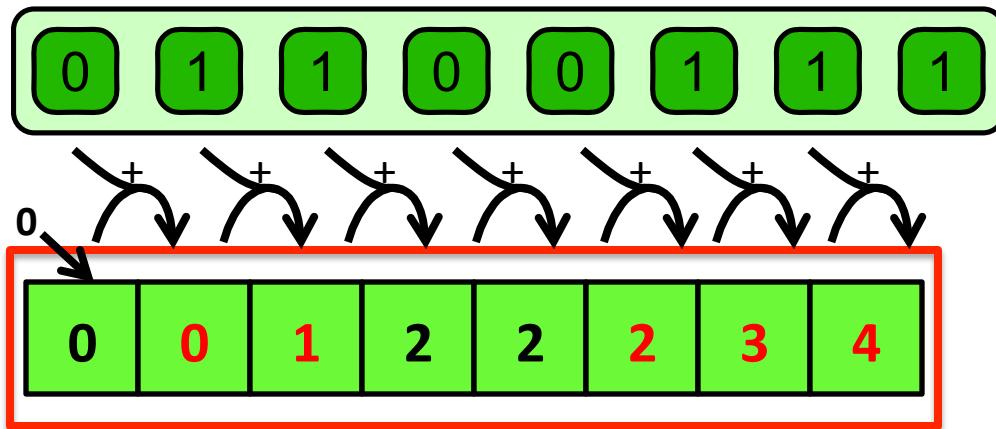
- Used to eliminate unused elements from a collection
- Retained elements are moved so they are contiguous in memory → performance improvement

# Pack Algorithm



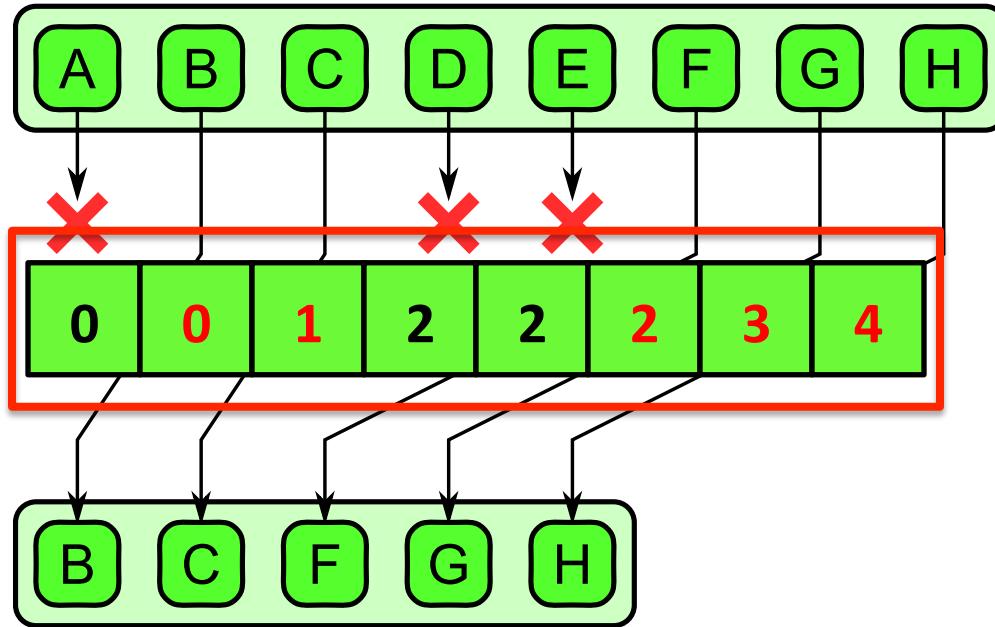
1. Convert input array of Booleans into integer 0's and 1's

# Pack Algorithm



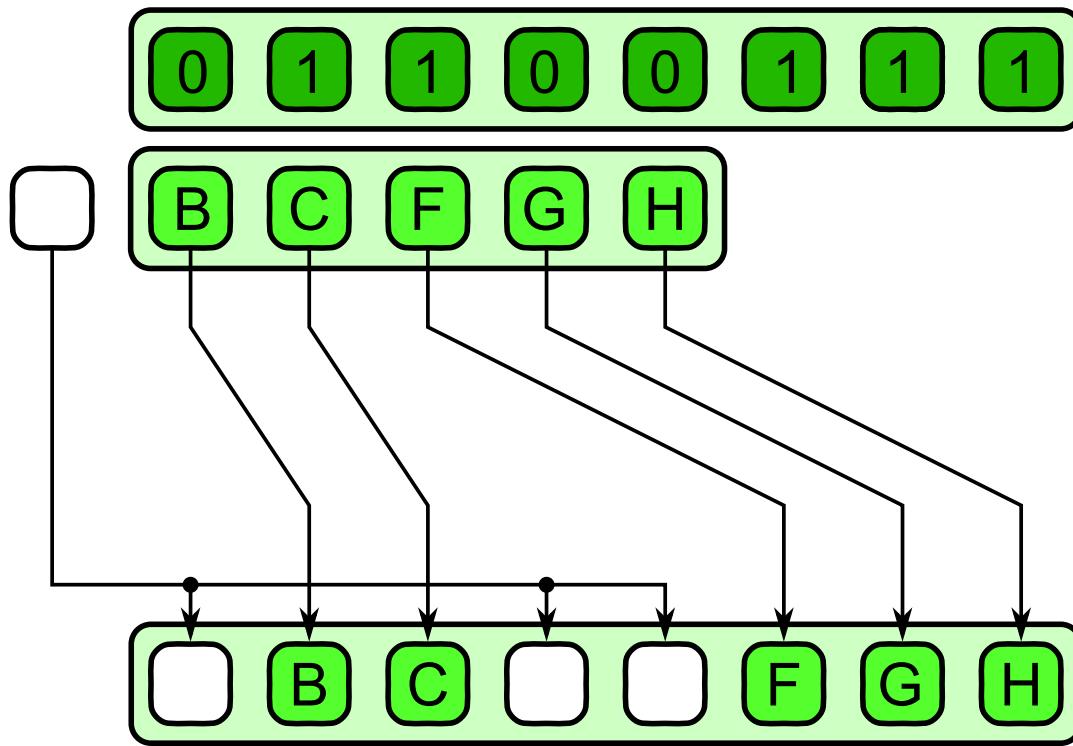
1. Convert input array of Booleans into integer 0's and 1's
2. Exclusive scan of this array with the addition operation

# Pack Algorithm



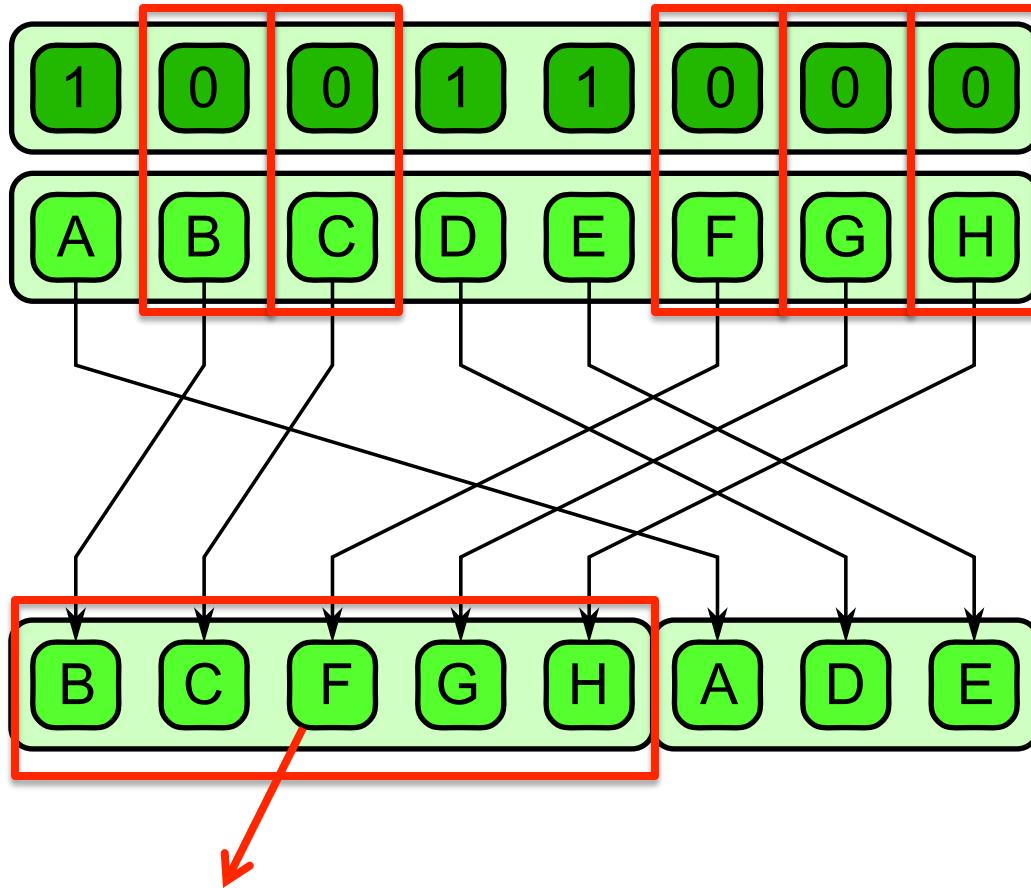
1. Convert input array of Booleans into integer 0's and 1's
2. Exclusive scan of this array with the addition operation
3. Write values to output array based on offsets

# *Unpack Defined*



- Inverse of pack operation
- Given the same data on which elements were kept and which were discarded, can place elements back in their original locations

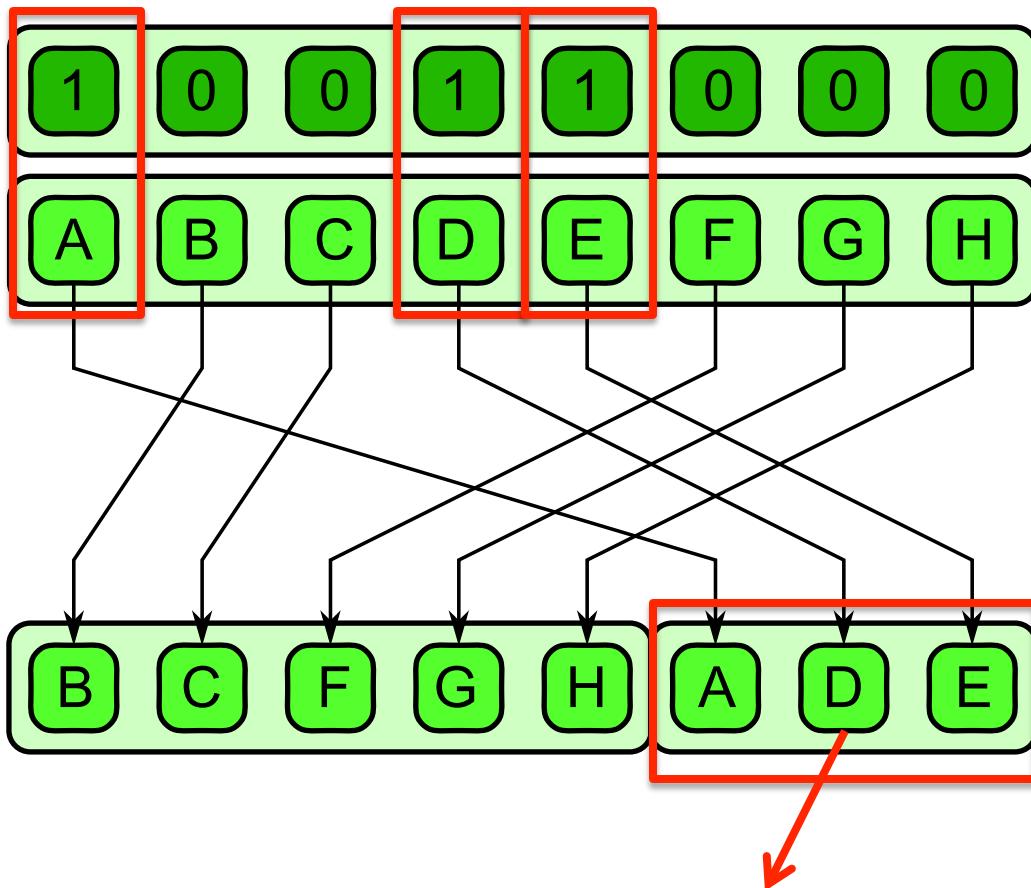
# *Generalization of Pack: Split*



**Upper half of output collection: values equal to 0**

- Generalization of pack pattern
- Elements are moved to upper or lower half of output collection based on some state
- Does not lose information like pack

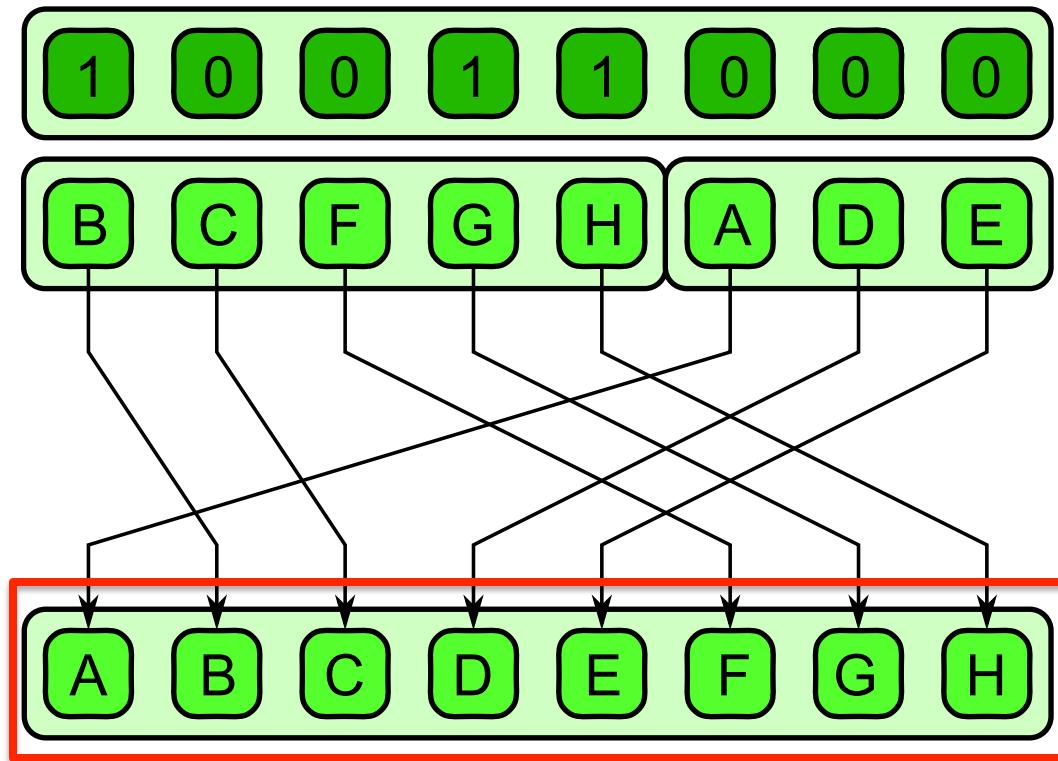
# *Generalization of Pack: Split*



**Lower half of output collection: values equal to 1**

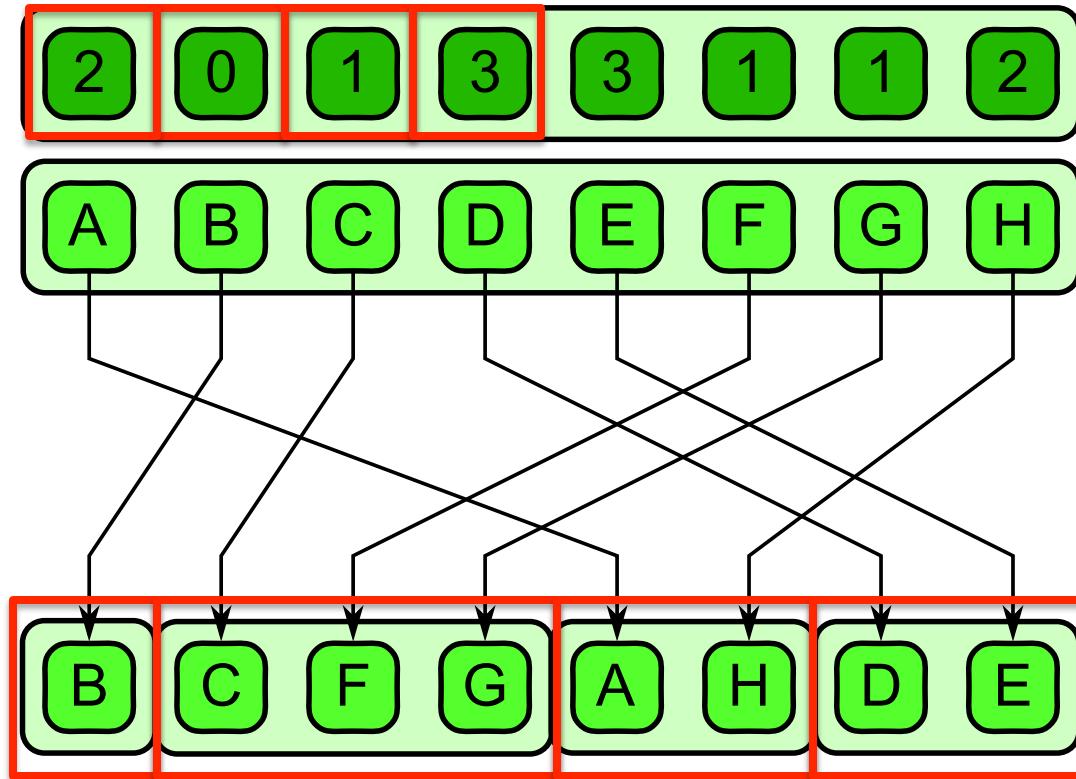
- Generalization of pack pattern
- Elements are moved to upper or lower half of output collection based on some state
- Does not lose information like pack

# *Generalization of Pack: Unsplit*



- Inverse of split
- Creates **output collection** based on original input collection

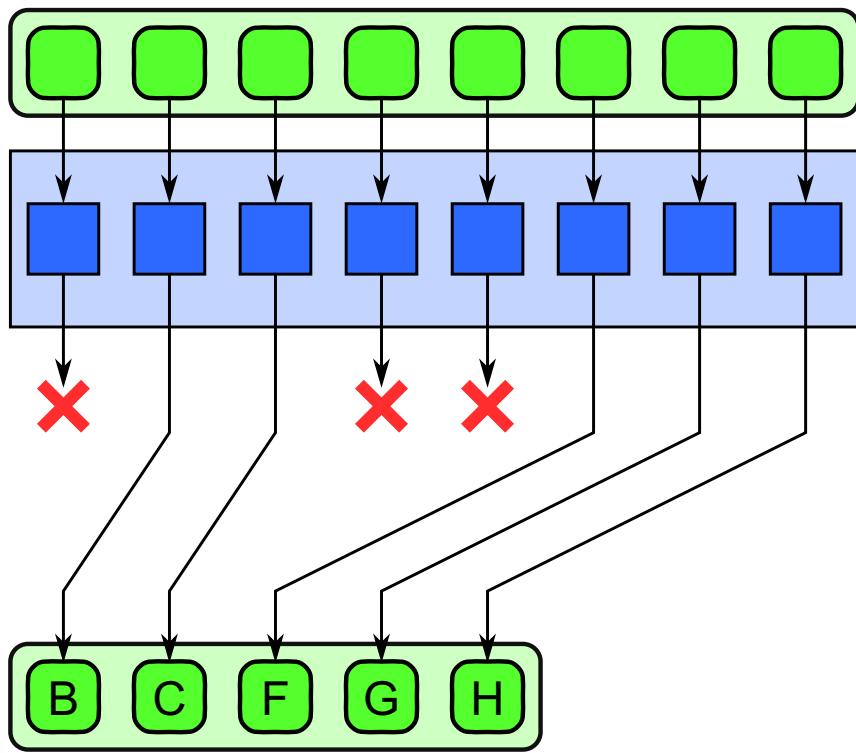
# *Generalization of Pack: Bin*



- Generalized split to support more than two categories
- Ex: radix sort, pattern classification

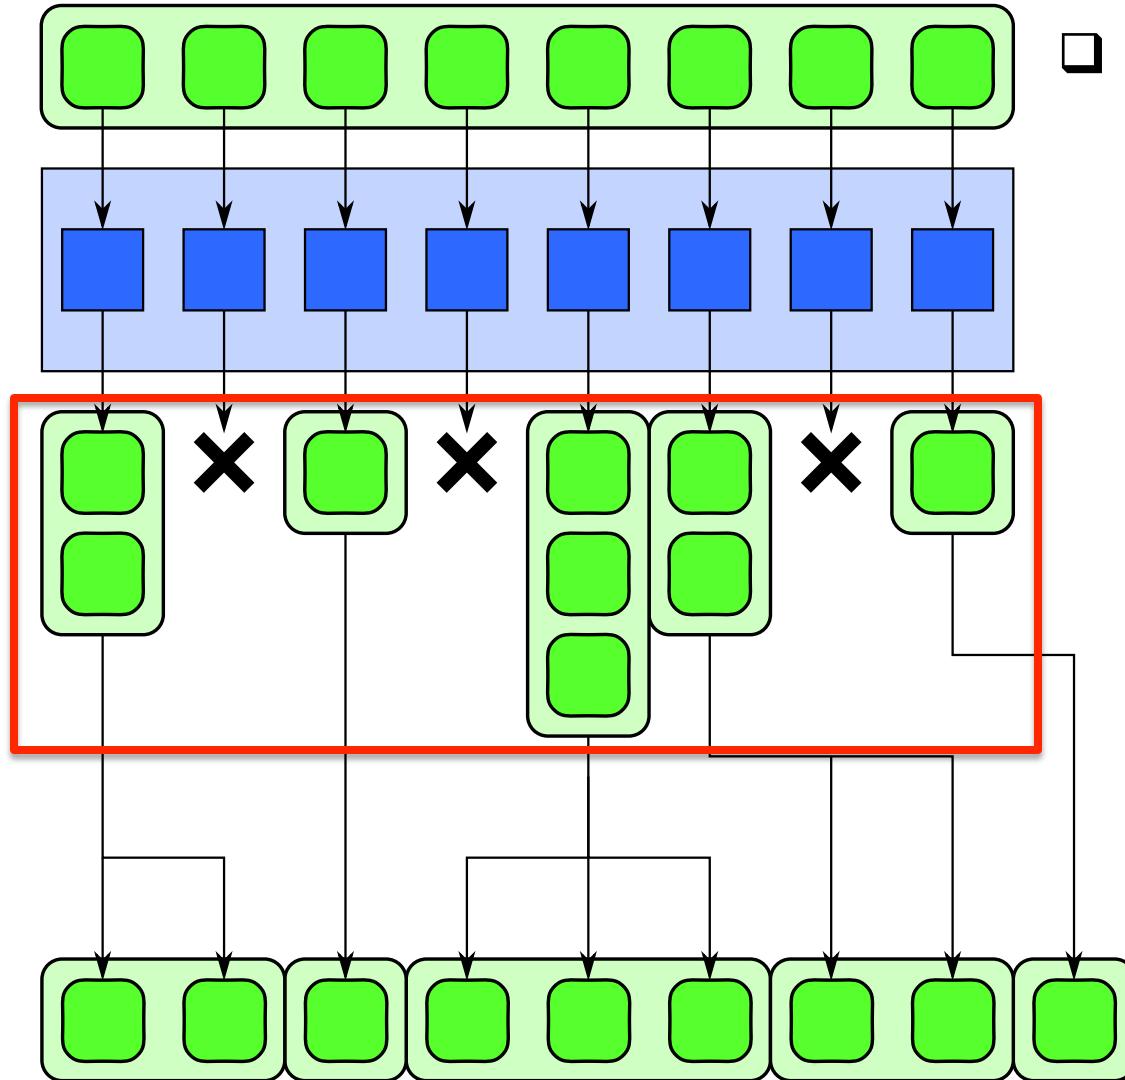
**4 different categories = 4 bins**

# *Fusion of Map and Pack*



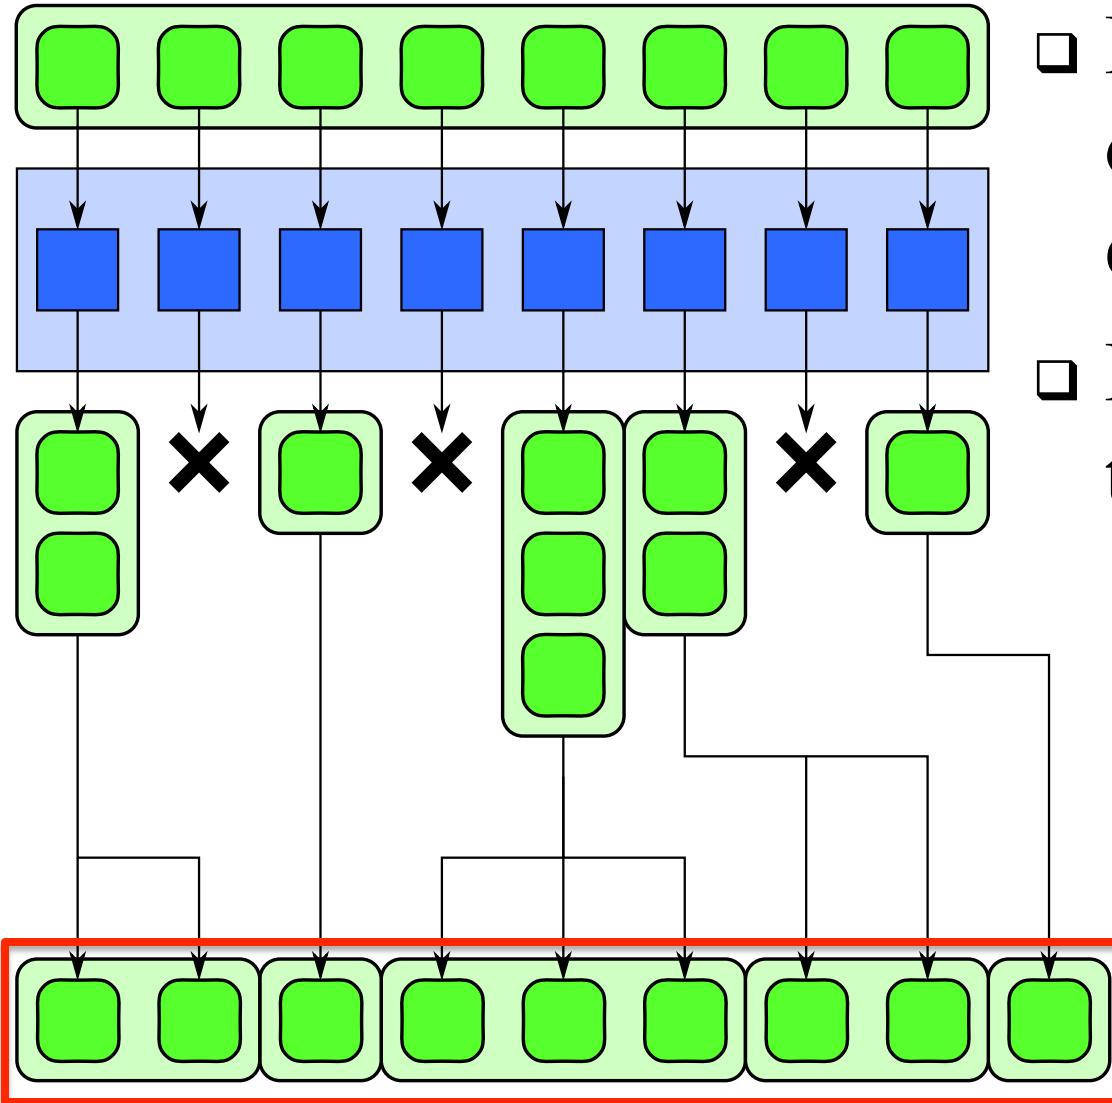
- Advantageous if most of the elements of a map are discarded
- **Map** checks pairs for collision
- **Pack** stores only actual collisions
- Output BW  $\sim$  results reported, not number of pairs tested
- Each element can output 0 or 1 element

# *Generalization of Pack: Expand*



- Each element can output any number of elements

# *Generalization of Pack: Expand*



- Each element can output any number of elements
- **Results are fused together in order**

# *Parallelizing Algorithms*

## □ Common strategy:

1. Divide up the computational domain into sections
2. Work on the sections individually
3. Combine the results

Discussed in ch 8



Methods: divide-and-conquer, fork-join, geometric decomposition, partitions, segments

# *Parallelizing Algorithms*

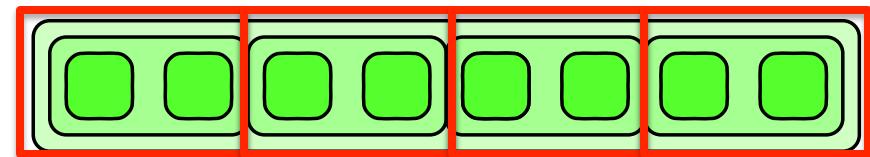
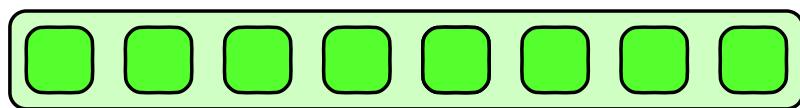
- Common strategy:
  1. Divide up the computational domain into sections
  2. Work on the sections individually
  3. Combine the results

Methods: divide-and-conquer, fork-join, geometric decomposition, **partitions**, segments

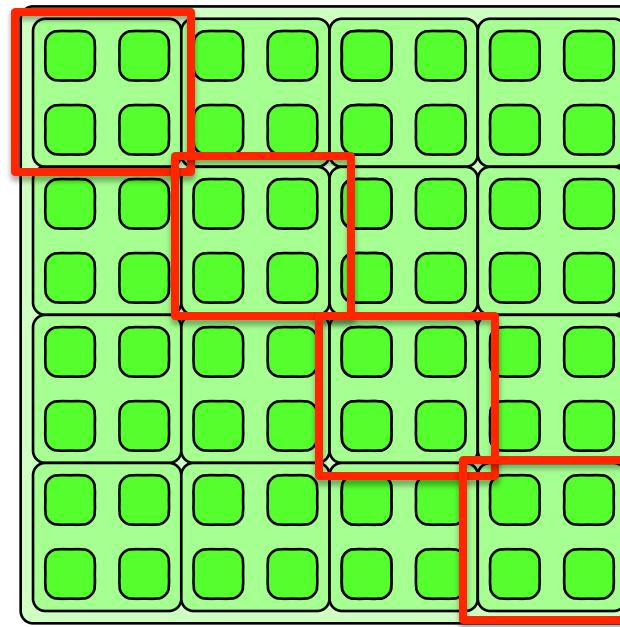
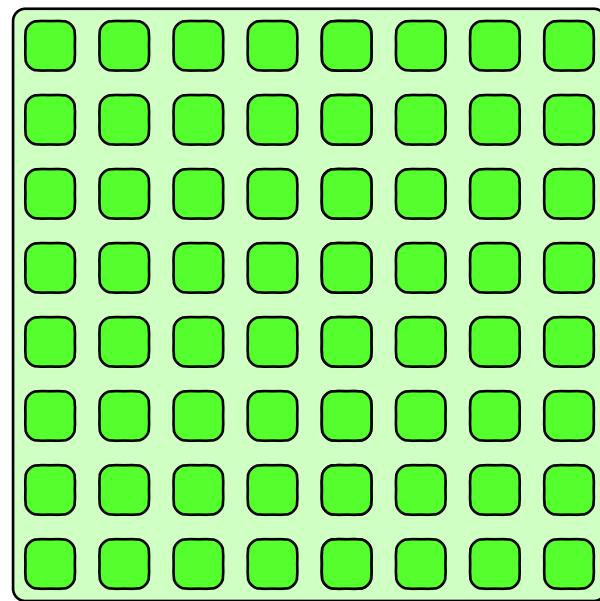


Discussed next...

# *Partitioning*



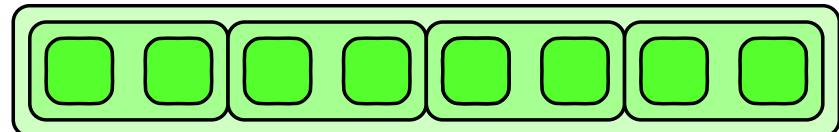
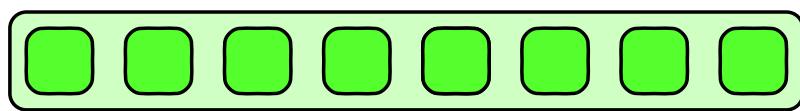
1D



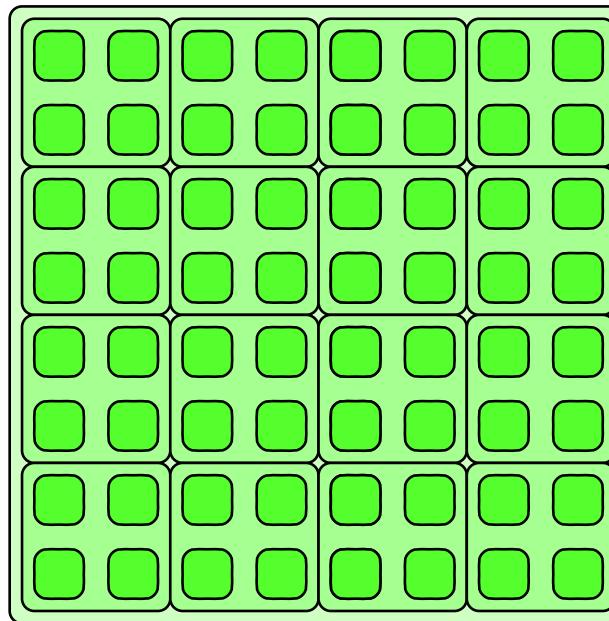
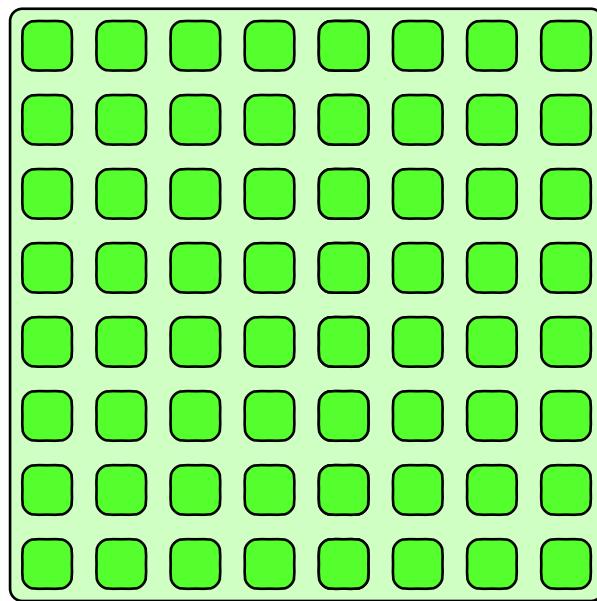
2D

- Data is divided into **non-overlapping, equal-sized** regions

# *Partitioning*



1D



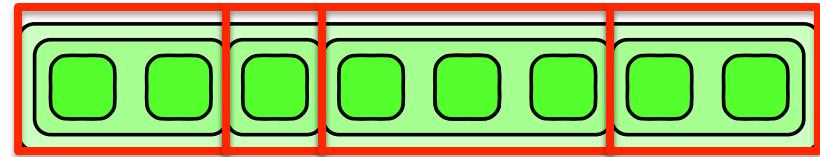
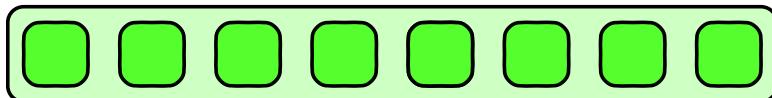
2D

- Data is divided into **non-overlapping**, equal-sized regions



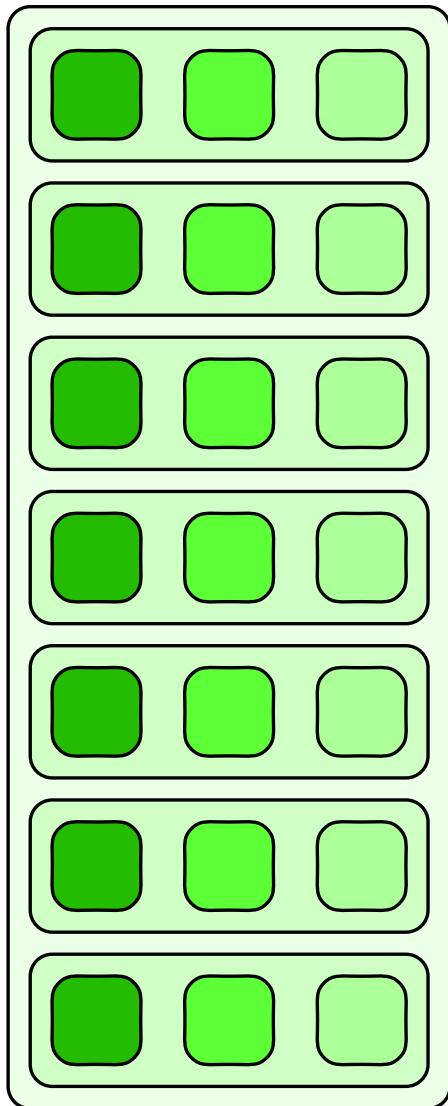
Avoid write conflicts and race conditions

# *Segmentation*



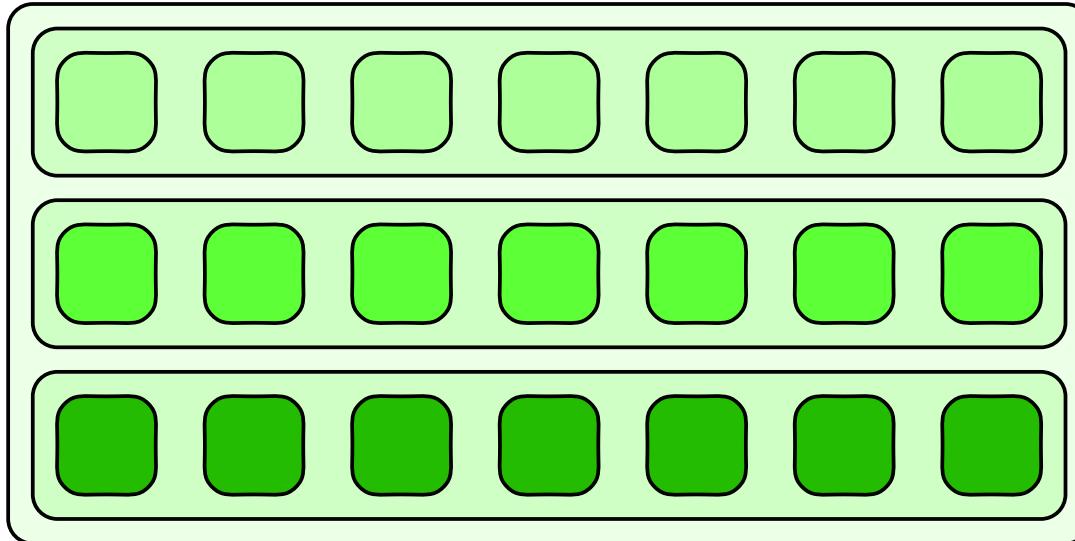
- Data is divided into **non-uniform** non-overlapping regions
- Start of each segment can be marked using:
  - Array of integers
  - Array of Boolean flags

# *Array of Structures (AoS)*



- May lead to better cache utilization if data is accessed randomly

# *Structures of Arrays (SoA)*



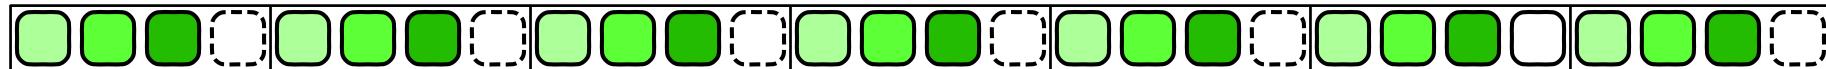
- Typically better for vectorization and avoidance of false sharing

# *Data Layout Options*

Array of Structures (AoS), padding at end



Array of Structures (AoS), padding after each structure



Structure of Arrays (SoA), padding at end



Structure of Arrays (SoA), padding after each component



# *Example Implementation*

## AoS Code

```
struct node {  
    float x, y, z;  
};  
struct node NODES[1024];  
  
float dist[1024];  
for(i=0;i<1024;i+=16){  
    float x[16],y[16],z[16],d[16];  
    x[:] = NODES[i:16].x;  
    y[:] = NODES[i:16].y;  
    z[:] = NODES[i:16].z;  
    d[:] = sqrtf(x[:]*x[:] + y[:]*y[:] + z[:]*z[:]);  
    dist[i:16] = d[:];  
}
```

## SoA Code

```
struct node1 {  
    float x[1024], y[1024], z[1024];  
}  
struct node1 NODES1;  
  
float dist[1024];  
for(i=0;i<1024;i+=16){  
    float x[16],y[16],z[16],d[16];  
    x[:] = NODES1.x[i:16];  
    y[:] = NODES1.y[i:16];  
    z[:] = NODES1.z[i:16];  
    d[:] = sqrtf(x[:]*x[:] + y[:]*y[:] + z[:]*z[:]);  
    dist[i:16] = d[:];  
}
```

# *AoS Code*

```
struct node {  
    float x, y, z;  
};  
struct node NODES[1024];  
  
float dist[1024];  
for(i=0;i<1024;i+=16){  
    float x[16],y[16],z[16],d[16];  
    x[:] = NODES[i:16].x;  
    y[:] = NODES[i:16].y;  
    z[:] = NODES[i:16].z;  
    d[:] = sqrtf(x[:]*x[:] + y[:]*y[:] + z[:]*z[:]);  
    dist[i:16] = d[:];  
}
```

- Most logical data organization layout
- Extremely difficult to access memory for reads (gathers) and writes (scatters)
- Prevents efficient vectorization

# *SoA Code*

```
struct node1 {  
    float x[1024], y[1024], z[1024];  
}  
struct node1 NODES1;  
  
float dist[1024];  
for(i=0;i<1024;i+=16){  
    float x[16],y[16],z[16],d[16];  
    x[:] = NODES1.x[i:16];  
    y[:] = NODES1.y[i:16];  
    z[:] = NODES1.z[i:16];  
    d[:] = sqrtf(x[:]*x[:] + y[:]*y[:] + z[:]*z[:]);  
    dist[i:16] = d[:];  
}
```

- Separate arrays for each structure-field keeps memory accesses contiguous when vectorization is performed over structure instances