

# Stencil Pattern

Parallel Computing

CIS 410/510

Department of Computer and Information Science



UNIVERSITY OF OREGON

# *Table of Contents*

- ❑ What is the stencil pattern?
- ❑ Implementing stencil with shift
- ❑ Stencil and cache optimizations
- ❑ Stencil and communication optimizations
- ❑ Recurrence


# *Table of Contents*

- ❑ What is the stencil pattern?
- ❑ Implementing stencil with shift
- ❑ Stencil and cache optimizations
- ❑ Stencil and communication optimizations
- ❑ Recurrence

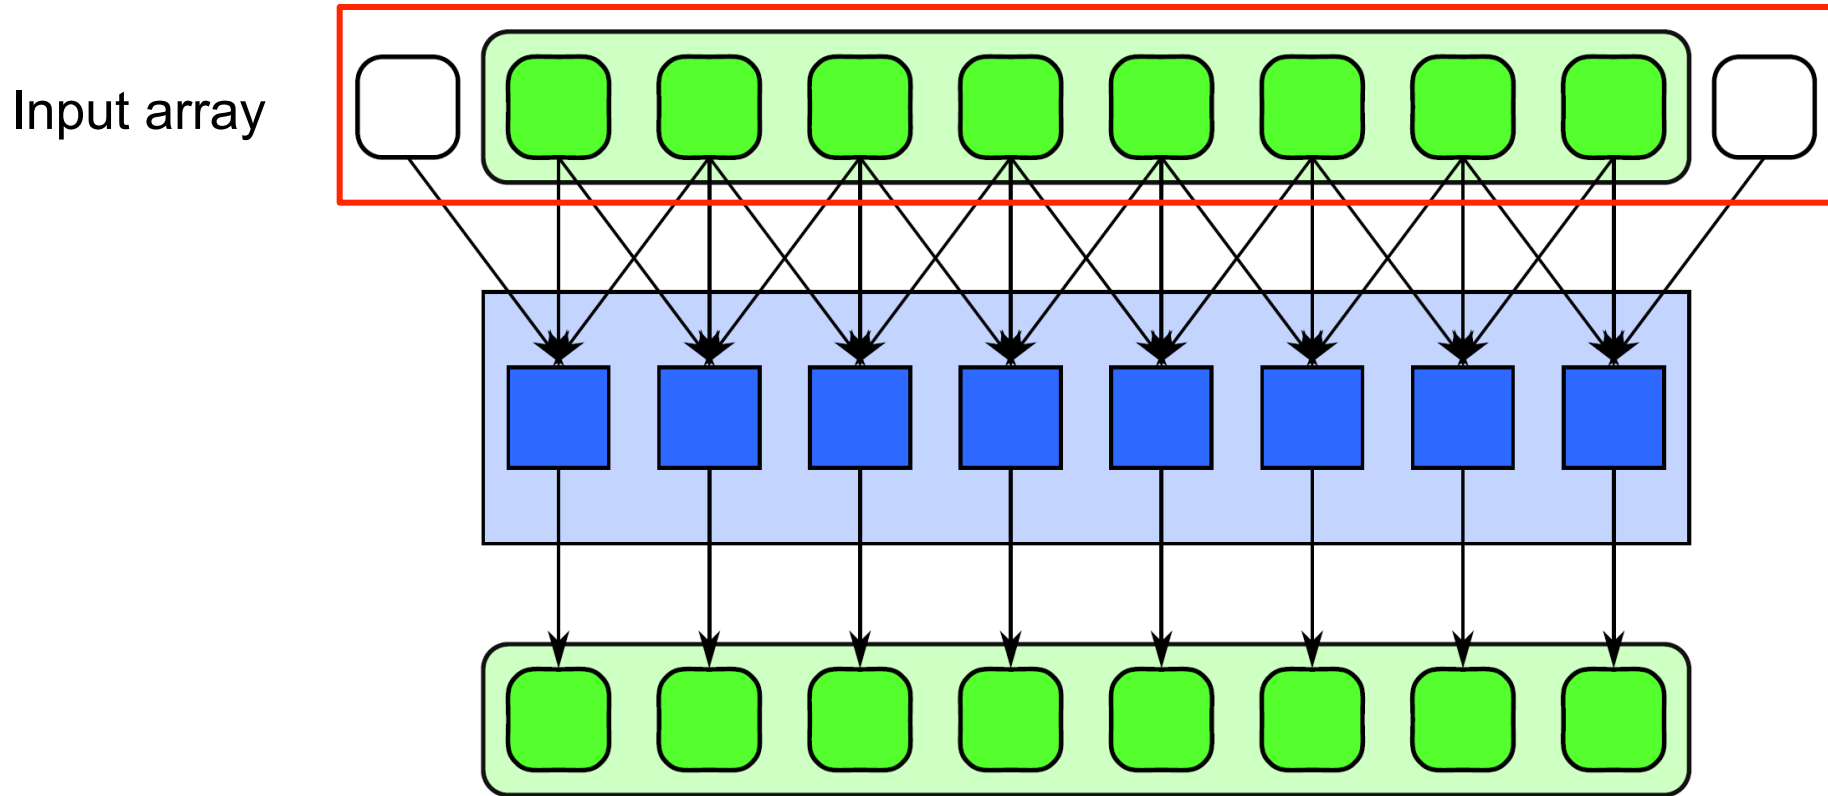
# *What is the stencil pattern?*

- ❑ **Stencil:** A map where each output depends on a “neighborhood” of inputs. These inputs are a set of fixed offsets relative to the output position.
- ❑ A stencil output is a function of a “neighborhood” of elements in an input collection


# *What is the stencil pattern?*




# *What is the stencil pattern?*



# *What is the stencil pattern?*

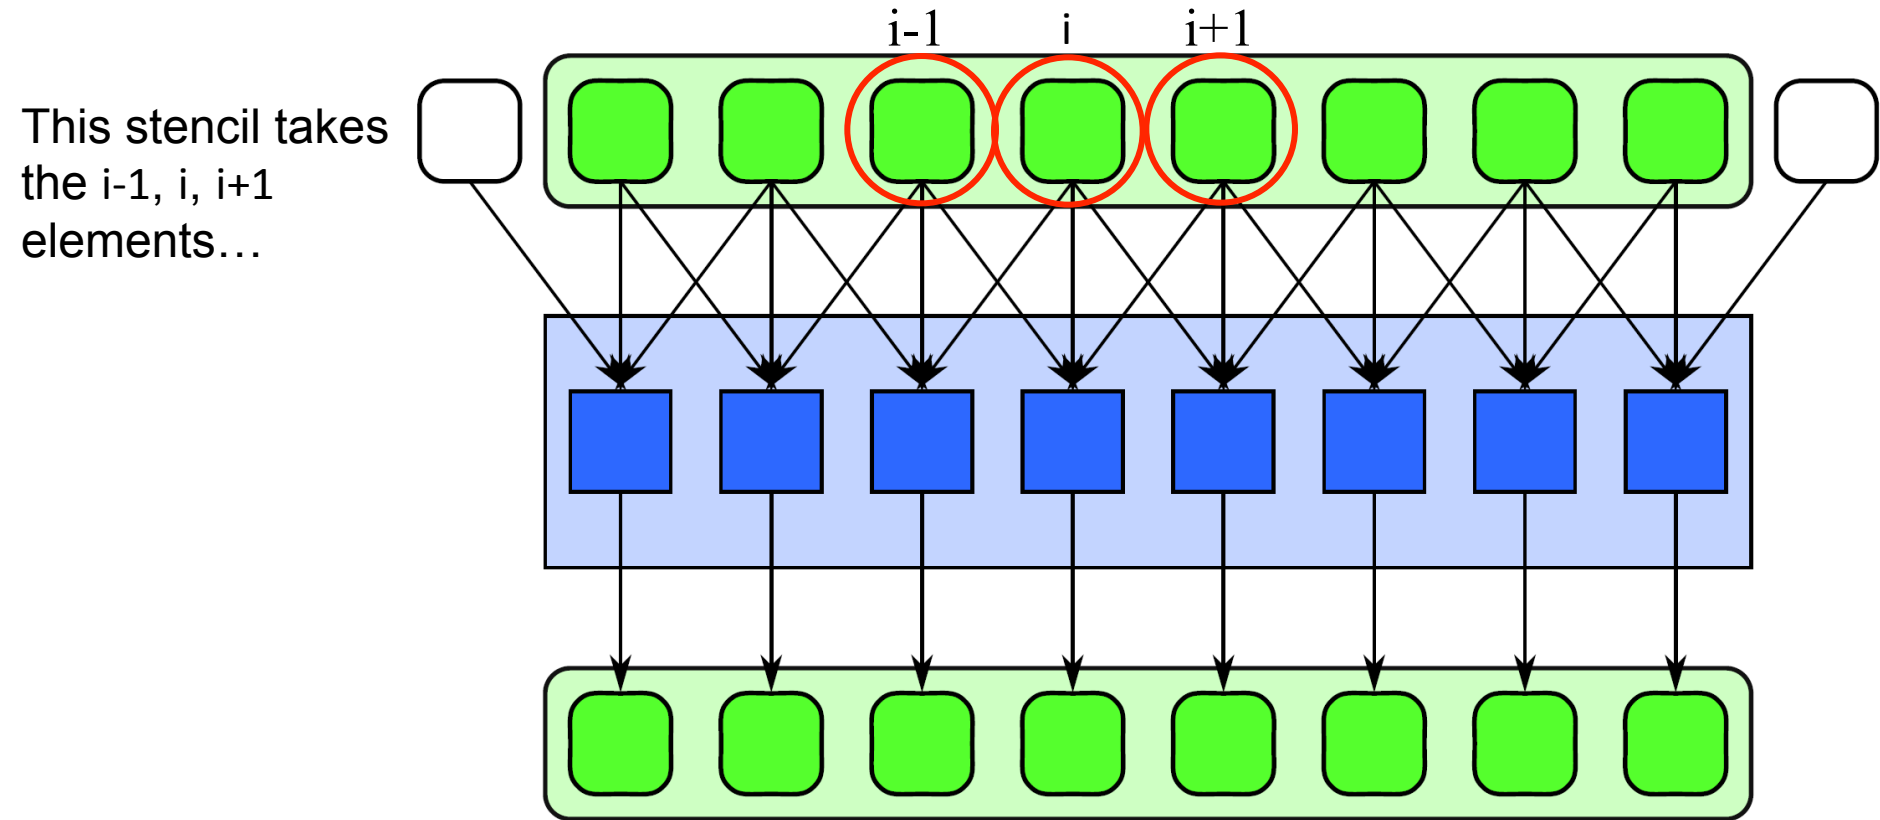


# *What is the stencil pattern?*







# *What is the stencil pattern?*



# *What is the stencil pattern?*



# *What is the stencil pattern?*



# *What is the stencil pattern?*

- ❑ Stencils can operate on one dimensional and multidimensional data
- ❑ Stencil neighborhoods can range from compact to sparse, square to cube, and anything else!

*What is the stencil pattern?*

Practice!

# What is the stencil pattern?

- Here is our array,  $A$

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

# What is the stencil pattern?

- Here is our array,  $A$
- Apply a stencil operation to the inner square of the form:  
$$(i,j) = \text{avg}((i,j), (i-1,j), (i+1,j), (i,j-1), (i,j+1))$$

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0

# What is the stencil pattern?


1) Average all blue squares

0	0	0	0
0	9	7	0
0	6	4	0
0	0	0	0




# What is the stencil pattern?

- 1) Average all blue squares
- 2) Store result in red square




# What is the stencil pattern?

- 1) Average all blue squares
- 2) Store result in red square
- 3) Repeat 1 and 2 for all green squares



# What is the stencil pattern?

- 1) Average all blue squares
- 2) Store result in red square
- 3) Repeat 1 and 2 for all green squares



# What is the stencil pattern?

- Output!
- How would we parallelize this?

4.4	4
3.8	3.4

# *Serial Stencil Example (part 1)*

```
1  template<
2      int NumOff,      // number of offsets
3      typename In,     // type of input locations
4      typename Out,    // type of output locations
5      typename F       // type of function / functor
6  >
7  void stencil(
8      int n,           // number of elements in data collection
9      const In a[],    // input data collection (n elements)
10     Out r[],          // output data collection (n elements)
11     In b,             // boundary value
12     F func,           // function / functor from neighborhood inputs to output
13     const int offsets[] // offsets (NumOffsets elements)
14 ) {
```

# *Serial Stencil Example (part 2)*

```
15      // array to hold neighbors
16      In neighborhood[NumOff];
17      // loop over all output locations
18      for (int i = 0; i < n; ++i) {
19          // loop over all offsets and gather neighborhood
20          for (int j = 0; j < NumOff; ++j) {
21              // get index of jth input location
22              int k = i+offsets[j];
23              if (0 <= k && k < n) {
24                  // read input location
25                  neighborhood[j] = a[k];
26              } else {
27                  // handle boundary case
28                  neighborhood[j] = b;
29              }
30          }
31          // compute output value from input neighborhood
32          r[i] = func(neighborhood);
33      }
34 }
```

# *Parallel Stencil Example*

# *Table of Contents*

- ❑ What is the stencil pattern?
- ❑ Implementing stencil with shift
- ❑ Stencil and cache optimizations
- ❑ Stencil and communication optimizations
- ❑ Recurrence




# *Implementing Stencil with Shift*

- ❑ One possible implementation of the stencil pattern includes shifting the input data
- ❑ For each offset in the stencil, we gather a new input vector by **shifting** the original input by the offset amount

# Implementing Stencil with Shift

All input arrays are derived from the same original input array



# *Implementing Stencil with Shift*

- ❑ This implementation is only beneficial for one dimensional stencils or the memory-contiguous dimension of a multidimensional stencil
- ❑ Memory traffic to external memory is not reduced with shifts
- ❑ But, shifts allow vectorization of the data reads, which may reduce the total number of instructions

# *Table of Contents*

- ❑ What is the stencil pattern?
- ❑ Implementing stencil with shift
- ❑ Stencil and cache optimizations
- ❑ Stencil and communication optimizations
- ❑ Recurrence

# *Stencil and Cache Optimizations*

- ❑ Assuming 2D array where rows are contiguous in memory...
  - Horizontally related data will tend to belong to the same cache line
  - Vertical offset accesses will most likely result in cache misses

# *Stencil and Cache Optimizations*

- ❑ Assigning rows to cores:
  - Maximizes horizontal data locality
  - Assuming vertical offsets in stencil, this will create redundant reads of adjacent rows from each core
- ❑ Assigning columns to cores:
  - Redundantly read data from same cache line
  - Create false sharing as cores write to same cache line

# *Stencil and Cache Optimizations*


- ❑ Assigning “strips” to each core can be a better solution
- ❑ **Strip-mining:** an optimization in a stencil computation that groups elements in a way that avoids redundant memory accesses and aligns memory accesses with cache lines

# *Stencil and Cache Optimizations*

- ❑ A strip's size is a multiple of a cache line in width, and the height of the 2D array
- ❑ Strip widths are in increments of the cache line size so as to avoid false sharing and redundant reads
- ❑ Each strip is processed serially from top to bottom within each core



# *Stencil and Cache Optimizations*



# *Table of Contents*

- ❑ What is the stencil pattern?
- ❑ Implementing stencil with shift
- ❑ Stencil and cache optimizations
- ❑ Stencil and communication optimizations
- ❑ Recurrence

# Stencil and Communication Optimizations

- ❑ When data is distributed, overlapping regions (or “**ghost cells**”) must be explicitly communicated between nodes between loop iterations
- ❑ Darker cells are PE 0’s ghost cells
- ❑ After first iteration of stencil computation
  - PE 1 & PE 2 must send their stencil results to PE 0
  - PE 0 can perform another iteration of stencil

0	0	0	0
0	PE 0 9	PE 1 7	0
0	PE 2 6	PE 3 4	0
0	0	0	0

# *Stencil and Communication Optimizations*

- ❑ Generally better to replicate ghost cells in each local memory and swap after each iteration than to share memory
  - Fine-grained sharing can lead to increased communication cost

# *Stencil and Communication Optimizations*

- ❑ **Halo:** set of all ghost cells
- ❑ Halo must contain all neighbors needed for one iteration
- ❑ Larger halo (**deep halo**)
  - Less communications
  - More redundant local computation
  
- ❑ **Latency Hiding:** Compute interior of stencil while waiting for ghost cell updates

# *Table of Contents*

- ❑ What is the stencil pattern?
- ❑ Implementing stencil with shift
- ❑ Stencil and cache optimizations
- ❑ Stencil and communication optimizations
- ❑ Recurrence

# *Recurrence*

- ❑ What if we have several nested loops with data dependencies between them when doing a stencil computation?

# Recurrence

```
1 void my_recurrence(  
2     size_t v,           // number of elements vertically  
3     size_t h,           // number of elements horizontally  
4     const float a[v][h], // input 2D array  
5     float b[v][h]       // output 2D array (boundaries already initialized )  
6 ) {  
7     for (int i=1; i<v; ++i)  
8         for (int j=1; j<h; ++j)  
9             b[i][j] = f(b[i-1][j], b[i][j-1], a[i][j]);  
10 }
```



# Recurrence

```
1 void my_recurrence(  
2     size_t v,          // number of elements vertically  
3     size_t h,          // number of elements horizontally  
4     const float a[v][h], // input 2D array  
5     float b[v][h]      // output 2D array (boundaries already initialized )  
6 ) {  
7     for (int i=1; i<v; ++i)  
8         for (int j=1; j<h; ++j)  
9             b[i][j] = f(b[i-1][j], b[i][j-1], a[i][j]);  
10 }
```




Data dependencies between loops

# *Recurrence*


- ❑ This can still be parallelized!
- ❑ Trick: find a plane that cuts through grid of intermediate results
  - Previously computed values on one side of plane
  - Values to still be computed on other side of plane
  - Computation proceeds perpendicular to plane through time (this is known as a sweep)
- ❑ This plane is called a *separating hyperplane*

# Recurrence



# Recurrence

- Same grid of intermediate results
- Each level corresponds to a loop iteration
- Computation proceeds downward



# *Example Implementation*

# *Conclusion*

- ❑ Examined the stencil and recurrence pattern
  - Both have a regular pattern of communication and data access
- ❑ In both patterns we can convert a set of offset memory accesses to shifts
- ❑ Stencils can use strip-mining to optimize cache use
- ❑ Ghost cells should be considered when stencil data is distributed across different memory spaces

