

# **Lecture 3: Technologies**

Allen D. Malony

Department of Computer and Information Science



UNIVERSITY OF OREGON

# *Performance Tools and Technologies*

- ❑ It is never the case that performance tools are developed from scratch
- ❑ They depend on a range of technologies that can themselves be significant engineering efforts
  - Even simple conceptual things can be hard
- ❑ Most technologies deal with how to observe performance metrics or state

*"If I have seen further it is by standing  
on the shoulders of giants."*

- Sir Isaac Newton

# *Technologies*

- ❑ Timers
- ❑ Counters
- ❑ Instrumentation
  - Source (PDT)
  - PMPI
  - Compiler instrumentation
  - Binary
    - ◆ Dyninst
    - ◆ PEBIL
    - ◆ MAQAO
  - Runtime Interfaces
- ❑ Program Address Resolution: Binutils
- ❑ Stack Walking
  - StackwalkerAPI
  - Libunwind
  - Backtrace
- ❑ Heterogeneous (accelerator) timers and counters

# *Time*

- ❑ How is time measured in a computer system?
- ❑ How do we derive time from a clock?
- ❑ What clock/time technologies are available to a measurement system?
- ❑ How are clocks synchronized in a parallel computer in order to provide a “global time” common between nodes?
- ❑ Different technologies are available
  - Issues of resolution and accuracy

## *Timer: gettimeofday()*

- ❑ UNIX function
- ❑ Returns wall-clock time in seconds and microseconds
- ❑ Actual resolution is hardware-dependent
- ❑ Base value is 00:00 UTC, January 1, 1970
- ❑ Some implementations also return the timezone

```
#include <sys/time.h>

struct timeval tv;
double walltime;  /* seconds */

gettimeofday(&tv, NULL);
walltime = tv.tv_sec + tv.tv_usec * 1.0e-6;
```

# *Timer: clock\_gettime()*

- ❑ POSIX function
- ❑ For *clock\_id* CLOCK\_REALTIME it returns wall-clock time in seconds and nanoseconds
- ❑ More clocks may be implemented but are not standardized
- ❑ Actual resolution is hardware-dependent

```
#include <time.h>

struct timespec tv;
double walltime;  /* seconds */

clock_gettime(CLOCK_REALTIME, &tv);
walltime = tv.tv_sec + tv.tv_nsec * 1.0e-9;
```

# *Timer: getrusage()*

- ❑ UNIX function
- ❑ Provides a variety of different information
  - Including user time, system time, memory usage, page faults, and other *resource use* information
  - Information provided system-dependent!

```
#include <sys/resource.h>

struct rusage ru;
double usertime;  /* seconds */
int memused;

getrusage(RUSAGE_SELF, &ru);
usertime = ru.ru_utime.tv_sec +
           ru.ru_utime.tv_usec * 1.0e-6;
memused = ru.ru_maxrss;
```

## *Timer: Others*

- ❑ MPI provides portable MPI wall-clock timer

```
#include <mpi.h>
double walltime;  /* seconds */

walltime = MPI_Wtime();
```

- Not required to be consistent/synchronized across ranks!

- ❑ OpenMP 2.0 also provides a library function

```
#include <omp.h>
double walltime;  /* seconds */

walltime = omp_get_wtime();
```

- ❑ Hybrid MPI/OpenMP programming?

- Interactions between both standards (yet) undefined



# *Timer: Others*

- ❑ Fortran 90 intrinsic subroutines
  - `cpu_time()`
  - `system_clock()`
- ❑ Hardware counter libraries typically provide “timers” because underlying them are cycle counters
  - Vendor APIs
    - ◆ PMAPI, HWPC, libhpm, libpfm, libperf, ...
  - PAPI (Performance API)

# *What Are Performance Counters*

- ❑ Extra processor logic inserted to count specific events
- ❑ Updated at every cycle (or when some event occurs)
- ❑ Strengths
  - Non-intrusive
  - Very accurate
  - Low overhead
- ❑ Weaknesses
  - Provides only hard counts
  - Specific for each processor
  - Access is not appropriate for the end user
    - ◆ nor is it well documented
  - Lack of standard on what is counted


# Hardware Counter Issues

## ❑ Kernel level

- Handling of overflows
- Thread accumulation
- Thread migration
- State inheritance
- Multiplexing
- Overhead
- Atomicity

## ❑ Multi-platform interfaces

- Performance API (*PAPI*)
  - ◆ University of Tennessee, USA
- Lightweight Performance Tools (*LIKWID*)
  - ◆ University of Erlangen, Germany



# *Hardware Measurement*

- Typical measured events account for:
  - Functional units status
    - ◆ float point operations
    - ◆ fixed point operations
    - ◆ load/stores
  - Access to memory hierarchy
  - Cache coherence protocol events
  - Cycles and instructions counts
  - Speculative execution information
    - ◆ instructions dispatched
    - ◆ branches mispredicted

# *Hardware Metrics*

## ❑ Typical hardware counter

Cycles / Instructions

Floating point instructions

Integer instructions

Load/stores

Cache misses

Cache misses

Cache misses

TLB misses

## Useful derived metrics

IPC

FLOPS

computation intensity

instructions per load/store

load/stores per cache miss

cache hit rate

loads per load miss

loads per TLB miss

- ❑ Derived metrics allow users to correlate the behavior of the application to hardware components
- ❑ Define threshold values acceptable for metrics and take actions regarding optimization when below/above thresholds

# *Accuracy Issues*

- ❑ Granularity of the measured code
  - If not sufficiently large enough, overhead of the counter interfaces may dominate
  - Mainly applies to time
- ❑ Pay attention to what is not measured:
  - Out-of-order processors
  - Sometimes speculation is included
  - Lack of standard on what is counted
    - ◆ microbenchmarks can help determine accuracy of the hardware counters
- ❑ Impact of measurement on counters themselves
  - Typically less of an issue

# *Hardware Counters Access on Linux*

- ❑ Linux had not defined an out-of-the-box interface to access the hardware counters!
  - Linux Performance Monitoring Counters Driver (PerfCtr) by Mikael Pettersson from Uppsala X86 + X86-64
    - ◆ needs kernel patching!  
<http://user.it.uu.se/~mikpe/linux/perfctr/>
  - Perfmon by Stephane Eranian from HP – IA64
    - ◆ it was being evaluated to be added to Linux  
<http://www.hpl.hp.com/research/linux/perfmon/>
- ❑ Linux 2.6.31
  - Performance Counter subsystem provides an abstraction of special performance counter hardware registers

# *Utilities to Count Hardware Events*

- ❑ There are utilities that start a program and at the end of the execution provide overall event counts
  - *hpmcount* (IBM)
  - *CrayPat* (Cray)
  - *pfmon* from HP (part of Perfmon for AI64)
  - *psrun* (NCSA)
  - *cputrack*, *har* (Sun)
  - *perfex*, *ssrun* (SGI)
  - *perf* (Linux 2.6.31)



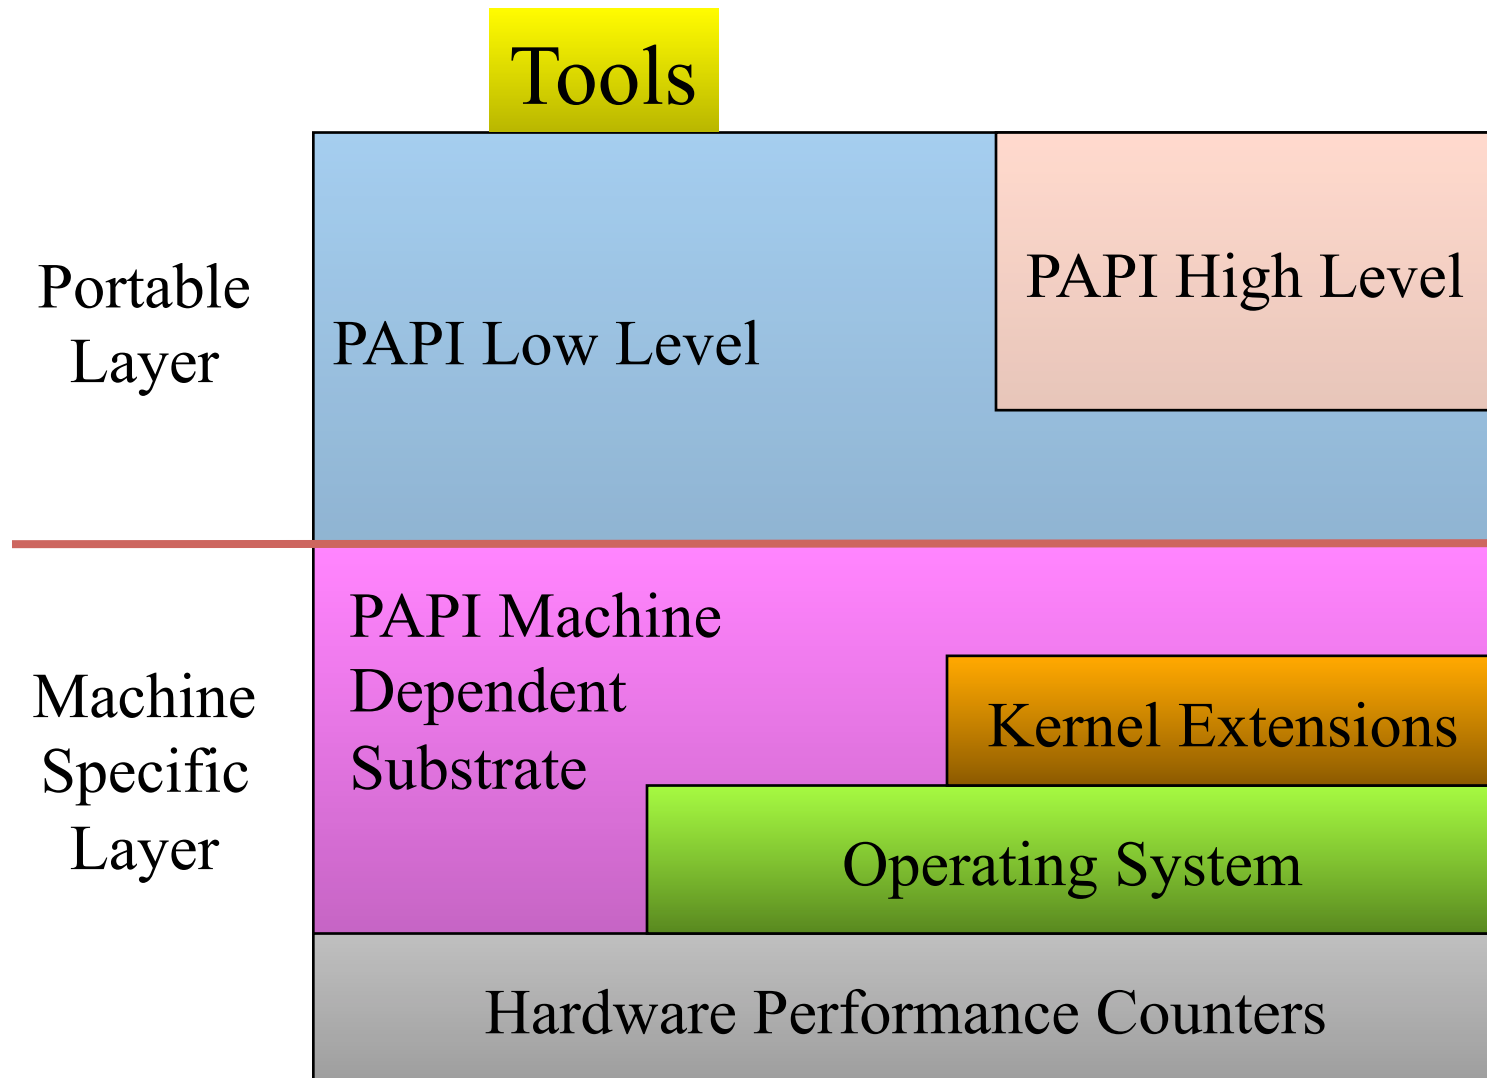


# *PAPI – Performance API*

- ❑ Middleware to provide a consistent and portable API for the performance counter hardware in microprocessors
- ❑ Countable events are defined in two ways:
  - Platform-neutral *preset* events
  - Platform-dependent *native* events
- ❑ Presets can be derived from multiple native events
- ❑ Two interfaces to the underlying counter hardware:
  - *High-level* interface simply provides the ability to start, stop and read the counters for a specified list of events
  - Low-level interface manages hardware events in user defined groups called *EventSets*
- ❑ Events can be multiplexed if counters are limited

<http://icl.cs.utk.edu/papi/>

# ***PAPI Architecture***



# *PAPI Predefined Events*

- ❑ Common set of events deemed relevant and useful for application performance tuning (wish list)
  - `papiStdEventDefs.h`
  - Accesses to the memory hierarchy, cache coherence protocol events, cycle and instruction counts, functional unit and pipeline status
  - Run PAPI `papi_avail` utility to determine which predefined events are available on a given platform
  - Semantics may differ on different platforms!
- ❑ PAPI also provides access to native events on all supported platforms through the low-level interface
  - Run PAPI `papi_native_avail` utility to determine which predefined events are available on a given platform

# *papi\_avail Utility*

- ❑ Provides information on what events are available on a particular hardware platform

```
% papi_avail -h
```

```
This is the PAPI avail program.
```

```
It provides availability and detail information  
for PAPI preset and native events.  Usage:
```

```
    papi_avail [options] [event name]
```

```
    papi_avail TESTS_QUIET
```

```
Options:
```

-a	display only available PAPI preset events
-d	display PAPI preset event info in detailed format
-e EVENTNAME	display full detail for named preset or native event
-h	print this help message
-t	display PAPI preset event info in tabular format (default)

# *High Level API*


- ❑ Meant for application programmers wanting simple but accurate measurements
- ❑ Calls the lower level API
- ❑ Allows only PAPI preset events
- ❑ Eight functions:
  - *PAPI\_num\_counters*
  - *PAPI\_start\_counters*, *PAPI\_stop\_counters*
  - *PAPI\_read\_counters*
  - *PAPI\_accum\_counters*
  - *PAPI\_flops*
  - *PAPI\_flips*, *PAPI\_ipc* (New in Version 3.x)
- ❑ Not thread-safe (Version 2.x)

# *Low Level API*

- ❑ Increased efficiency and functionality over the high level PAPI interface
- ❑ 54 functions
- ❑ Access to native events
- ❑ Obtain information about the executable, the hardware, and memory
- ❑ Set options for multiplexing and overflow handling
- ❑ System V style sampling (profil())
- ❑ Thread safe

# *Component PAPI*

- ❑ Developed for the purpose of extending counter sets while providing a common interface



# *Source Instrumentation with Timers*


- ❑ Measuring performance using timers requires instrumentation
  - Have to uniquely identify code region (name)
  - Have to add code for timer start and stop
  - Have to compute delta and accumulate statistics
- ❑ Hand-instrumenting becomes tedious very quickly, even for small software projects
- ❑ Also a requirement for enabling instrumentation only when wanted
  - Avoids unnecessary overheads when not needed



# *Program Database Toolkit (PDT)*

- ❑ University of Oregon, Research Center Juelich (FZJ Germany), Edison Design Group, Inc. (USA), LLNL (USA)
- ❑ Automated instrumentation of C/C++, Fortran source code
- ❑ Source code parser(s) identify blocks such as function boundaries, loop boundaries, generates a .PDB file for each source file
- ❑ Instrumentor uses .PDB file to insert API calls into source code files at block enter/exit, outputs an instrumented code file
- ❑ Instrumented source passed to compiler for compilation to object file
- ❑ Linker links application with measurement library providing definitions for API calls
- ❑ Free download: <http://tau.uoregon.edu>

# *PDT Architecture*



# *PMPI – MPI Standard Profiling Interface*

- ❑ The MPI (Message Passing Interface) standard defines a mechanism for instrumenting all API calls in an MPI implementation
- ❑ Each MPI\_\* function call is actually a weakly defined interface that can be re-defined by performance tools
- ❑ Each MPI\_\* function call eventually calls a corresponding PMPI\_\* function call which provides the expected MPI functionality
- ❑ Performance tools can redefine MPI\_\* calls

# *PMPI Example*

## ❑ Original MPI\_Send() definition:

```
int __attribute__((weak))
MPI_Send(void *buf, int count, MPI_Datatype datatype, int dest,
int tag, MPI_Comm comm) {
    PMPI_Send(buf, count, datatype, dest, tag, comm);
}
```

## ❑ *Possible* Performance tool definition:

```
int MPI_Send(void *buf, int count, MPI_Datatype datatype, int
dest, int tag, MPI_Comm comm) {
    MYTOOL_Timer_Start("MPI_Send");
    PMPI_Send(buf, count, datatype, dest, tag, comm);
    MYTOOL_Timer_Stop("MPI_Send");
    MYTOOL_Message_Size("MPI_Send", count * sizeof(datatype));
}
```

# *Compiler Instrumentation*

- ❑ Modern compilers provide the ability to instrument functions at compile time
- ❑ Can exclude files, functions
- ❑ GCC example:
  - `-finstrument-functions` parameter
  - Instruments function entry and exit(s)

```
void __cyg_profile_func_enter (void *this_fn, void *call_site);  
void __cyg_profile_func_exit  (void *this_fn, void *call_site);
```

# *Compiler Instrumentation – Tool Interface*

- ❑ Measurement libraries have to implement those two functions:

```
void __cyg_profile_func_enter (void *this_fn, void *call_site);  
void __cyg_profile_func_exit  (void *this_fn, void *call_site);
```

- ❑ The function and call site pointers are instruction addresses
- ❑ How to resolve those addresses to source code locations?
  - Binutils: libbfd, libiberty (discussed later)

# *Binary Instrumentation*

- ❑ Source Instrumentation not possible in all cases
  - Exotic / Domain Specific Languages (no parser support)
  - Pre-compiled system libraries
  - Utility libraries without source available
- ❑ Binary instrumentation modifies the existing executable and all libraries, adding user-specified function entry/exit API calls
- ❑ Can be done once, or as first step of execution

# *Binary Instrumentation: Dyninst API*

- ❑ University of Wisconsin, University of Maryland
- ❑ Provides binary instrumentation for runtime code patching:
  - Performance Measurement Tools
  - Correctness Debuggers (efficient data breakpoints)
  - Execution drive simulations
  - Computational Steering

<http://www.dyninst.org>



# *Binary Instrumentation: PEBIL*

- ❑ San Diego Supercomputing Center / PMaC group
- ❑ Static binary instrumentation for x86\_64 Linux
- ❑ PEBIL = PMaC's Efficient Binary Instrumentation for Linux/x86
- ❑ Lightweight binary instrumentation tool that can be used to capture information about the behavior of a running executable



<http://www.sdsc.edu/PMaC/projects/pebil.html>

# *PEBIL Design*

- ❑ Efficiency is priority #1
- ❑ Designed around a few use cases
  - Execution counting
  - Memory tracing
- ❑ Static binary rewriter
  - Write instrumented + runnable executable to disk
    - ◆ keep original behavior intact
    - ◆ gather information as a side-effect
  - Instrument once, run many times
  - No instrumentation cost at runtime
  - Code patching (not just-in-time compiled!)

# How Binary Instrumentation Works

## Original

	0000c000 <foo>:
Basic Block 1	c000: 48 89 7d f8 mov %rdi,-0x8(%rbp)
Basic Block 2	c004: 5e pop %rsi
	c005: 75 f8 jne 0xc004
Basic Block 3	c007: c9 leaveq
	c008: c3 retq

## Instrumented

```
0000d000 <foo>:
d000: e9 de ad be ef jmp 0x1000 # to instrumentation
d005: 48 89 7d f8 mov %rdi,-0x8(%rbp)
d000: e9 de ad be ef jmp 0x1010 # to instrumentation
d00a: 5e pop %rsi
d00b: 75 00 00 00 f8 jne 0xd009
d000: e9 de ad be ef jmp 0x1020 # to instrumentation
d00a: c9 leaveq
d00b: c3 retq
```

// do stuff  
// jump back

# *Use case: Memory Address Collection*

- ❑ Collect the address of every load/store issued by the application
  - Put addresses in a buffer, process addresses in batch
    - ◆ Fewer function calls
    - ◆ Less cache pollution

```
for (i = 0; i < n; i++) {  
    if (cur + 2 > BUF_SIZE) clear_buf();  
    buffer[cur + 0] = &(A[i]);  
    buffer[cur + 1] = &(B[i]);  
    A[i] = B[i];  
}
```

# *Binary Instrumentation: MAQAO*

- ❑ Modular Assembly Quality Analyzer and Optimizer
- ❑ Tool for analyzing and optimizing binary code
- ❑ Intel64 and Xeon Phi architectures supported
- ❑ Binary release only (for now)

<http://maqao.org>



# ***MAQAO: Introduction***

## ❑ Easy install

- Packaging : ONE (static) standalone binary
- Easy to embed

## ❑ Audience

- User/Tool developer: analysis and optimization tool
- Performance tool developer: framework services
- TAU: tau\_rewrite (MIL)
- ScoreP: on-going effort (MIL)

# MAQAO: Architecture

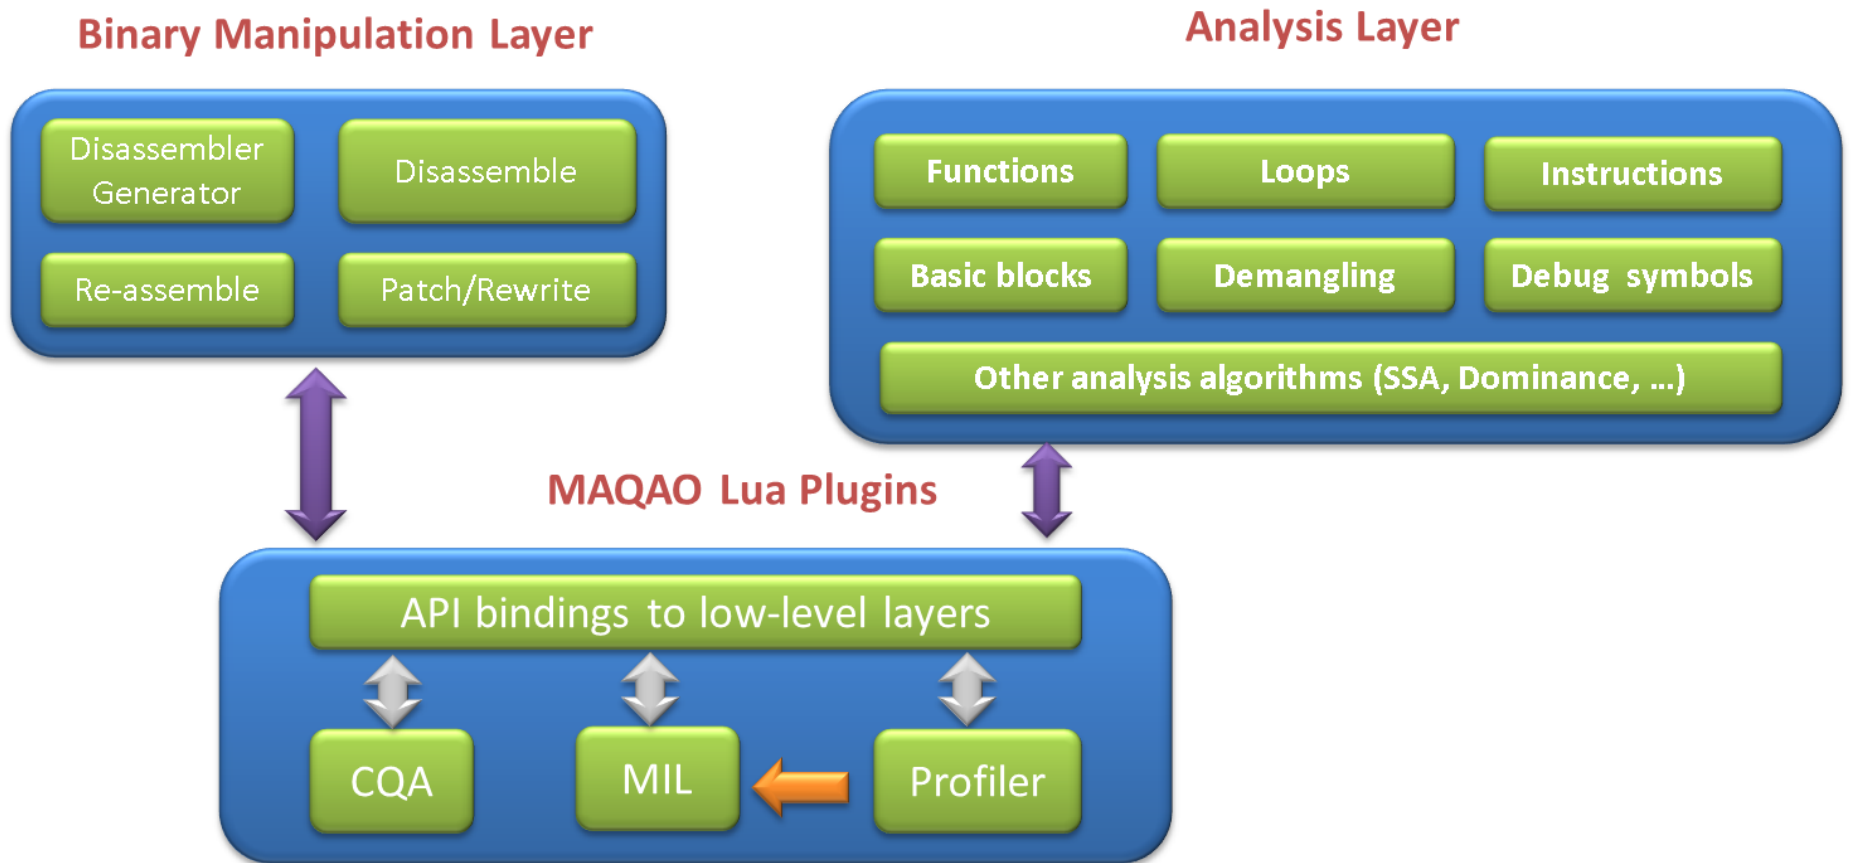


Image source: <http://maqao.org>

# ***MAQAO: Measurement and Analysis Tool***

- ❑ Scripting language
  - Lua language : simplicity and productivity
  - Fast prototyping
  - MAQAO Lua API : Access to services
- ❑ Built on top of the Framework
- ❑ Loop-centric approach
- ❑ Output: reports
  - System deals with low level details
  - User gets high level reports



# *Runtime Measurement Support*

- ❑ Some runtime systems provide callback mechanisms for function entry / exit or state transitions
  - Java JVM
  - Python
  - Some OpenMP runtimes (Collector API, OMPT)
    - ◆ Sun/Oracle, OpenUH, Intel (in development)
- ❑ Measurement tools / libraries:
  - implement event handlers for callback functions
  - register with the runtime, are notified when events happen

## *Periodic Sampling – what is it?*

- ❑ The application is interrupted after a specified period of time
- ❑ Interruption handler queries the program state
- ❑ The timer is reset and the process repeats until program termination
- ❑ Either at termination or during handler, a statistical profile is constructed
- ❑ Sampling theory states that the state (function) sampled the most frequently is the most time consuming state (function)

# *Periodic Sampling – how to do it?*

- ❑ ANSI C / POSIX signaling and signal handling
- ❑ `sigaction()`
  - Specify a handler for when a signal is raised
  - Handler has to be signal-safe\*
  - Handler gets program state context pointer, including current instruction pointer address and full program stack
- ❑ `setitimer()`
  - Portable (POSIX) interval timer
  - A signal is raised when the timer expires
  - Timers: real time (process-level only), user CPU time, or user CPU + system CPU time counters
  - Undefined behavior for threaded applications
- ❑ `timer_create()` / `timer_settime()`
  - POSIX function like `setitimer()`, but with a Linux-specific interval timer with threaded support for real time counter

\*POSIX.1-2004 lists the functions that are signal-safe

# *Address Resolution: GNU Binutils*


- ❑ Compiler instrumentation and signal handling deal with instruction pointer addresses
- ❑ Binutils provides utilities and libraries for looking up addresses and getting function name, source code file and line number
  - Source info available if code compiled with `-g`
- ❑ Iterates over executable and any shared object libraries (if applicable) to find address
- ❑ Command line version:
  - `addr2line -f -e <executable> <address_1> ... <address_n>`

<http://www.gnu.org/software/binutils/>


# *Stack Walking – how did I get here?*

- ❑ Periodic sampling freezes the program state
- ❑ “Walking the stack” answers the question: how did I get here?

```
int func_c(int c) {  
    return (c+3);  
}  
int func_b(int b) {  
    return func_c(b+2);  
}  
int func_a(int a) {  
    return func_b(a+1);  
}  
int main(int argc, char* argv[]) {  
    int in = atoi(argv[1]);  
    printf("%d: %d\n", in, func_a(in));  
}
```



Current program stack:



# *Stack Walking: libunwind*

- ❑ Libunwind defines a portable and efficient C programming interface (API) to determine the call-chain of a program
- ❑ Provides the means to manipulate the preserved (callee-saved) state of each call-frame and to resume execution at any point in the call-chain (non-local goto)
- ❑ Supports both local (same-process) and remote (across-process) operation.
- ❑ Developed and supported by the Free Software Foundation (FSF)
- ❑ <https://savannah.nongnu.org/projects/libunwind/>

# *Stack Walking: StackWalkerAPI*

- ❑ University of Wisconsin, University of Maryland
- ❑ An API that allows users to collect a call stack and access information about its stack frames
- ❑ Support for Linux (x86\_64, AMD-64, Power, Power-64), BlueGene/L and BlueGene/P
- ❑ <http://www.dyninst.org/stackwalker>

# *Stack Walking: Linux Backtrace (libc)*

- ❑ A *backtrace* is a list of the function calls that are currently active in a thread
- ❑ 2 steps: get array of  $n$  addresses, resolve them to symbols
- ❑ Warning! By default, address resolution is not signal-safe
- ❑ Signal-safe option writes to file descriptor (no malloc)

[http://www.gnu.org/software/libc/manual/html\\_node/Backtraces.html](http://www.gnu.org/software/libc/manual/html_node/Backtraces.html)

<http://man7.org/linux/man-pages/man3/backtrace.3.html>

```
#include <execinfo.h>
```

```
int backtrace(void **buffer, int size);
```

```
char **backtrace_symbols(void *const *buffer, int size);
```

```
void backtrace_symbols_fd(void *const *buffer, int size, int fd);
```



# *Backtrace example*

```
#include <execinfo.h>
#include <stdio.h>
#include <stdlib.h>

/* Obtain a backtrace and print it to stdout. */
void print_trace (void) {
    void *array[10];
    size_t size;
    char **strings;
    size_t i;

    size = backtrace (array, 10);
    // not signal safe! backtrace_symbols calls "malloc"
    strings = backtrace_symbols (array, size);
    printf ("Obtained %zd stack frames.\n", size);
    for (i = 0; i < size; i++)
        printf ("%s\n", strings[i]);
    free (strings);
}
```

# *NVIDIA CUDA Performance Tool Interface*

- ❑ Use of accelerator and coprocessor hardware also requires access to timer and counting information
- ❑ NVIDIA is developing CUDA Performance Tool Interface (CUPTI) to enable the creation of profiling and tracing tools
  - CUPTI support was released with CUDA 4.0
  - Capabilities have steadily improved
  - Current version is released with CUDA 5.x and the just announced CUDA 6
- ❑ CUPTI is delivered as a dynamic library

# ***NVIDIA CUPTI APIs***

## ❑ *Callback API*

- Interject tool callback code at the entry and exist to each CUDA runtime and driver API call
- Registered tools are invoked for selected events

## ❑ *Counter API*

- Query, configure, start, stop, read counters on CUDA devices
- Device-level counter access

## ❑ *Activity API*

- GPU kernel and memory copy timing information is stored in a buffer until a synchronization point is encounter and these timings are recorded by the CPU
- Synchronization can be either be within a device, stream or occur during some synchronous memory copies and event synchronizations

## ❑ Can also get information on kernel registers and instructions

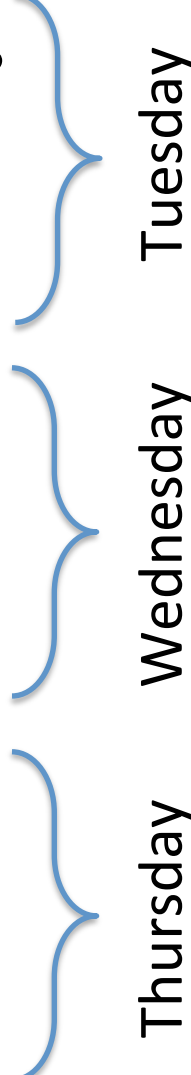
# *PAPI CUDA Component*

- ❑ HW performance counter measurement technology for NVIDIA CUDA platform
- ❑ Access to HW counters inside the GPUs
- ❑ Based on CUPTI (CUDA Performance Tool Interface)
- ❑ PAPI CUDA component can provide detailed performance counter info regarding execution of GPU kernel
  - Initialization, device management and context management are enabled by CUDA driver API
  - Domain and event management are enabled by CUPTI
- ❑ Names of events are established by the following hierarchy:  
    Component.Device.Domain.Event

# Portion of CUDA Events on Tesla C870

Event Code	Symbol	Long Description
0x44000000	CUDA.GeForce_GTX_480.gpc0.local_load	# executed local load instructions per warp on a multiprocessor
0x44000001	CUDA.GeForce_GTX_480.gpc0.local_store	# executed local store instructions per warp on a multiprocessor
0x44000002	CUDA.GeForce_GTX_480.gpc0.gld_request	# executed global load instructions per warp on a multiprocessor
0x44000003	CUDA.GeForce_GTX_480.gpc0.gst_request	# executed global store instructions per warp on a multiprocessor
0x44000004	CUDA.GeForce_GTX_480.gpc0.shared_load	# executed shared load instructions per warp on a multiprocessor
0x44000005	CUDA.GeForce_GTX_480.gpc0.shared_store	# executed shared store instructions per warp on a multiprocessor
0x44000006	CUDA.GeForce_GTX_480.gpc0.branch	# branches taken by threads executing a kernel
0x44000007	CUDA.GeForce_GTX_480.gpc0.divergent_branch	# divergent branches within a warp
0x4400000b	CUDA.GeForce_GTX_480.gpc0.active_cycles	# cycles a multiprocessor has at least one active warp
0x4400000c	CUDA.GeForce_GTX_480.gpc0.sm_cta_launched	# thread blocks launched on a multiprocessor
0x4400000d	CUDA.GeForce_GTX_480.gpc0.l1_local_load_hit	# local load hits in L1 cache
0x4400000e	CUDA.GeForce_GTX_480.gpc0.l1_local_load_miss	# local load misses in L1 cache
0x44000011	CUDA.GeForce_GTX_480.gpc0.l1_global_load_hit	# global load hits in L1 cache
0x4400002e	CUDA.Tesla_C870.domain_a.tex_cache_hit	# texture cache misses
0x4400002f	CUDA.Tesla_C870.domain_a.tex_cache_miss	# texture cache hits
0x44000034	CUDA.Tesla_C870.domain_b.local_load	# local memory load transactions
0x44000037	CUDA.Tesla_C870.domain_b.branch	# branches taken by threads executing a kernel
0x44000038	CUDA.Tesla_C870.domain_b.divergent_branch	# divergent branches within a warp
0x44000039	CUDA.Tesla_C870.domain_b.instructions	# instructions executed

# *Short Course Outline*

- ❑ Lecture 1: Introduction and Fundamentals
  - ❑ Lecture 2: Methodology
  - ❑ Lecture 3: Tools Technology
  - ❑ Lecture 4: Tools Landscape – Part 1
  - ❑ Lecture 5: Tools Landscape – Part 2
  - ❑ Lecture 6: TAU Performance System
  - ❑ Lecture 7: TAU Applications
  - ❑ Lecture 8: Advances in TAU
  - ❑ Lecture 9: Future Directions
- 
- Tuesday
- Wednesday
- Thursday