

# **Lecture 9:**

# **Parallel Performance Tools Future**

Allen D. Malony

Department of Computer and Information Science



UNIVERSITY OF OREGON

# *Perspective – Parallel Tools Experience*

- ❑ *Performance* has been the fundamental driving concern for parallel computing ➡ high-performance computing
  - *Performance observation, modeling, and engineering* are key methodologies to deliver the potential of parallel, HPC machines
- ❑ However, there has always been a strained relationship between *performance tools* and the parallel computing
  - Performance considered necessary, but an afterthought
- ❑ There are compelling reasons for performance methods and technologies to be integrated throughout parallel computing
  - To obtain, learn, and carry forward performance knowledge
  - To address increasing scalability and complexity issues
  - To bridge between programming semantics (computation model) and execution model operation

# *Parallel Performance Engineering*

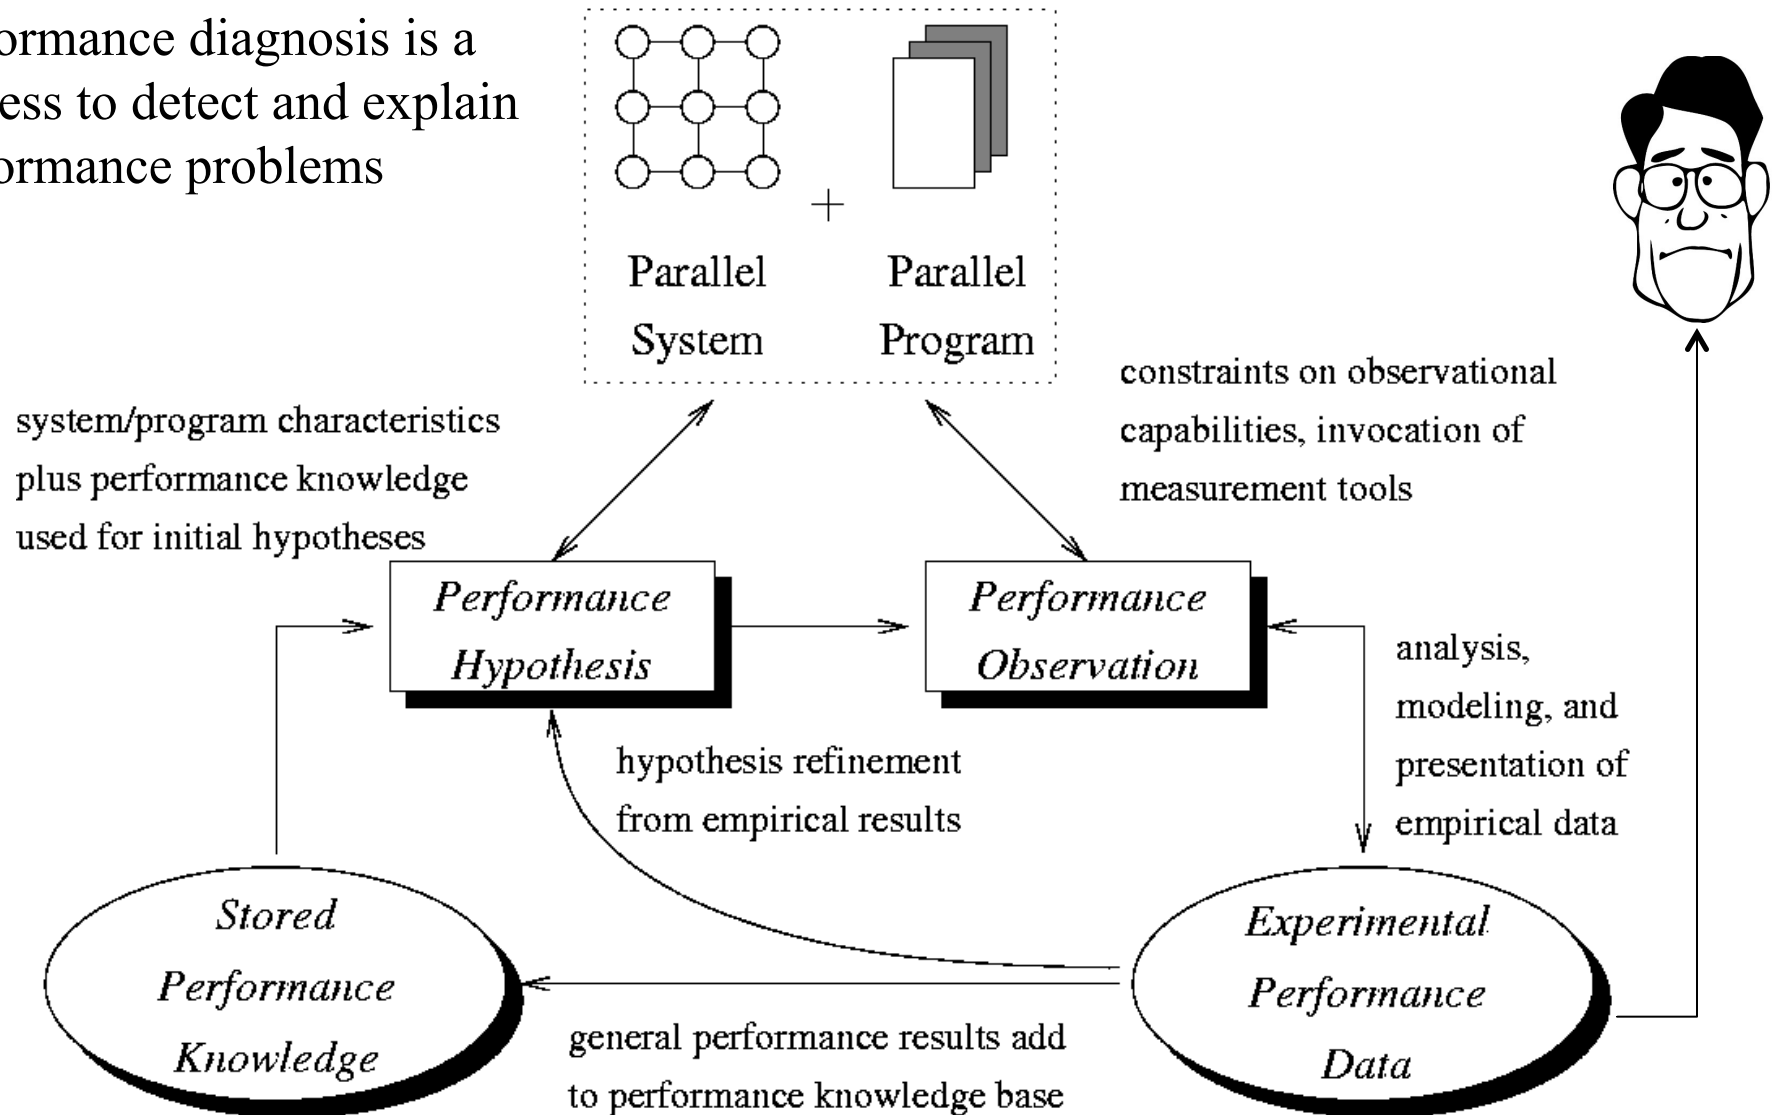
- ❑ Scalable, optimized applications deliver HPC promise
- ❑ Optimization through *performance engineering process*
  - Understand performance complexity and inefficiencies
  - Tune application to run optimally on high-end machines
- ❑ How to make the process more effective and productive?
  - What is the nature of the performance problem solving?
  - What is the performance technology to be applied?
- ❑ Performance tool efforts have been focused on performance observation, analysis, problem diagnosis
  - Application development and optimization productivity
  - Programmability, reusability, portability, robustness
  - Performance technology part of larger programming system
- ❑ Parallel systems evolution will change process, technology, and use of tools

# *Retrospective (1991) – Performance Observability*

- ❑ Performance evaluation problems define the requirements for performance measurement and analysis methods
- ❑ *Performance observability* is the ability to “accurately” capture, analyze, and present understand (collectively *observe*) information about parallel software and system
- ❑ Tools for performance observability must balance the *need* for performance data against the *cost* of obtaining it (environment complexity, performance intrusion)
  - Too little performance data makes analysis difficult
  - Too much data perturbs the measured system
- ❑ Important to understand performance observability complexity and develop technology to address it

# Retrospective (1998-2007) – Performance Diagnosis

Performance diagnosis is a process to detect and explain performance problems



# *Performance Diagnosis Projects*

- ❑ ***APART*** – Automatic Performance Analysis - Real Tools
  - Problem specification and identification
- ❑ ***Poirot*** – theory of performance diagnosis processes
  - Compare and analyze performance diagnosis systems
  - Use theory to create system that is automated / adaptable
    - ◆ *Heuristic classification*: match to characteristics
    - ◆ *Heuristic search*: look up solutions with problem knowledge
  - Problems: low-level feedback, lack of explanation power
- ❑ ***Hercule*** – knowledge-based (model-based) diagnosis
  - Capture knowledge about performance problems
  - Capture knowledge about how to detect and explain them
  - Knowledge comes from *parallel computational models*
    - ◆ associate computational models with performance models

## *Retrospective (2008-2012) – Performance Complexity*

- ❑ Performance tools have evolved incrementally to serve the dominant architectures and programming models
  - Reasonably stable, static parallel execution models
  - Allowed application-level observation focus
- ❑ Observation by *first person* measurement model:
  - Performance measurement can be made locally (per thread)
  - Performance data collected at the end of the execution
  - Post-mortem analysis and presentation of performance results
  - Offline performance engineering
- ❑ Increasing performance complexity
  - Factors: core counts, hierarchical memory architecture, interconnection technology, heterogeneity, and scale
- ❑ Focus on performance technology integration

# *Evolution*

- ❑ Increased performance complexity and scale forces the engineering process to be more intelligent and automated
  - Automate performance data analysis / mining / learning
  - Automated performance problem identification
- ❑ Even with intelligent and application-specific tools, the decisions of what to analyze are difficult
  - Performance engineering tools and practice must incorporate a performance knowledge discovery process
- ❑ Model-oriented knowledge
  - Computational semantics of the application
  - Symbolic models for algorithms
  - Performance models for system architectures / components
- ❑ Application developers can be more directly involved in the performance engineering process

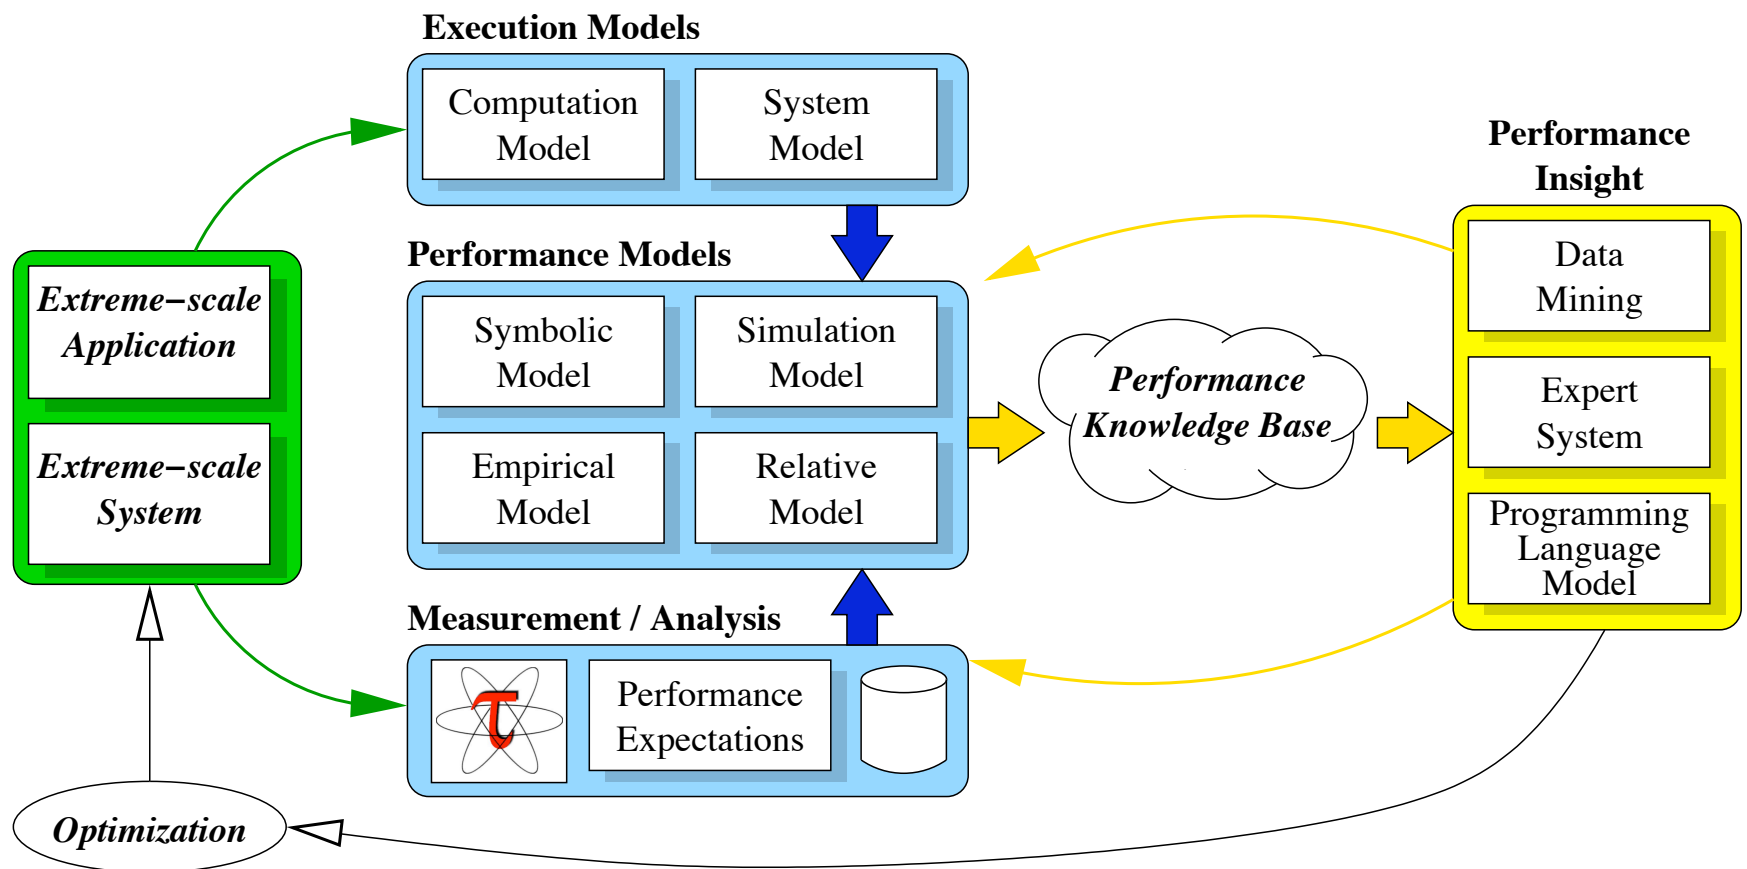


# *Need for Whole Performance Evaluation*

- ❑ Extreme scale performance is an optimized orchestration
  - Application, processor, memory, network, I/O
- ❑ Reductionist approaches to performance will be unable to support optimization and productivity objectives
- ❑ Application-level only performance view is myopic
  - Interplay of hardware, software, and system components
  - Ultimately determines how performance is delivered
- ❑ Performance should be evaluated *in toto*
  - Application and system components
  - Understand effects of performance interactions
  - Identify opportunities for optimization across levels
- ❑ Need *whole performance evaluation practice*

# *“Extreme” Performance Engineering*

- Empirical performance data evaluated with respect to performance expectations at levels of abstraction



# *Revolution (for Exascale)*

- ❑ First person measurement (post-mortem analysis) not viable
  - Highly concurrent and dynamic execution model
    - ◆ post-mortem analysis of low-level data prohibitive
  - Interactions with scarce and shared resources
    - ◆ introduces bottlenecks and queues on chip/node and between nodes
  - Multiple objectives (performance, energy, resilience, ...)
  - Dynamic variability motivates need for runtime adaptation
- ❑ ***Third person*** measurement model required (in addition)
  - Focus is on system activity characterization at different levels
    - ◆ system resource usage is a primary concern
  - Measurements are analyzed relative to contributors
  - Online analysis and availability of performance allows introspective adaptation for objective evaluation and tuning

# *A New “Performance” Observability*

- ❑ Exascale requires a fundamentally different “performance” observability paradigm
  - Designed specifically to support introspective adaptation
  - Reflects computation model mapped to execution model
  - Aware of multiple objectives (“performance”)
  - In-situ analysis of performance state and objectives
- ❑ Key parallel “performance” abstraction
  - Inherent state of exascale execution is dynamic
  - Embodies non-stationarity of “performance”
  - Constantly shaped by the adaptation of resources to meet computational needs and optimize execution objectives

# *Needs Integration in Exascale Software Stack*

- ❑ Exascale programming methodology can include observability awareness and adaptability
  - Programming system exposes alternatives
    - ◆ parameters, algorithms, parallelism control, ...
  - Runtime state awareness can be coupled with application knowledge for self-adaptive, closed-loop runtime tuning
    - ◆ richer contextualization and attribution of performance state
- ❑ Enables *top-down* application transformations to be optimized for runtime and system layers
- ❑ Feeding back dynamic information about HW/SW software resources from *bottom-up*
- ❑ Performance introspection and dynamic adaptivity
- ❑ Requires support in OS and runtime environment

## *Willing Suspension of (Observability) Disbelief*

- ❑ Suppose that full knowledge of the state of the exascale system were available
  - Can observe it at any time
  - No cost to produce it, access it, or analyze it
  - No cost to provide it to wherever it might be used
- ❑ How/where could it be used in an exascale system?
- ❑ Would it fundamentally change how an exascale system is programmed?
- ❑ How would you incorporate such a capability in the exascale software stack?

# ***XPRESS Project (DOE X-Stack)***

- ❑ Design and development of exascale software stack to support the ParalleX execution model
  - Highly concurrent
  - Asynchronous
  - Message driven
  - Global address space
- ❑ OpenX
  - XPI programming API
  - HPX runtime system
  - RIOS interface to OS
  - APEX performance system
- ❑ Team
  - Universities: IU, LSU, UH, UNC/RENCI, UO
  - Laboratories: SNL, LBNL, ORNL

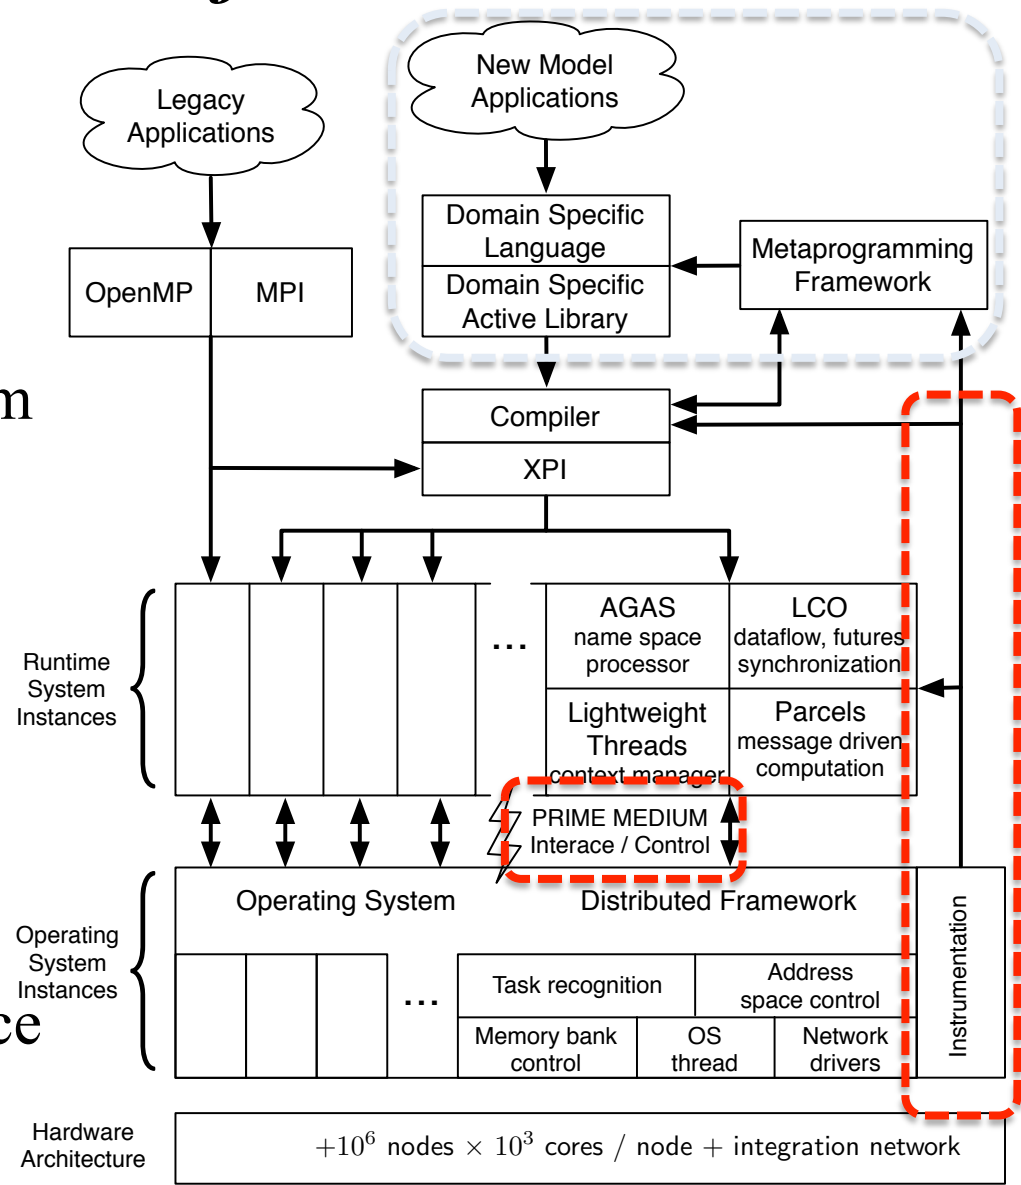
# Integrated Software Stack for ParallelX

## ❑ OpenX

- XPI programming API
- HPX runtime system
- RIOS interface to OS
- APEX performance system

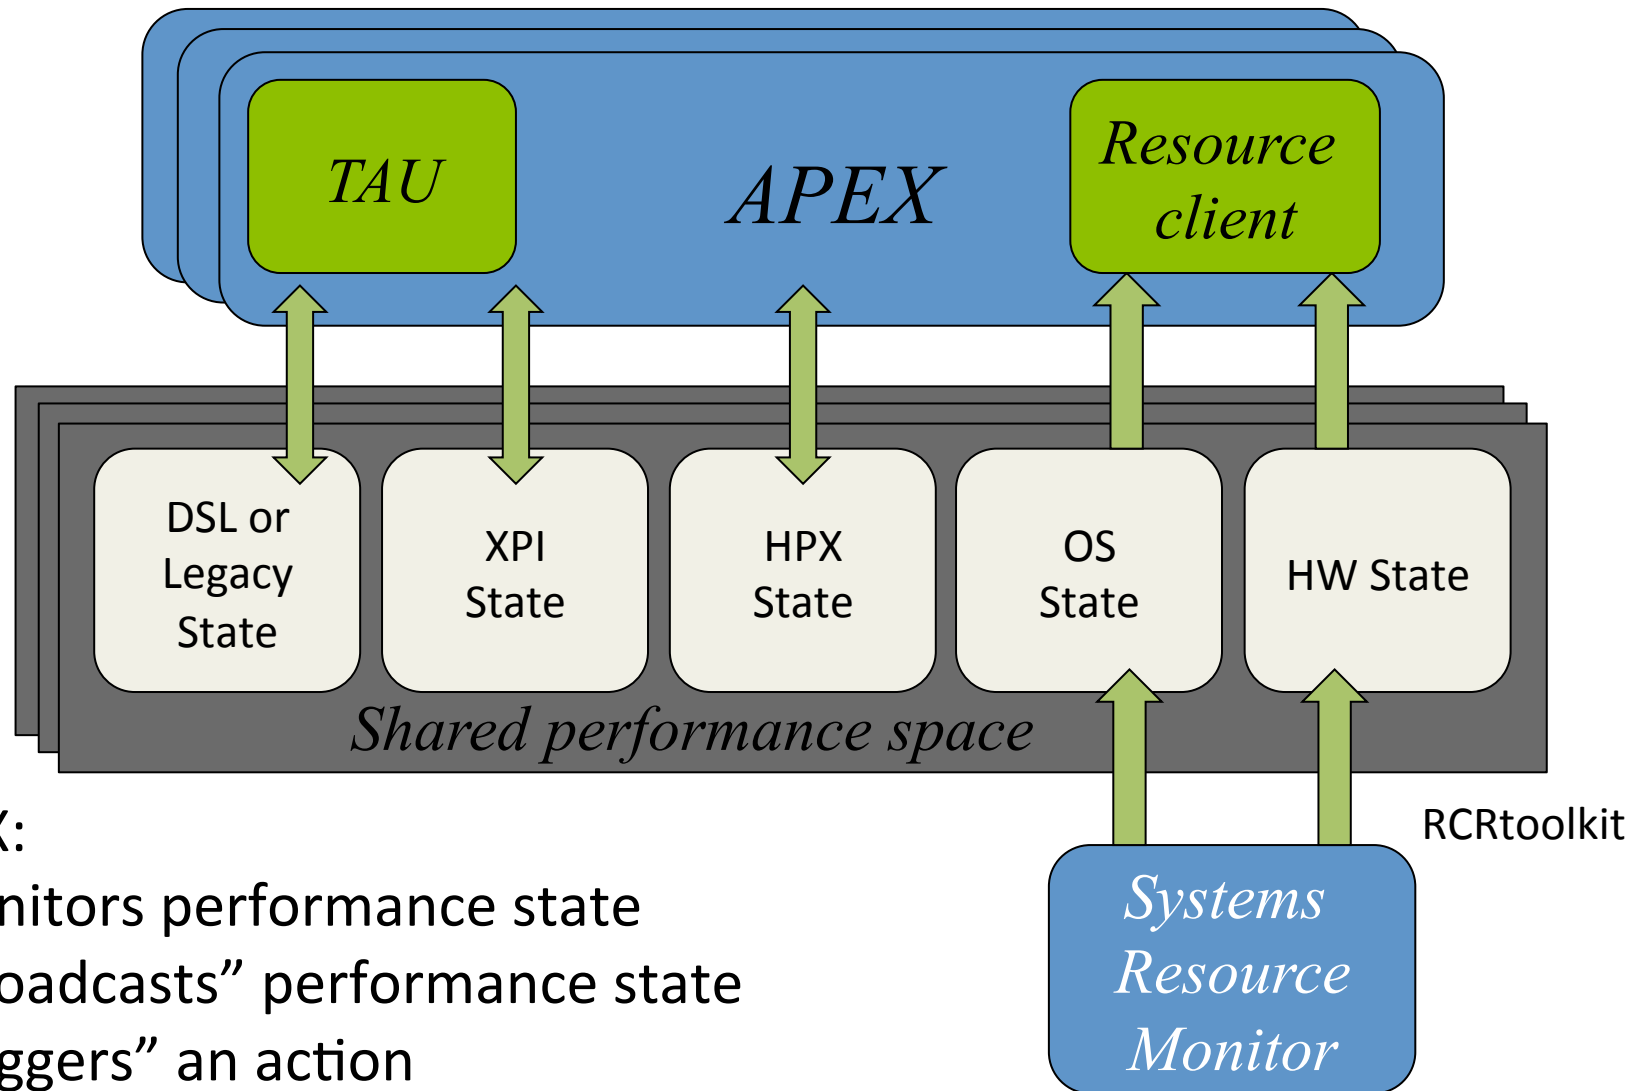
## ❑ APEX

- OS (LXK) tracks system-level resources
- Runtime (HPX) tracks threads, queues, parcels, remote ops, memory, concurrency
- XPI allows, allow language-level performance semantics to be measured






# *APEX Prototype Approach*



APEX:

- monitors performance state
- “broadcasts” performance state
- “triggers” an action

# *APEX Prototyping with HPX-3 and TAU*



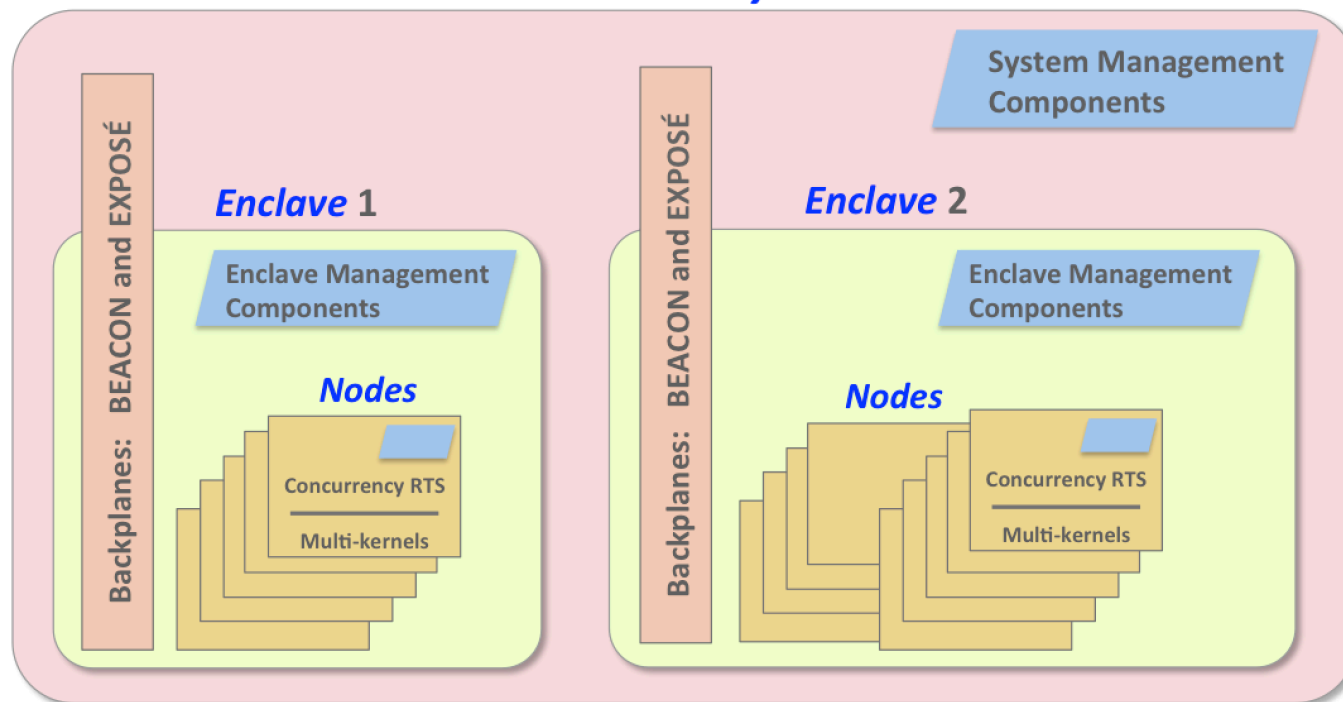
# *Argo DOE ExaOSR Project*

- ❑ Exascale OS and runtime research project
- ❑ Team
  - Labs: ANL, LLNL, PNL
  - Universities: BU, UC, UIUC, UO, UTK
- ❑ Philosophy
  - Whole-system view
    - ◆ dynamic user environment (functionality, dynamism, flexibility)
    - ◆ first-class managed resources (performance, power, ...)
    - ◆ hierarchical response to faults
  - Massive concurrency support
- ❑ Key ideas relevant to ESPT
  - Hierarchical (control, communication, goals, data resolution)
  - Embedded (performance, power, ...) feedback and response
  - Global system support

# Argo Exascale System View

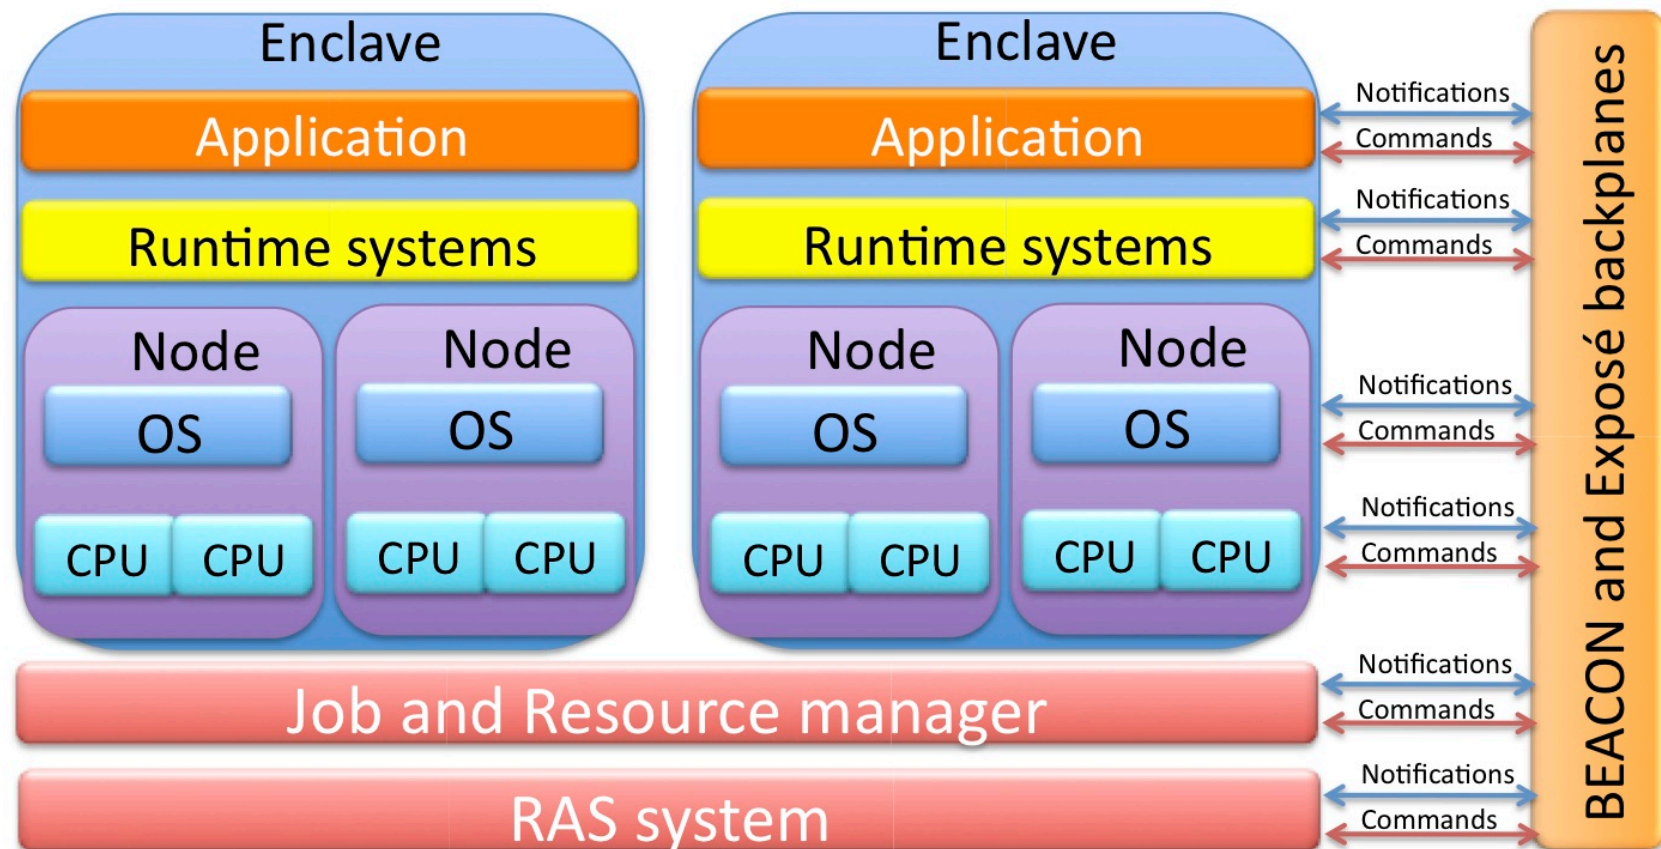
- ❑ Node OSR
- ❑ Lightweight runtime for concurrency
- ❑ Event, control, and performance backplane
- ❑ Global optimization / global view

Exascale *System*

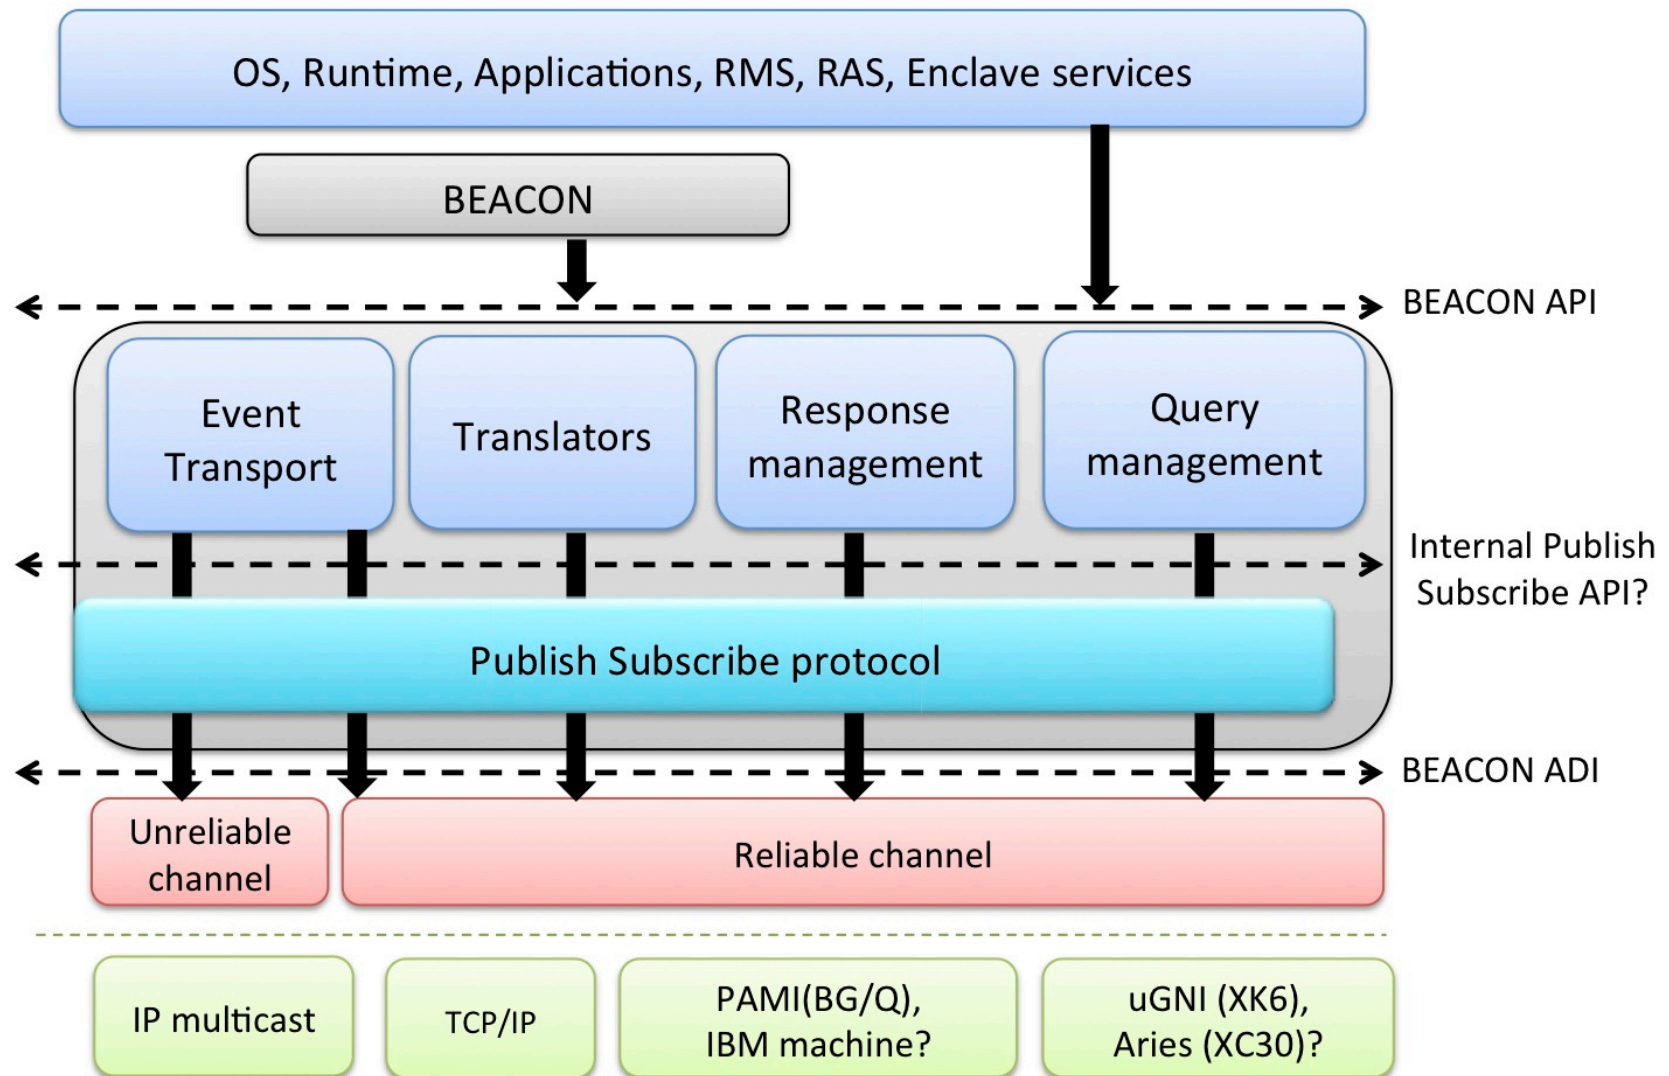


# *Argo Global Information Backplane*

- ❑ BEACON: event/action/control notification
- ❑ Exposé: performance observability system

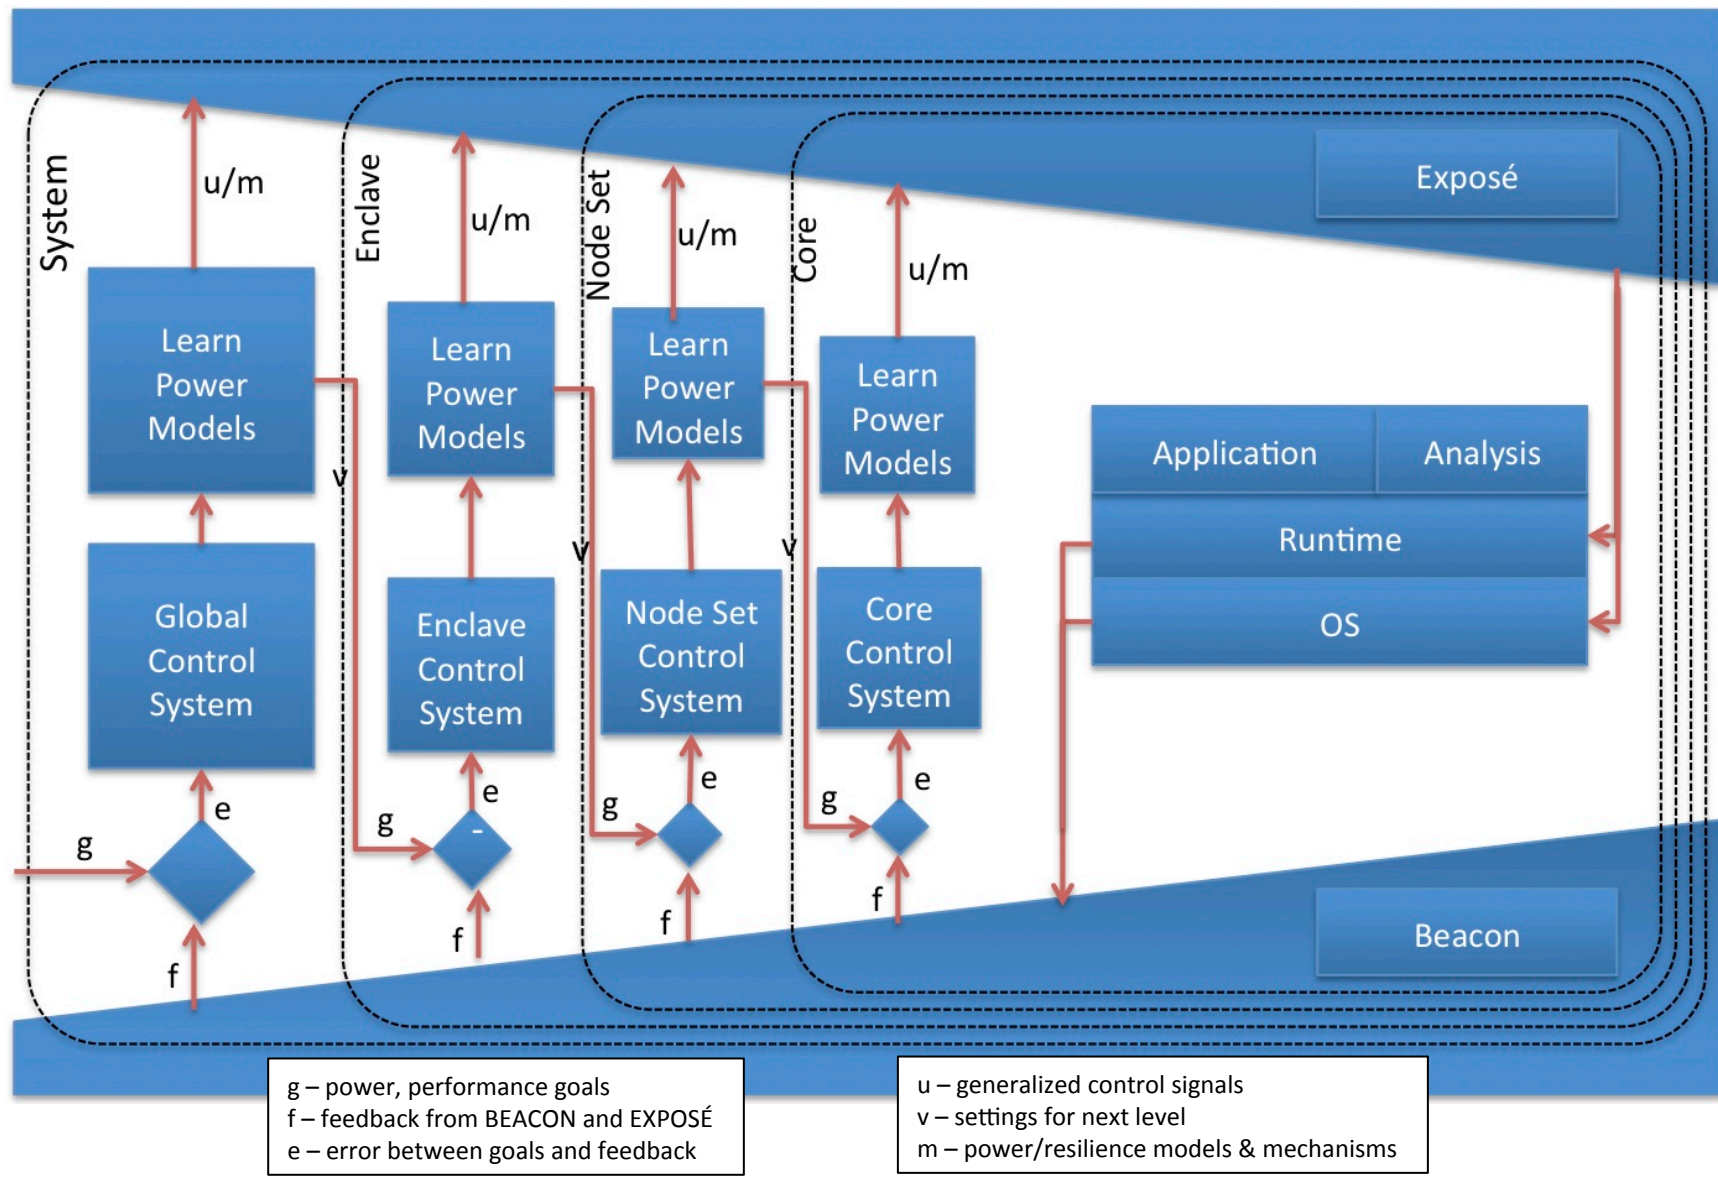


# BEACON Architecture





# Global Optimization View



# Research Projects

- ❑ *PRIMA-X* (DOE ASCR, UO, Research Centre Juelich, Aachen) E (R)
  - Refactoring core performance measurement infrastructure
- ❑ *Vancouver* (DOE X-Stack, ORNL, UO, GT, UIUC) (in renewal) E
  - Heterogeneous performance measurement (with GPUs)
- ❑ *SUPER* (DOE SciDAC Institute, many partners) E
  - Performance tools integration for end-to-end analysis, autotuning, energy, and resilience
- ❑ *GlassBox* (NSF OCI SI<sup>2</sup> program, GT, UH, UO) E (R)
  - Open interactions in HPC toolchain (compiler, runtime, I/O)
  - Integrated performance observation and knowledge sharing
- ❑ *XPRESS* (DOE X-Stack, SNL, IU, UNC, UH, UO, LSU) R
  - Integrated software ecosystem for ParalleX
  - Online, adaptive optimization and control
- ❑ *Argo* (DOE ExaOSR, ANL, PNNL, LLNL, UO, UC, UIUC) R
  - OS and runtime technology for exascale
  - Scalable signaling, introspection, feedback

Evolution  
Revolution



# *Performance as Collective Behavior*

- ❑ New performance observability paradigm can be thought of as providing *collective awareness* for exascale
- ❑ It allows the exascale system to be *performance-aware* and *performance-reactive*, able to observe performance state wholistically and holistically and to couple it with application knowledge for self-adaptive, runtime control
- ❑ Exascale systems might be considered to then be functioning as an self-organizing performance collective whose behavior is a natural consequence of seeking highly efficient operational forms optimizing exascale objectives

