

Collectives Pattern

Parallel Computing

CIS 410/510

Department of Computer and Information Science



UNIVERSITY OF OREGON

Reduce

Summing an Array

```
long sum = 0;
for (int n = 0, n < size; ++n) {
    sum += arr[n];
}
```

Fully sequential

How do we do it in parallel?

Divide and conquer!

❑ Basic algorithm:

- Split array into two parts
- Sum each part (in parallel)
- Return the sum of the two parts

❑ Why does this work?

❑ How can we generalize?

Summing an Array Recursively

```
long sum(int *arr, size_t size) {  
    if (size == 0) return 0;  
    else return arr[0] + sum(&arr[1], size-1);  
}
```

Back to being sequential...

Right to Left Reduction

We can generalize the sum example!

- Work over a triple (T, R, h, z)
 - T is the type of elements in the array
 - R is the type we want to return
 - h is a function $T * R \rightarrow R$...that is, it takes a T and an R and returns an R
 - z is an element of type R
- We apply h to the first element of the array, together with the reduction of the rest of the array
- Right to left – computes the sum over the rest of the array first

C++ Right Reduction

```
template<class T, class R, class H> R reduce(H h, R z, T* arr, size_t size)
{
    if (size == 0) return z;
    else return h(arr[0], reduce(h, z&arr[1], size-1));
}
```

Generalization!

```
sum(arr, size) = reduce([](int n, long m) {return n + m}, 0, arr, size)
```

Note that we need the z: it is how we handle empty arrays.

What about the parallel sum?

We could make sum happen in parallel, can we do this in general?

$$\text{reduce}(h, z, [1, 2, 3, 4, 5], 4) = h(1, h(2, h(3, h(4, h(5, z))))))$$



The parallel sum would be more like:

$$h(h(h(1, 2), h(3, 4)), h(5, z))$$

This might compute something totally different. It might not even have the same types (requires $T = R$).

What is going on?

- ❑ Right reduce is *list* structured
- ❑ A “parallel reduce” has to be *tree* structured



Sum Simple Algebra

□ Associativity Property

- $(x + y) + z = x + (y + z)$
- Examples: addition, multiplication

□ Commutative Property

- $x + y = y + z$
- Examples: addition, multiplication of integers
- Counter Examples: multiplication of matrices

Monoids

A monoid is a triple $(S, +, 0)$, where S is a set

- S is a set
- $+$ is an operation on S (not necessarily addition)
- 0 is an element of S (not necessarily 1)

such that

- Associativity: for all x, y, z in S

$$(x + y) + z = x + (y + z)$$

- Identity: for all x in S

$$0 + x = x + 0 = x$$

- Not necessarily commutative

Many Many Monoids

- ❑ $(\mathbb{Z}, +, 0)$ where \mathbb{Z} is the integers and $+$, 0 are the usual meaning of those
- ❑ $(\mathbb{R}, +, 0)$ where \mathbb{R} is the reals
- ❑ $(\mathbb{Z}, *, 1)$, that is, multiplication of integers
- ❑ (Real Valued 2-2 matrices, matrix multiplication, the identity matrix)
→
- ❑ For any set S : $(\{f \mid f : S \rightarrow S\}, \text{function composition, the identity function})$
- ❑ (Strings over an alphabet, string concatenation, empty string)
- ❑

Many Many Monoids

- ❑ $(\mathbb{Z}, +, 0)$ where \mathbb{Z} is the integers and $+$, 0 are the usual meaning of those
- ❑ $(\mathbb{R}, +, 0)$ where \mathbb{R} is the reals
- ❑ $(\mathbb{Z}, *, 1)$, that is, multiplication of integers
- ❑ (Real Valued 2-2 matrices, matrix multiplication, the identity matrix)
- ❑ For any set S : $(\{f \mid f : S \rightarrow S\}, \text{function composition, the identity function})$
- ❑ (Strings over an alphabet, string concatenation, empty string)
- ❑

(If you have taken abstract algebra: every **group** is a monoid)

Arbitrary Order Reduction

Performance of Parallel Reduction

Looks like:

- Work: $O(n)$
- Span: $O(n \cdot \log(n))$


But slight nuance...

These numbers are not *the number of calls to worker function* (+) **not** the total time

We will come back to this!

Map/Reduce Fusion

- ❑ Map then fuse pattern very common
- ❑ Fuse them!



Example: Dot Product

- 2 vectors of same length
- Map (*) to multiply the components
- Then reduce with (+) to get the final answer

$$\mathbf{a} \cdot \mathbf{b} = \sum_{i=0}^{n-1} a_i b_i.$$

Dot Product in TBB

```
1 float tbb_sprod(
2     size_t n,
3     const float *a,
4     const float *b
5 ) {
6     return tbb::parallel_reduce(
7         tbb::blocked_range<size_t>(0,n),
8         float(0),
9         [=]( // lambda expression
10             tbb::blocked_range<size_t>& r,
11             float in
12         ) {
13             return std::inner_product(
14                 a+r.begin(), a+r.end(),
15                 b+r.begin(), in );
16         },
17         std::plus<float>()
18     );
19 }
```

Merge Sort as a reduction

Define the monoid $(S, \langle \rangle, [])$ where

S is the set of in order vectors over some type

$\langle \rangle$ is the merge operation: $[1,3,5,7] \langle \rangle [2,6,15] = [1,2,3,5,6,7,15]$

$[]$ is the empty list

We can sort an array via a pair of a map and a reduce

Map each element into a vector containing just that element

Reduce using the monoid above

How fast is this?

Right Biased Sort

Start with [14,3,4,8,7,52,1]

Map to [[14],[3],[4],[8],[7],[52],[1]]

Reduce:

$$\begin{aligned} & [14] \triangleleft ([3] \triangleleft ([4] \triangleleft ([8] \triangleleft ([7] \triangleleft ([52] \triangleleft [1]))))) \\ &= [14] \triangleleft ([3] \triangleleft ([4] \triangleleft ([8] \triangleleft ([7] \triangleleft [1,52])))) \\ &= [14] \triangleleft ([3] \triangleleft ([4] \triangleleft ([8] \triangleleft [1,7,52]))) \\ &= [14] \triangleleft ([3] \triangleleft ([4] \triangleleft [1,7,8,52])) \\ &= [14] \triangleleft ([3] \triangleleft [1,4,7,8,52]) \\ &= [14] \triangleleft [1,3,4,7,8,52] \\ &= [1,3,4,7,8,14,52] \end{aligned}$$

Right Biased Sort Cont

- ❑ How long did that take?
- ❑ Well we did $O(n)$ merges...but each one took $O(n)$ time
- ❑ $O(n^2)$
- ❑ We wanted merge sort, but instead we got insertion sort!

Tree Shape Sort

Start with [14,3,4,8,7,52,1]

Map to [[14],[3],[4],[8],[7],[52],[1]]

Reduce:

$$\begin{aligned} & (([14] \triangleleft [3]) \triangleleft ([4] \triangleleft [8])) \triangleleft ((([7] \triangleleft [52]) \triangleleft [1])) \\ &= ([3,14] \triangleleft [4,8]) \triangleleft ([7,52] \triangleleft [1]) \\ &= [3,4,8,14] \triangleleft [1,7,52] \\ &= [1,3,4,7,8,14,52] \end{aligned}$$

Tree Shaped Sort Performance

- ❑ Even if we only had a single processor this is better
 - We do $O(\log n)$ merges
 - Each one is $O(n)$
 - So $O(n * \log(n))$
- ❑ But opportunity for parallelism is not so great
 - $O(n)$ assuming sequential merge
- ❑ We will explore parallel merging later in the class
- ❑ Takeaway: the shape of reduction matters!

Shape Matters

- ❑ Often tree based reductions are slower
- ❑ In fact, any reduction can be rewritten using monoids...but the cost of the operation **won't be a constant**
- ❑ Parallel reduction is suitable for problems where you have
 - Cost free associativity (like adding number)
 - Or tree based reduction is better (like merging)

Scan

Prefix Scans

What if instead of summing a list we wanted to compute all the **prefix sums**?

```
int * results = malloc (sizeof(int)*size);
```

```
int temp = 0;
```

```
for (int n = 0, n < size, ++n) {
```

```
    temp += arr[n];
```

```
    results[n] = temp;
```

```
}
```

$[1,2,3,4,5] \rightarrow [1,3,6,10,15]$

Semi-Groups

A pair $(S, +)$ where

S is a set



$+$ is a binary operation on S

- Is a semi-group if $+$ is associative
- Unlike monoids, we don't require identity
- Every monoid is a semi-group, but not vice versa!

Scans

- ❑ Scan operation generalizes prefix sum
- ❑ Parameterized by a semi-group
- ❑ Scan is often the trick to turn algorithm that seems to have **sequential data dependencies** into a parallel algorithm

Scans



Binary Addition

Given two bit vectors of length n , compute their sum

$$1010 + 0111 = 10001$$

Standard algorithm you learned in elementary school

- Add from least significant digit to most significant digit
- Keep track of carrying
- Very sequential: $O(n)$

Can we do better?

Binary Addition of Single Bits

Single bit addition

$$0 + 0 \rightarrow 00$$

$$0 + 1 \rightarrow 01$$

$$1 + 0 \rightarrow 01$$

$$1 + 1 \rightarrow 10$$

Note that we produce carries.

Binary Addition of Single Bits

Single bit addition

$$0 + 0 \rightarrow 00$$

$$0 + 1 \rightarrow 01$$

$$1 + 0 \rightarrow 01$$

$$1 + 1 \rightarrow 10$$

Clever Carrying

Single bit addition assumes we don't have a carry

What if we do?

Single bit addition assuming a carry

$$0 + 0 \rightarrow 01$$

$$0 + 1 \rightarrow 10$$

$$1 + 0 \rightarrow 10$$

$$1 + 1 \rightarrow 11$$

Note that we produce carries.

Clever Carrying Part 2

Single bit addition behavior depends on if we had a carry!

Single bit with and without carries

$$0 + 0 \rightarrow 00, 01$$

$$0 + 1 \rightarrow 01, 10$$

$$1 + 0 \rightarrow 01, 10$$

$$1 + 1 \rightarrow 10, 11$$

Clever Carrying part 3

We can classify these pairs in terms of

- “generates a carry” (G)
- “preserves the carry given to it, but does not produce one” (P)
- “does not produce a carry” (N)

Single bit carrying behavior:

$$0 + 0 \rightarrow N$$

$$0 + 1 \rightarrow P$$

$$1 + 0 \rightarrow P$$

$$1 + 1 \rightarrow G$$

The Carrying Semi-Group

The set $\{P, N, G\}$ forms a semi-group under the operation $\langle \rangle$

$$\circ P \langle \rangle x \rightarrow x$$

$$\circ G \langle \rangle _ \rightarrow G$$

$$\circ N \langle \rangle _ \rightarrow N$$

Here $_$ means “any element”

Multi Bit Carrying

We can expand this to multiple bits

$11 + 01$ always generates a carry (G)

$01 + 01$ never generates a carry (N)

$01 + 10$ preserves a carry (P)

$AB + CD$ works in the following way:

$A + B \rightarrow G$ means $AB + CD \rightarrow G$

$A + B \rightarrow N$ means $AB + CN \rightarrow N$

$A + B \rightarrow P$ means $AB + CD$ does the same thing as
ever $B + D$

$AB + CD$

Computing the Carry

If you want to compute the carry behavior of the sum of binary numbers

$$1010100101 + 0101011100$$

We can do it first computing the carries pairwise

$$PPPPPPGNP$$

And then using the semi-group structure

$$\begin{aligned} & ((P \triangleleft P) \triangleleft (P \triangleleft P)) \triangleleft (((P \triangleleft P) \triangleleft (P \triangleleft G)) \triangleleft (N \triangleleft P)) \\ &= (P \triangleleft P) \triangleleft ((P \triangleleft G) \triangleleft N) = P \triangleleft (G \triangleleft N) = P \triangleleft G \\ &= G \end{aligned}$$

Thus we know this example produces a carry.

Look Ahead Carry Addition

- ❑ We can take this idea to produce an algorithm for binary addition
- ❑ We write all numbers “backwards” (least significant digit to most significant)
- ❑ Use map to compute a {P, N, G} array
- ❑ Scan over this array with $\langle \rangle$ producing an array C
- ❑ Use map to compute the result

```
parallel_for(int n = 1, n < len, ++n) {  
    if (R[n-1] == G) R[n] = A[n]^B[n];  
    else R[n] = ! (A[n]^B[n]);  
}
```

Scan Performance

- ❑ Sequential scan calls the function $O(n)$ times
- ❑ Parallel scan is calls the function
 - Span: $O(\log n)$
 - Work: $O(n * \log n)$
- ❑ Unlike other patterns we have to do extra work to be parallel
- ❑ But, scan is broadly applicable. If you can't make an algorithm parallel, scan is often a good place to look