

Fork-Join Pattern

Parallel Computing

CIS 410/510

Department of Computer and Information Science



UNIVERSITY OF OREGON

Fibonacci

- Recursive Fibonacci is simple and inefficient

```
long fib ( int n ) {  
    if (n < 2) return 1;  
    else {  
        long x = fib (n-1);  
        long y = fib (n-2);  
        return x + y;  
    }  
}
```

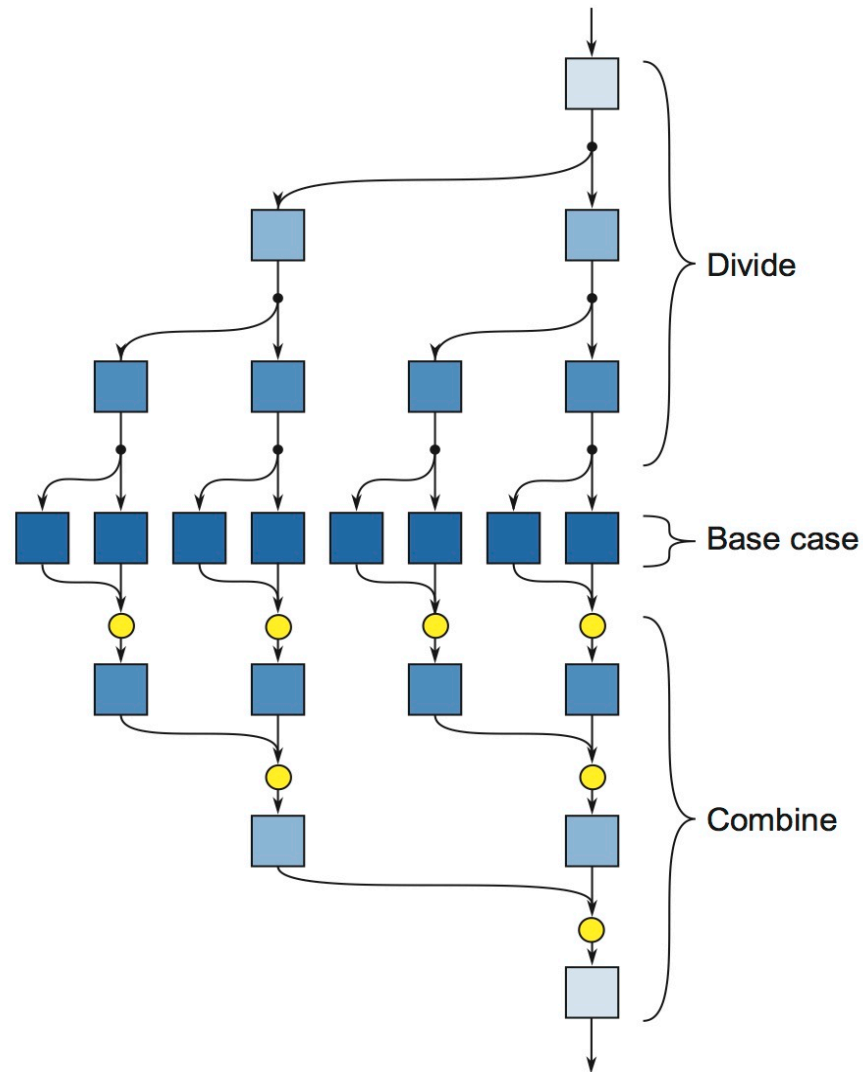
Fibonacci...

- ❑ Recursive Fibonacci is simple and inefficient
- ❑ But it does have the property that the sub-calls are independent
- ❑ Can we parallelize it?

Fibonacci...in Parallel?

```
long fib ( int n ) {  
    if (n < 2) return 1;  
    else {  
        long x = fork fib (n-1);  
        long y = fib(n-2);  
        join;  
        return x + y;  
    }  
}
```

Recursive Fork Join



Fork Join


- ❑ Simple Idea for Concurrency
 - “fork” new tasks
 - “join to delay execution until forked tasks have finished

Fork Join Control Flow

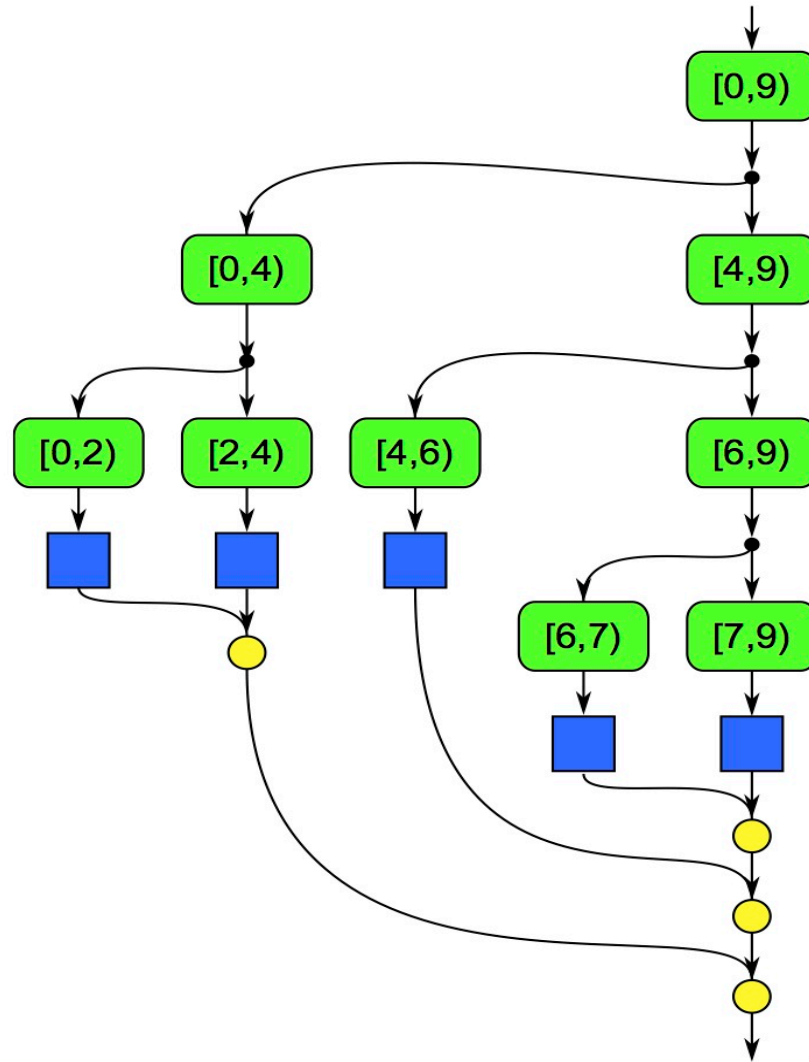
```
fork B ( ) ;
```

```
C ( ) ;
```

```
join
```



Executing Map as Fork/Join



Work Stealing

- ❑ The runtimes for TBB/CilkPlus do something known as work stealing
- ❑ Each **worker thread** has a queue of tasks
- ❑ When a call to fork is executed, the thread puts the task on its queue
- ❑ This provides good locality...but can cause starvation
- ❑ So, if a thread runs out of work, it “steals” some tasks from a different queue

Steal Continuation vs. Steal Child

- ❑ TBB and CilkPlus handle fork in different ways
- ❑ Given code of the form:

```
fork f ( ) ;
```


```
g ( ) ;
```

```
join ;
```


- ❑ A TBB thread would put `f()` on its queue and then execute `g()`. It would only start working on the queue when it got to the join.
- ❑ A Cilk thread would put **both** `f()` and the remainder of the program (`g()`; `join`; etc) on its queue. It is probable that it will execute `f()` before `g()`.

Steal Continuation vs. Steal Child

Steal Continuation




Steal Child



Extra Forking

- ❑ Simple Idea for Concurrency
 - “fork” new tasks
 - “join” to delay execution until forked tasks have finished



Don't do this!

Performance of Fork/Join

Let $A\|B$ be interpreted as “fork A, do B, and join”

$$\text{Work: } T(A\|B)_1 = T(A)_1 + T(B)_1$$

$$\text{Span: } T(A\|B)_\infty = \max(T(A)_\infty, T(B)_\infty)$$

From these you can figure out the Work/Span of algorithms using the asymptotic analysis technique you learned in CIS 315/621